The background of the slide features a complex network graph overlaid on a satellite image of North America. The network consists of numerous small white dots representing nodes, connected by thin white lines representing edges, forming a dense web of connections across the continent.

# Practical Machine Learning with Python

@MatthewRenze

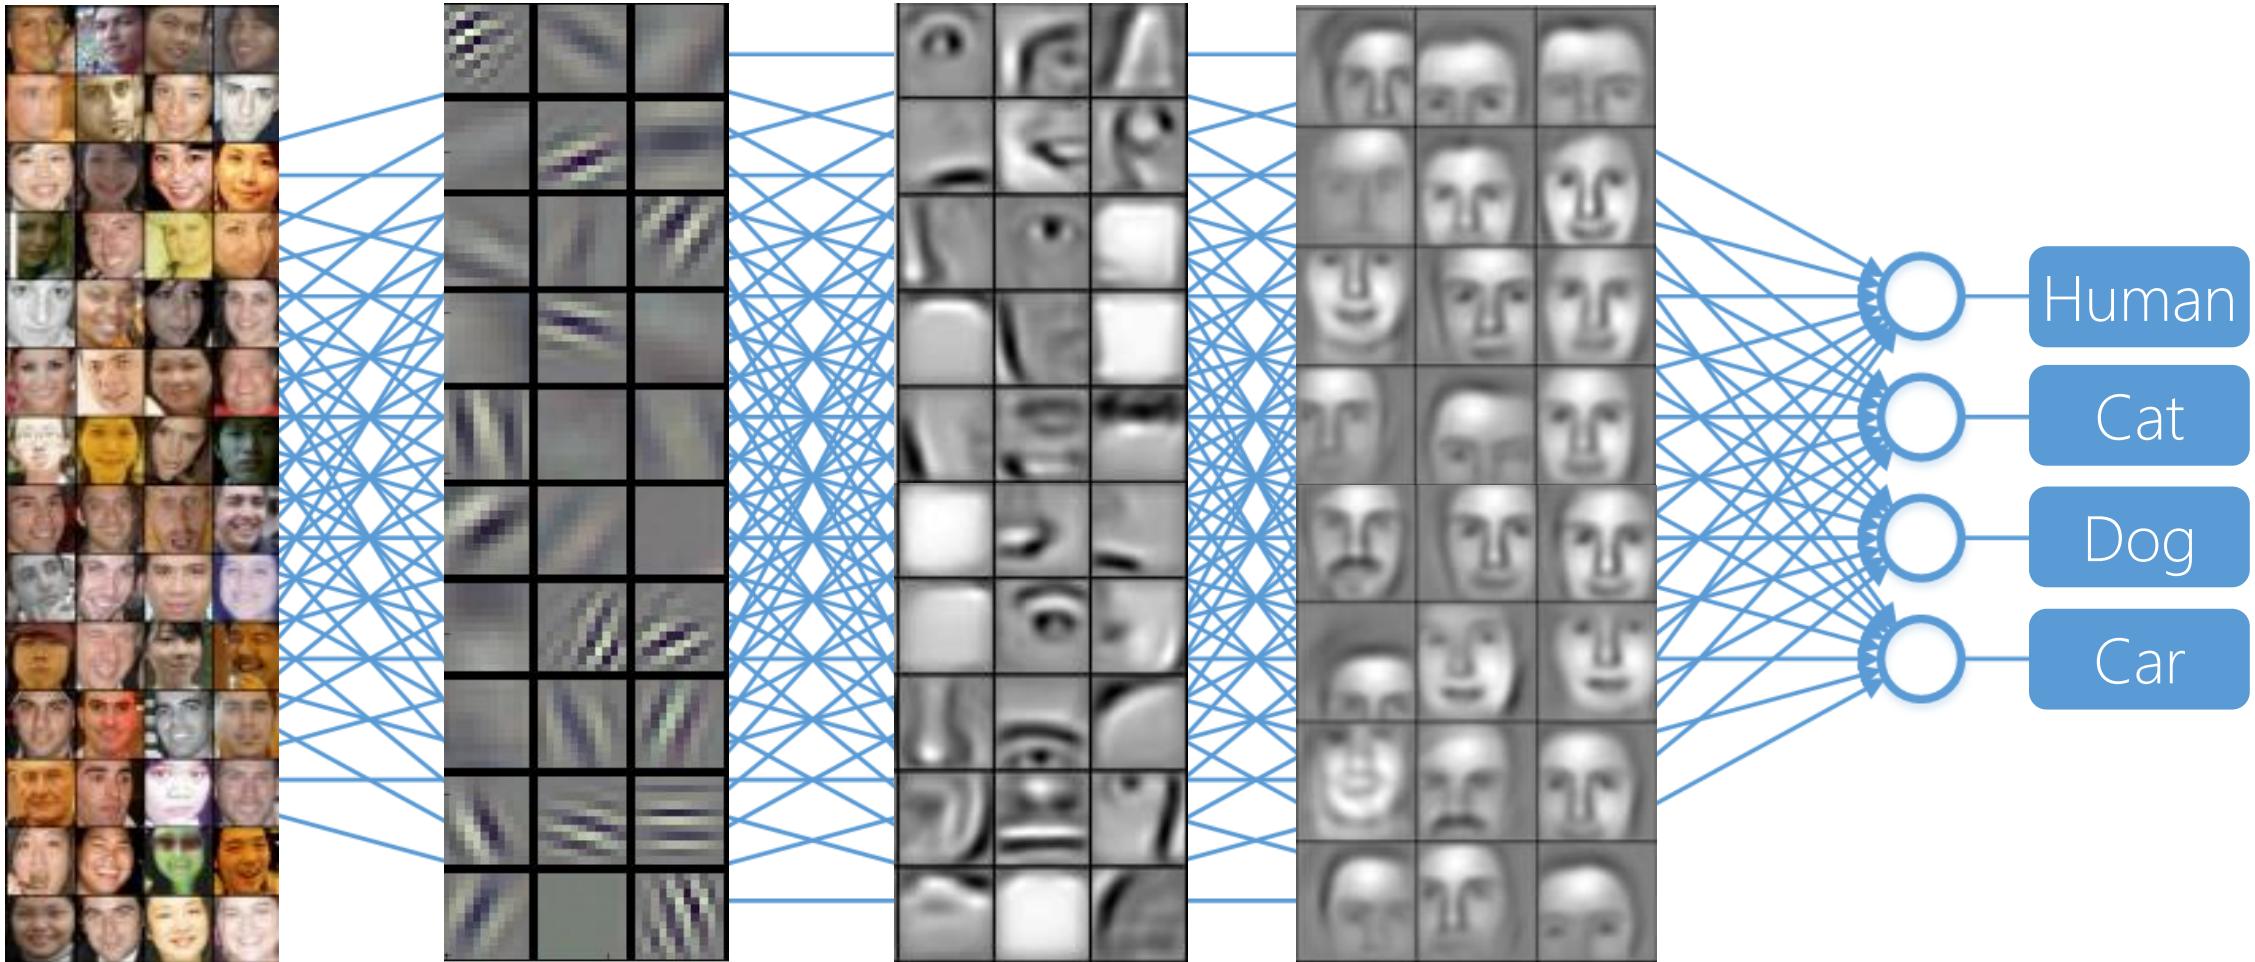
#IndyCode



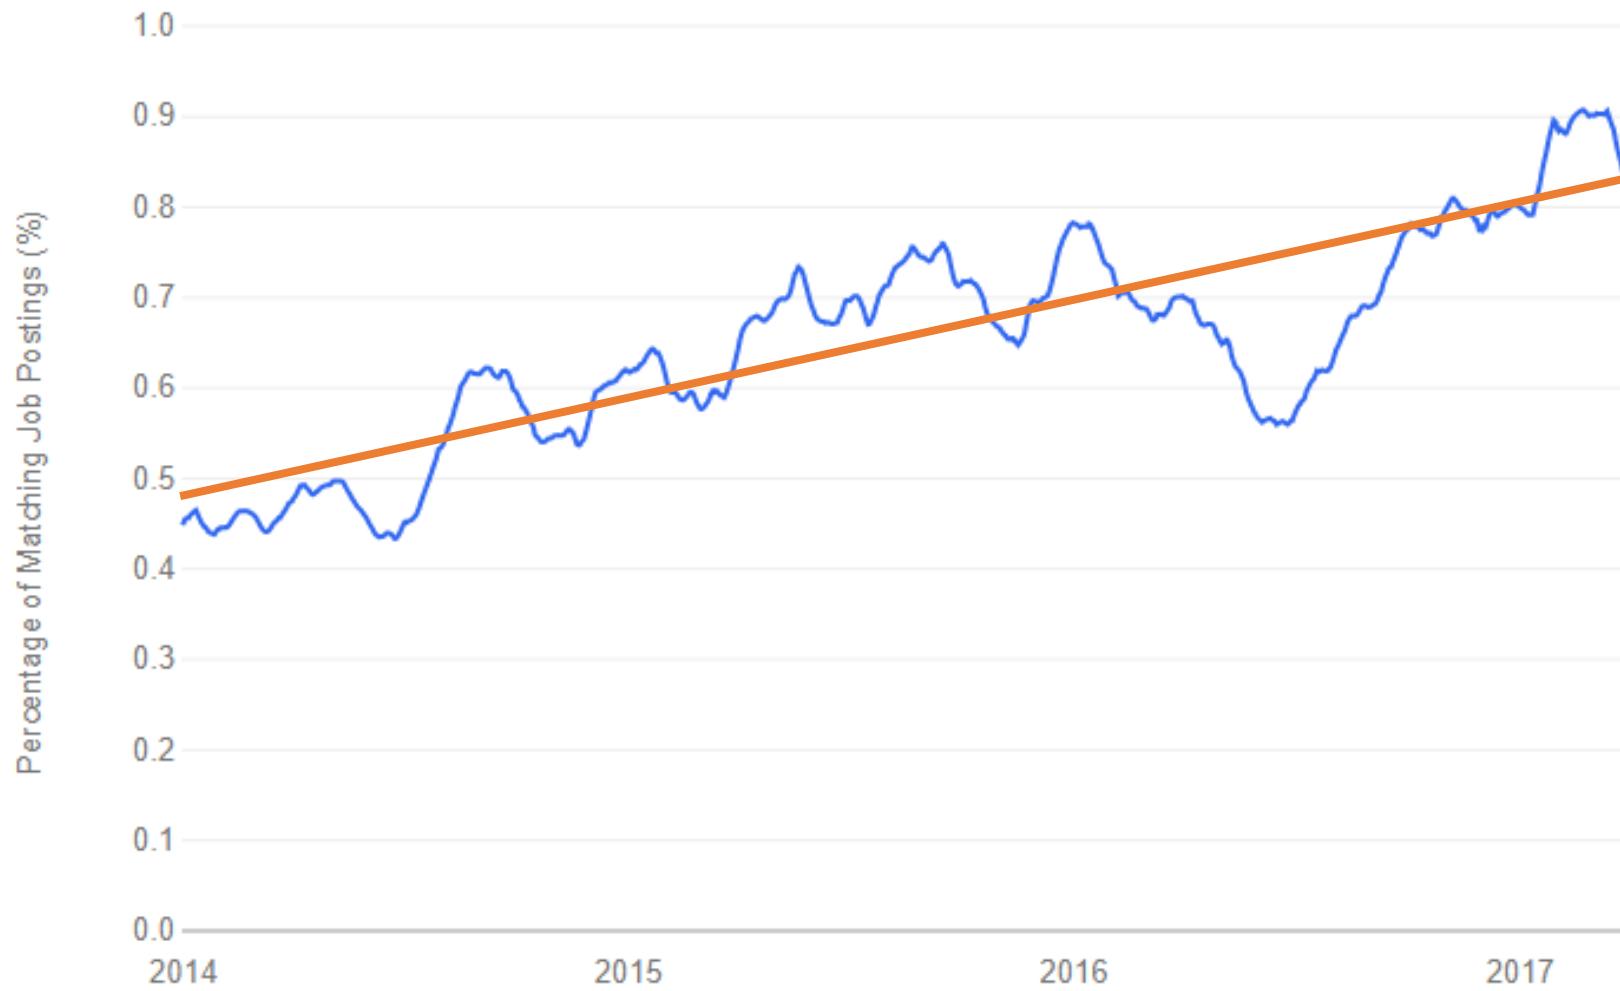


```
function updatePhotoDescription() {
    if (descriptions.length > (page * 9) + (currentImage - 1)) {
        document.getElementById('bigImageDesc').innerHTML = descriptions[currentImage - 1];
    }
}

function updateAllImages() {
    var i = 1;
    while (i < 10) {
        var elementId = 'foto' + i;
        var elementIdBig = 'bigImage' + i;
        if (page * 9 + i - 1 < photos.length) {
            document.getElementById(elementId).src = 'image/min/' + photos[i - 1];
            document.getElementById(elementIdBig).src = 'image/big/' + photos[i - 1];
        } else {
            document.getElementById(elementId).src = '';
        }
        i++;
    }
}
```

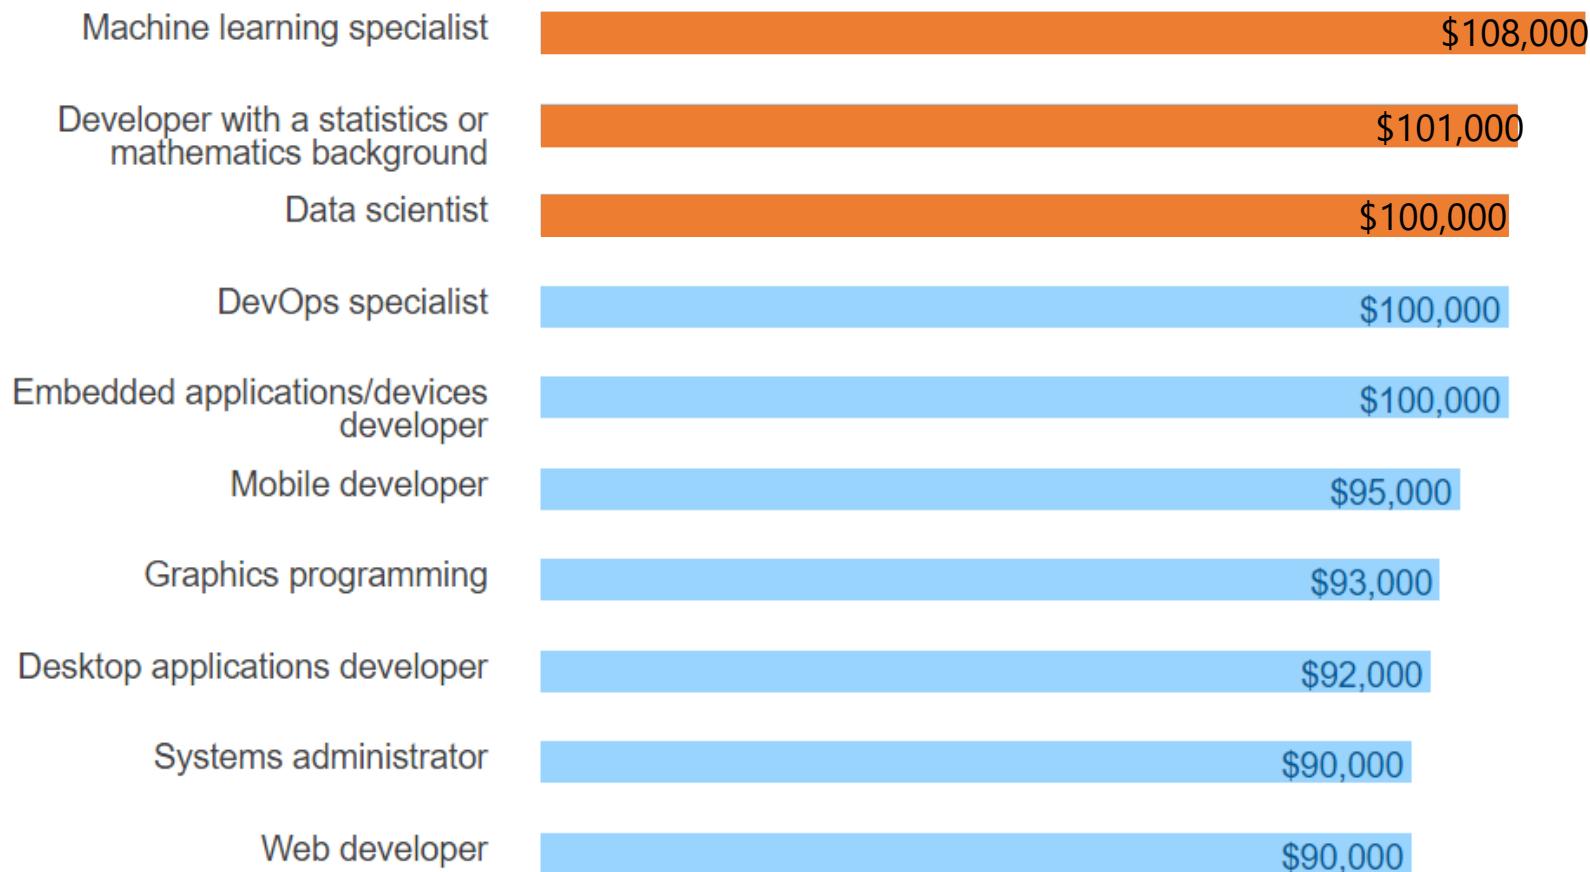


# Job Postings for Machine Learning



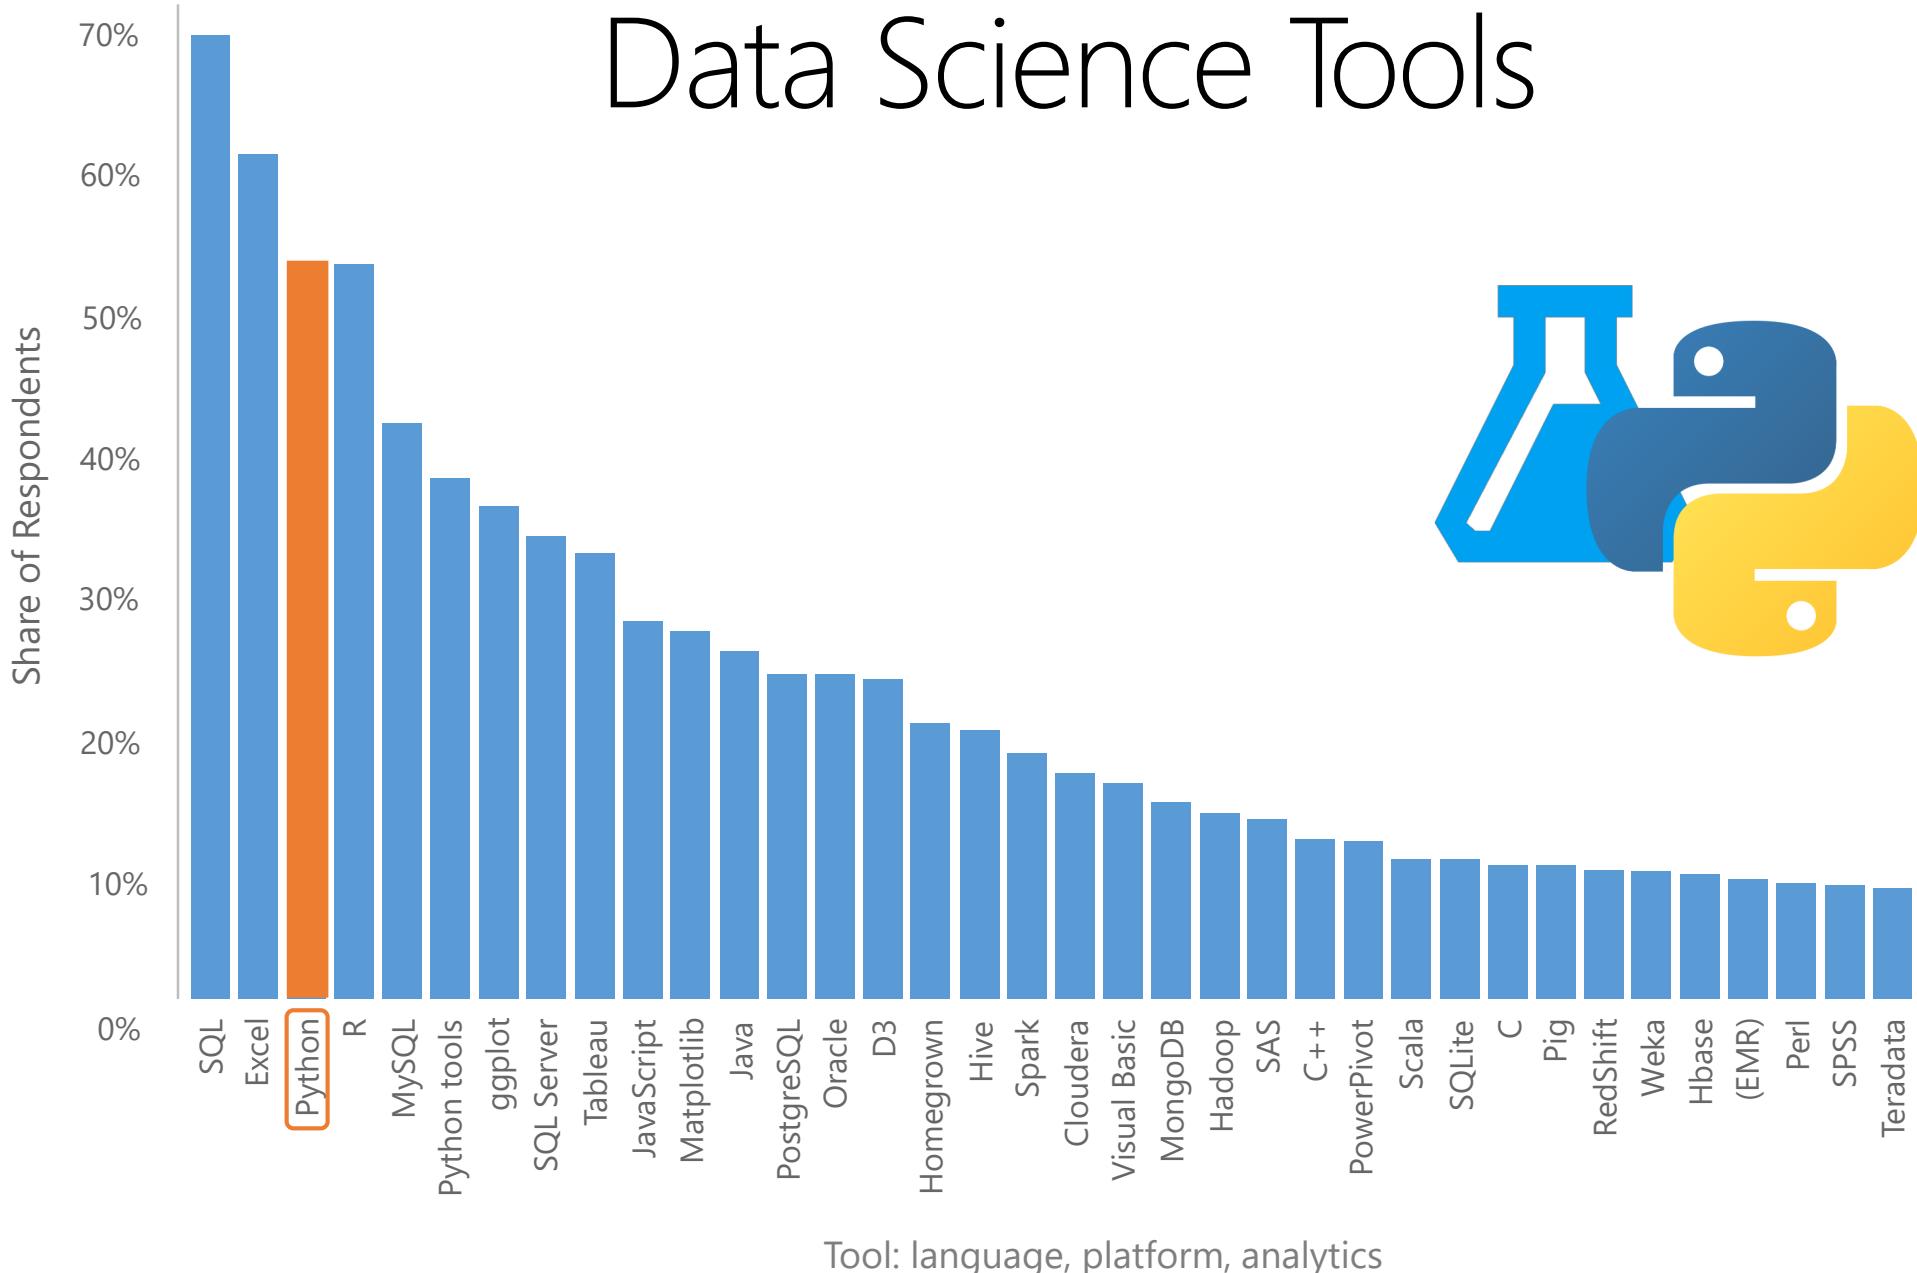
Source: Indeed.com

# Average Salary by Job Type (USA)

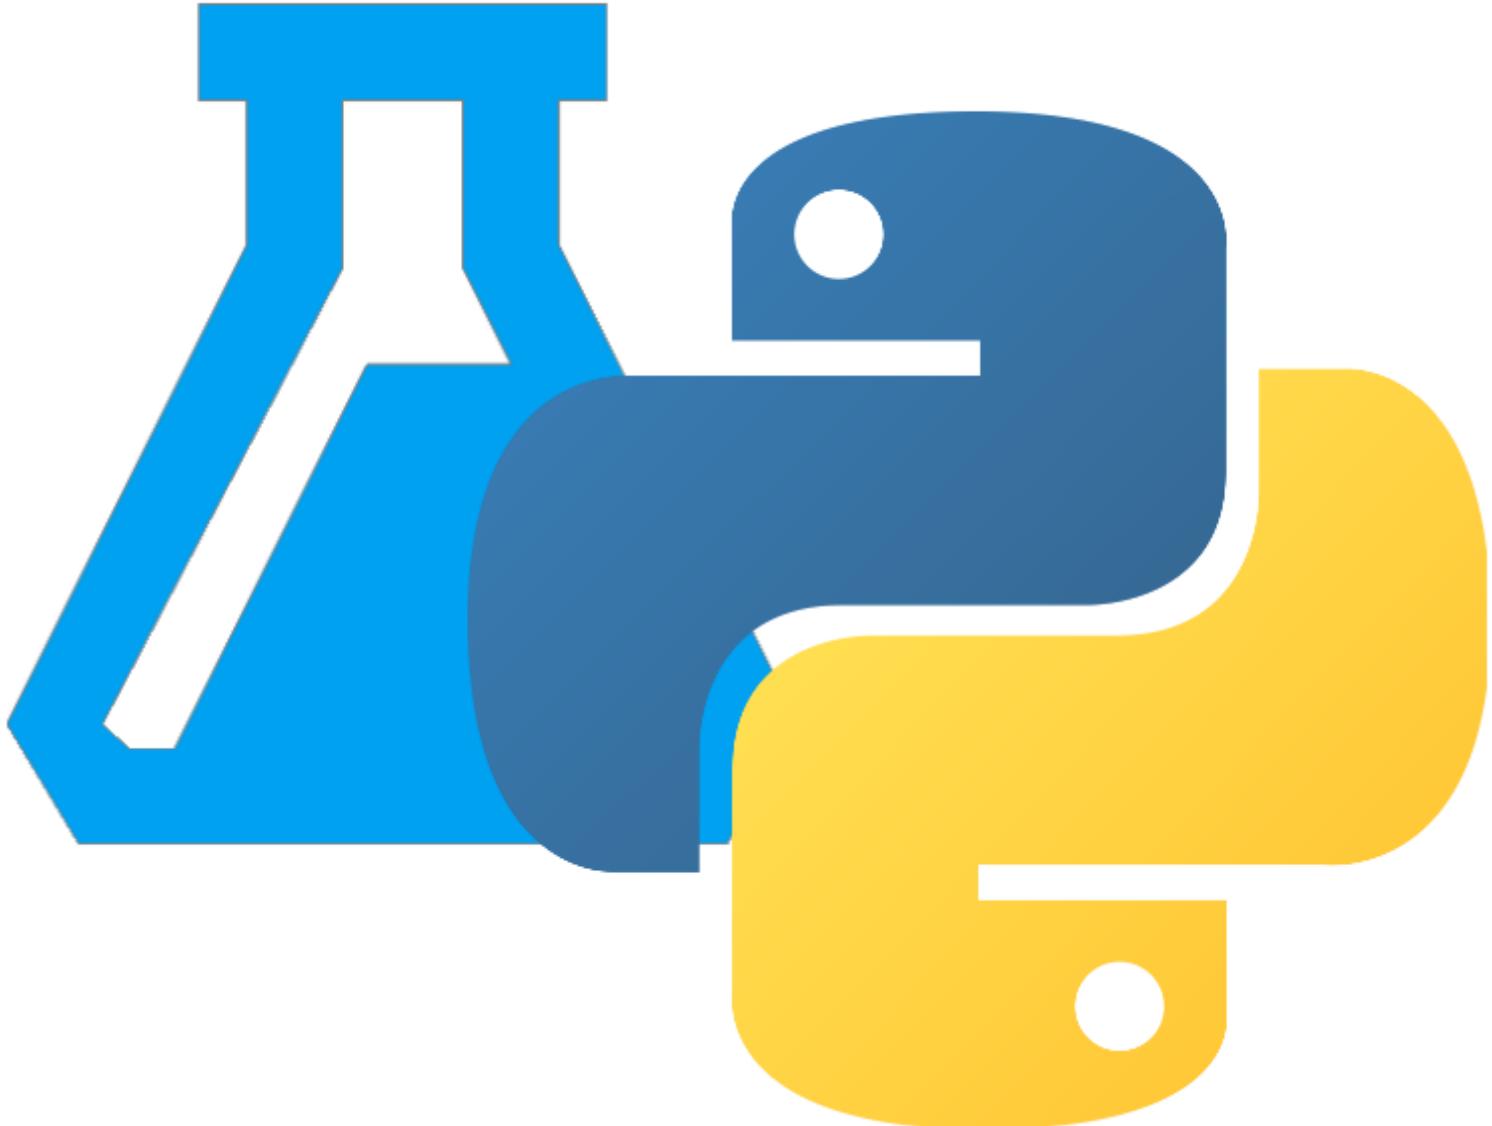


Source: Stack Overflow 2017

# Data Science Tools

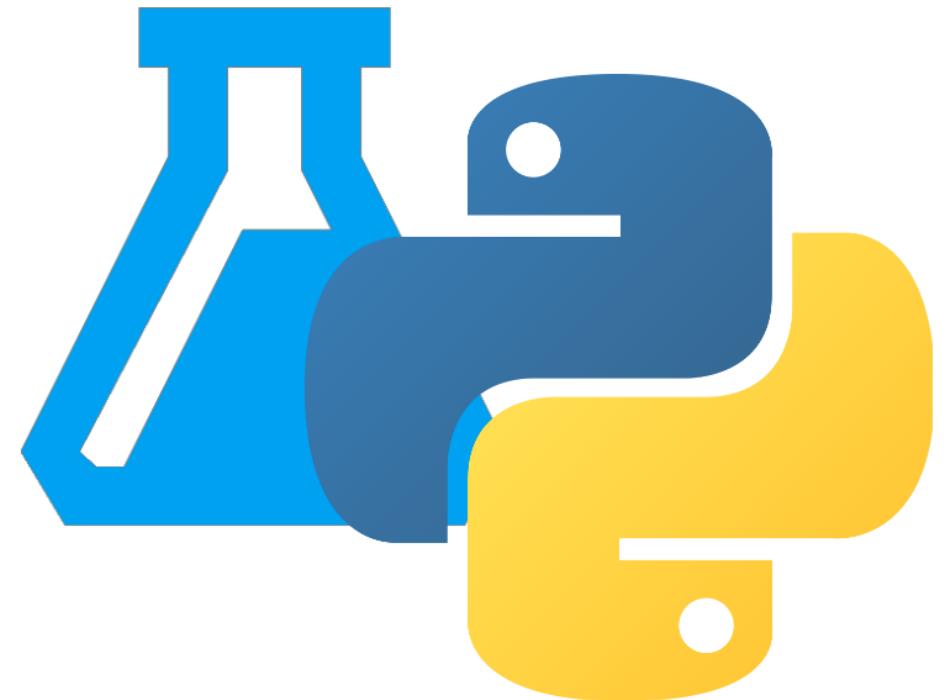


Source: O'Reilly 2015 Data Science Salary Survey



# Overview

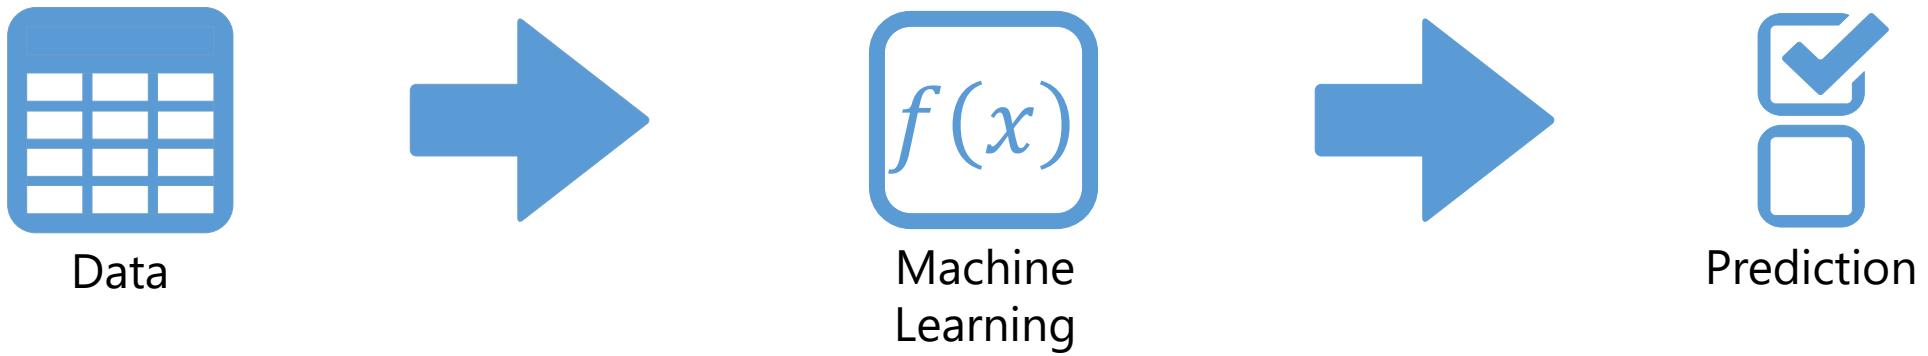
1. Intro to ML and Python
2. Classification
3. Regression
4. Clustering
5. ML in Practice



# How Does This Apply to Me?

- Make decisions using data
- Make predictions using data
- Make recommendations using data
- Automate these with code

# Conceptual Model









# About Me

Data Science Consultant  
Education

B.S. in Computer Science

B.A. in Philosophy

Data Science specializations

Community

Public speaker

Pluralsight author

Microsoft MVP

Open source

IOWA STATE  
UNIVERSITY



# Schedule

Lectures (15 min)

Demos (10 min)

Labs (30 min)

Breaks (5 min)

# Logistics

Pairing for labs is optional

Ask questions if needed

Come and go as needed

Feedback at the end

# Labs

# Labs

A  
(Easy)

# Labs

A

(Easy)

B

(Hard)

# Labs

A  
(Easy)

B  
(Hard)

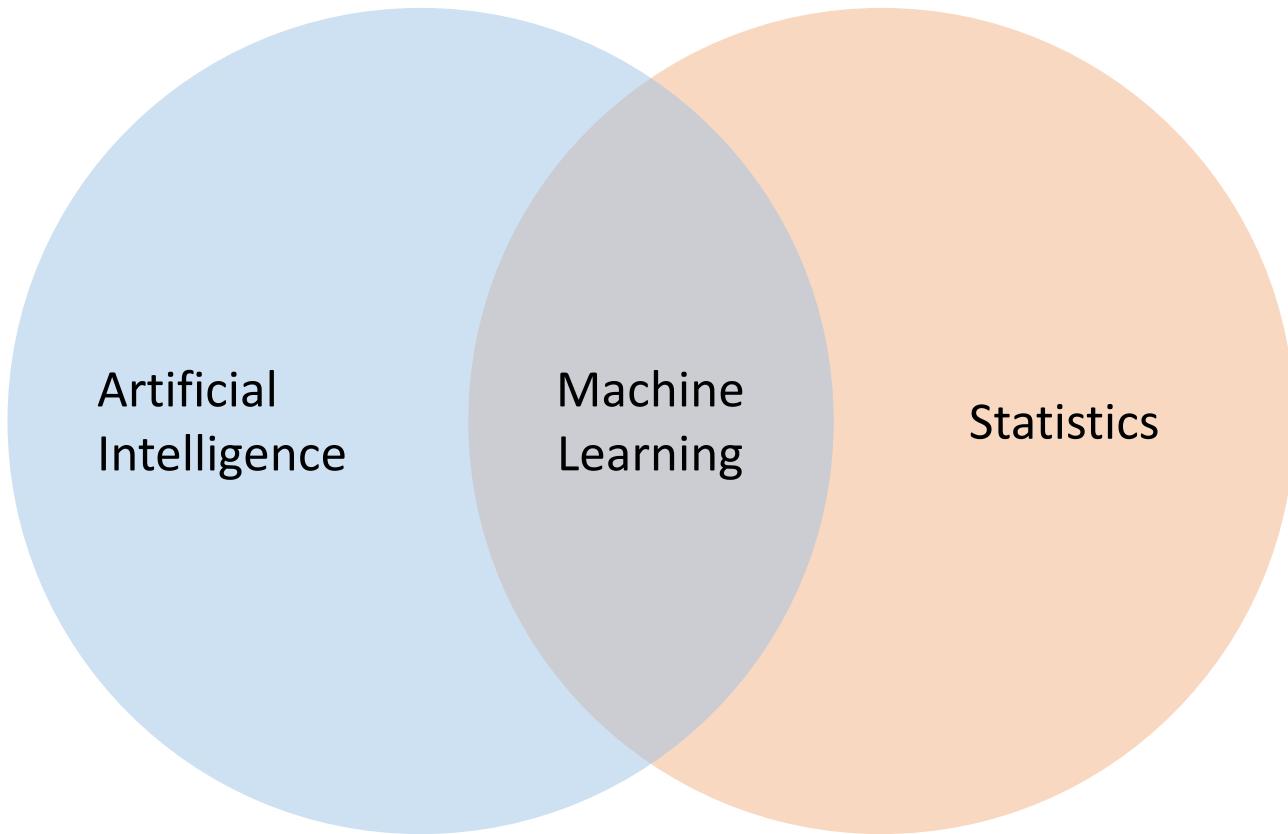


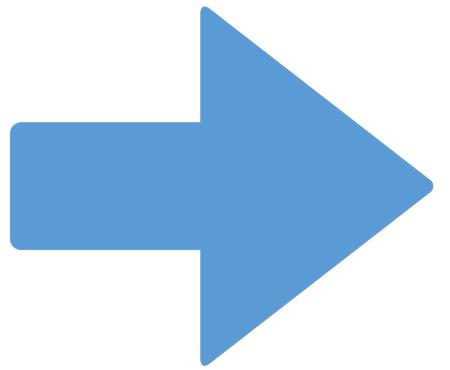
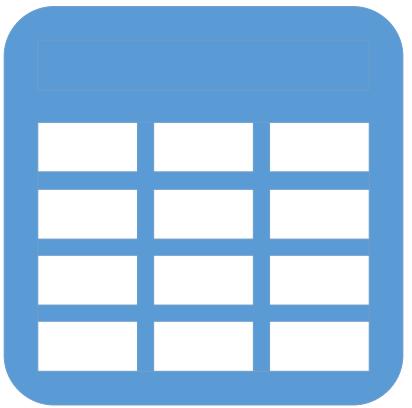
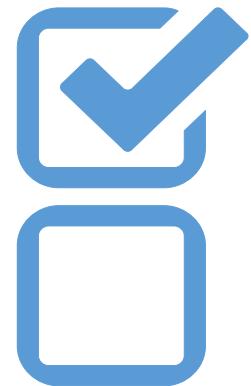
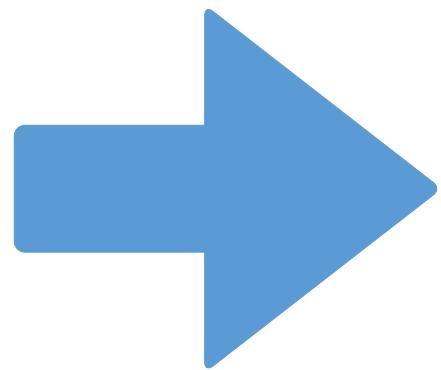
# Workshop URL

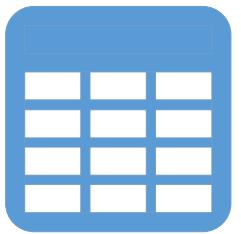
<http://www.matthewrenze.com/workshops/practical-machine-learning-with-python/>

# Introduction to Machine Learning

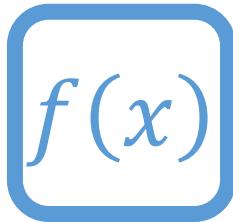
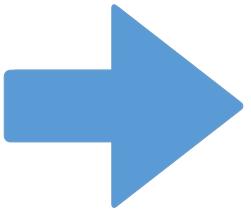
What is machine learning?



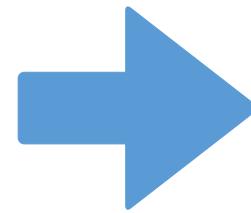
 $f(x)$ 



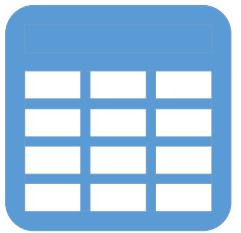
Data



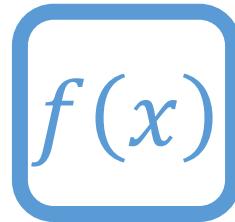
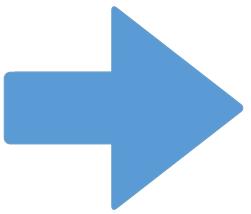
Function



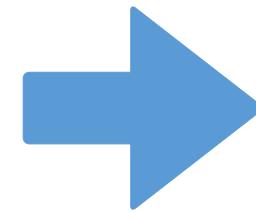
Prediction



Data

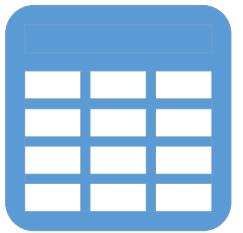


Function

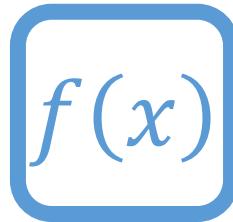
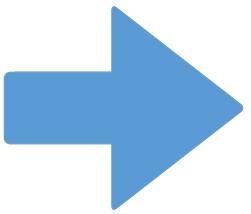


Prediction

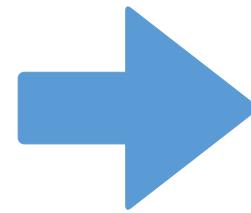




Data



Function



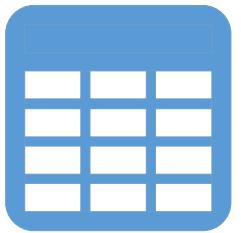
Prediction



Cat



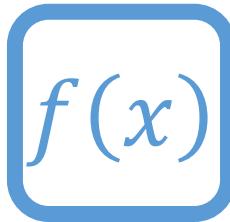
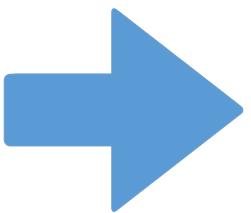
Not cat



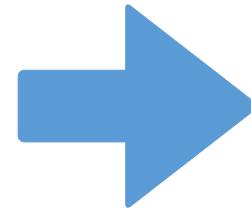
Data



Cat



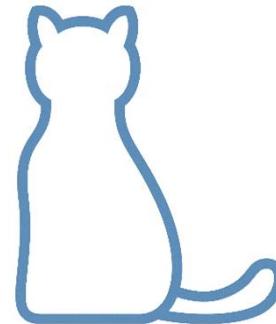
Function

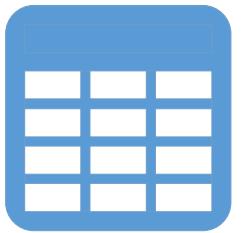


Prediction



Not cat

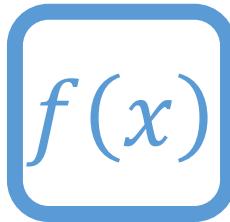
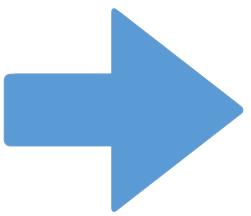




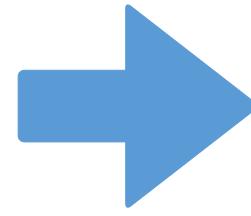
Data



Cat



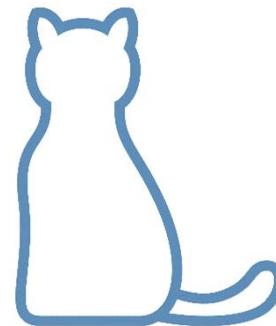
Function



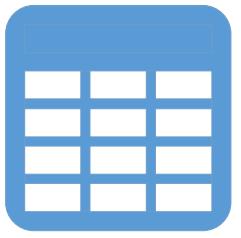
Prediction



Not cat



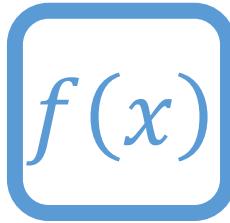
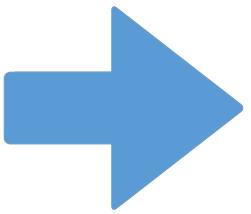
Is cat?



Data



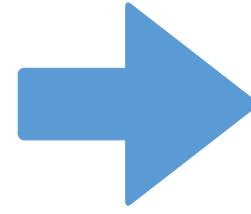
Cat



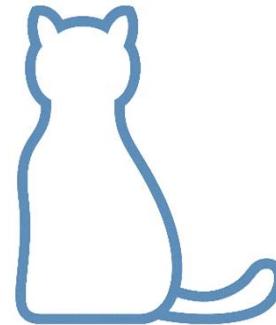
Function



Not cat

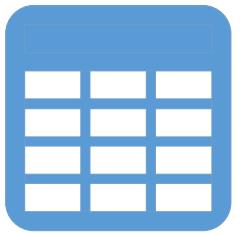


Prediction



Is cat?

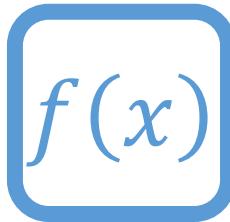
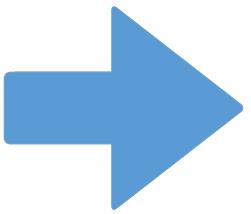




Data



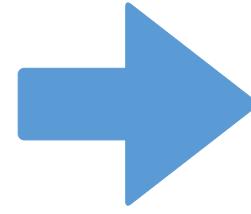
Cat



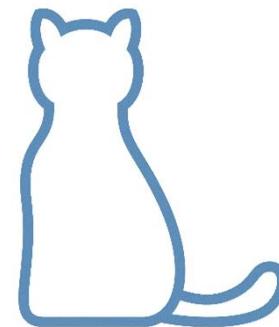
Function



Not cat



Prediction



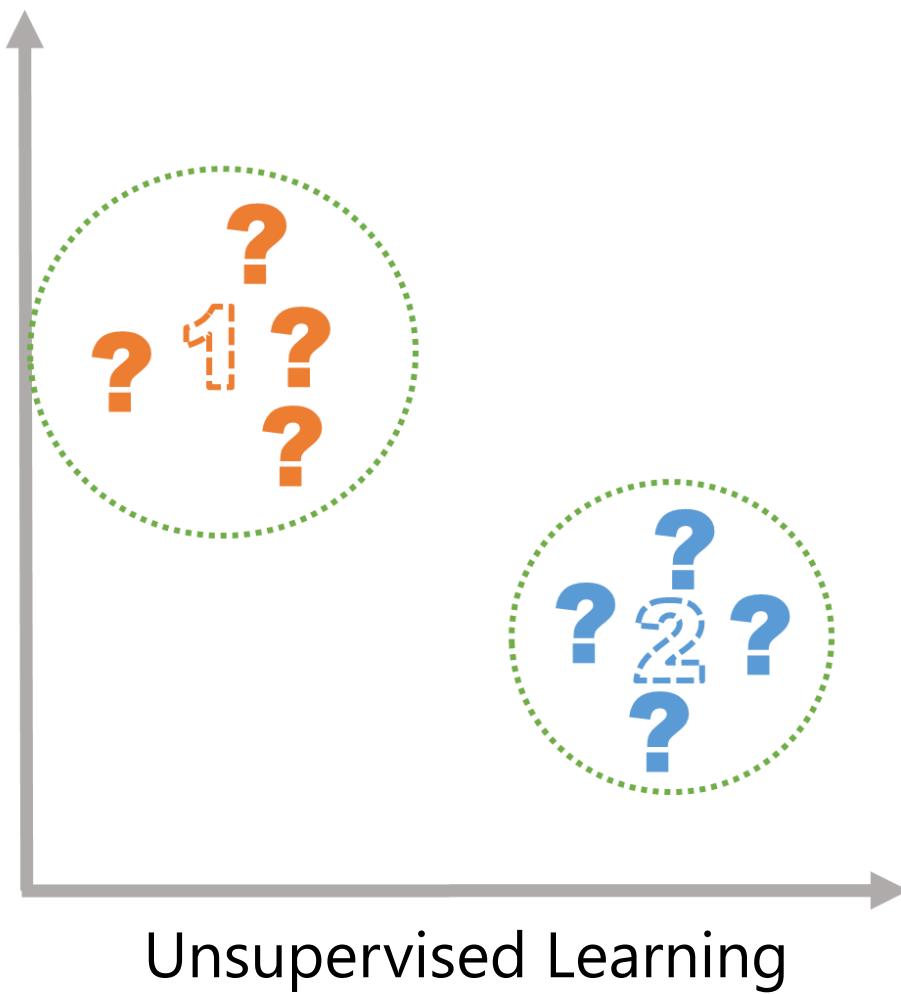
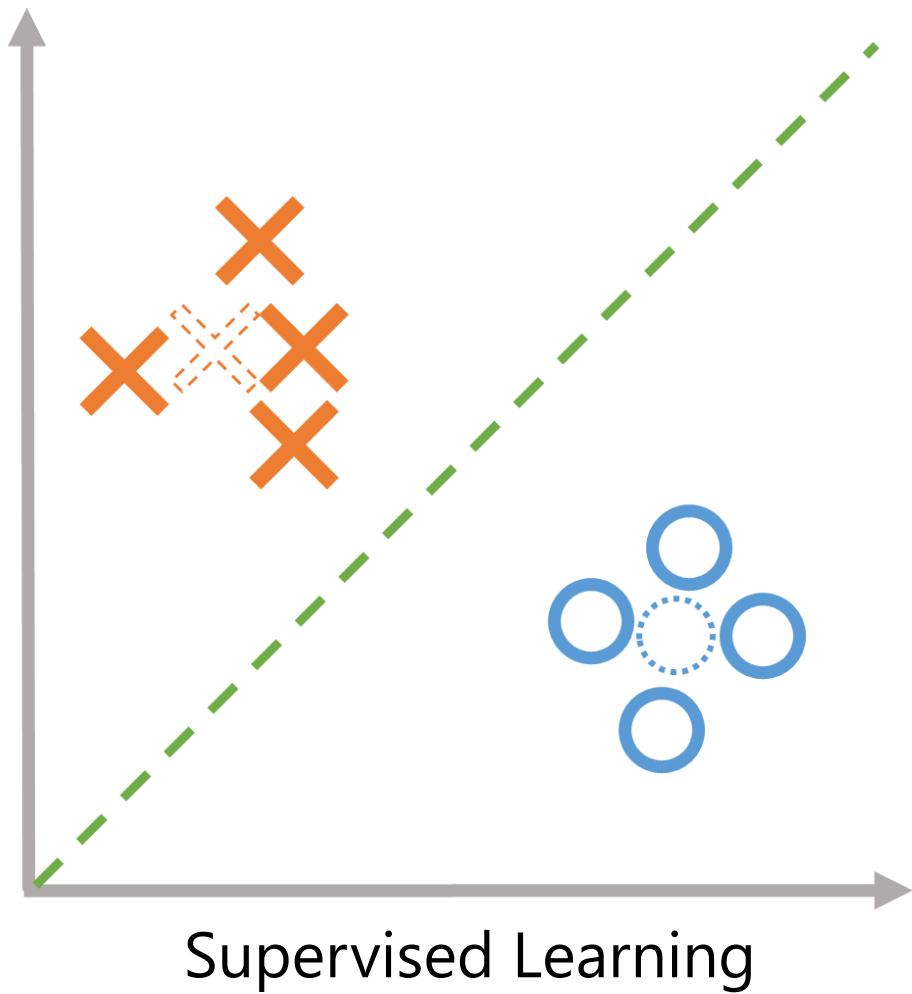
Is cat?



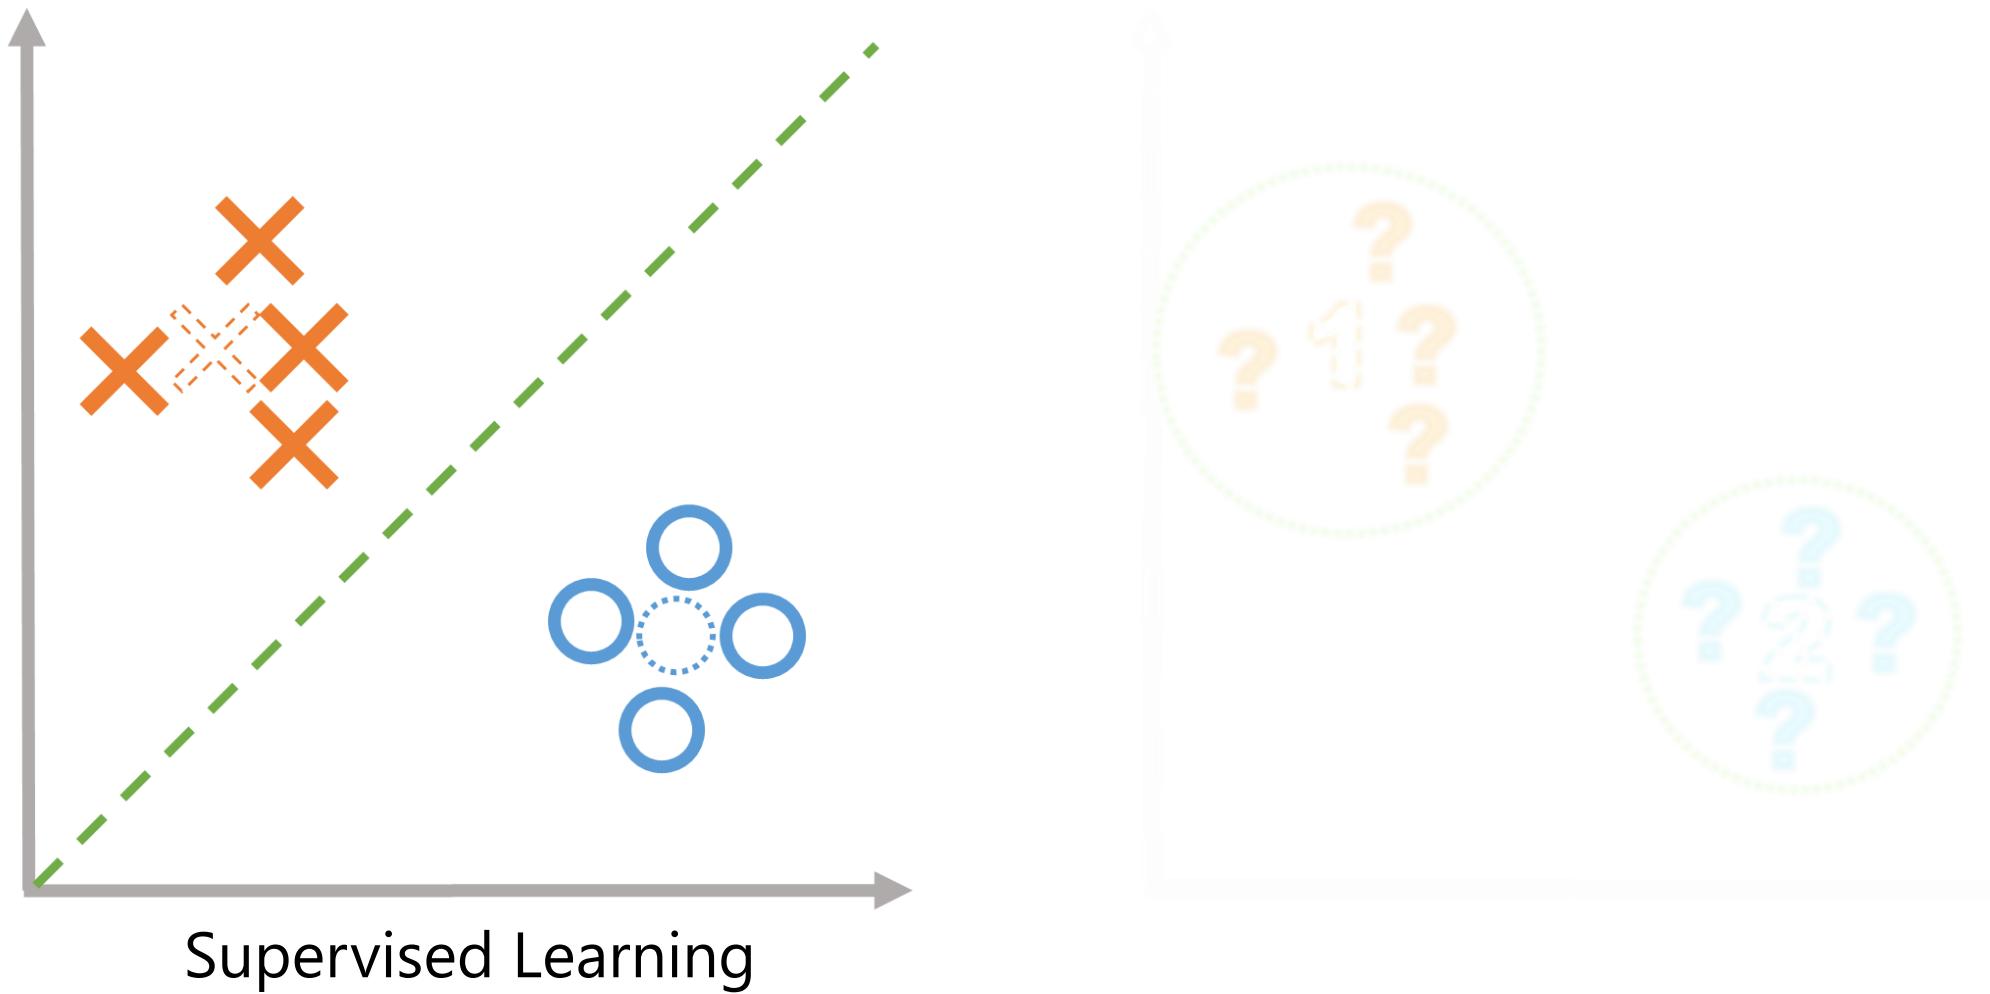
Yes

What types of machine learning exist?

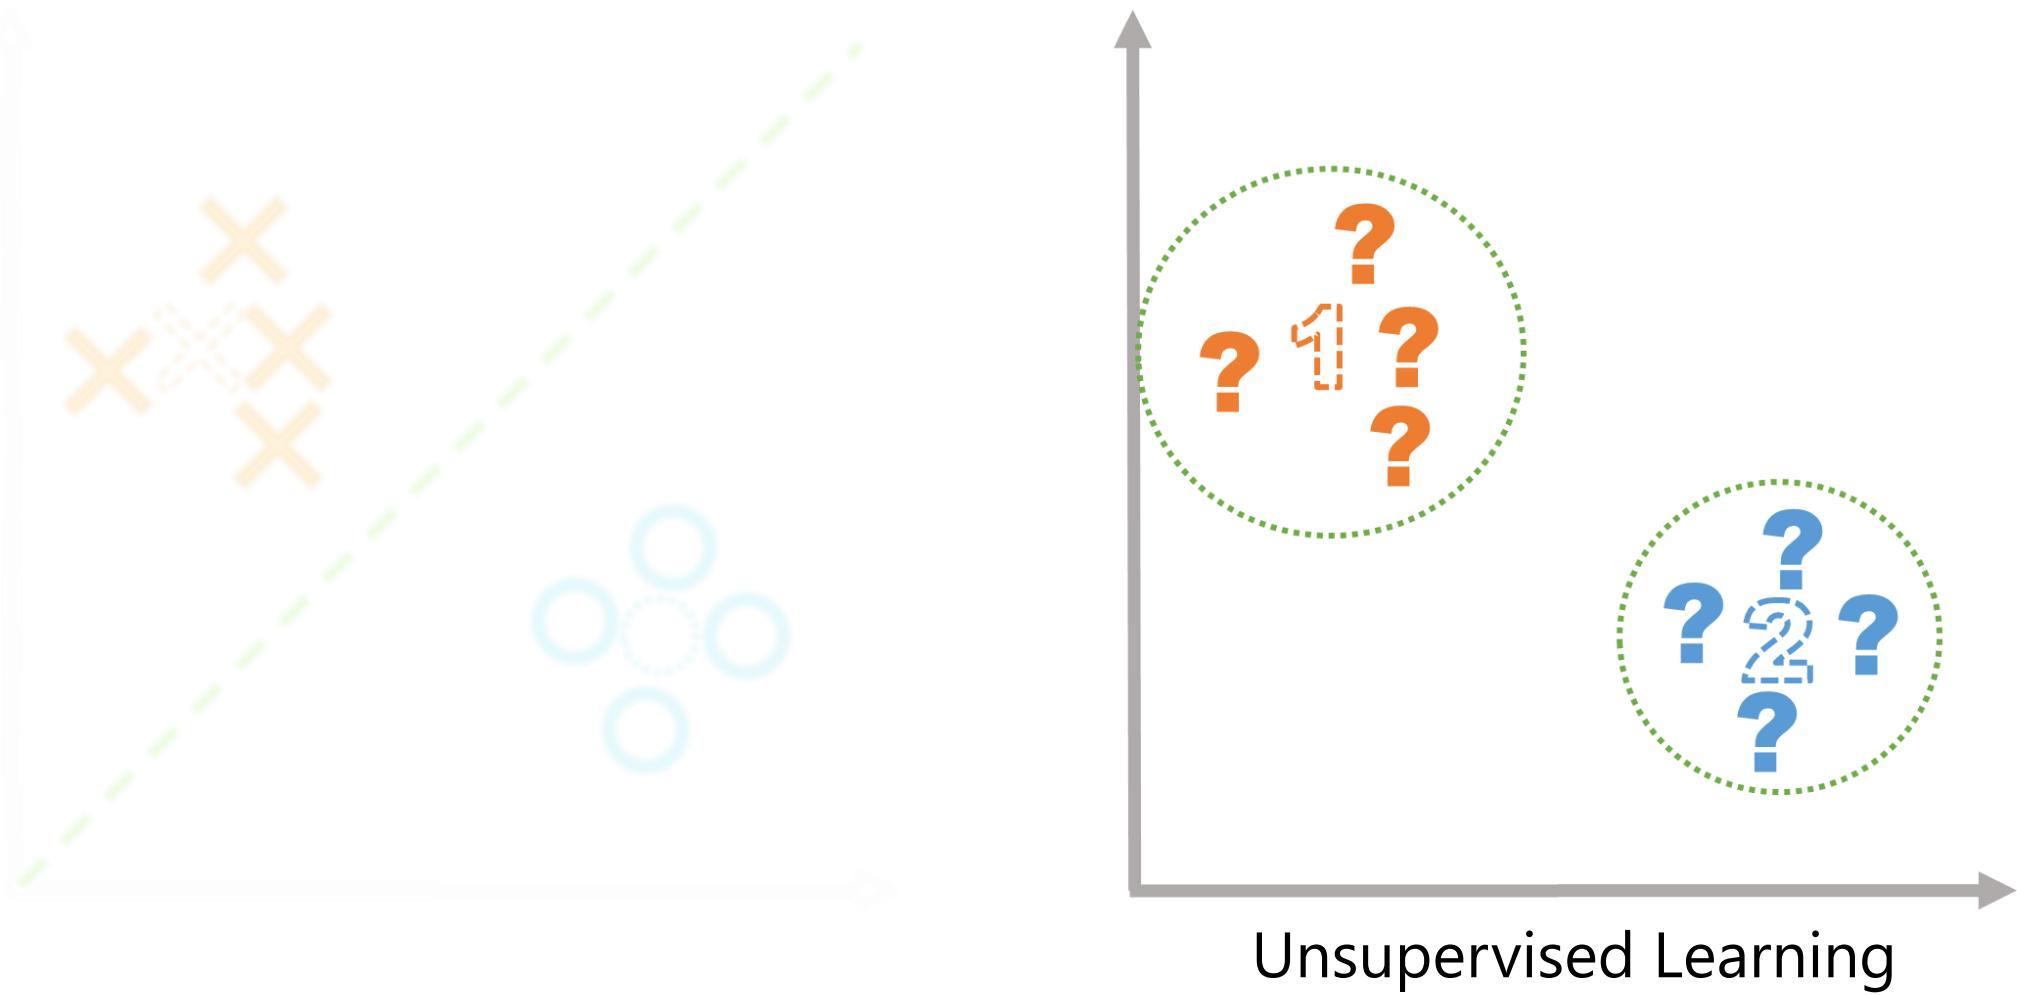
# Types of Machine Learning



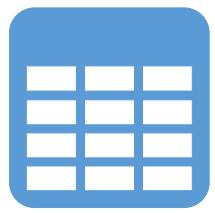
# Types of Machine Learning



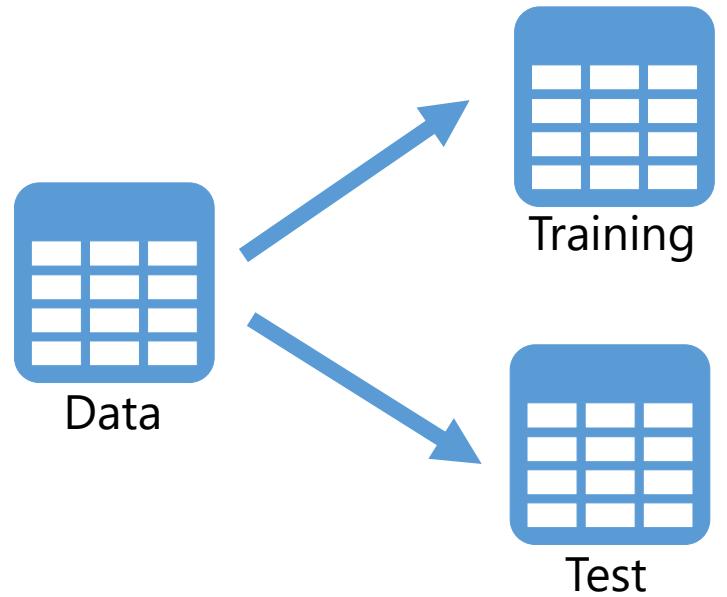
# Types of Machine Learning

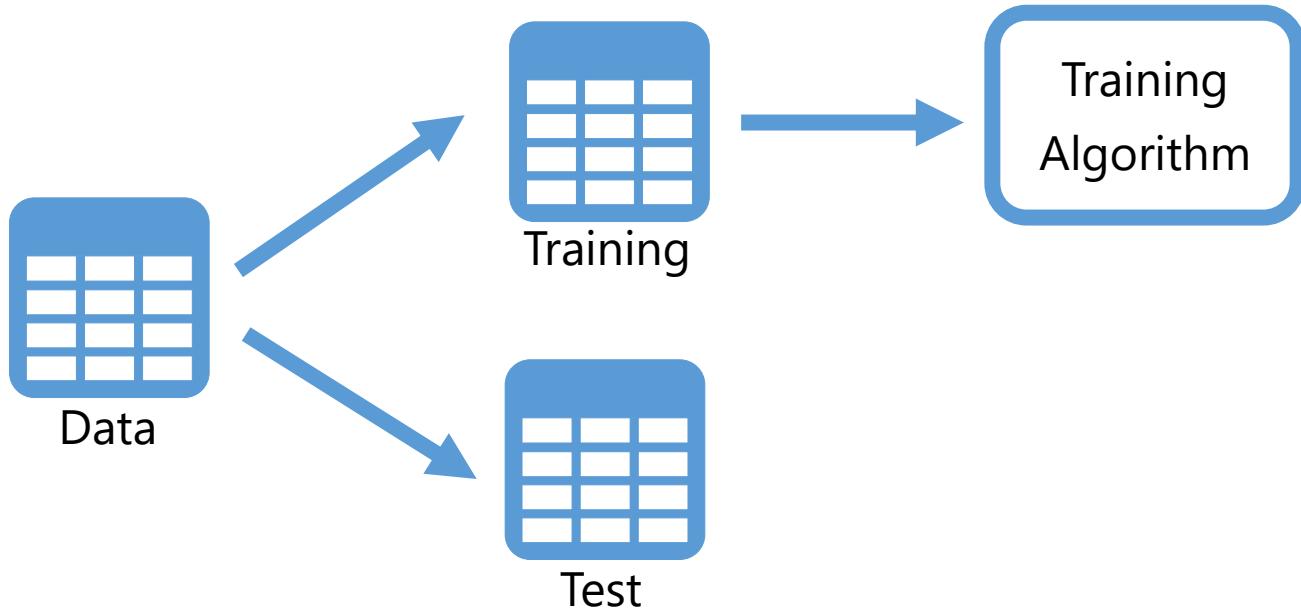


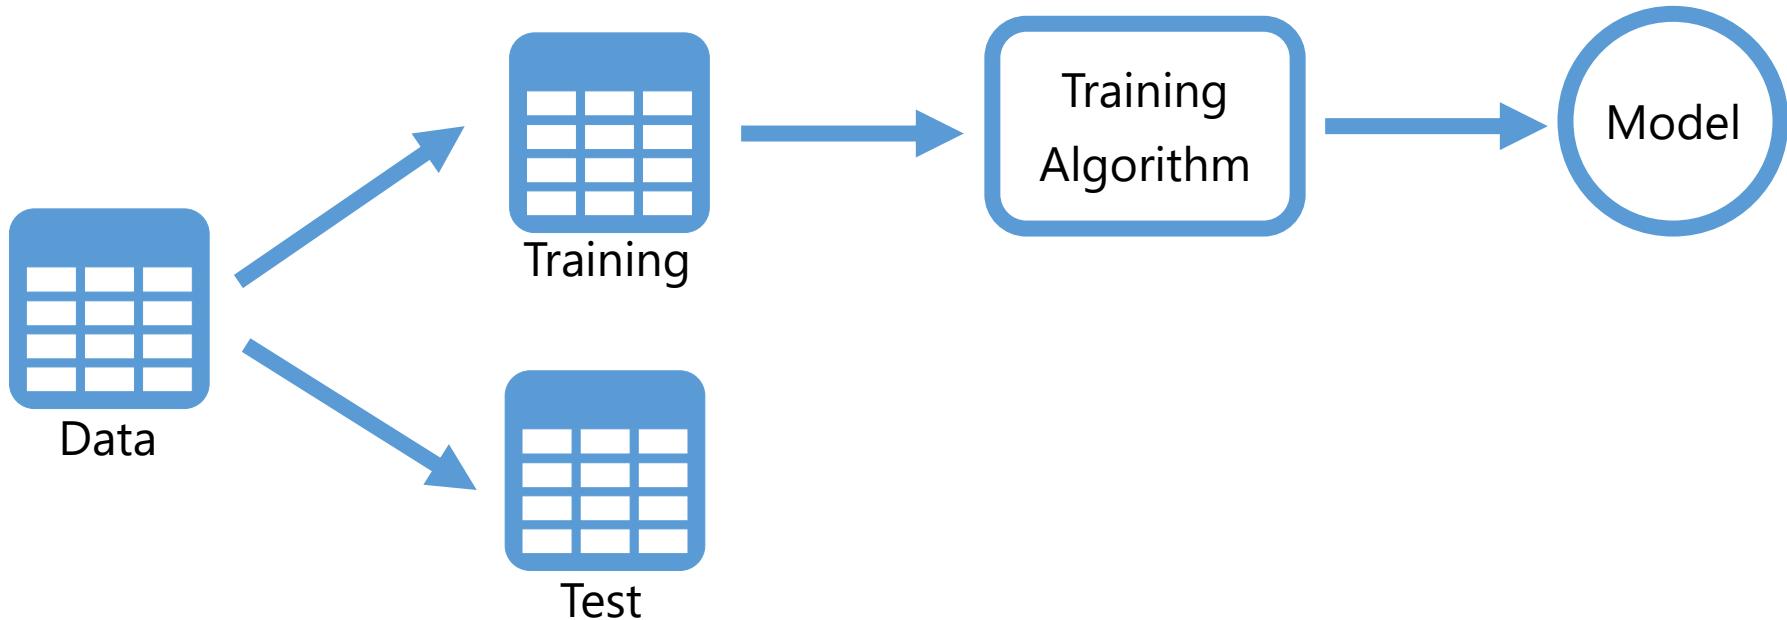
How does machine learning work?

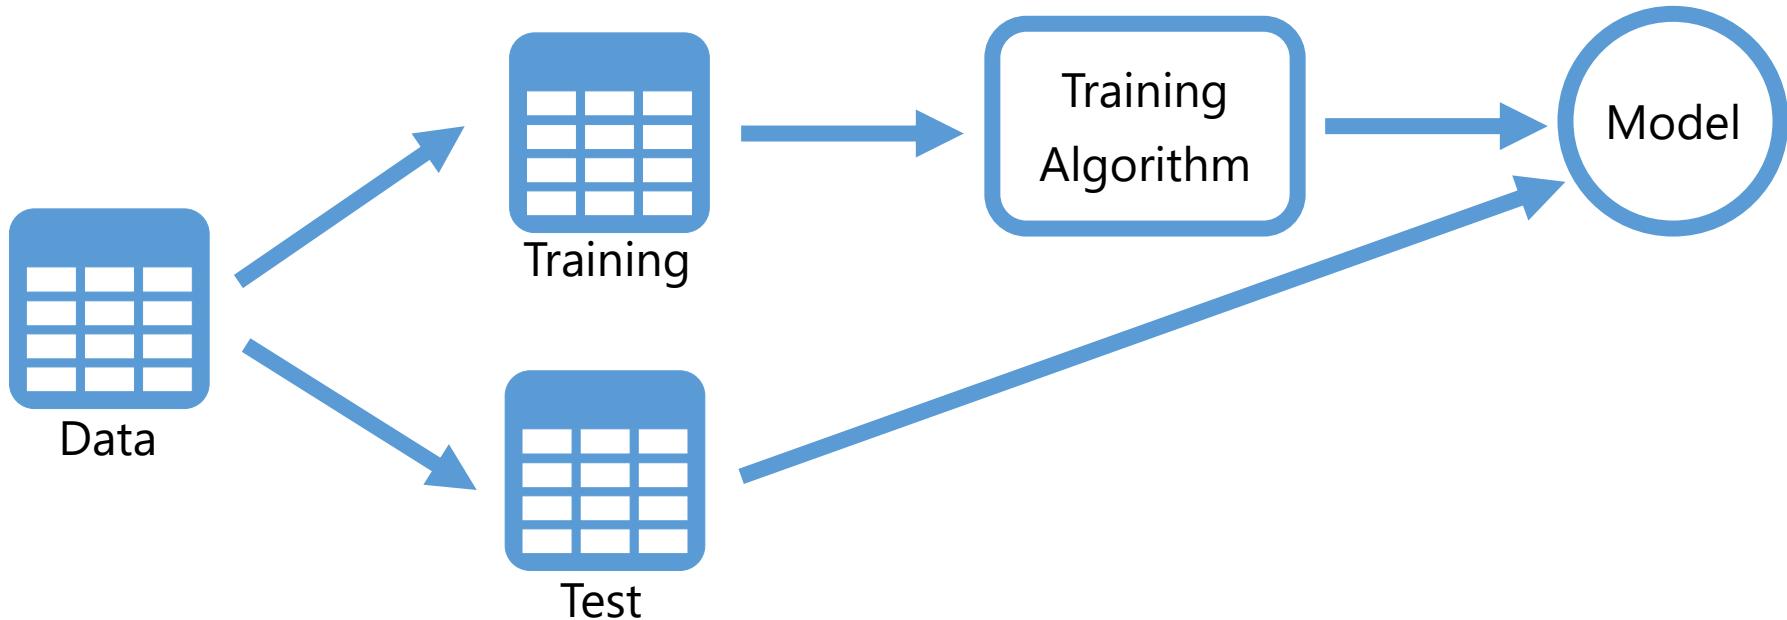


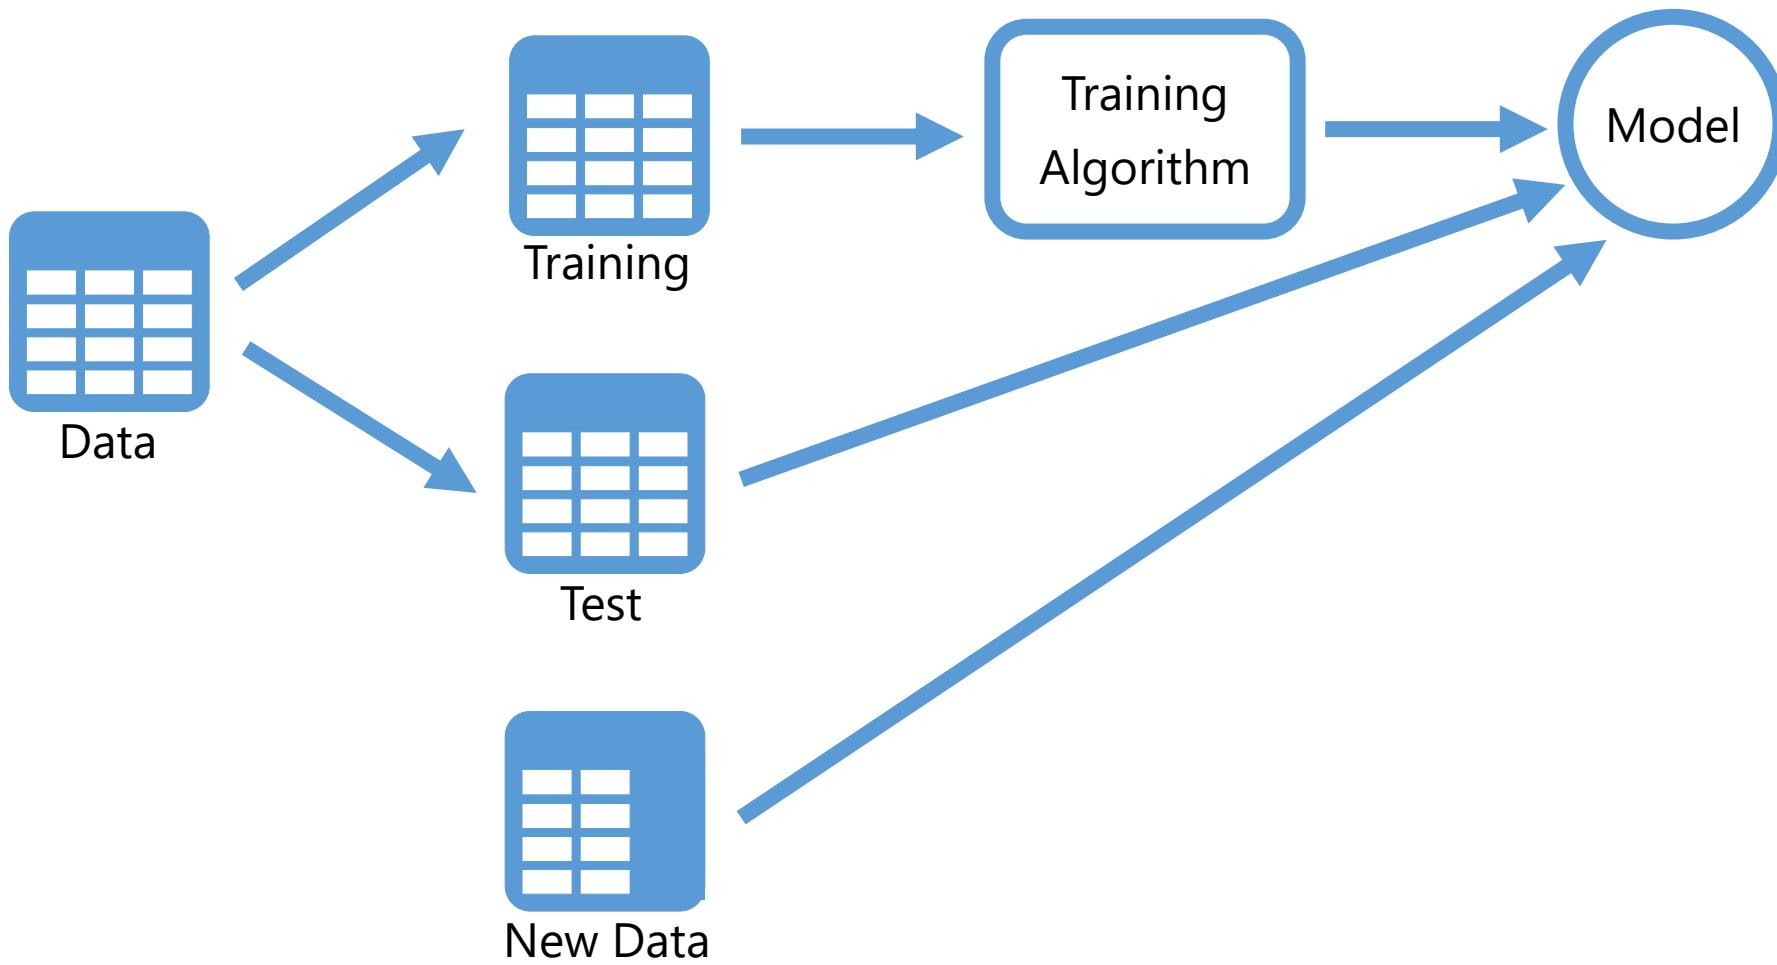
Data

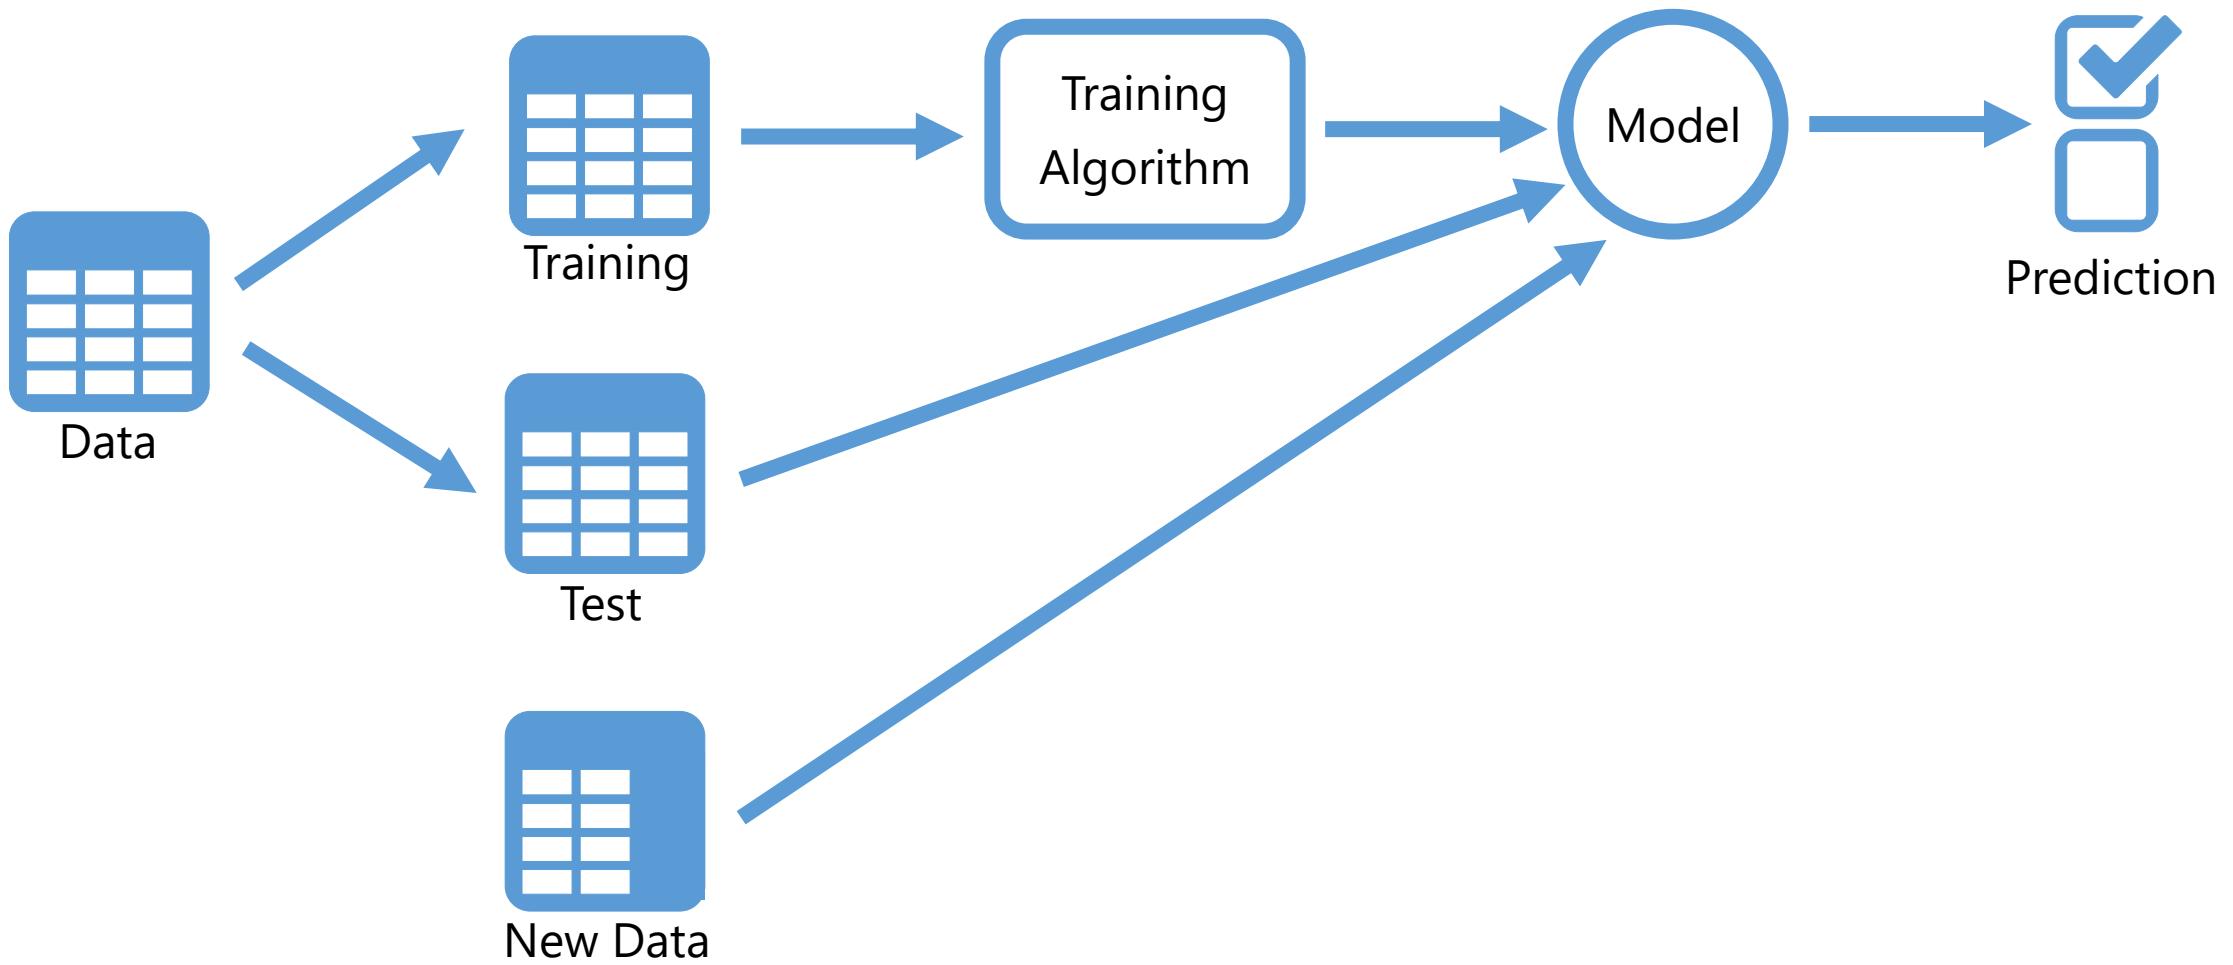






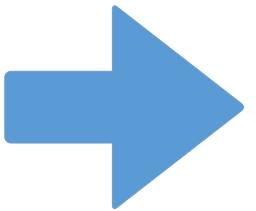
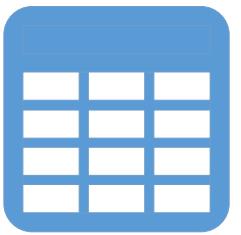
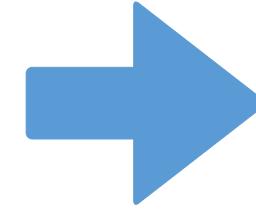






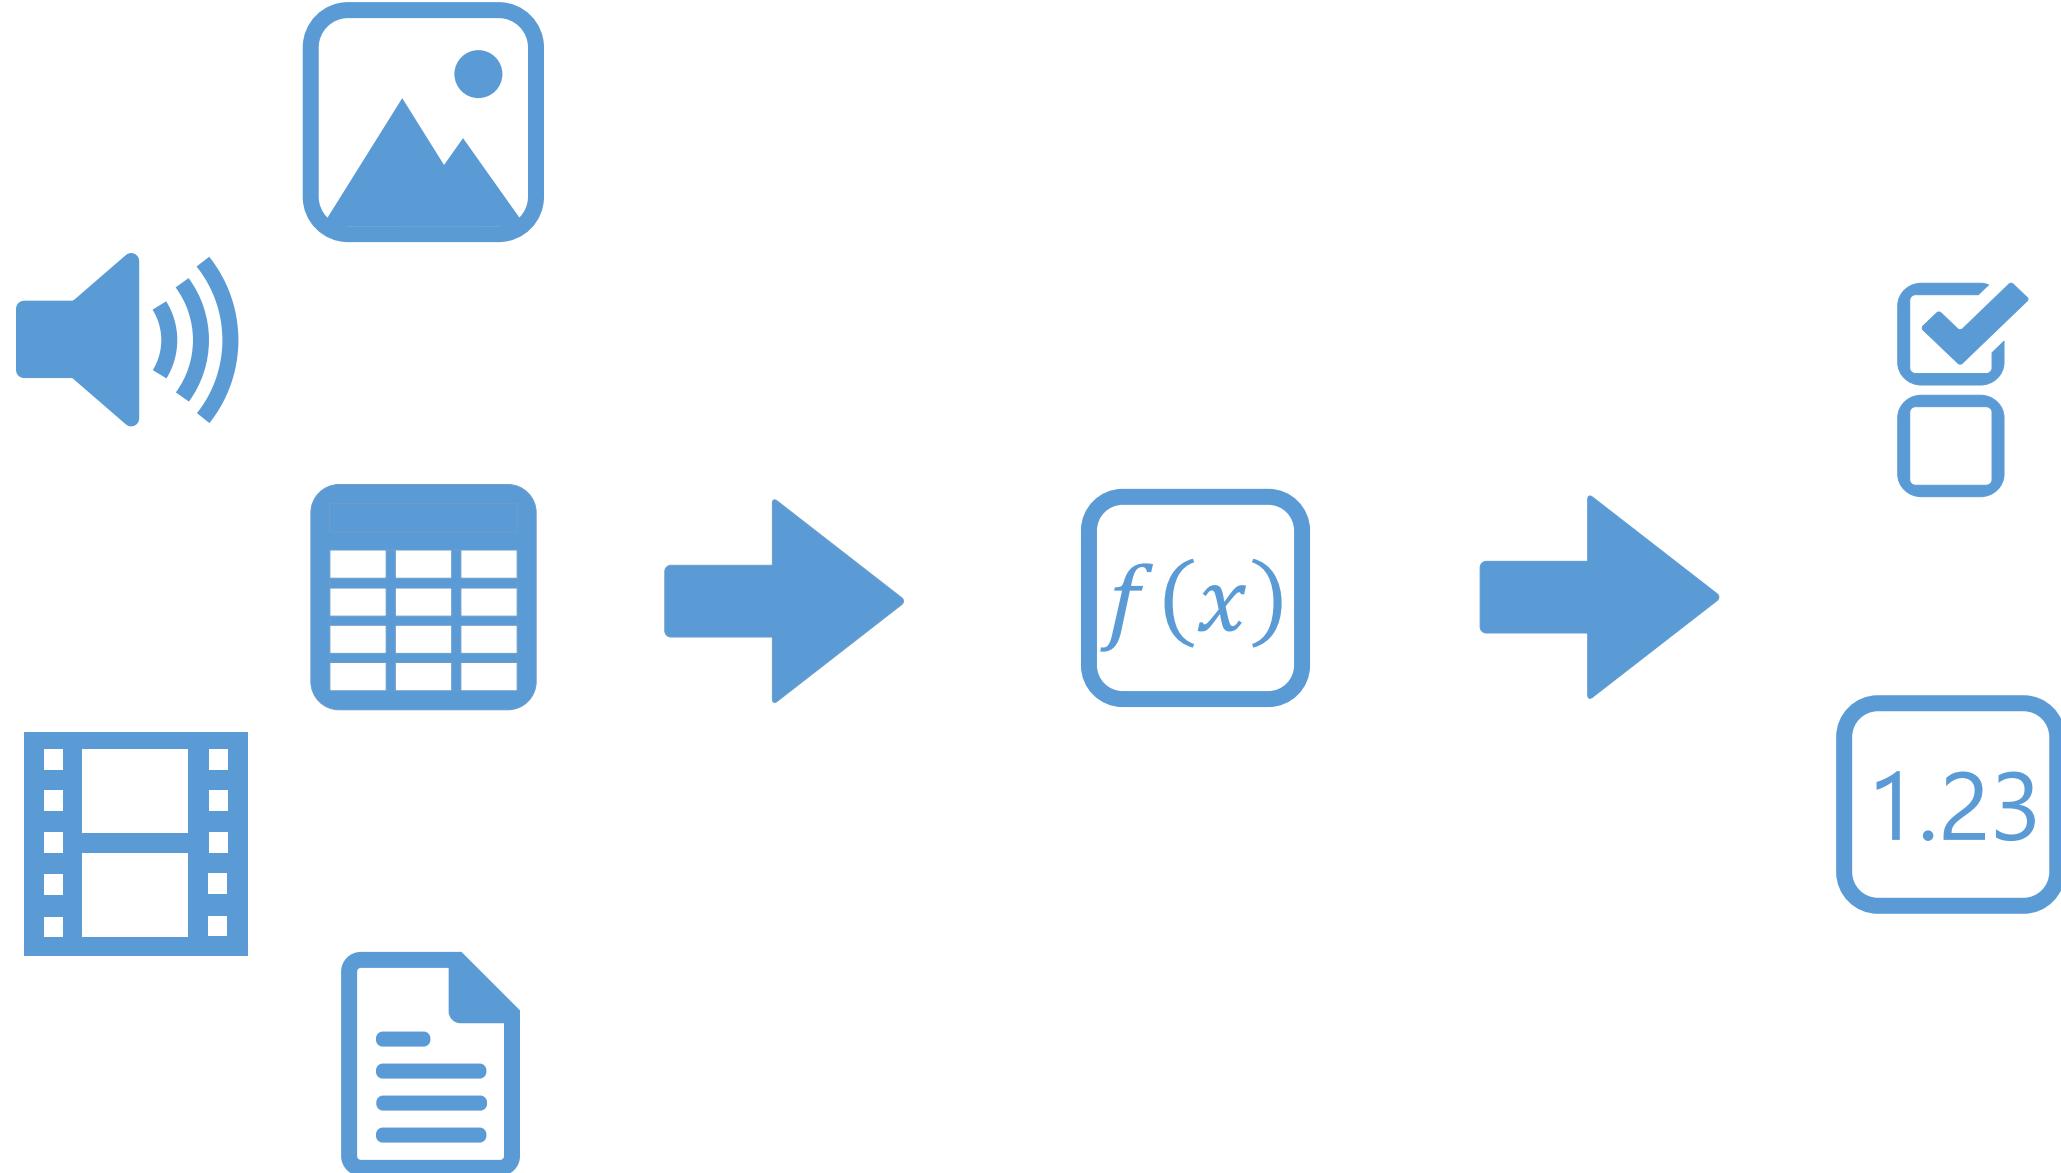


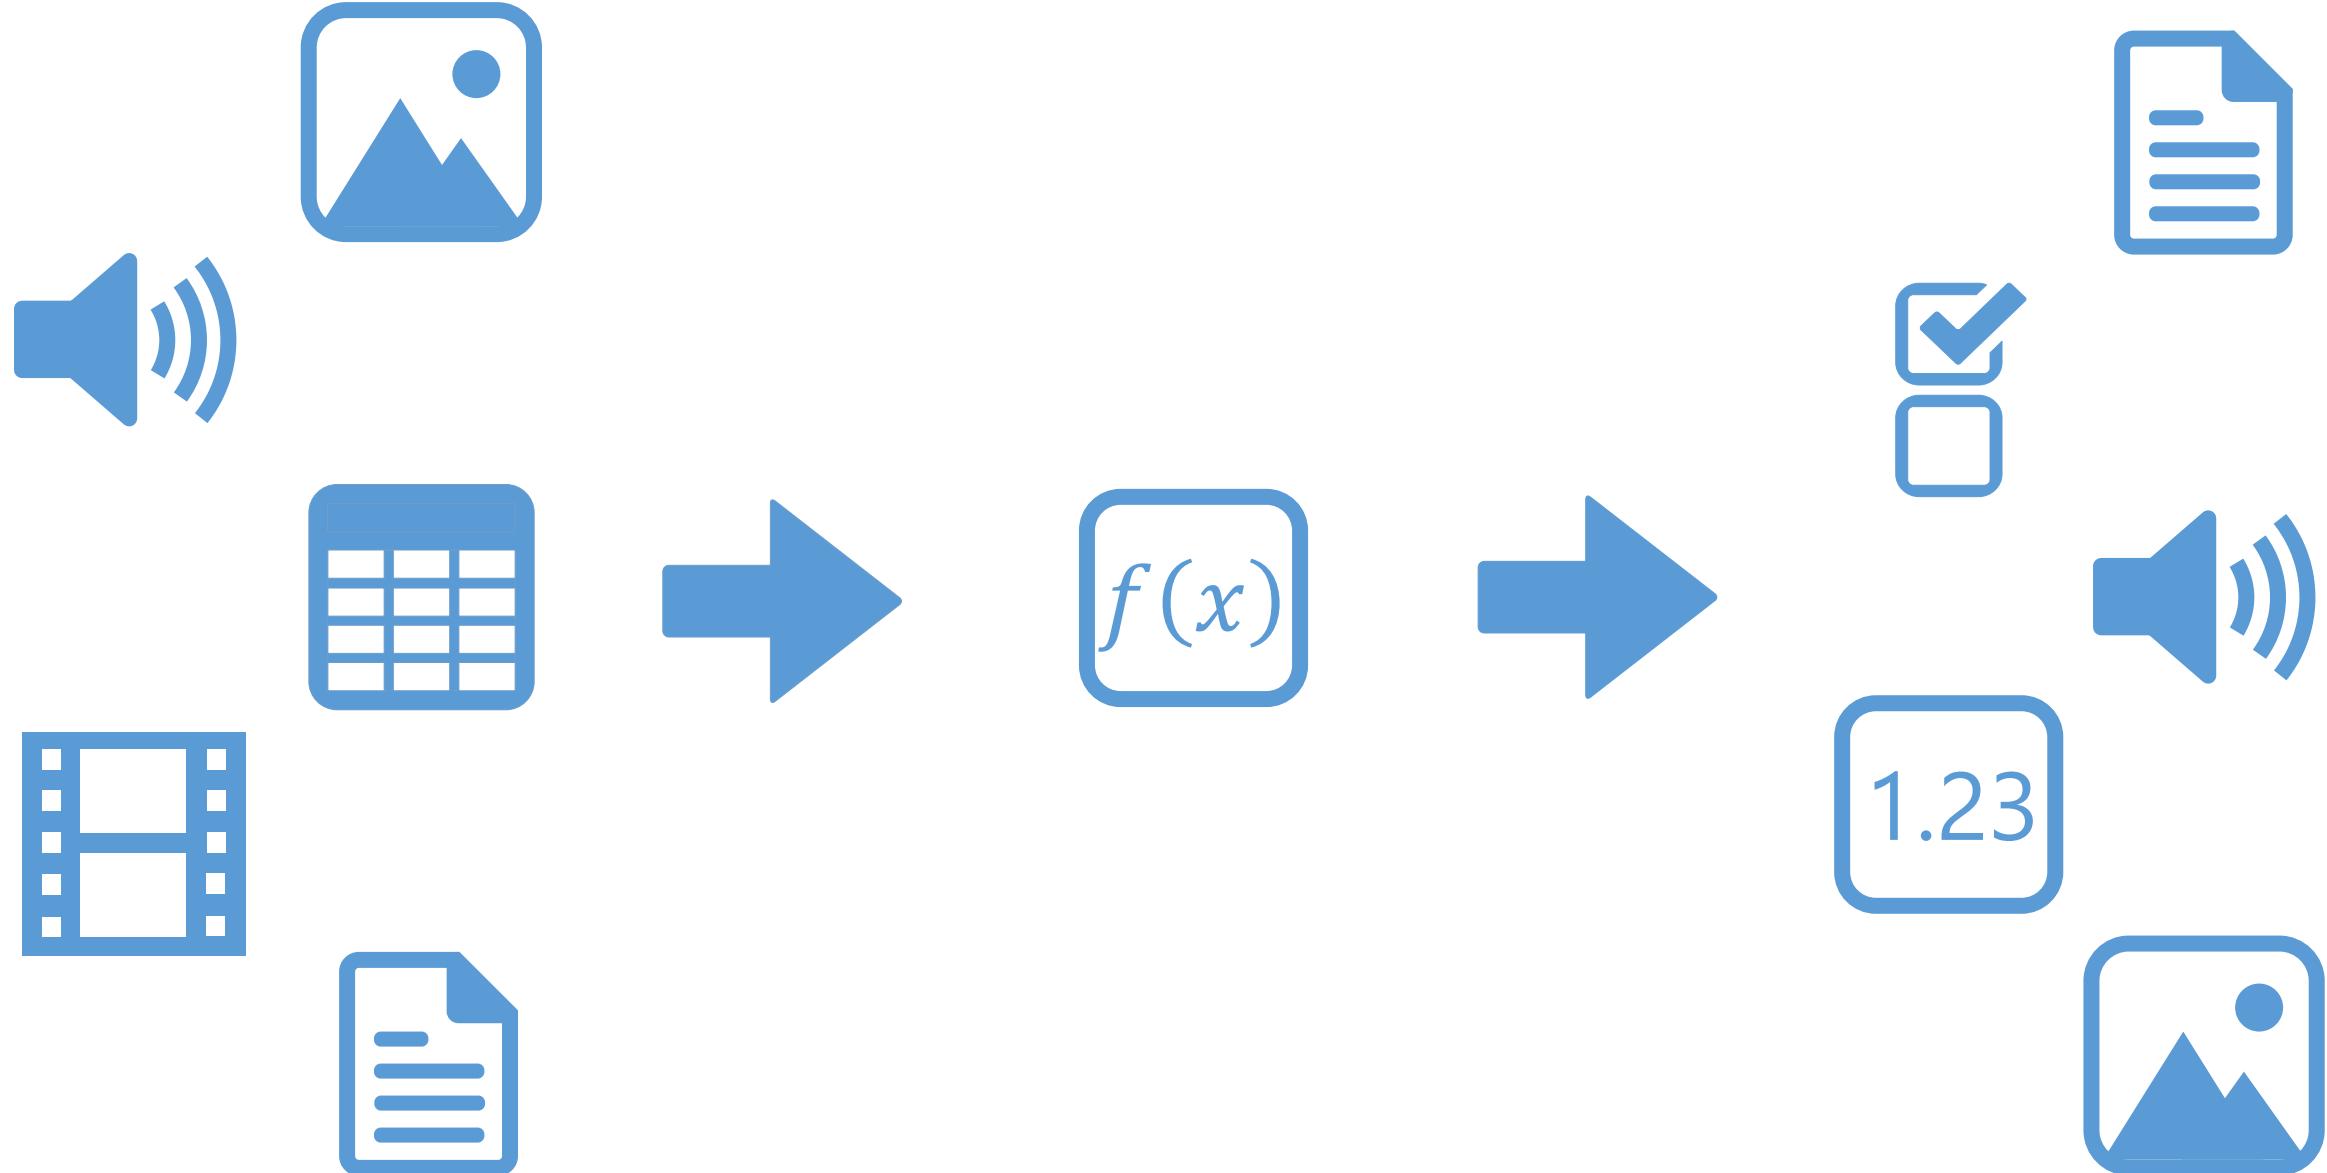
What can machine learning do?

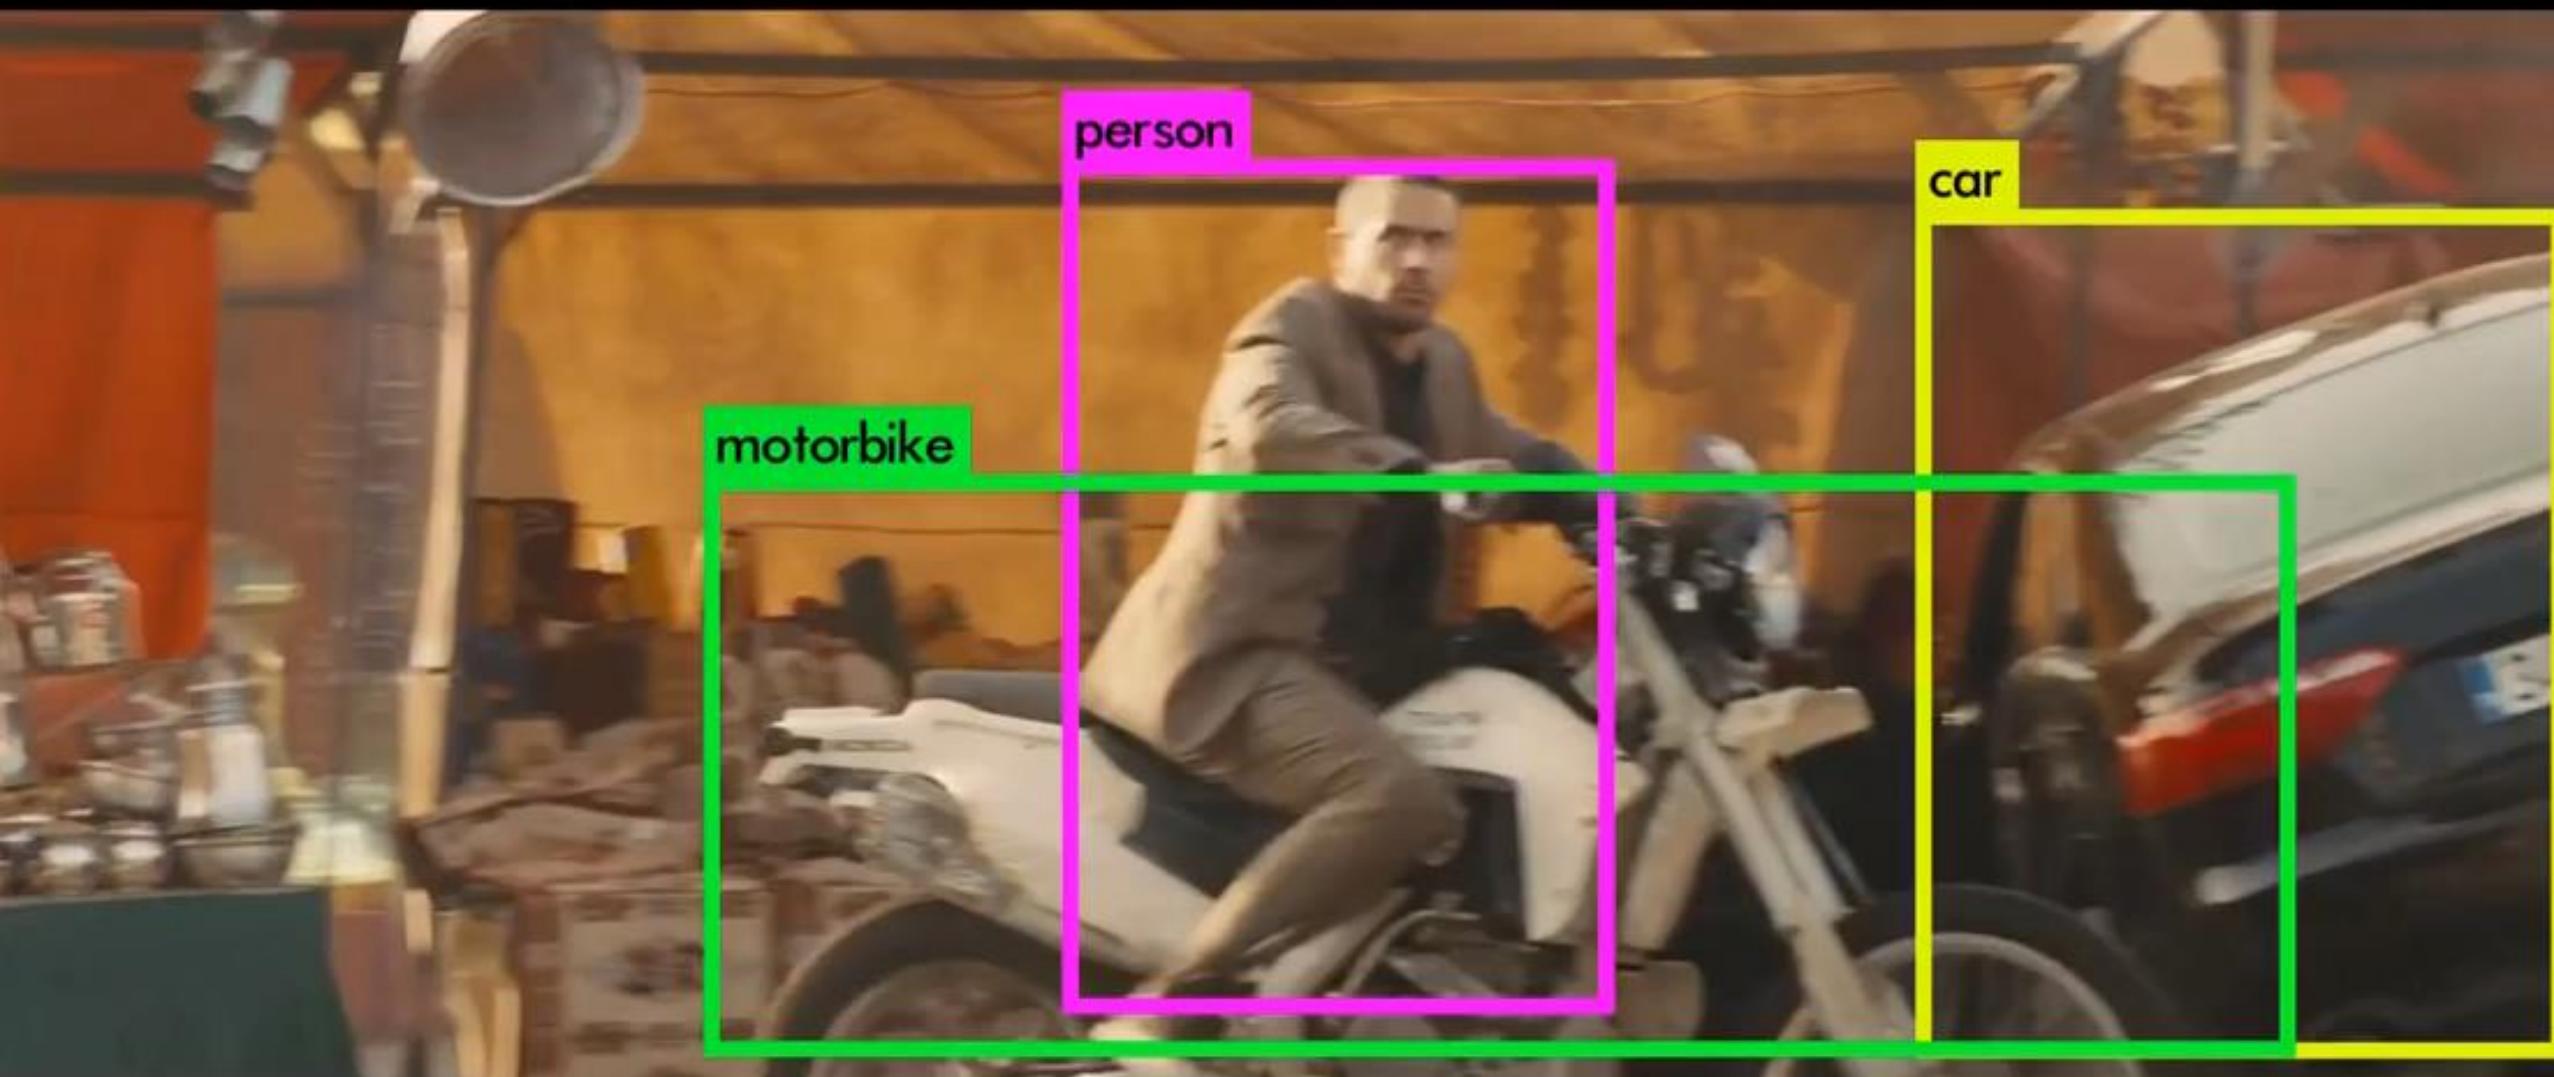
 $f(x)$ 

1.23











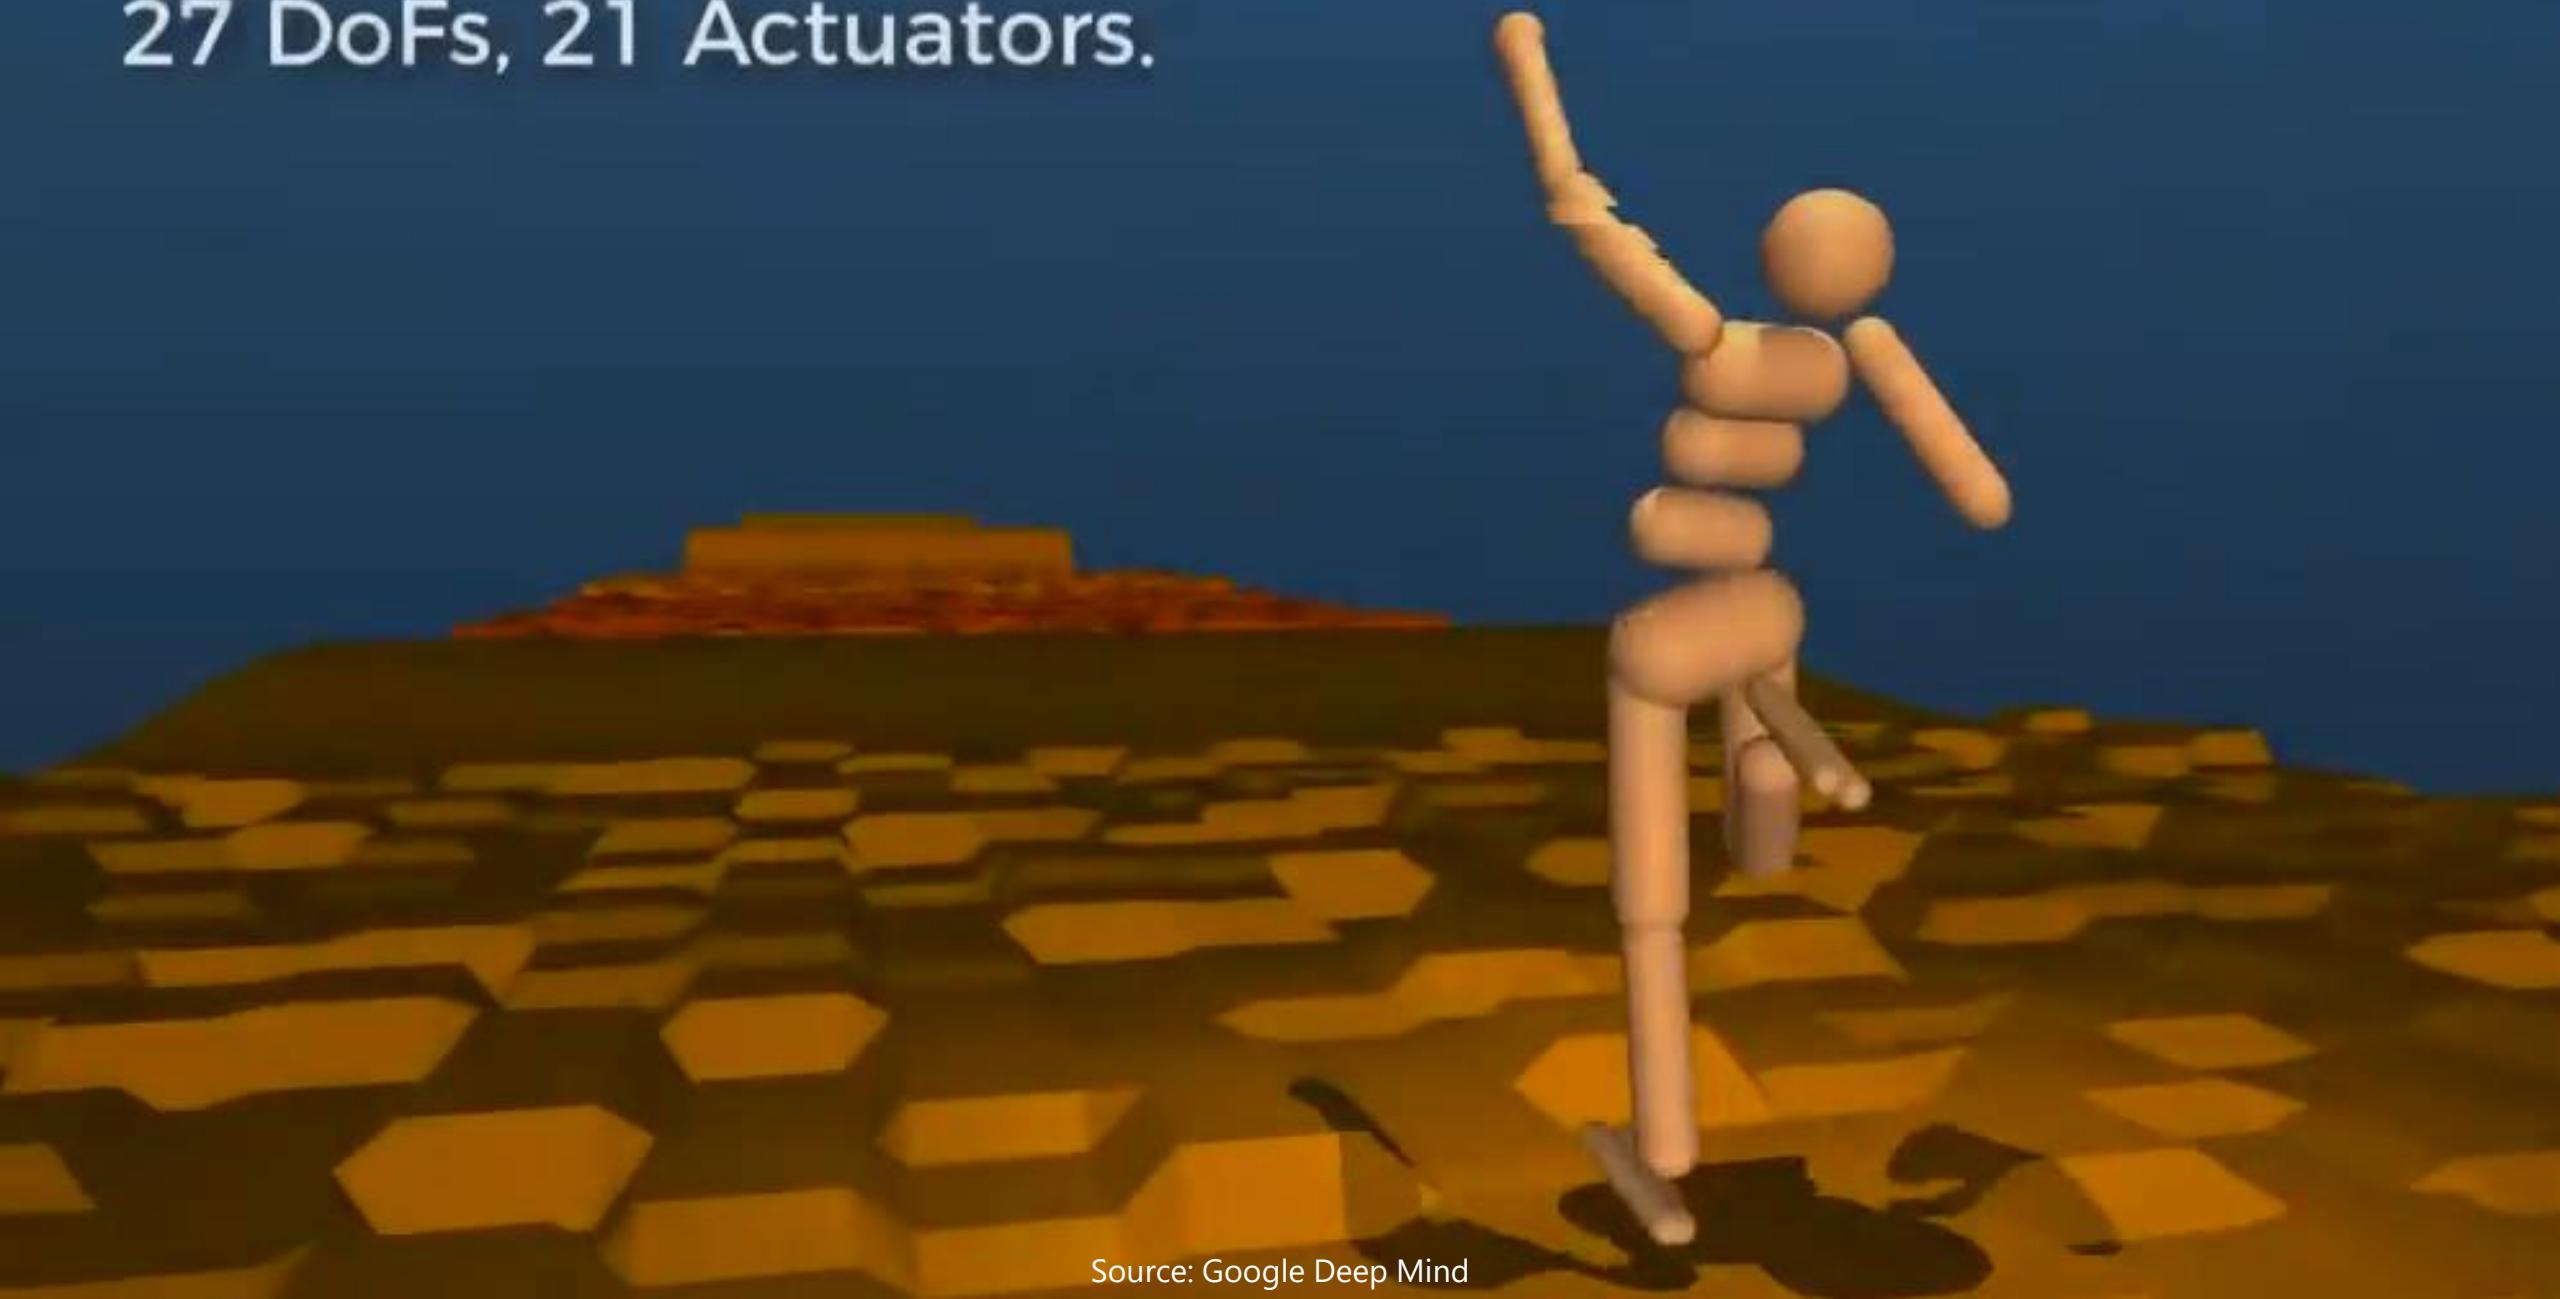
Source: Nvidia

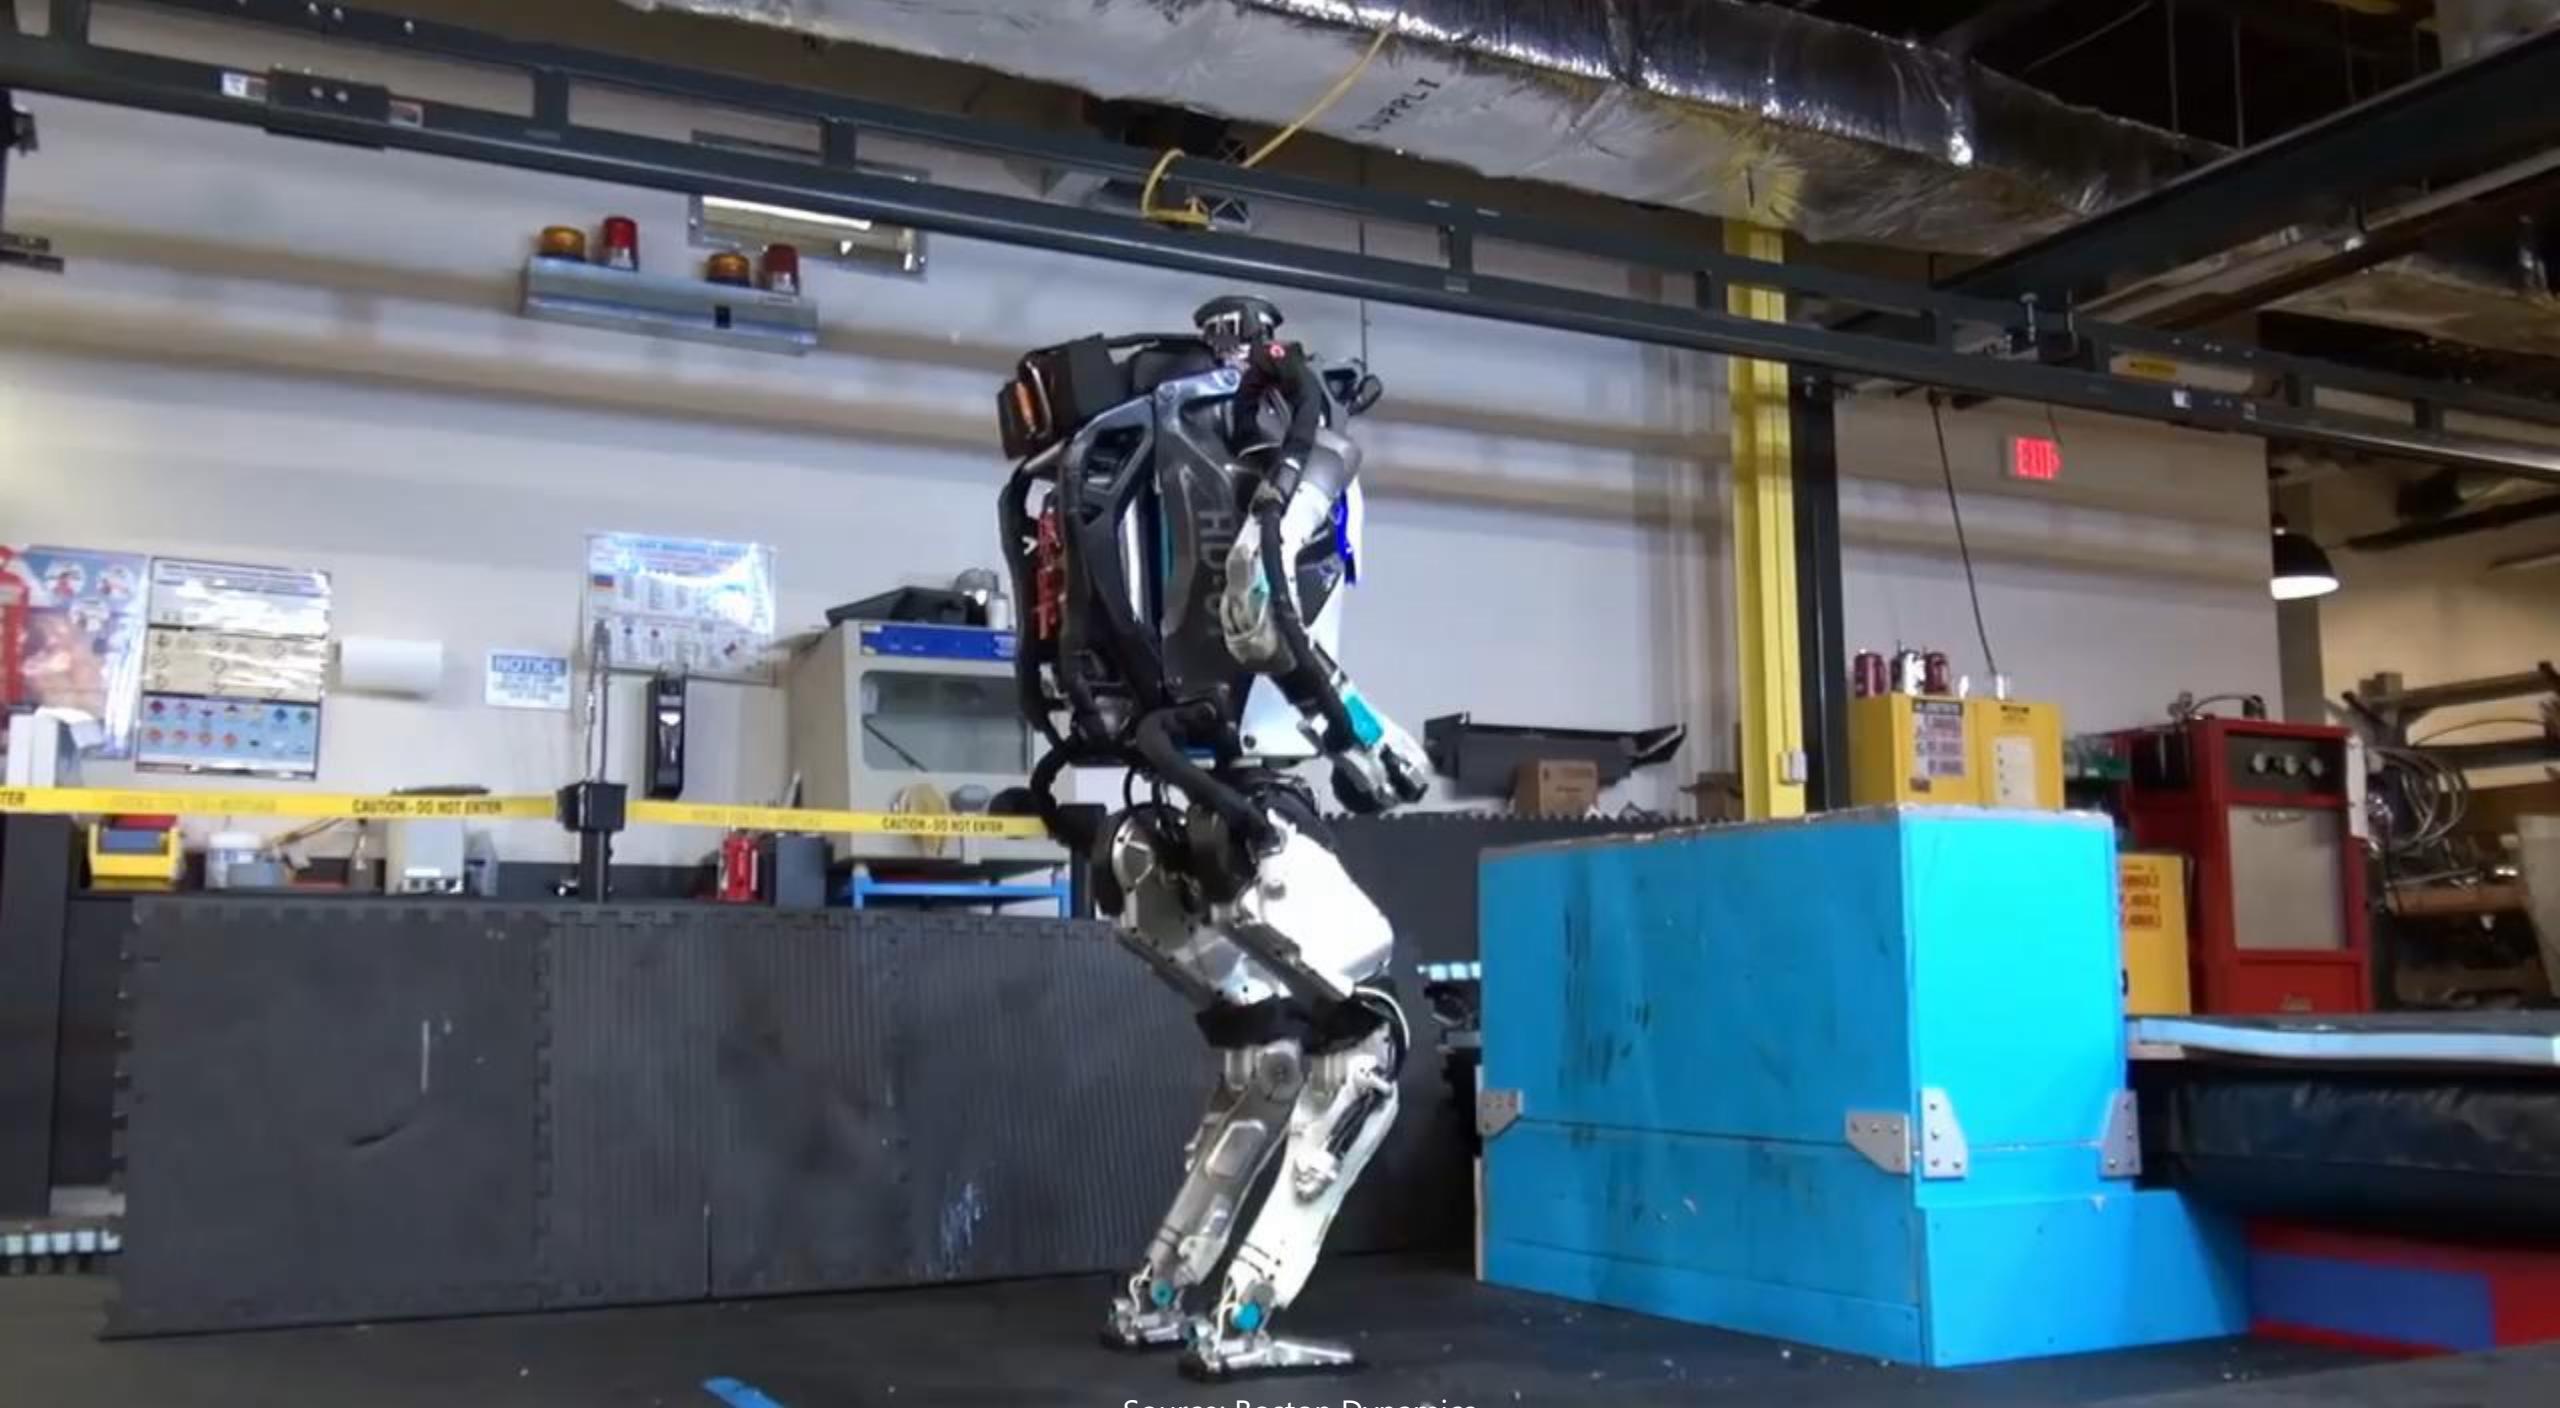




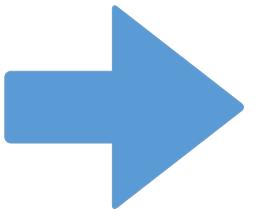
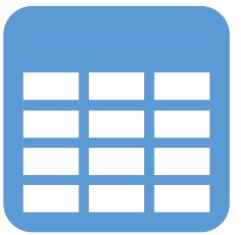
Source: <http://grail.cs.washington.edu/projects/AudioToObama/>

Humanoid:  
27 DoFs, 21 Actuators.

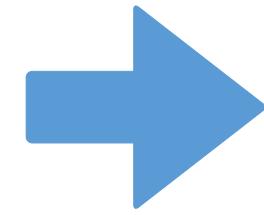




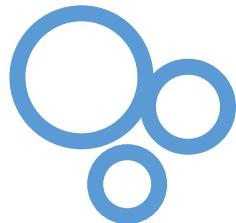
Source: Boston Dynamics

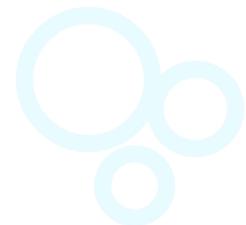
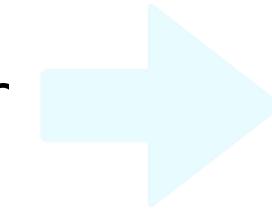
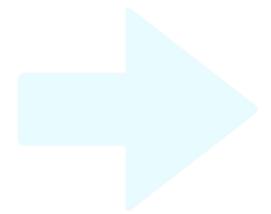
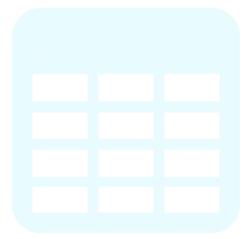


$f(x)$



1.23





# Introduction to Python

# What is Python?

- Open source
- Interpreted
- Readability
- Cross platform



# What is Python?

Active development

Large user community

Modular and extensible

135,000+ packages



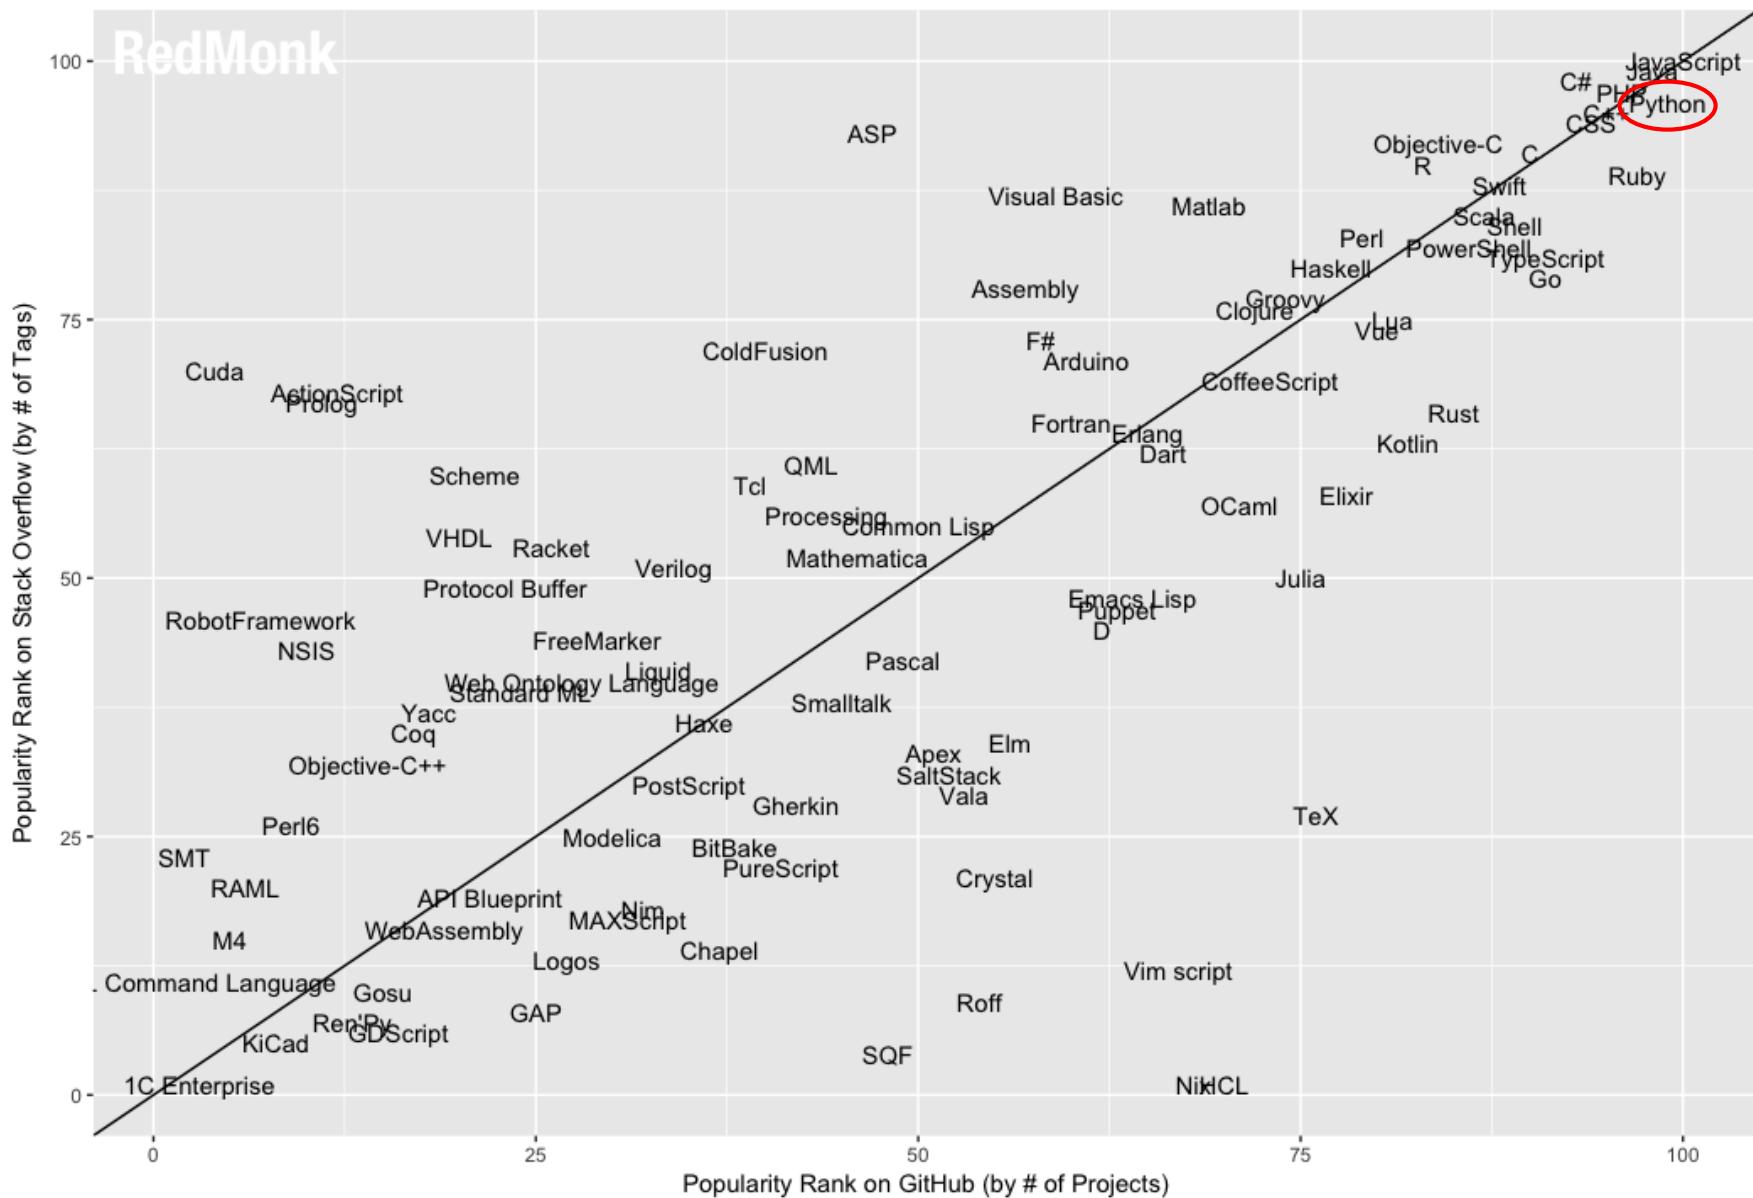
# FREE



A low-angle photograph of the Statue of Liberty against a clear blue sky. Her right arm is raised high, holding the torch aloft. Her left arm is bent, holding a tablet or smartphone that displays the word "FREE".

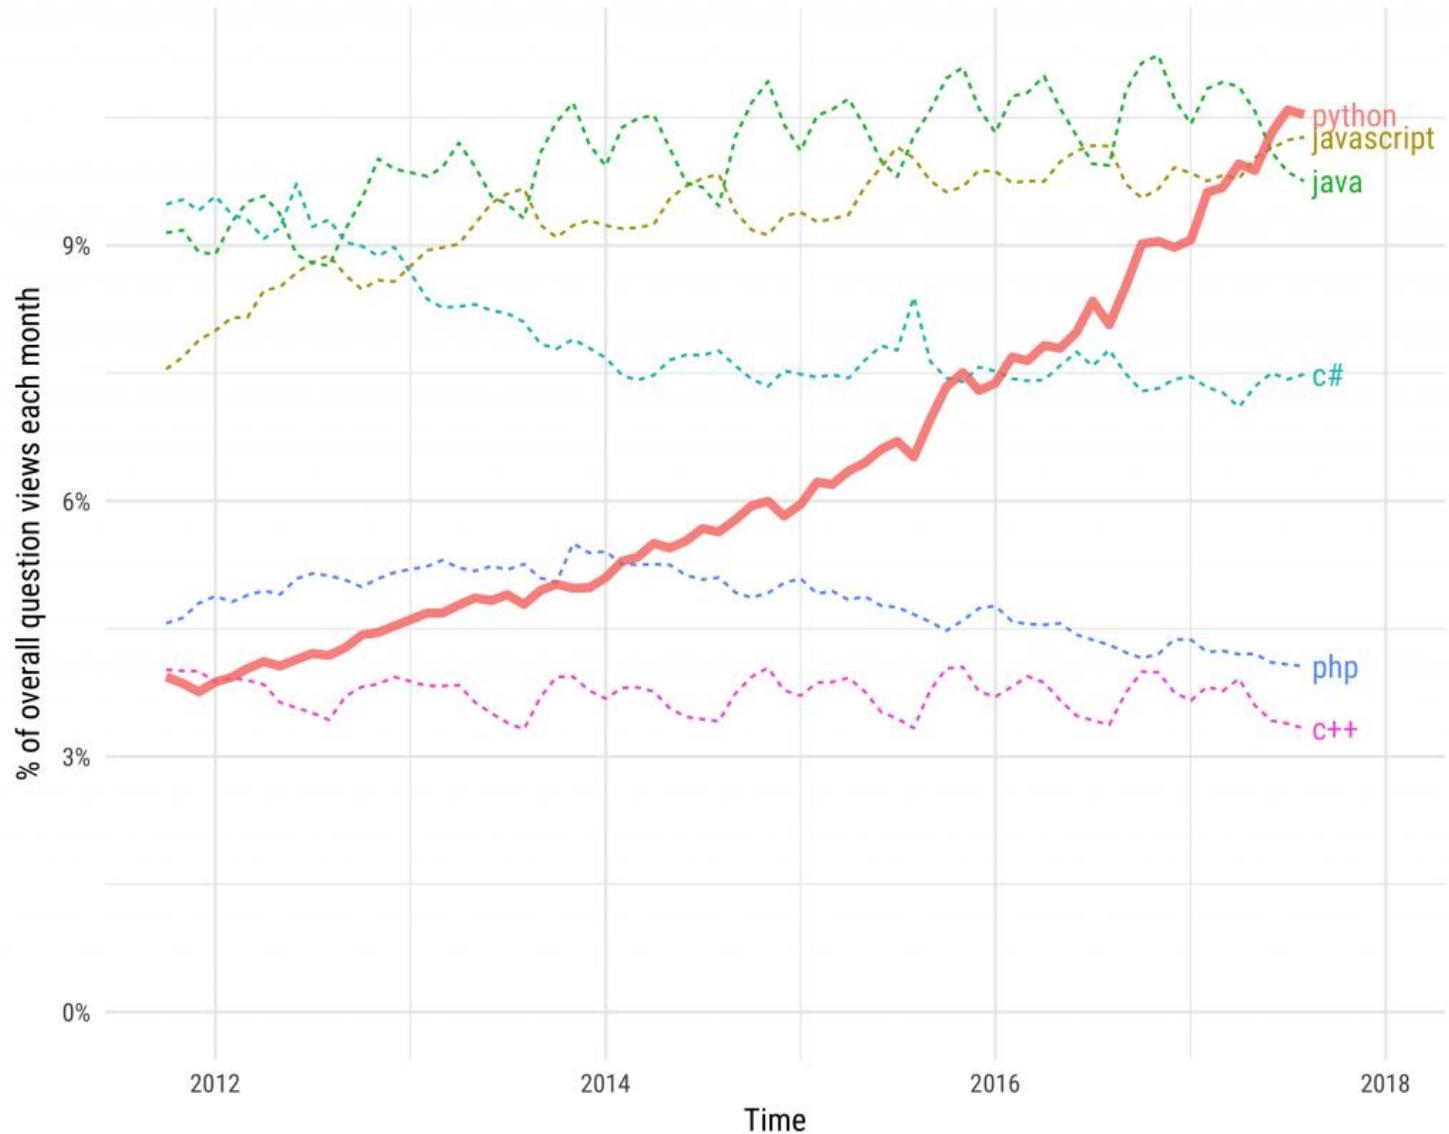
FREE

## RedMonk Q118 Programming Language Rankings



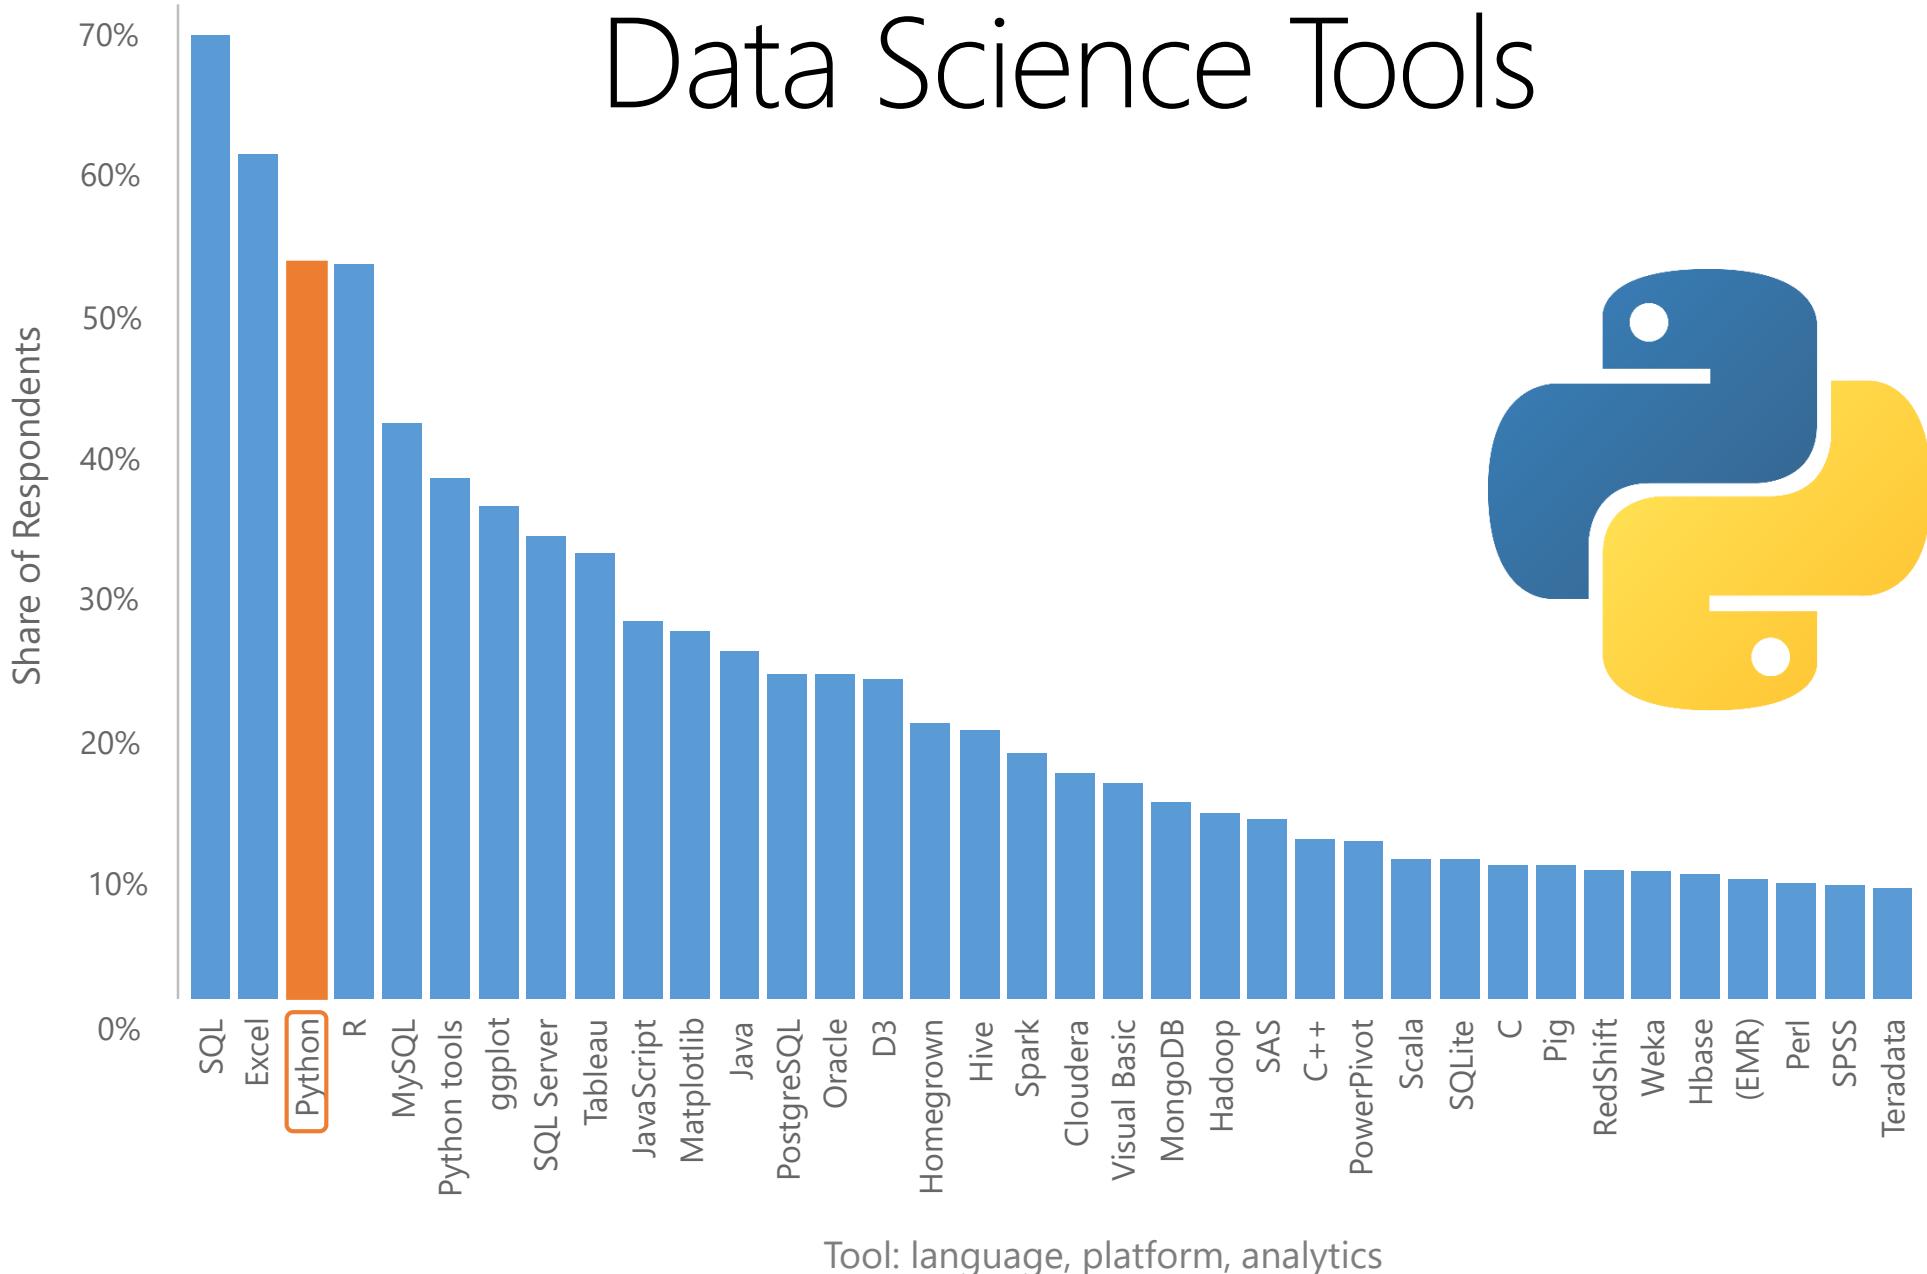
## Growth of major programming languages

Based on Stack Overflow question views in World Bank high-income countries

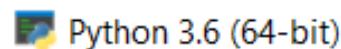


Source: Stack Overflow

# Data Science Tools



Source: O'Reilly 2015 Data Science Salary Survey



Python 3.6.3 (v3.6.3:2c5fed8, Oct 3 2017, 18:11:49) [MSC v.1900 64 bit (AMD64)] on win32  
Type "help", "copyright", "credits" or "license" for more information.

```
>>> x = "Hello World"
```

```
>>> print(x)
```

```
Hello World
```

```
>>> ■
```

Spyder (Python 3.5)

File Edit Search Source Run Debug Consoles Tools View Help

Editor - C:\temp\plot\_bayesian\_ridge.py

plot\_bayesian\_ridge.py

```
1 """
2 -----
3 Bayesian Ridge Regression
4 -----
5
6 Computes a Bayesian Ridge Regression on a synthetic data
7
8 See :ref:`bayesian_ridge_regression` for more information
9
10 Compared to the OLS (ordinary least squares) estimator,
11 weights are slightly shifted toward zeros, which stabilizes
12
13 As the prior on the weights is a Gaussian prior, the histogram
14 of estimated weights is Gaussian.
15
16 The estimation of the model is done by iteratively maximizing
17 the marginal log-likelihood of the observations.
18 """
19 print(__doc__)
20
21 import numpy as np
22 import matplotlib.pyplot as plt
23 from scipy import stats
24
25 from sklearn.linear_model import BayesianRidge, LinearRegression
26
27 #####
28 # Generating simulated data with Gaussian weights
29 np.random.seed(0)
30 n_samples, n_features = 100, 100
31 X = np.random.randn(n_samples, n_features) # Create Gaussian data
32 # Create weights with a precision lambda_ of 4.
33 lambda_ = 4.
34 w = np.zeros(n_features)
35 # Only keep 10 weights of interest
36 relevant_features = np.random.randint(0, n_features, 10)
37 for i in relevant_features:
38     w[i] = np.random.normal(0, 1 / lambda_)
```

Variable explorer

Name	Type	Size	Value
X	float64	(100, 100)	array([[ 1.7640..., 0.1269...]])
alpha_	float	1	50.0
digits	datasets.base.Bunch	5	{'target': array([ 1.,  0.])}
i	int32	1	42
iris	datasets.base.Bunch	5	{'target': array([ 5.1,  1.8])}

Object inspector Variable explorer File explorer

IPython console

Console 1/A

Weights of the model

Values of the weights

Bayesian Ridge estimate  
Ground truth  
OLS estimate

Console History log IPython console

Permissions: RW End-of-lines: LF Encoding: UTF-8-GUESSED Line: 1 Column: 1 Memory: 40 %

15-Preview-of-Data-Science-Tools.py - Microsoft Visual Studio

File Edit View Project Debug Team Tools Test Analyze Window Help

Steve Dower

15-Preview-of-Data-Science-Tools.py

# In[14]:

```
import matplotlib.pyplot as plt
plt.style.use('ggplot') # make graphs in the
```

# Now let's create some data (as NumPy arrays)

# In[15]:

```
x = np.linspace(0, 10) # range of values from 0 to 10
y = np.sin(x) # sine of these values
plt.plot(x, y); # plot as a line
```

# If you run this code live, you will see an

# What we see is a smooth interpolation between the points.

# ## Other Data Science Packages

#

# Built on top of these tools are a host of c

Anaconda 4.3.0 Interactive

In [18]: # Now let's create some data (as NumPy arrays)
.....
....: # In[15]:
....:
....: x = np.linspace(0, 10) # range of values from 0 to 10
....: y = np.sin(x) # sine of these values
....: plt.plot(x, y); # plot as a line

In [19]:

# Demo 1

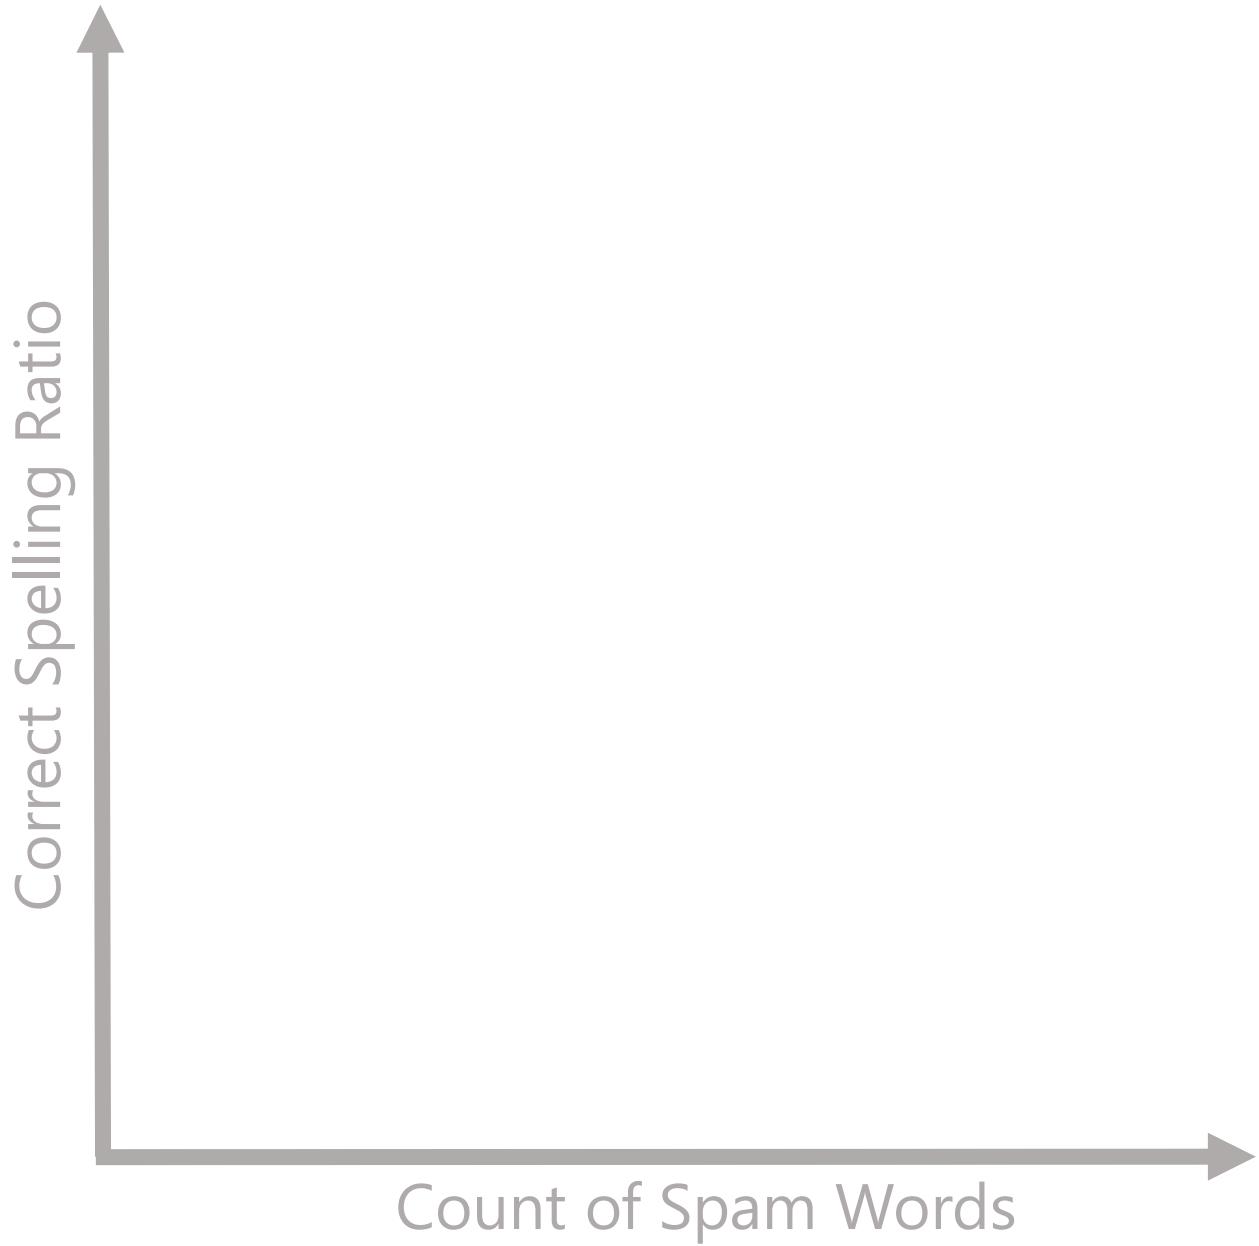
## Python Language Basics

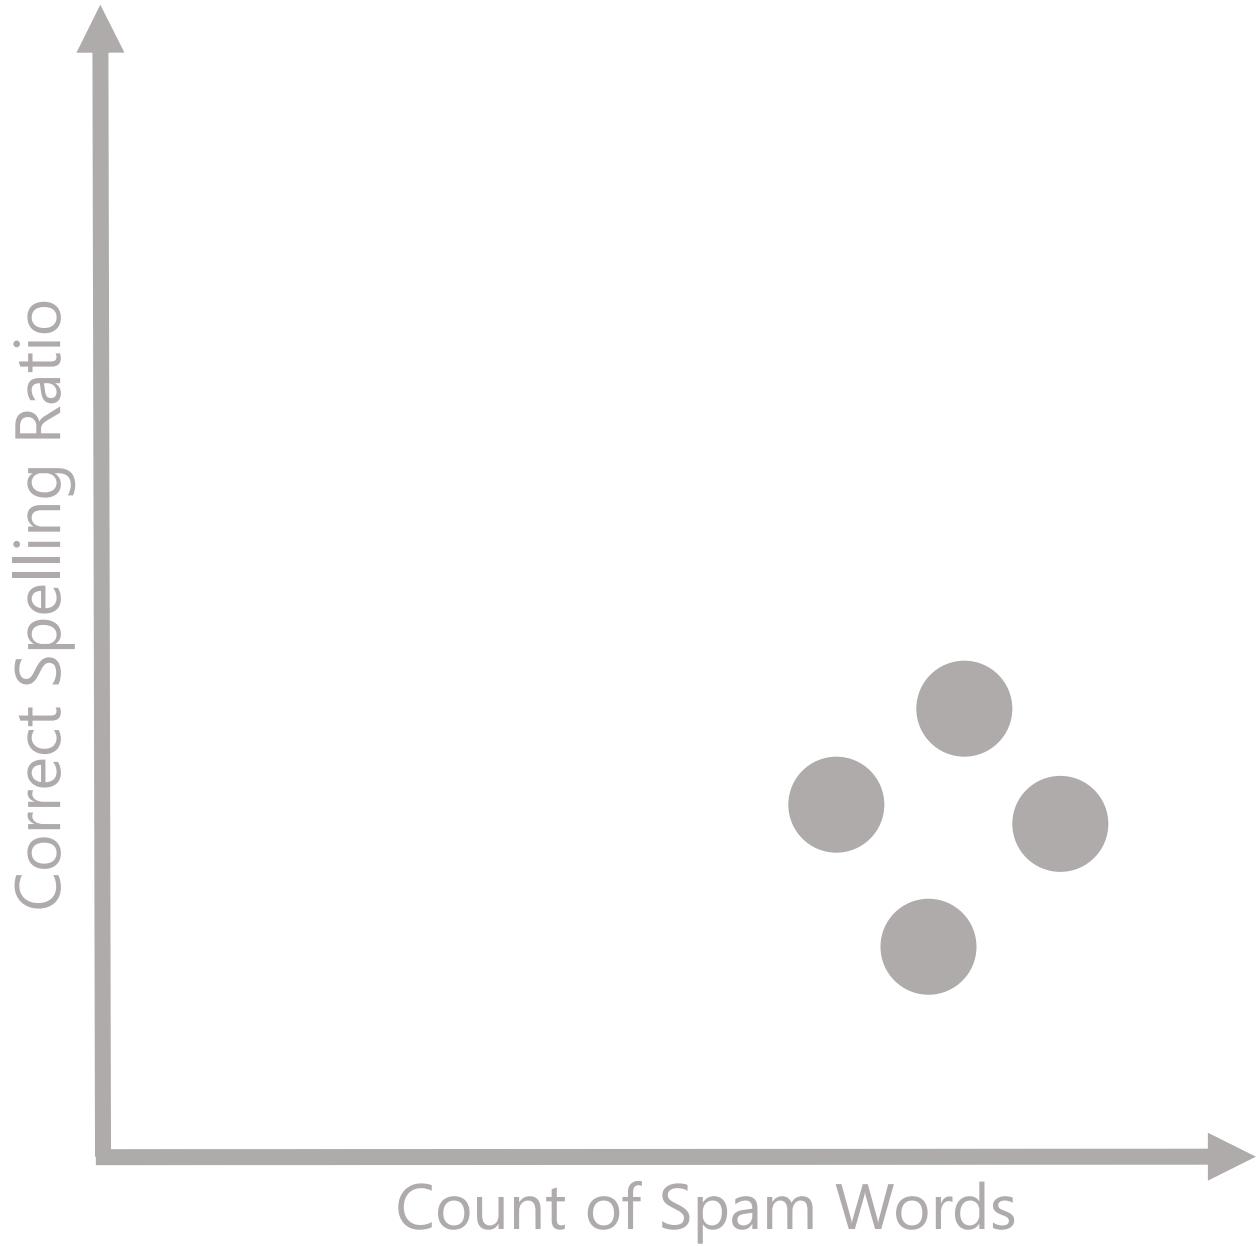
# Lab 1

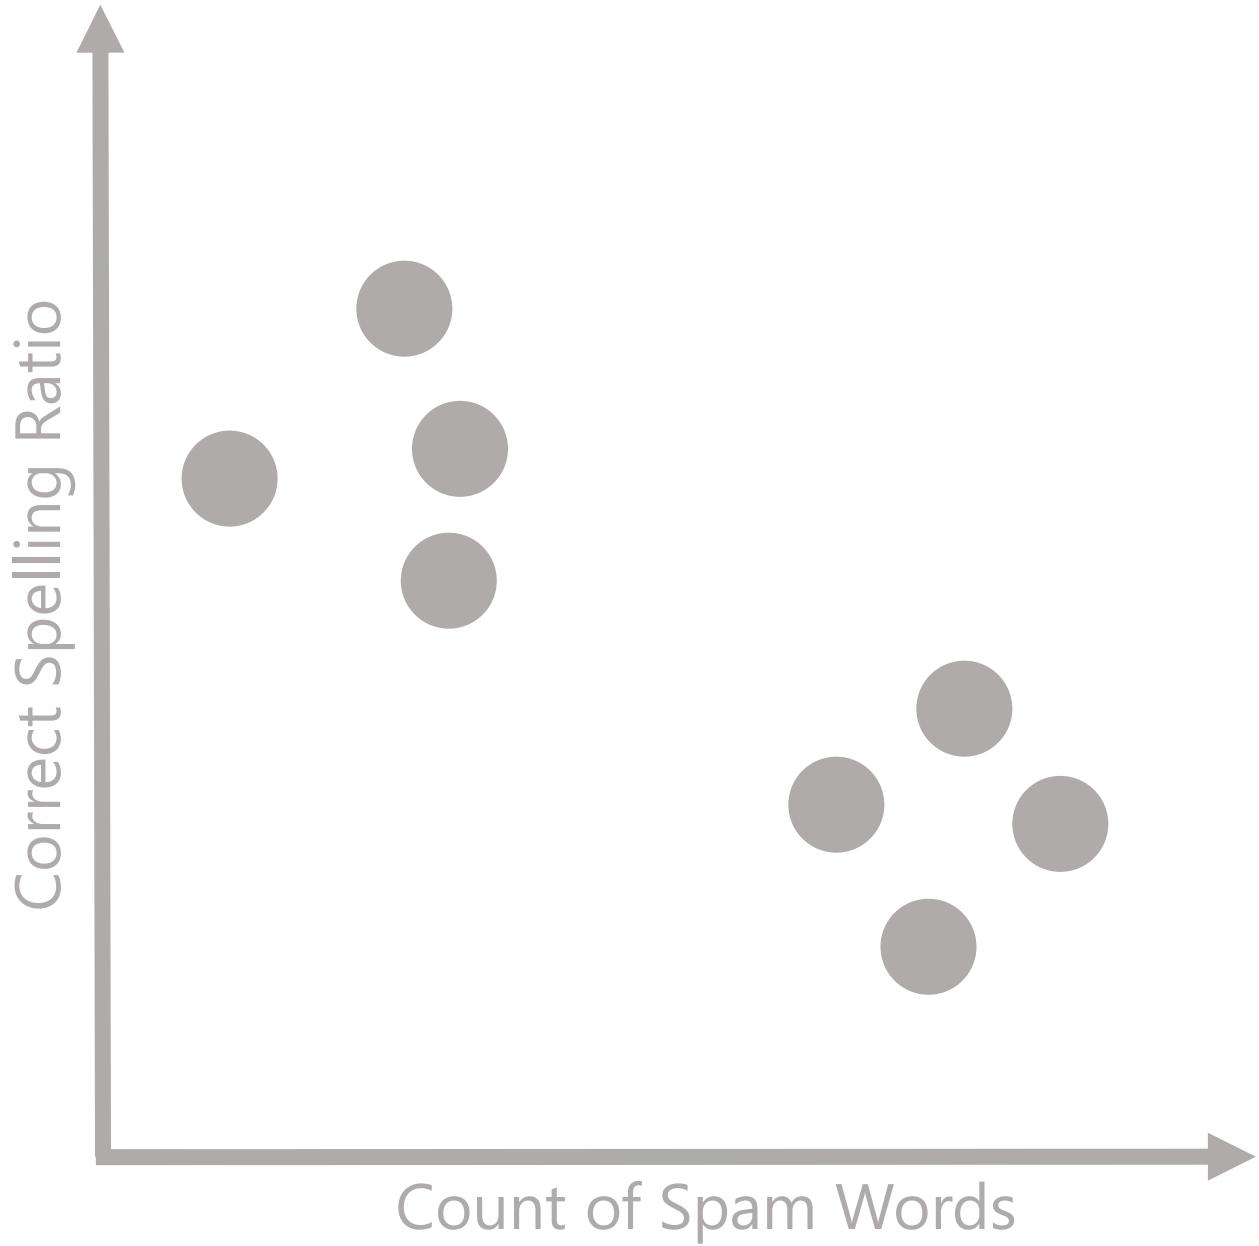
## Python Language Basics

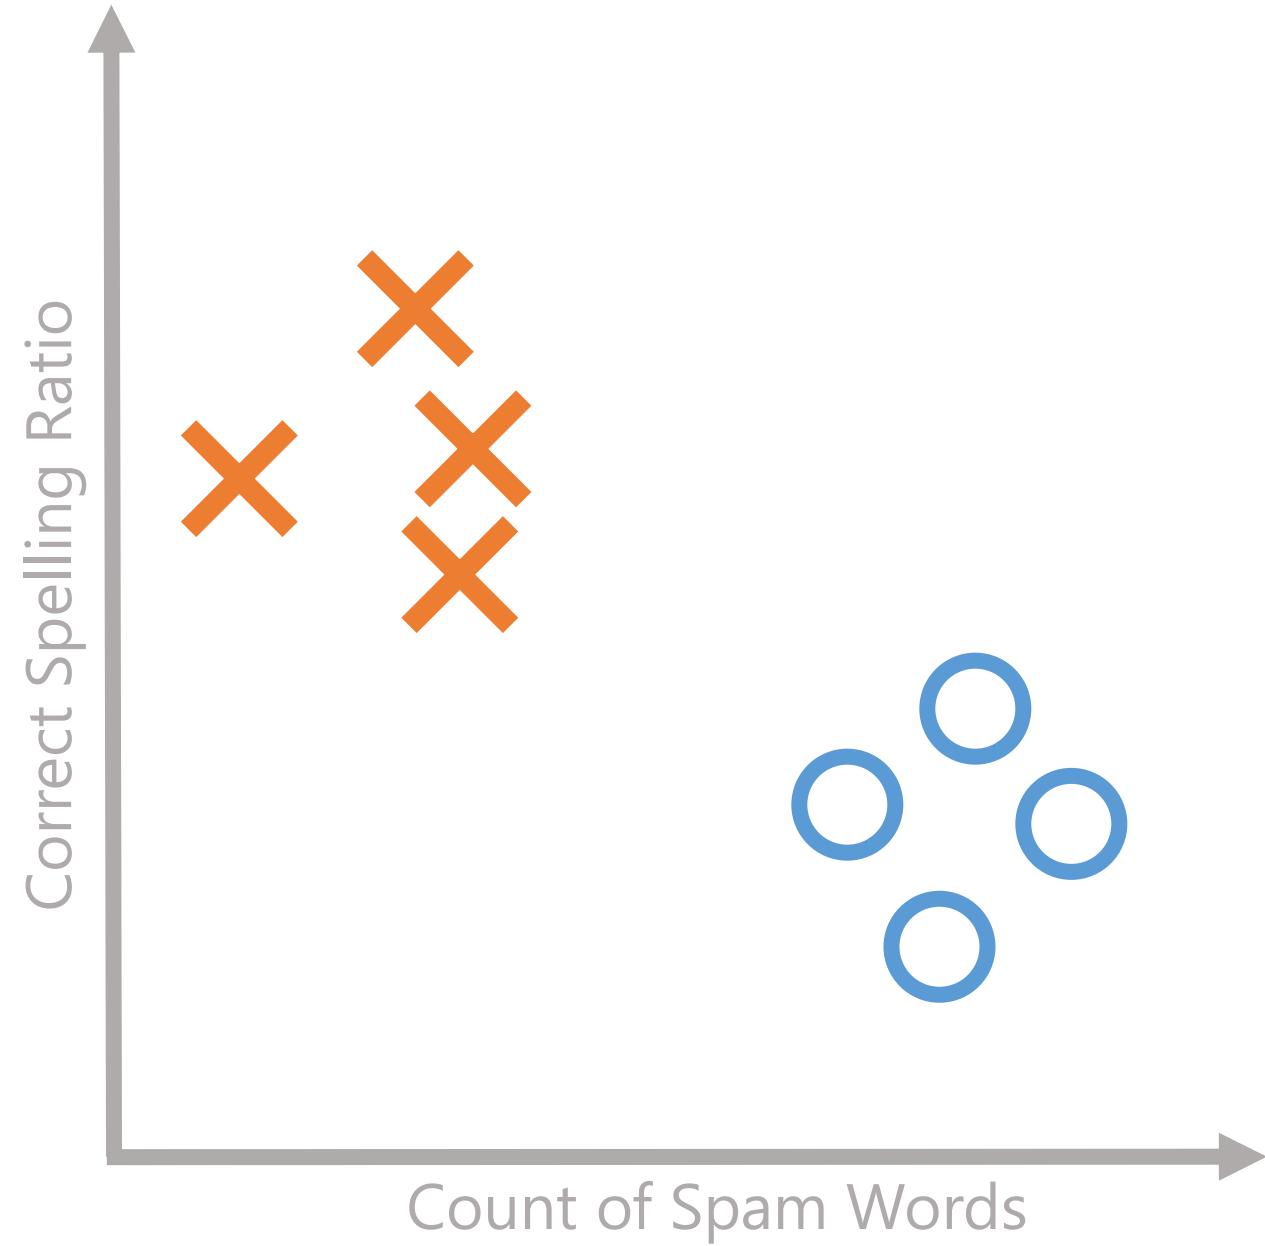
# Classification

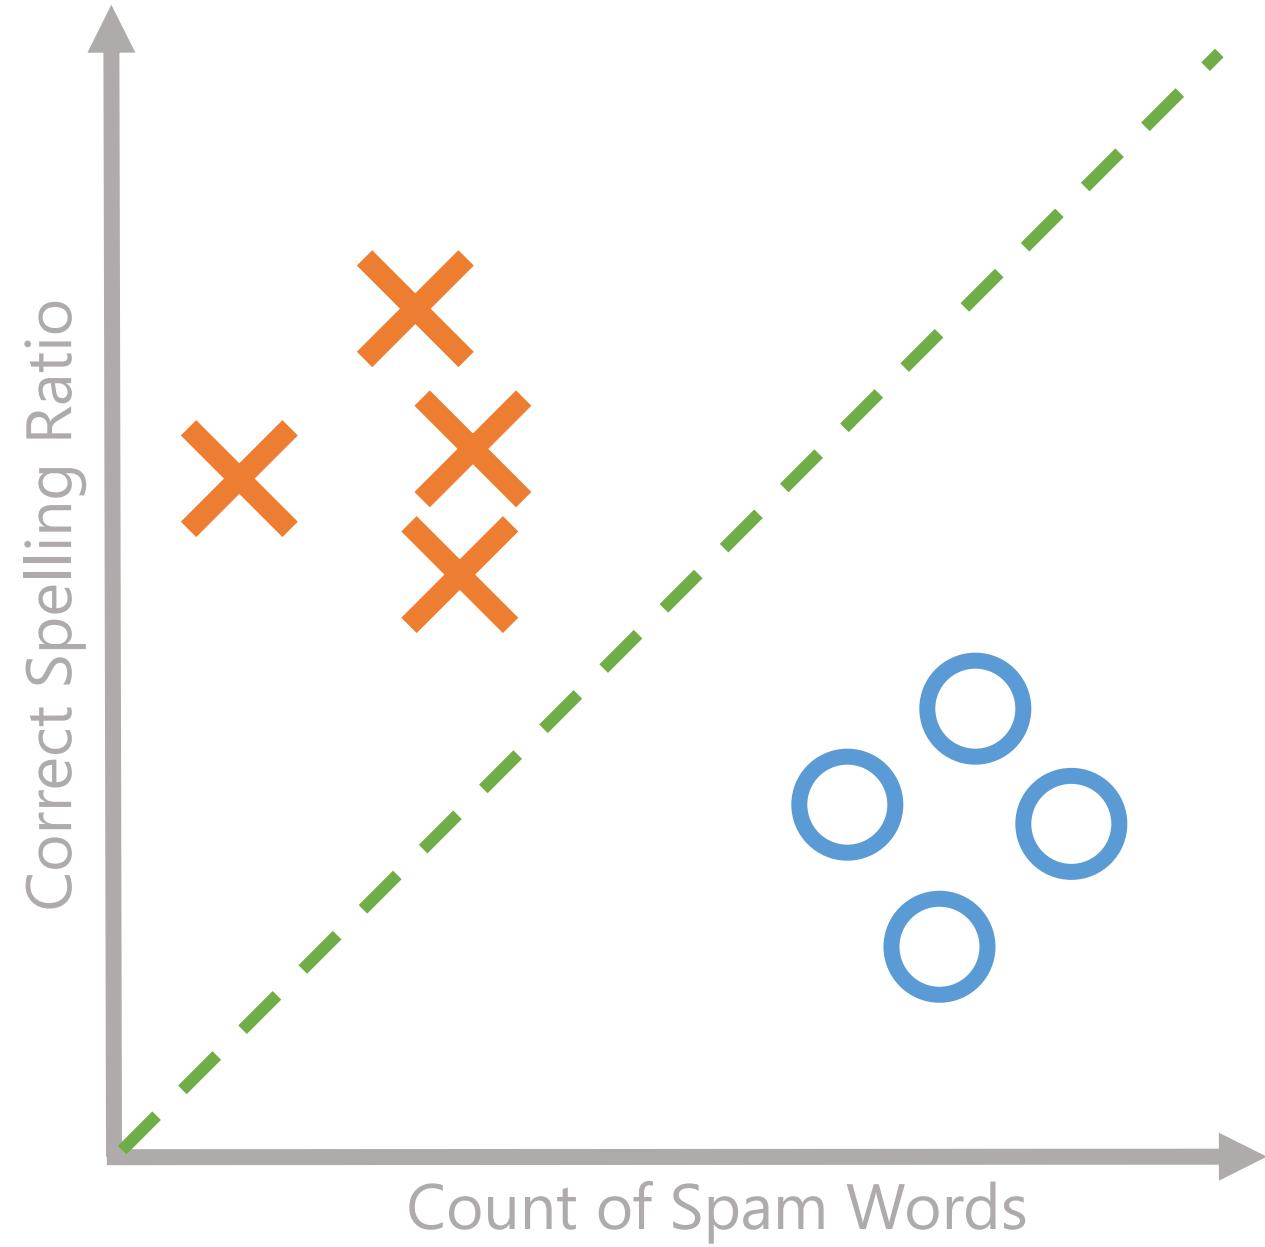


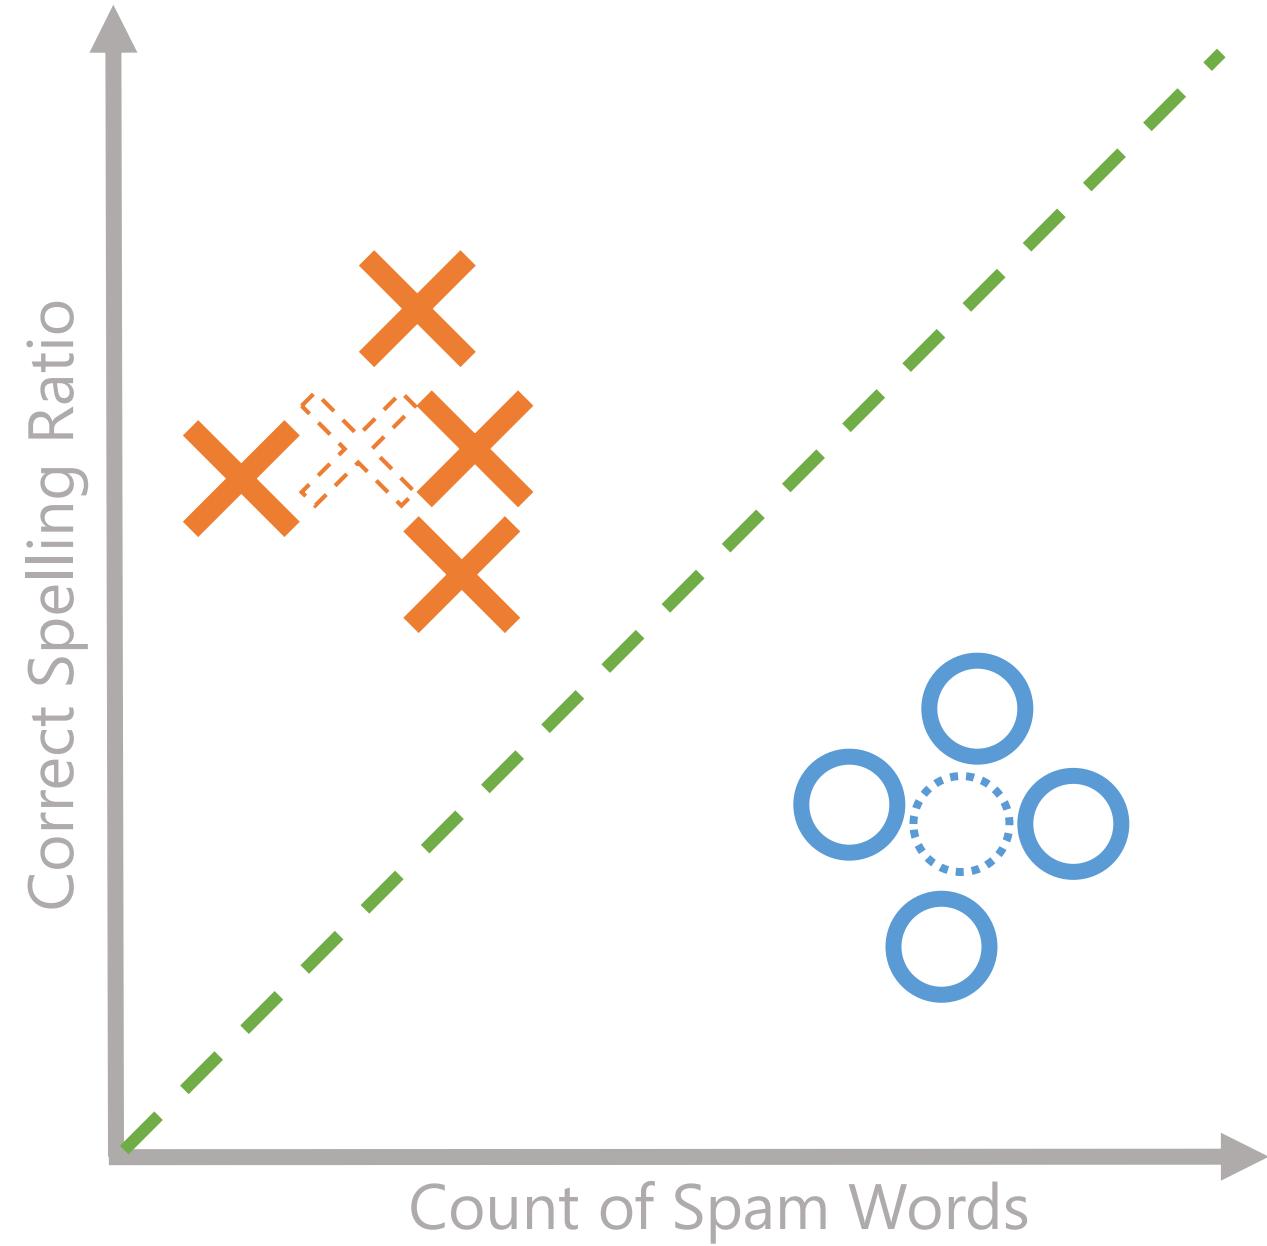


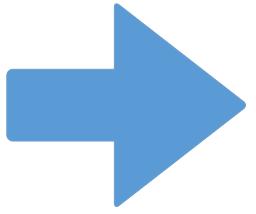
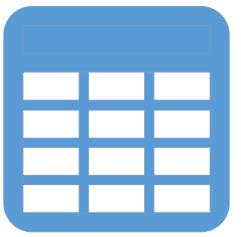
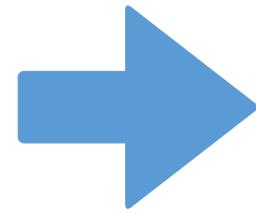






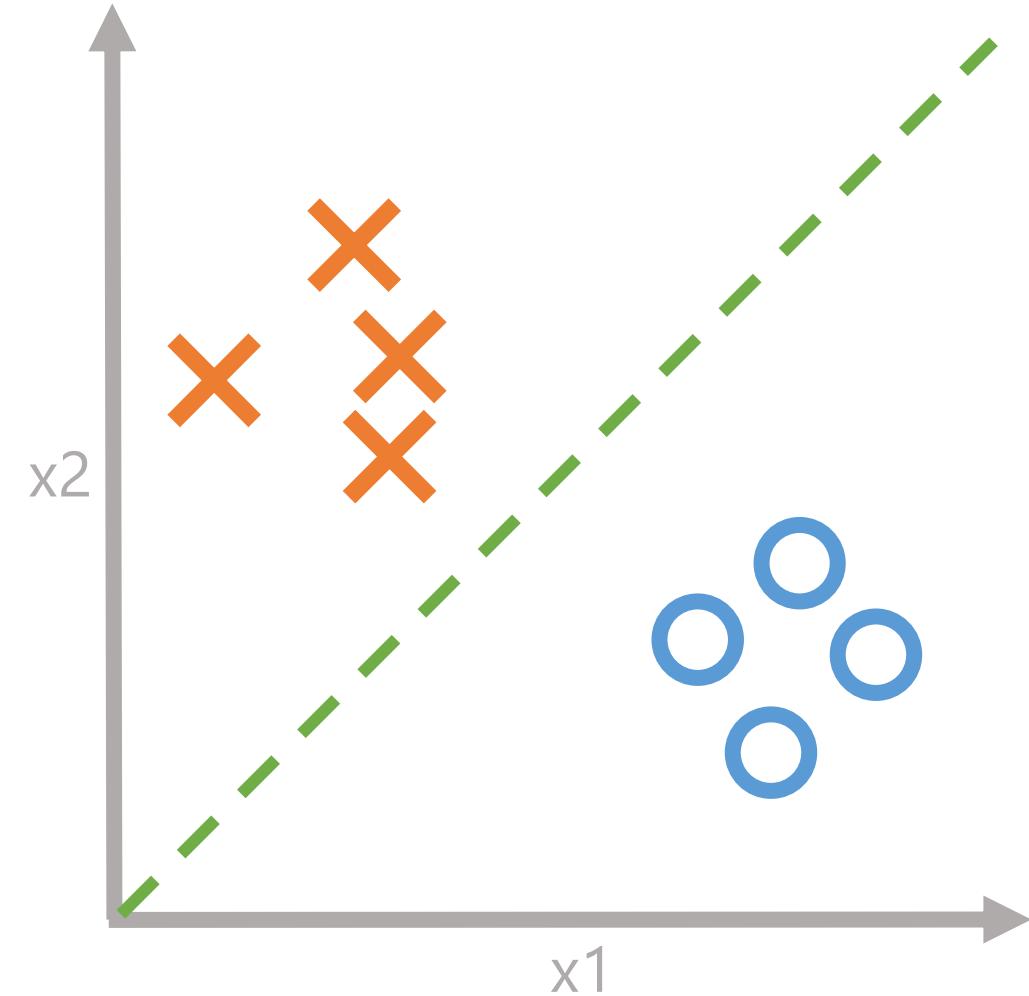




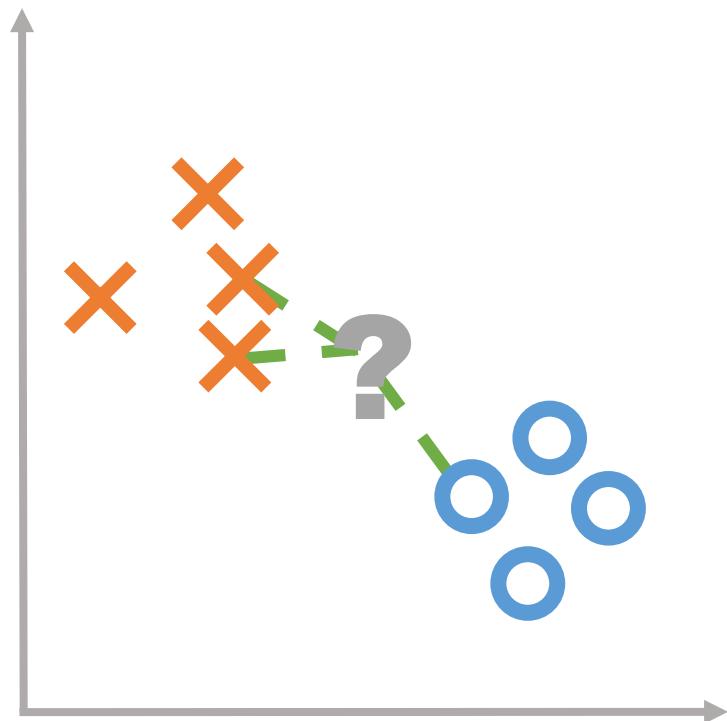
 $f(x)$ 

# Classification Algorithms

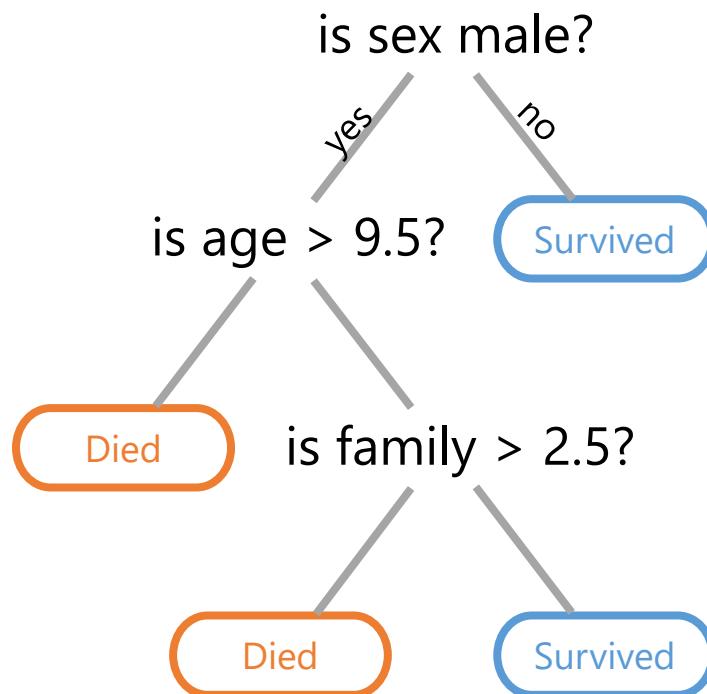
k-Nearest Neighbors  
Decision Tree Classifier  
Naïve Bayes Classifier  
Support Vector Machine  
Neural Network Classifier



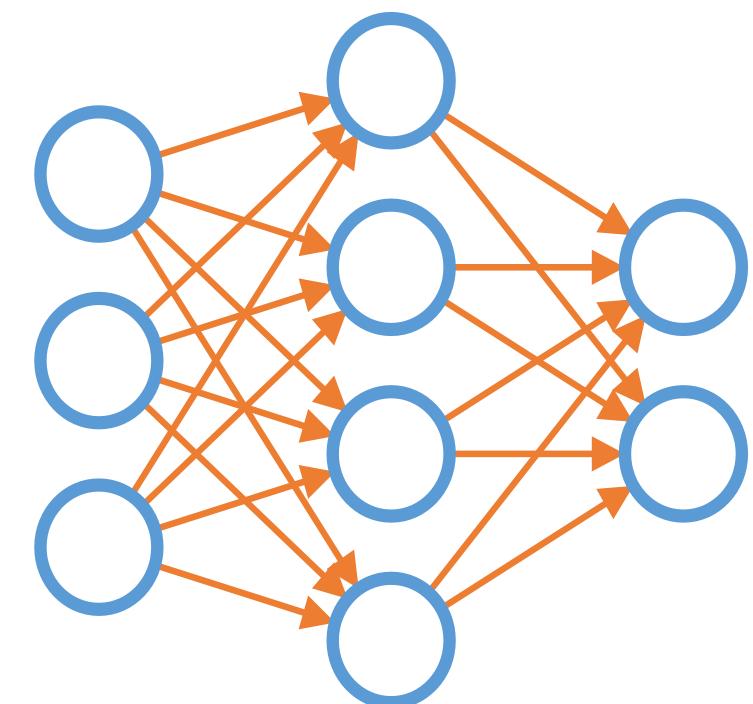
# Classification Algorithms



k-Nearest Neighbors

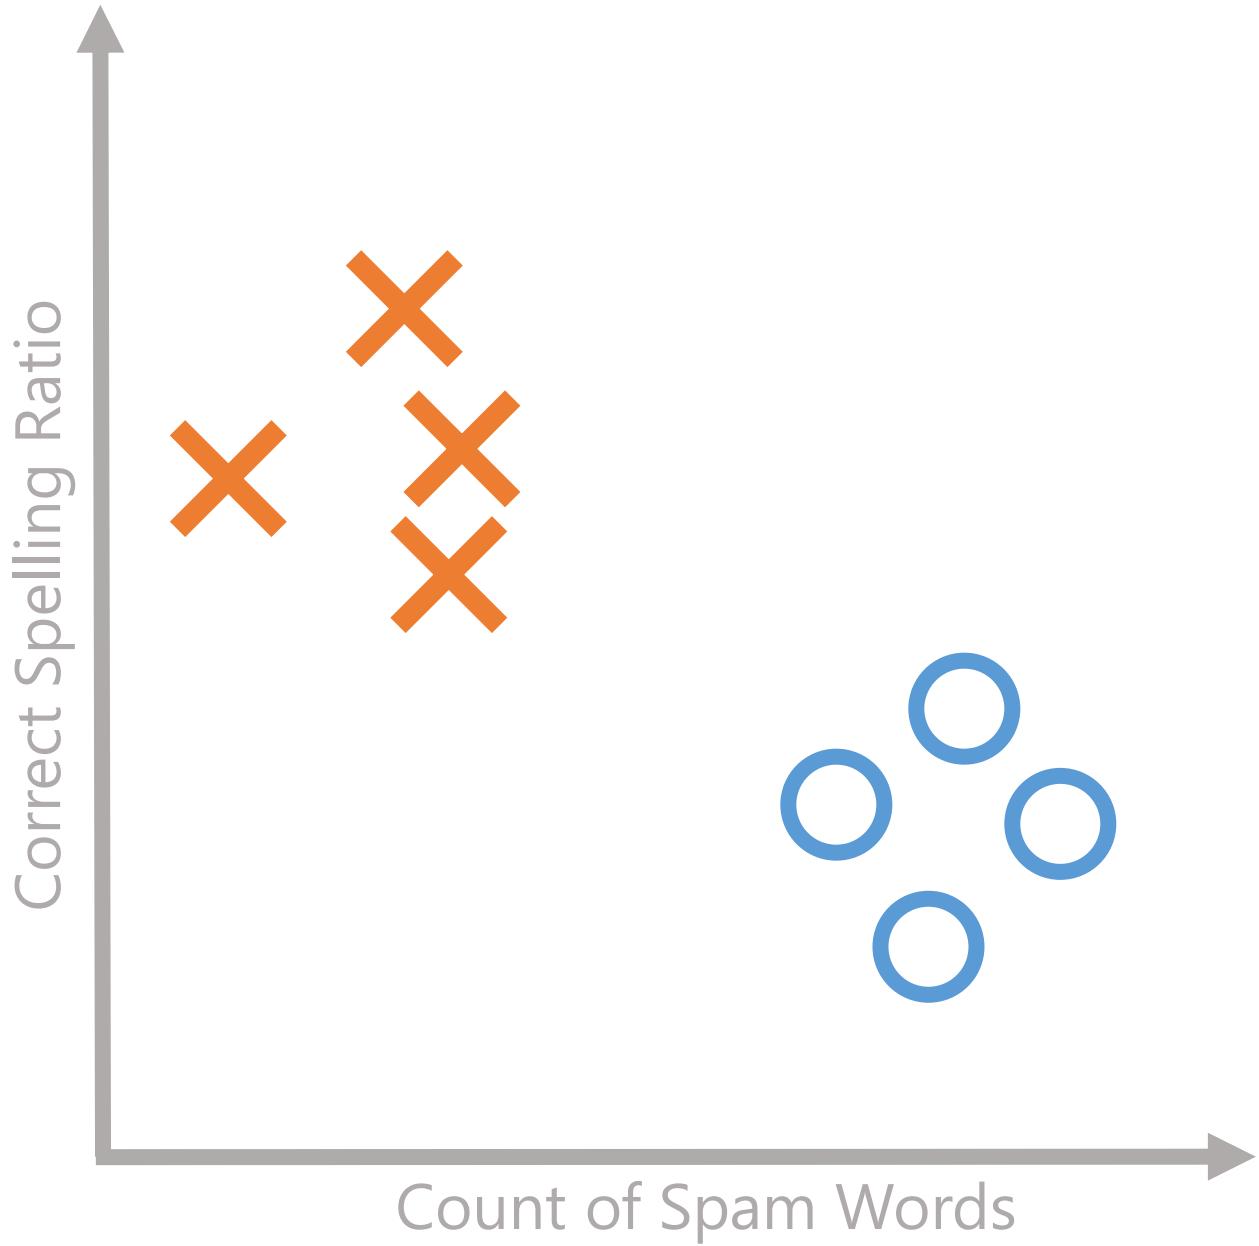


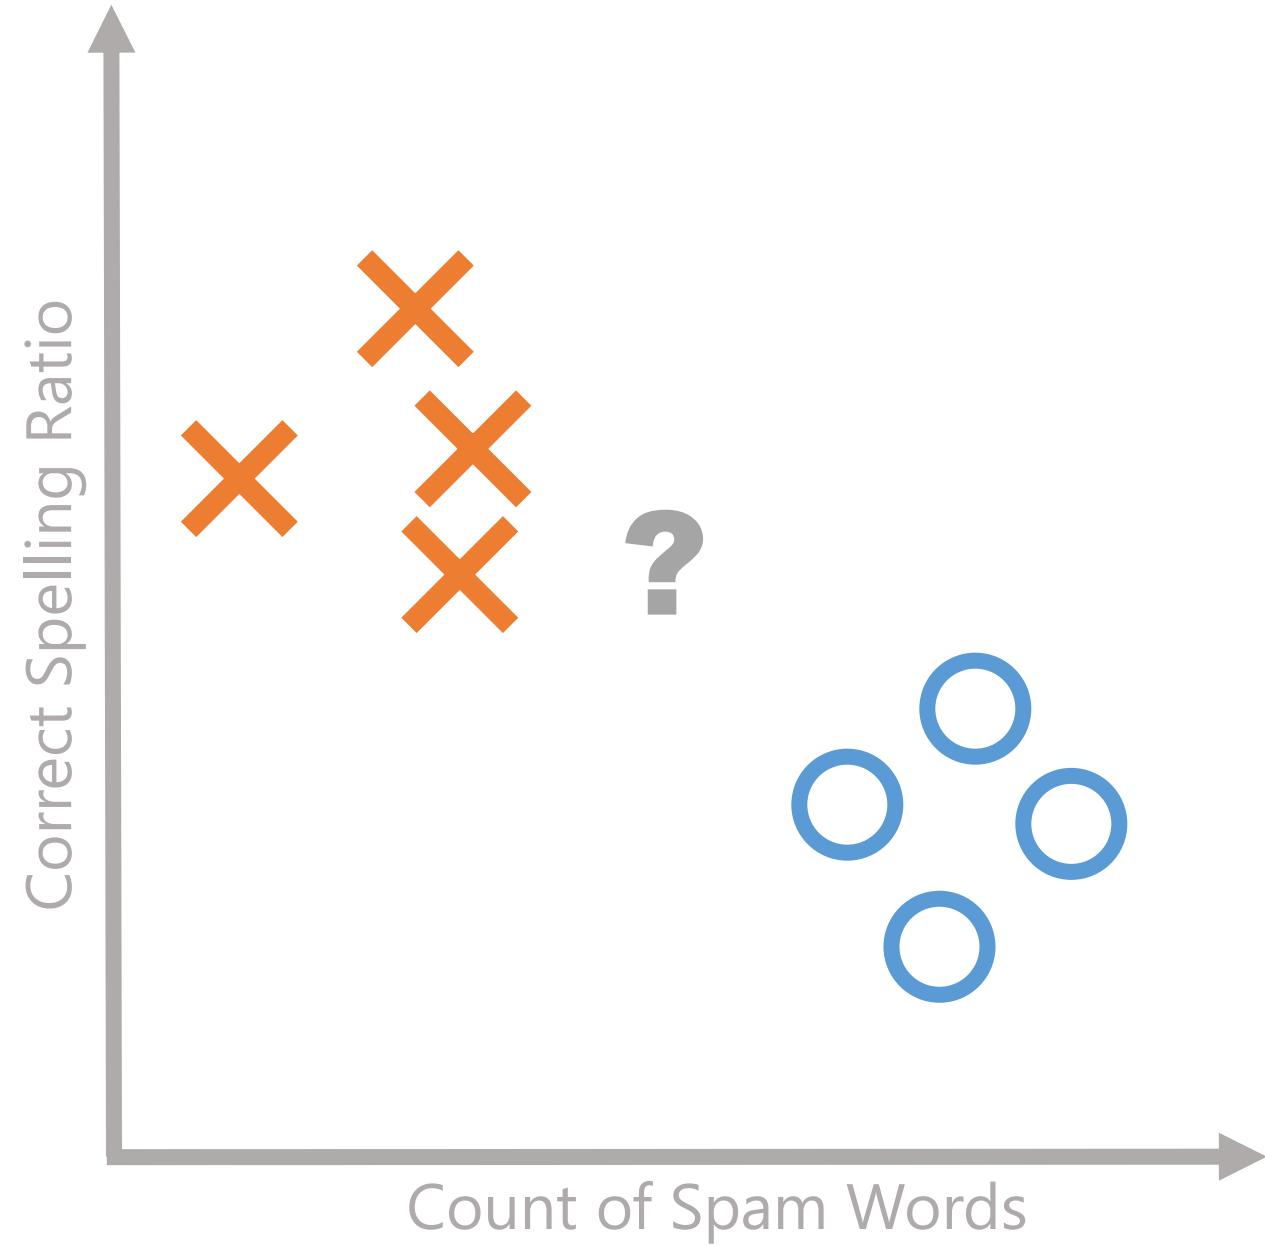
Decision Tree

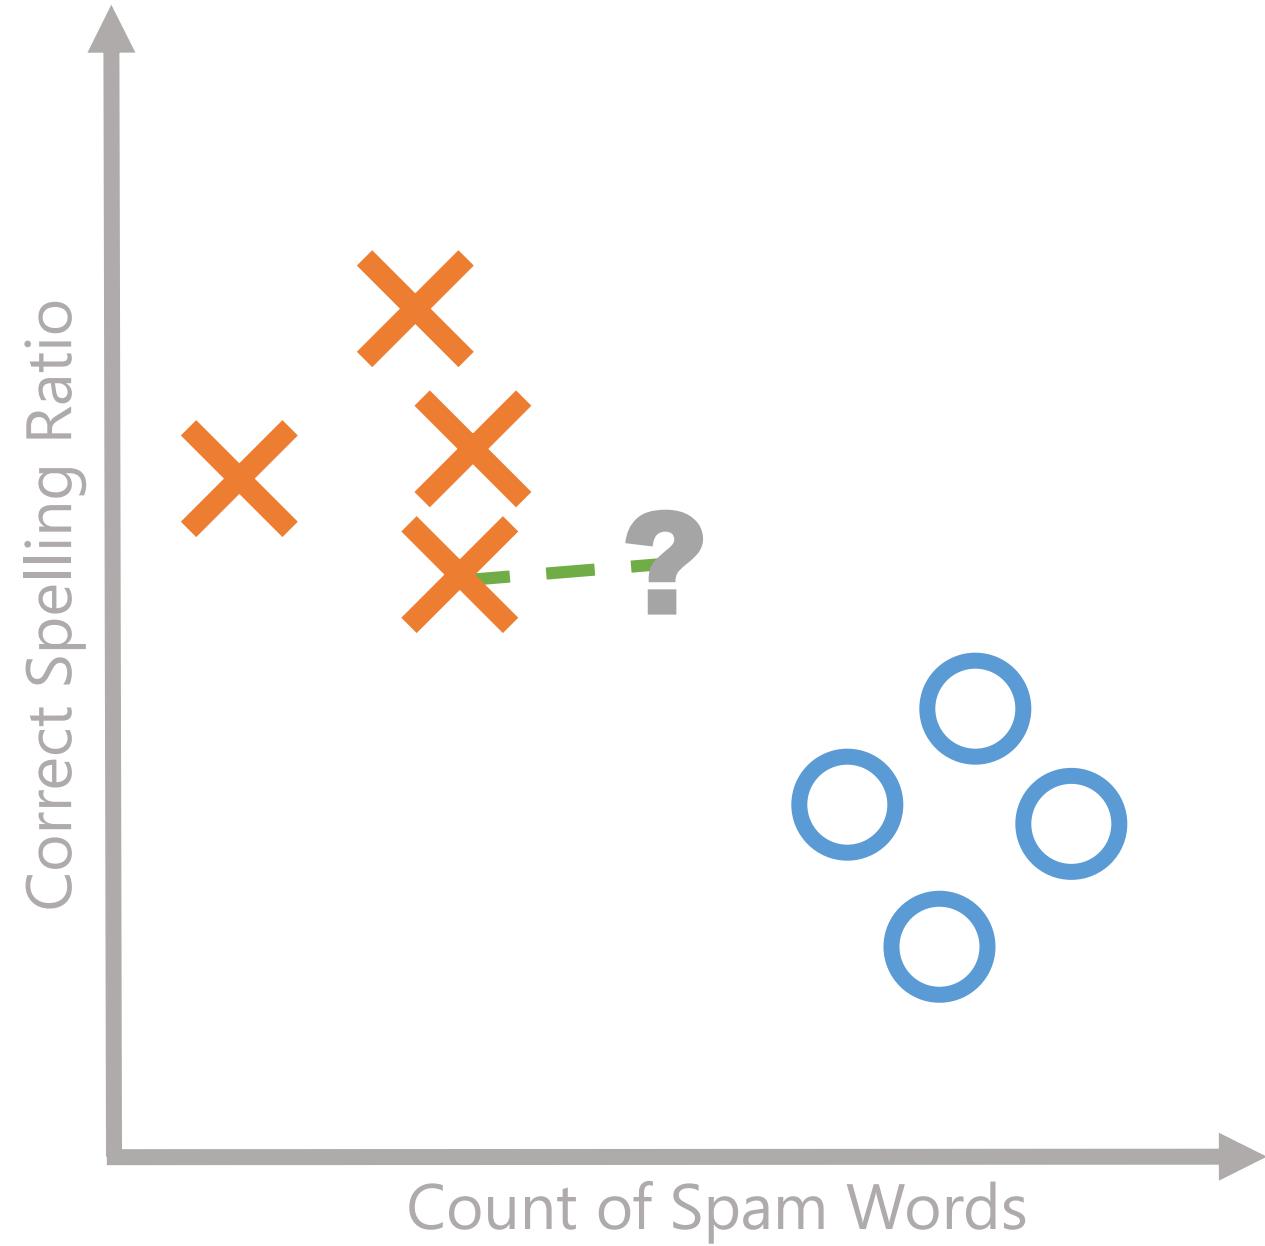


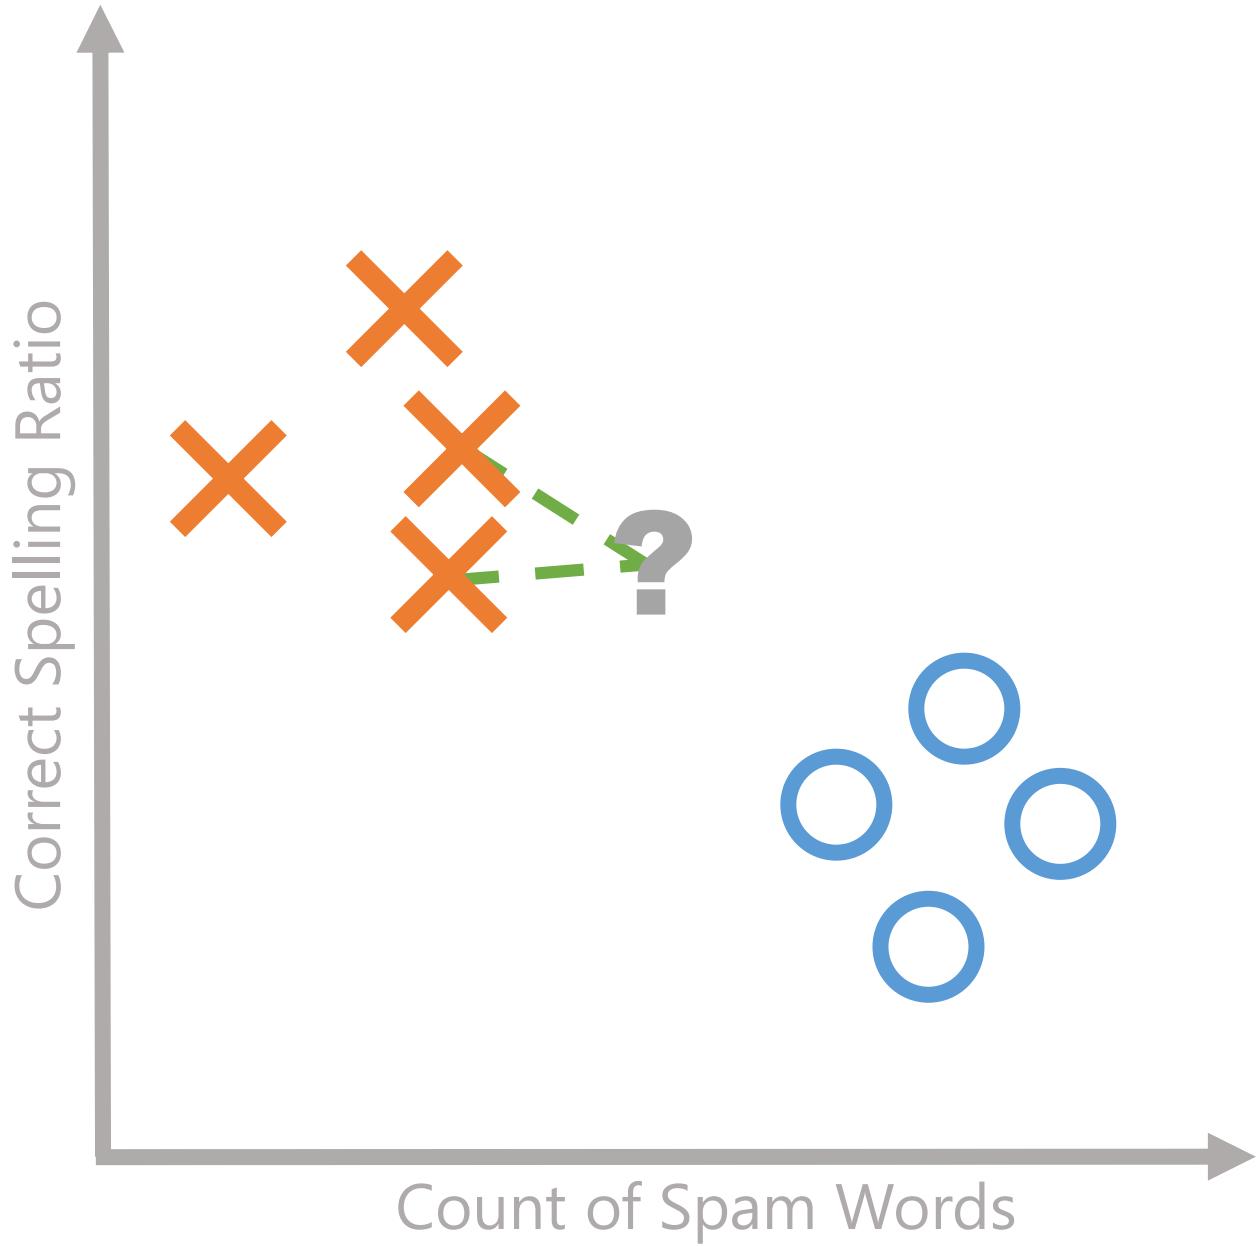
Neural Network

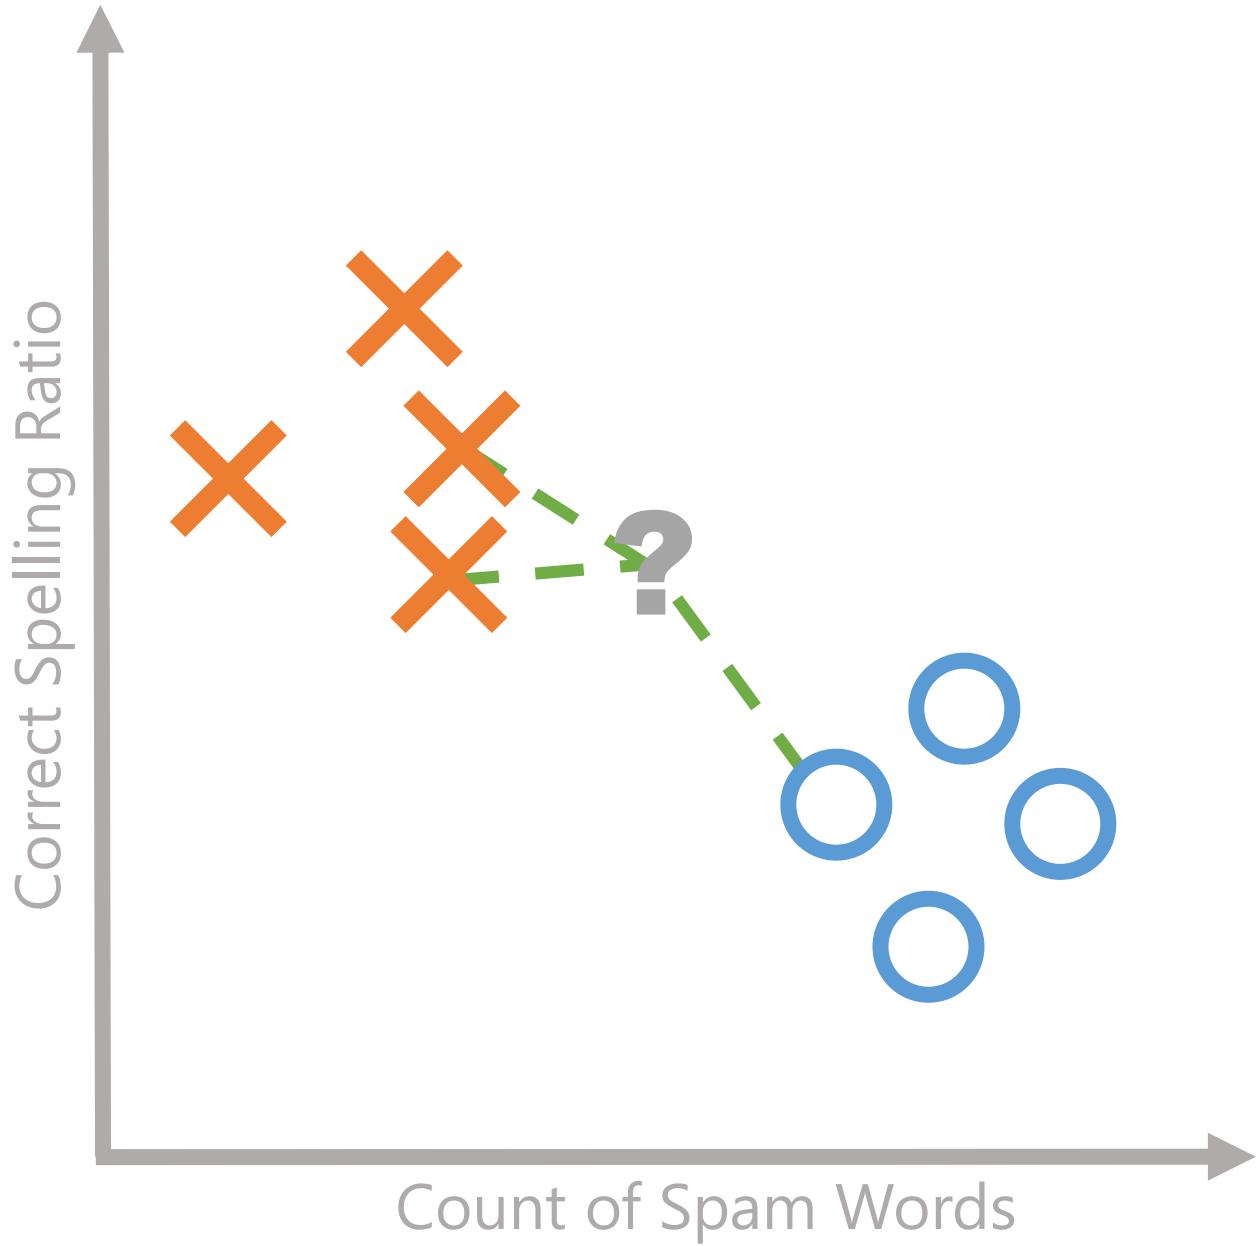
# k-Nearest Neighbors Classifier

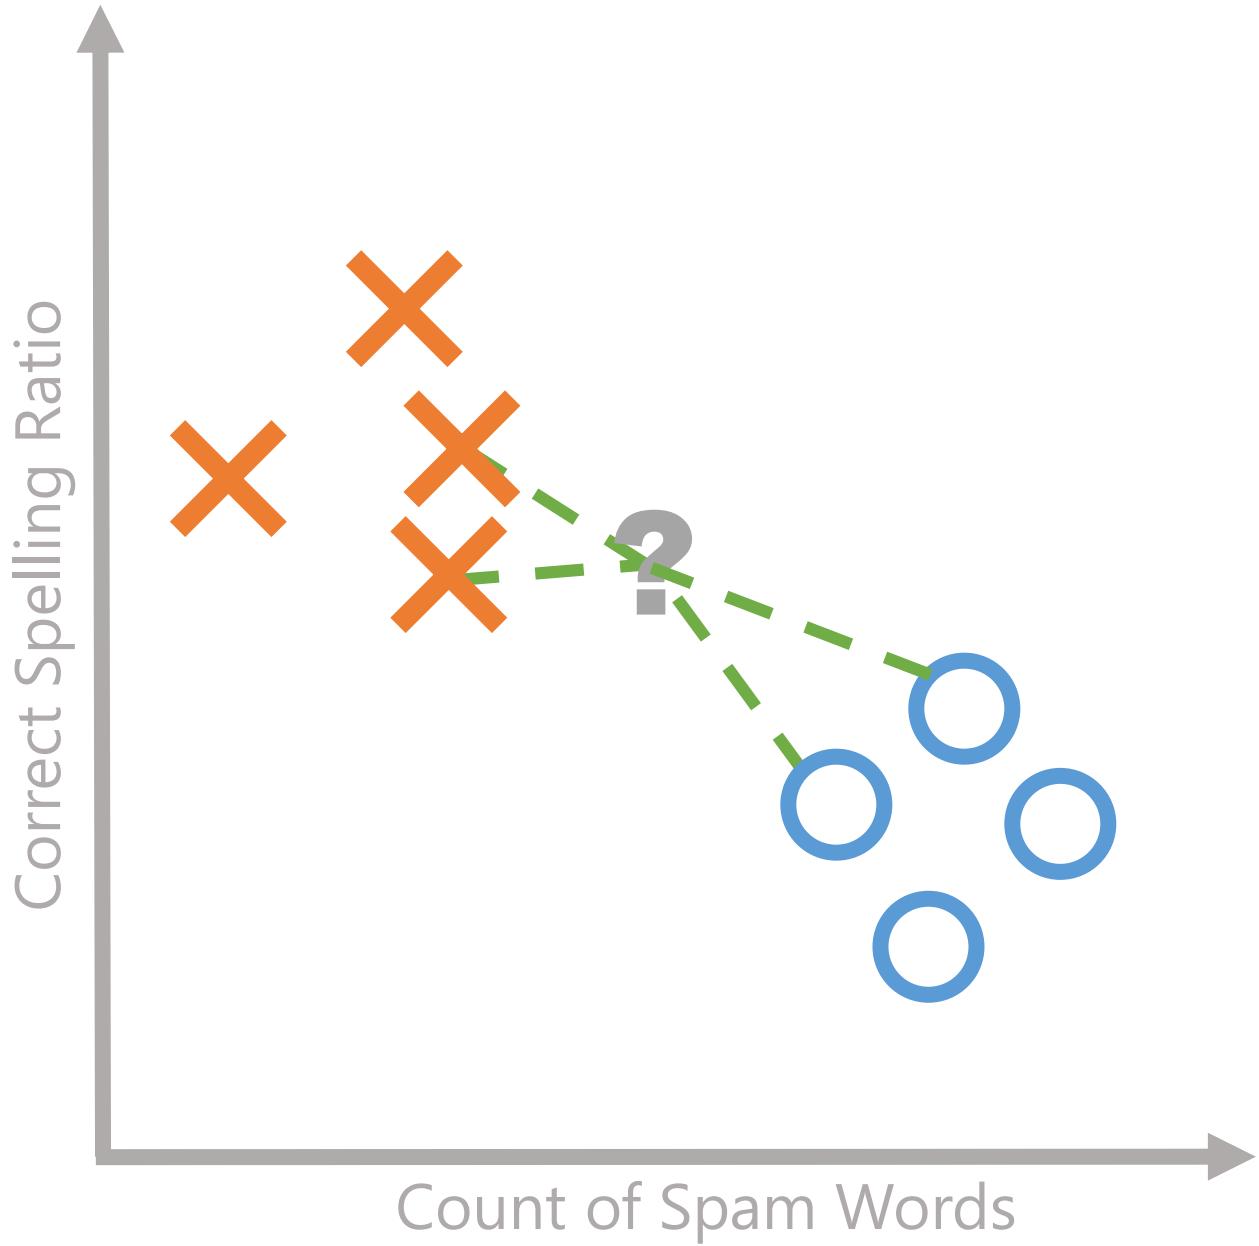






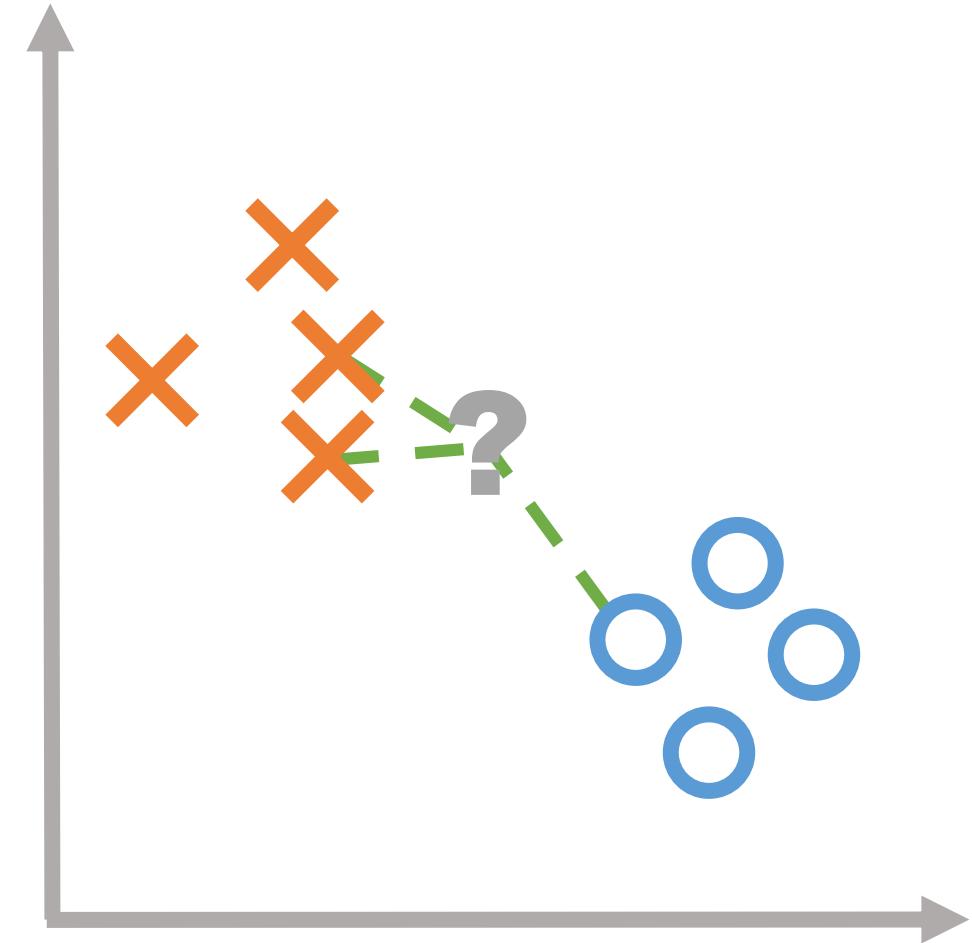






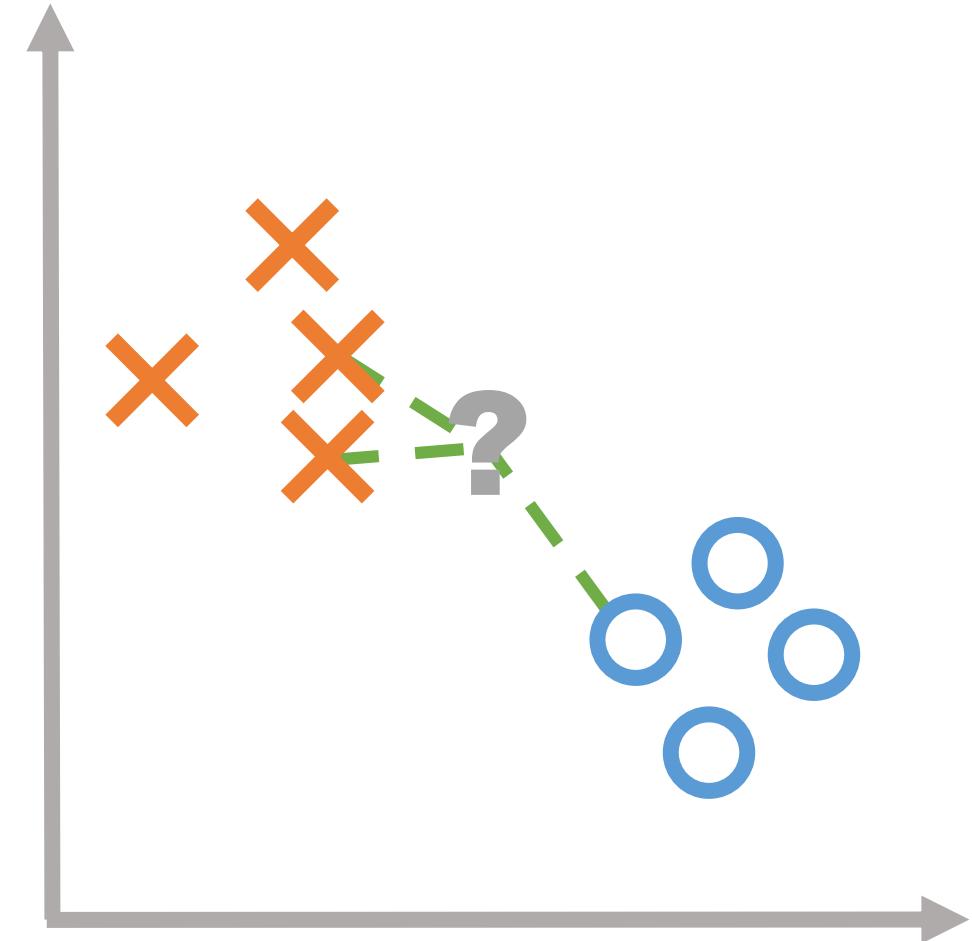
# K-Nearest Neighbors Classifier

Supervised learning



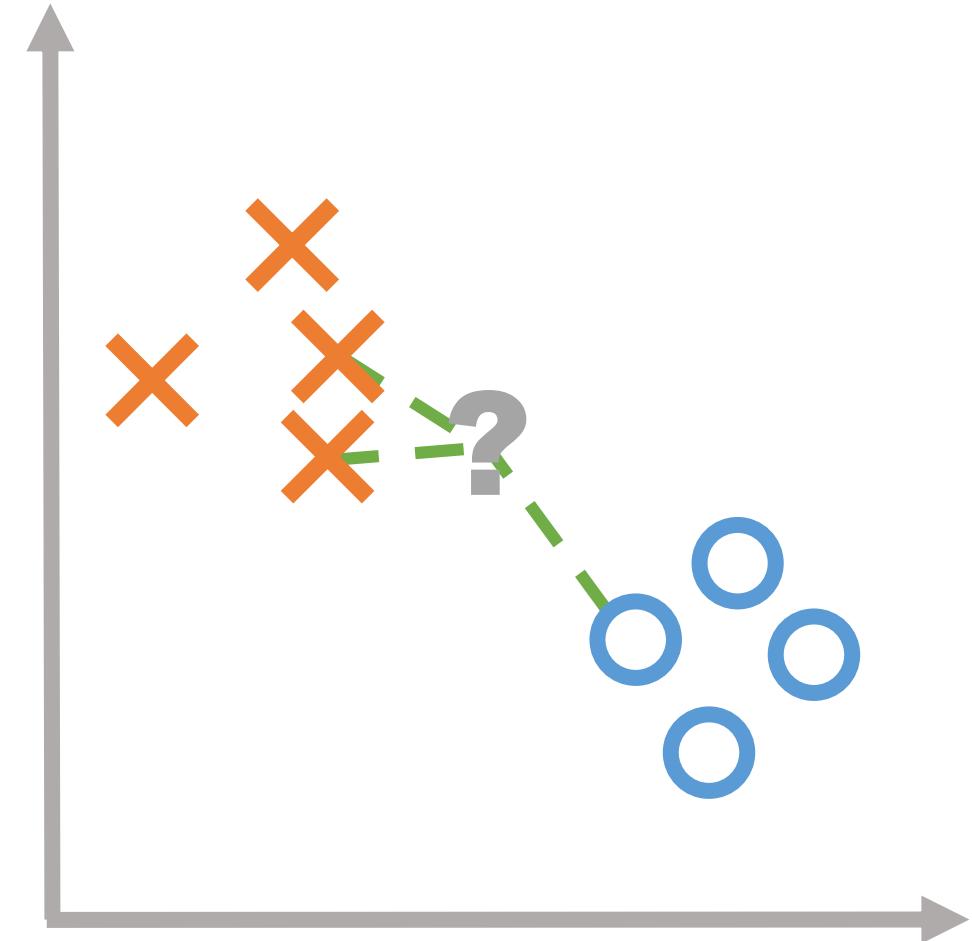
# K-Nearest Neighbors Classifier

Supervised learning  
Uses class of neighbors



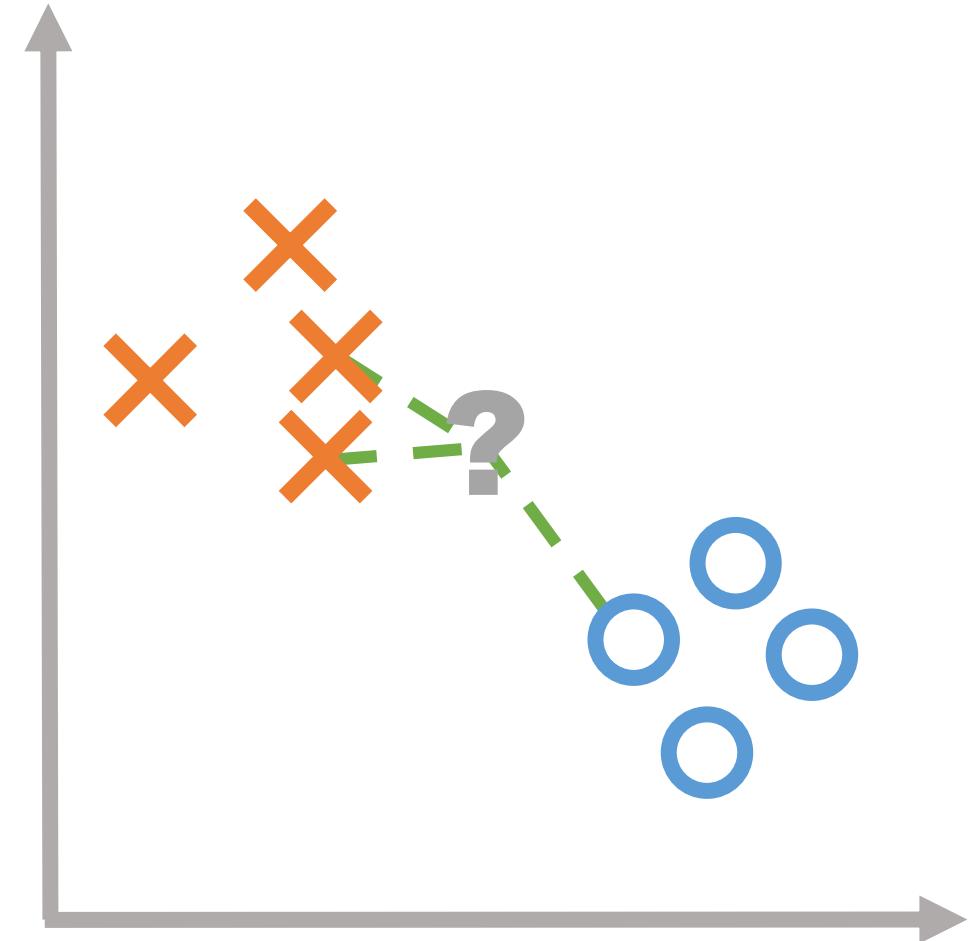
# K-Nearest Neighbors Classifier

Supervised learning  
Uses class of neighbors  
 $k$  specifies how many

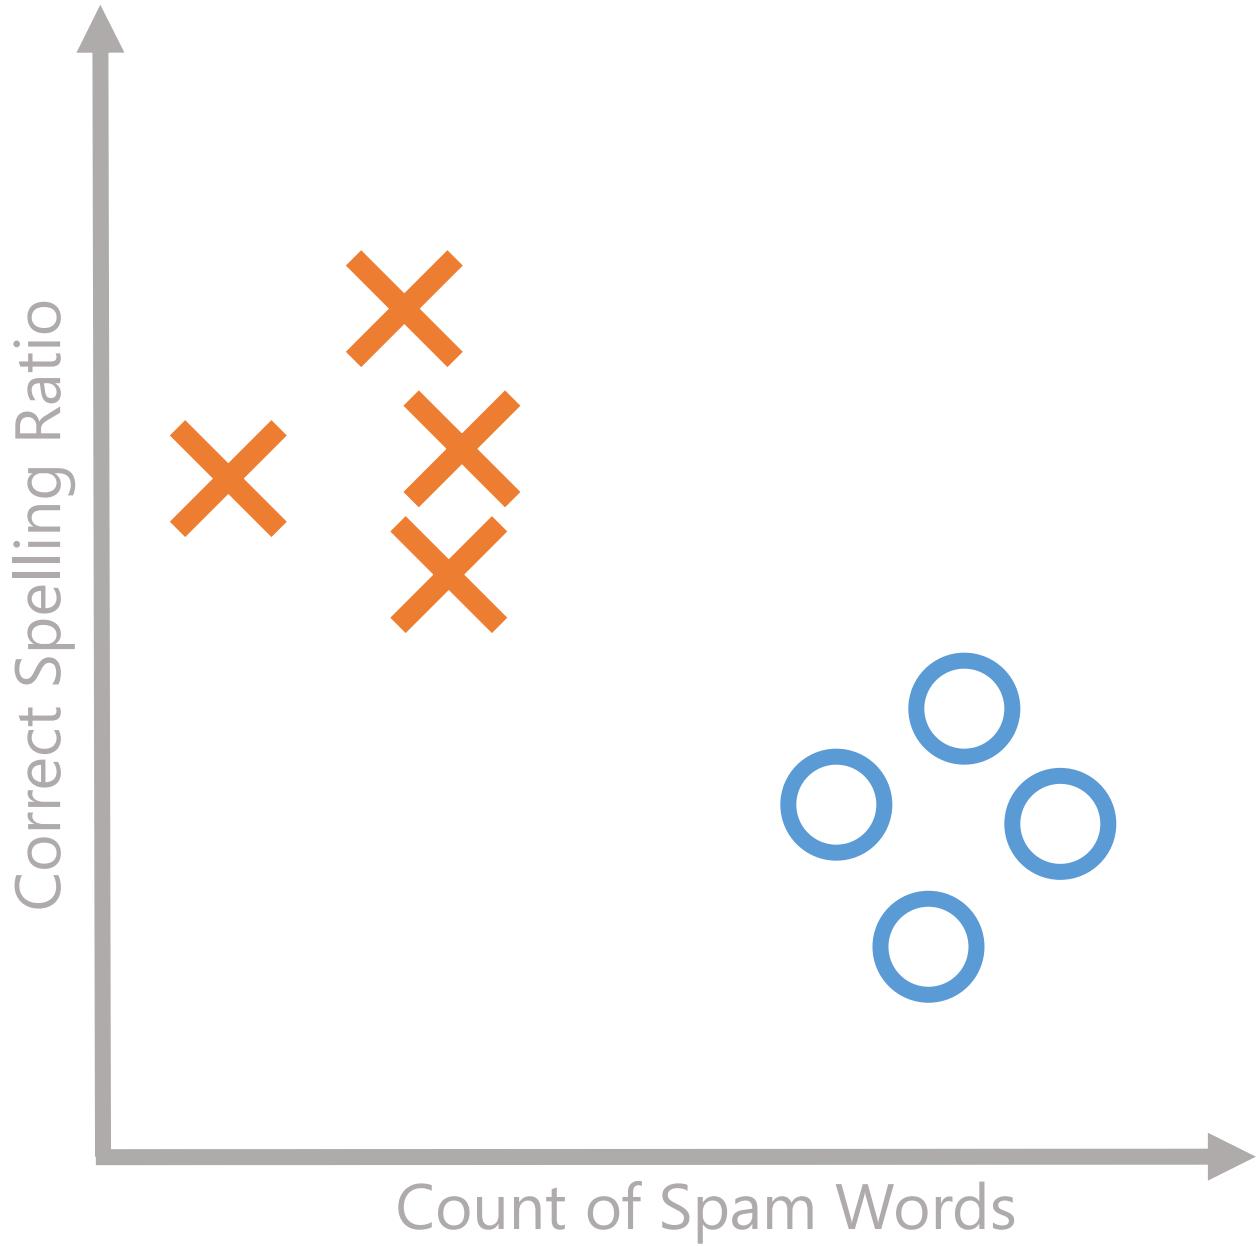


# K-Nearest Neighbors Classifier

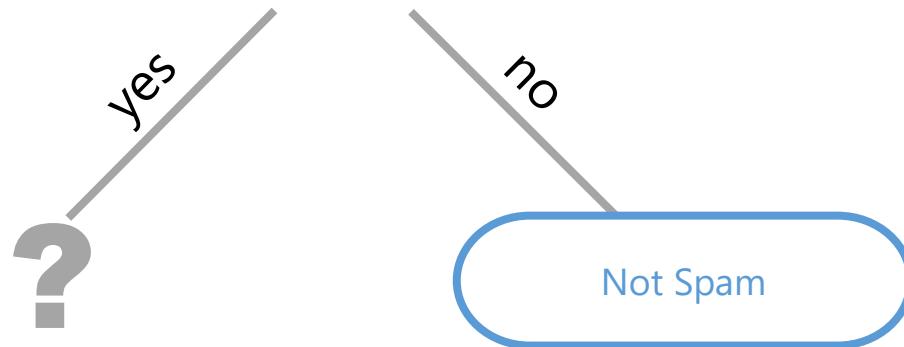
Supervised learning  
Uses class of neighbors  
 $k$  specifies how many  
Simple and easy



# Decision Tree Classifier



Is count of spam words > 5?



Is count of spam words > 5?

yes

Is correct-spelling ratio > 50%

yes

Not Spam

no

Not Spam

?

Is count of spam words > 5?

yes

Is correct-spelling ratio > 50%?

yes

Not Spam

no

Not Spam

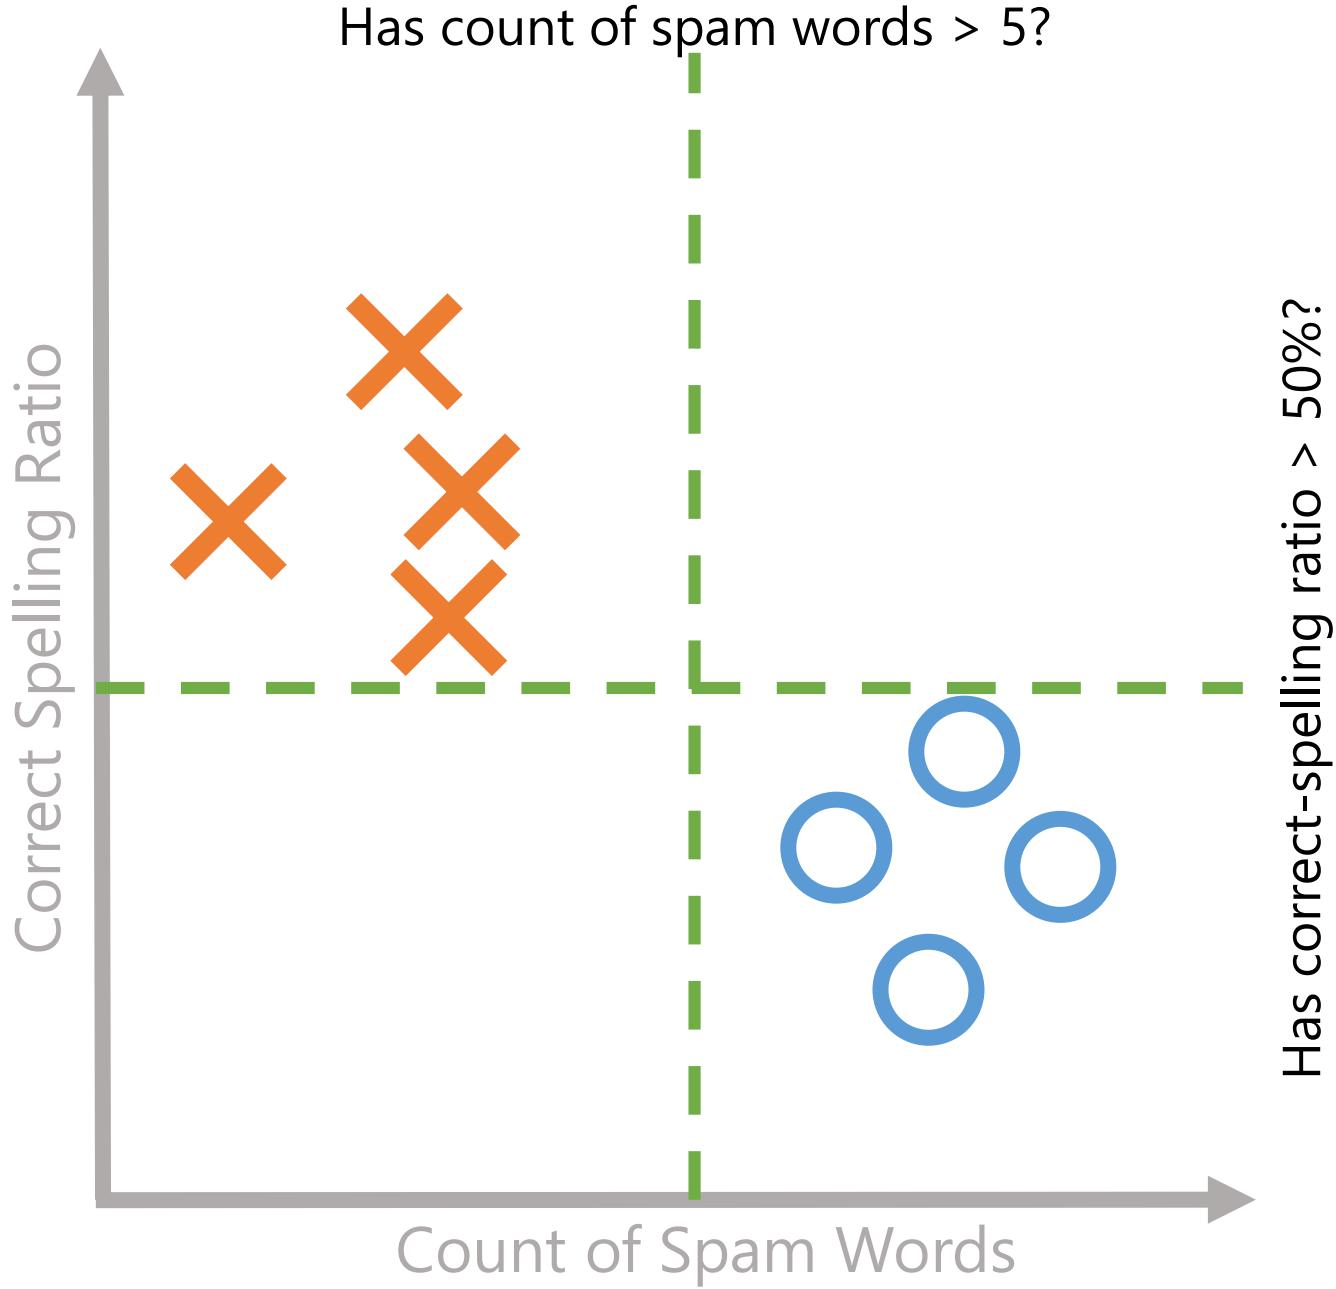
Is known contact?

yes

Not spam

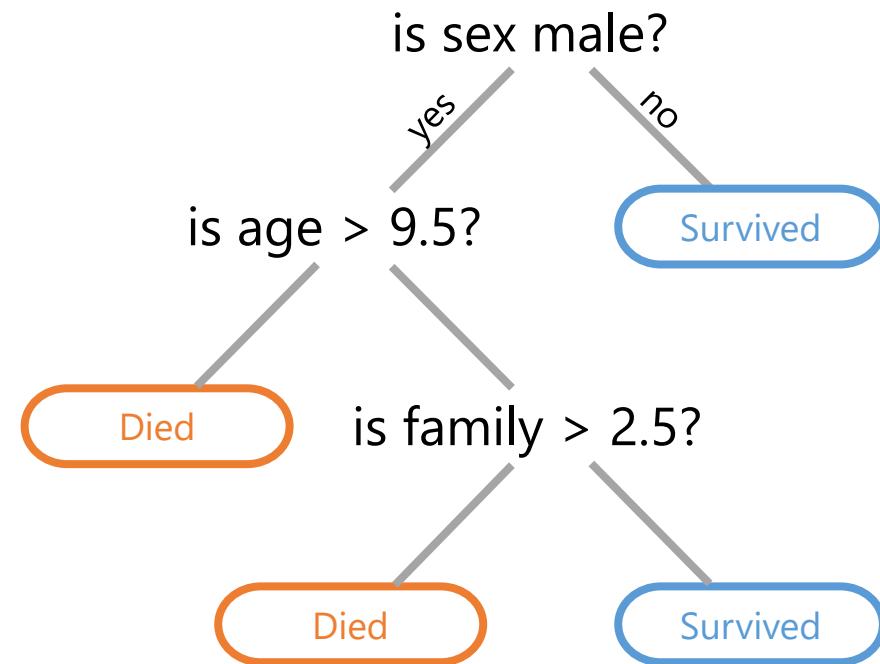
no

Spam



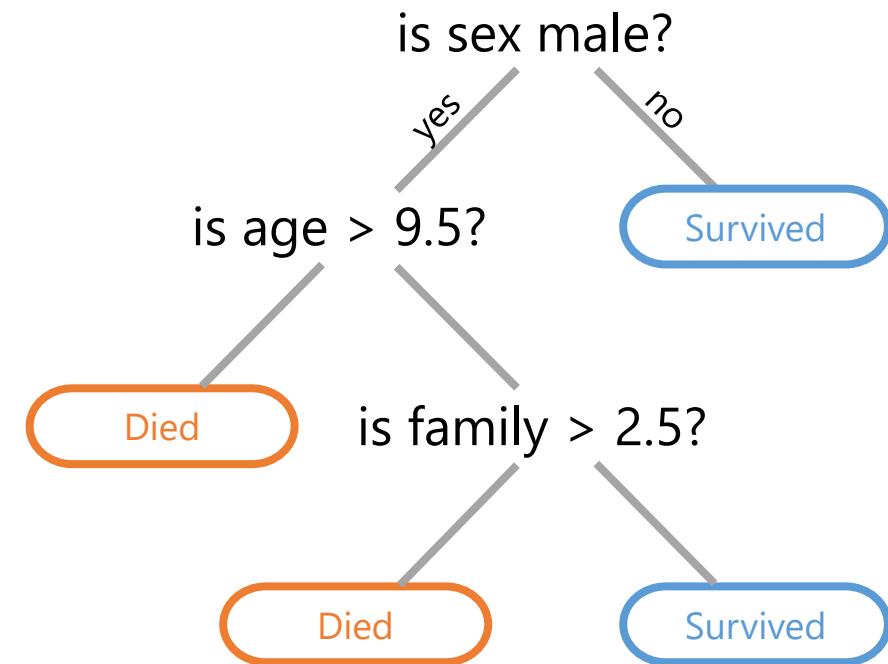
# Decision Tree Classifier

Supervised learning



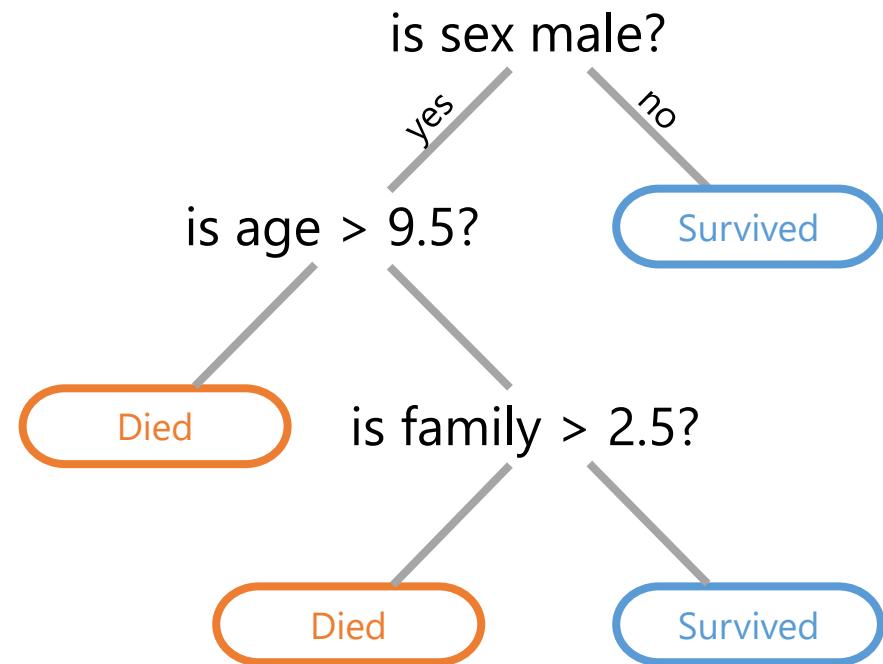
# Decision Tree Classifier

Supervised learning  
Tree of decisions



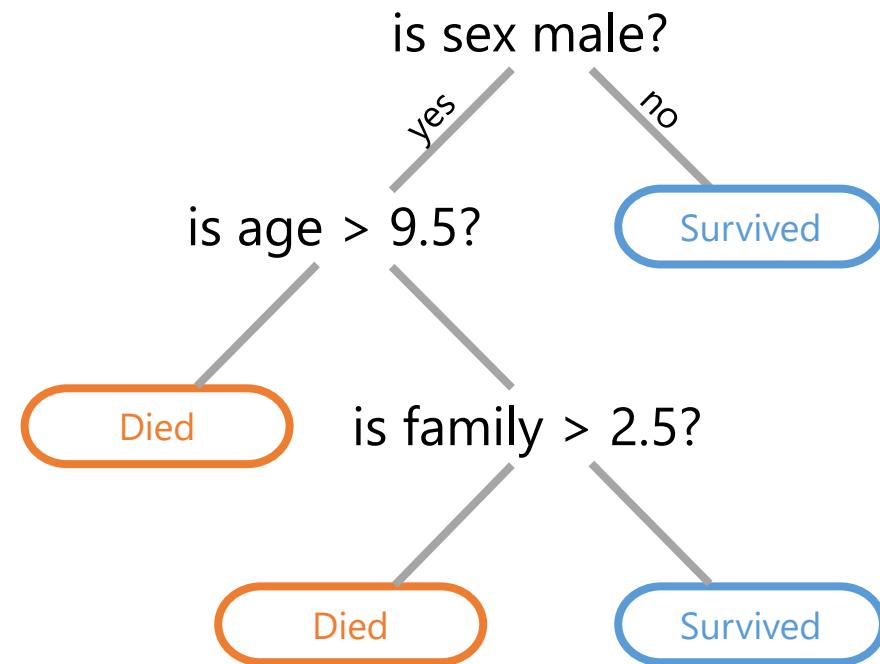
# Decision Tree Classifier

Supervised learning  
Tree of decisions  
Information gain

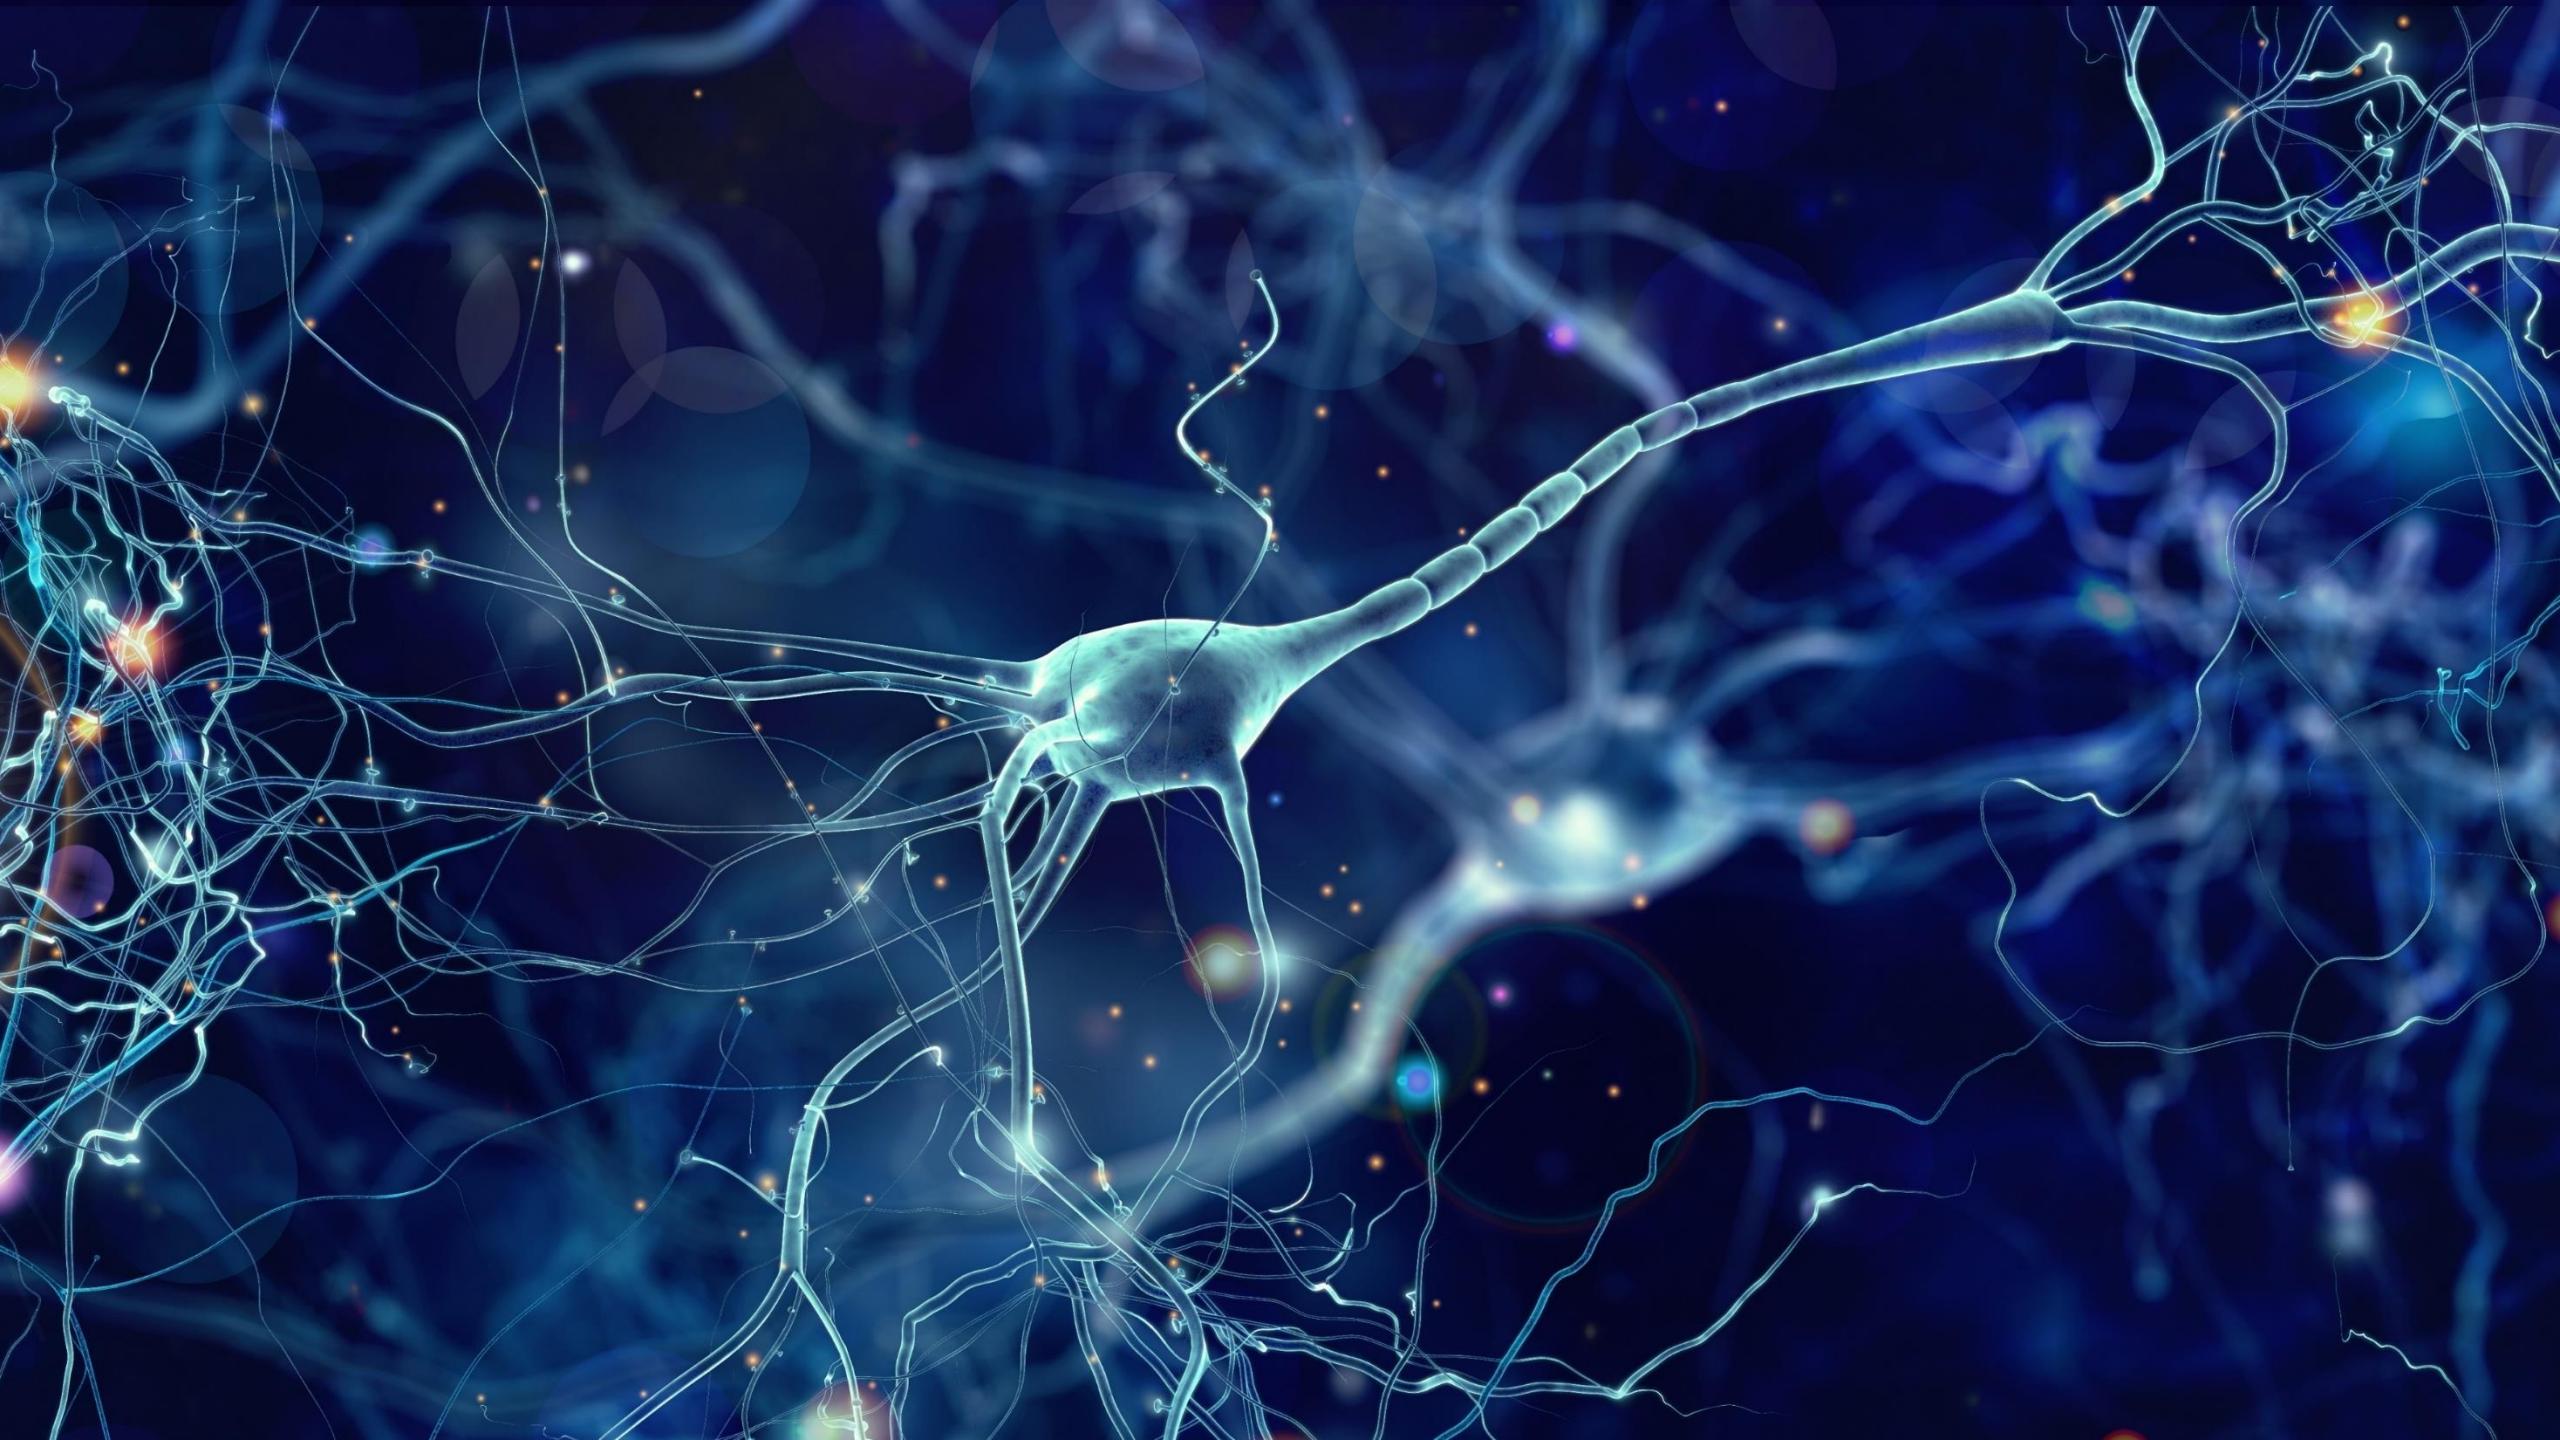


# Decision Tree Classifier

Supervised learning  
Tree of decisions  
Information gain  
Simple and easy

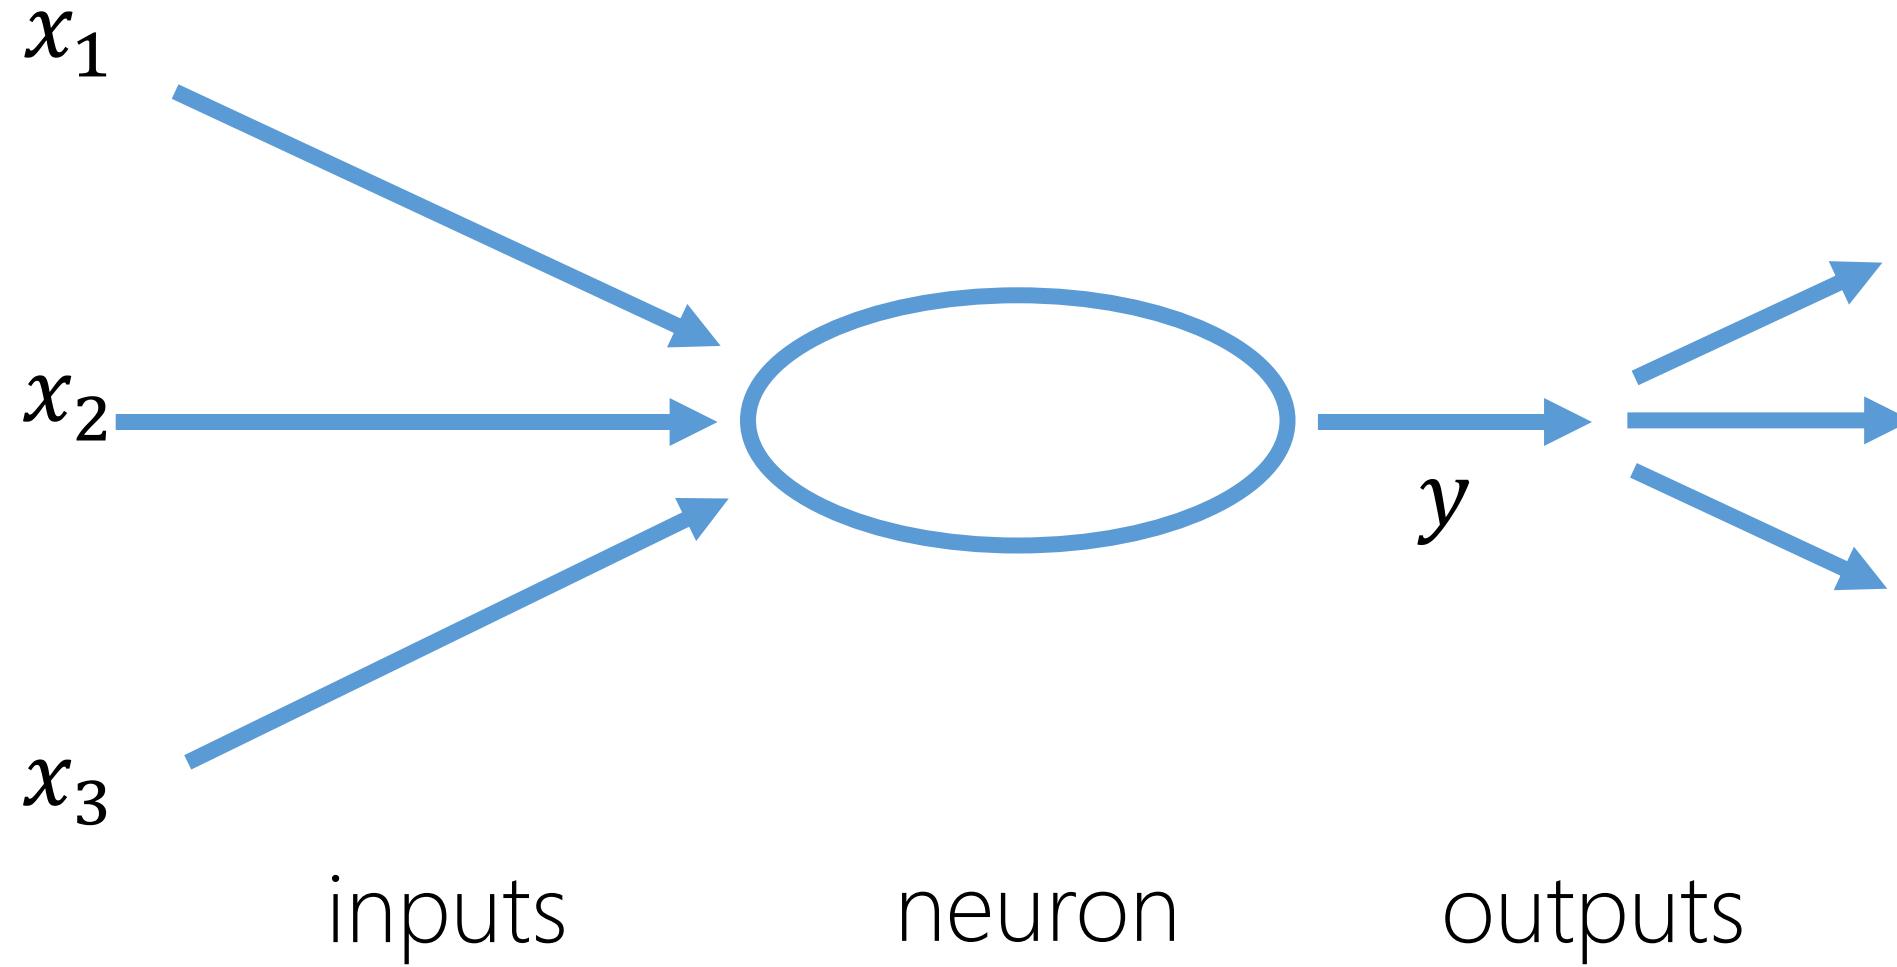


# Neural Network Classifier

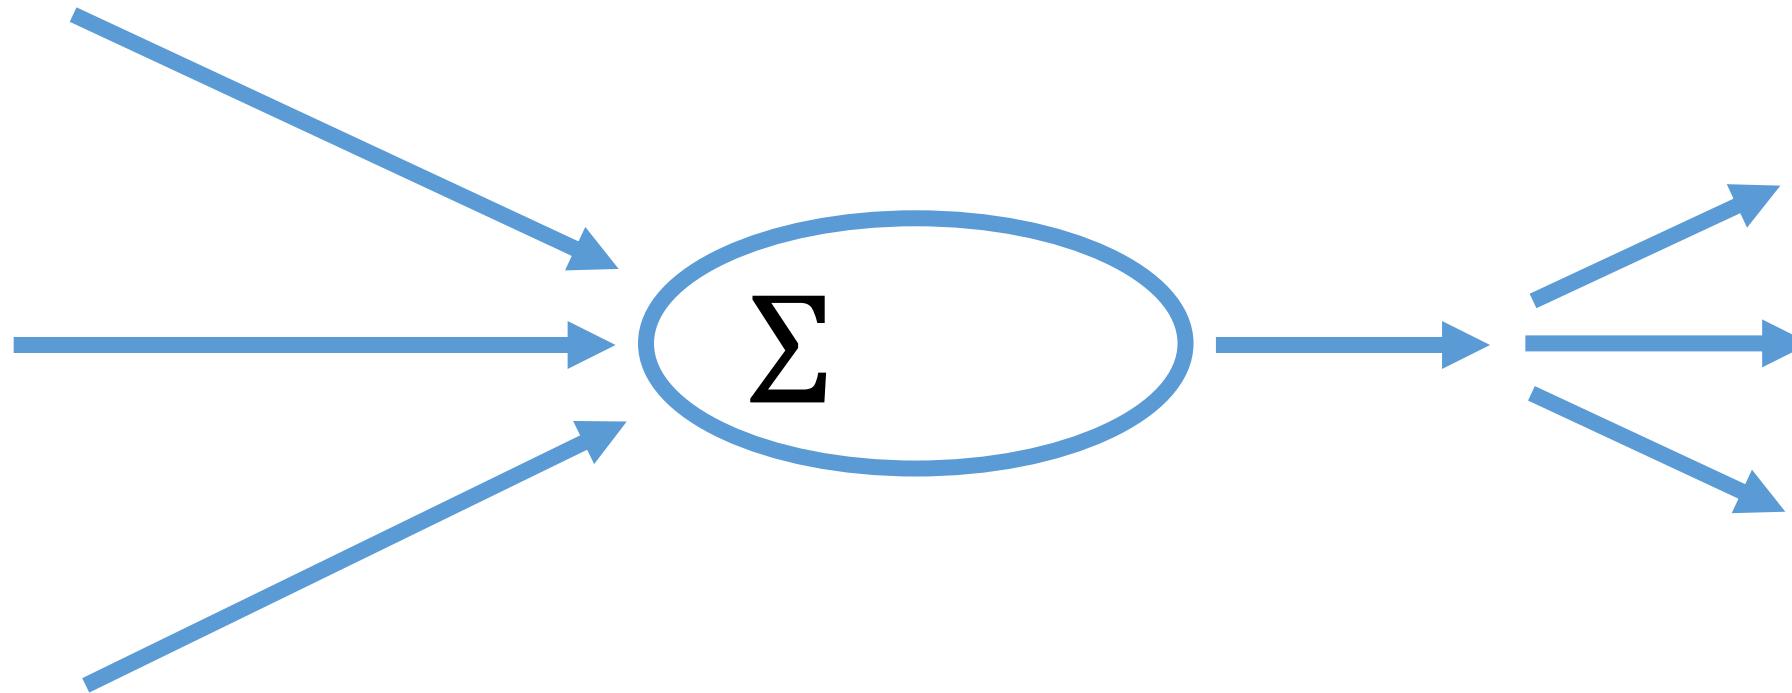




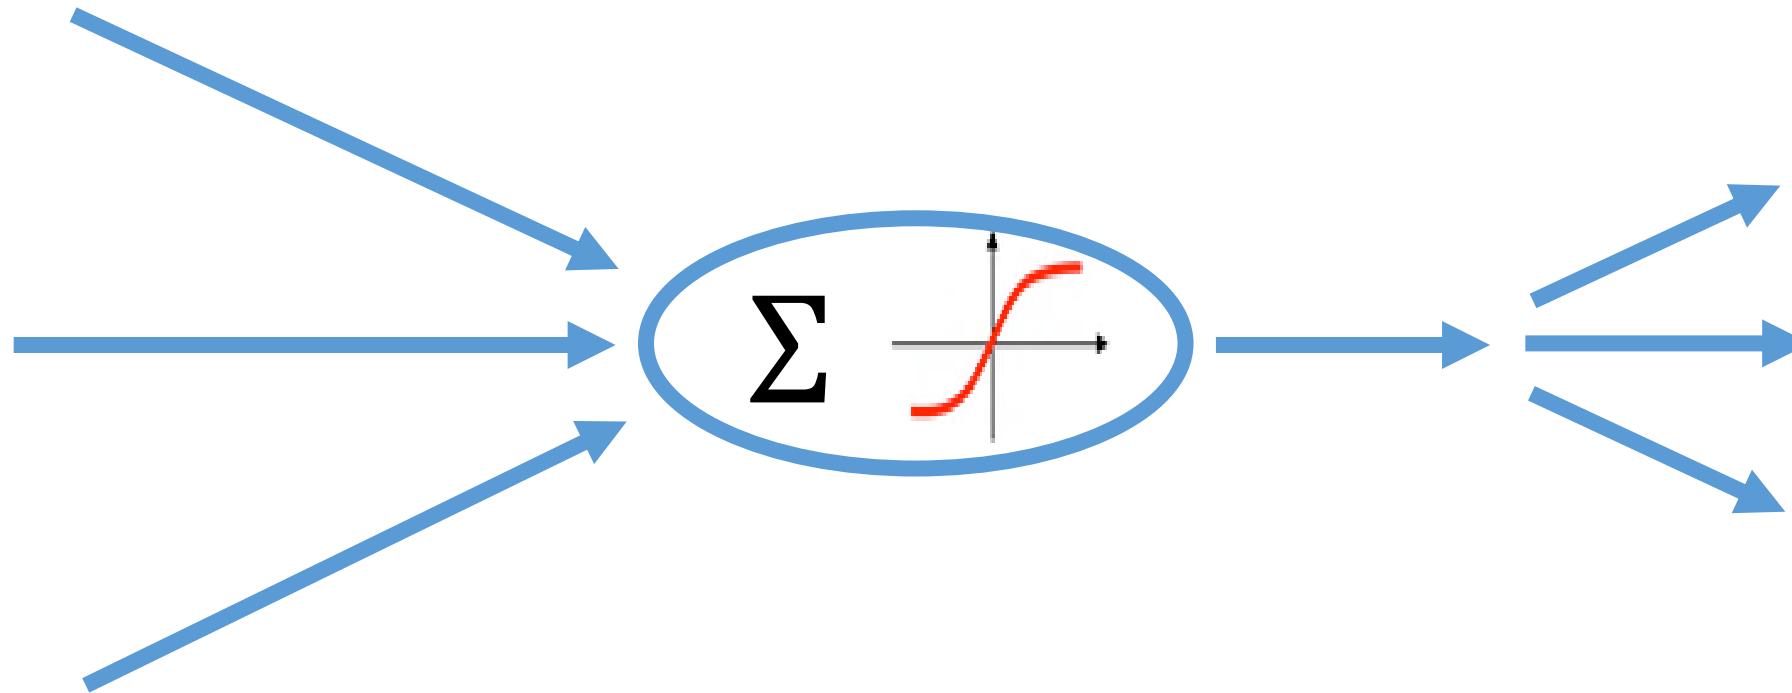
# Artificial Neuron



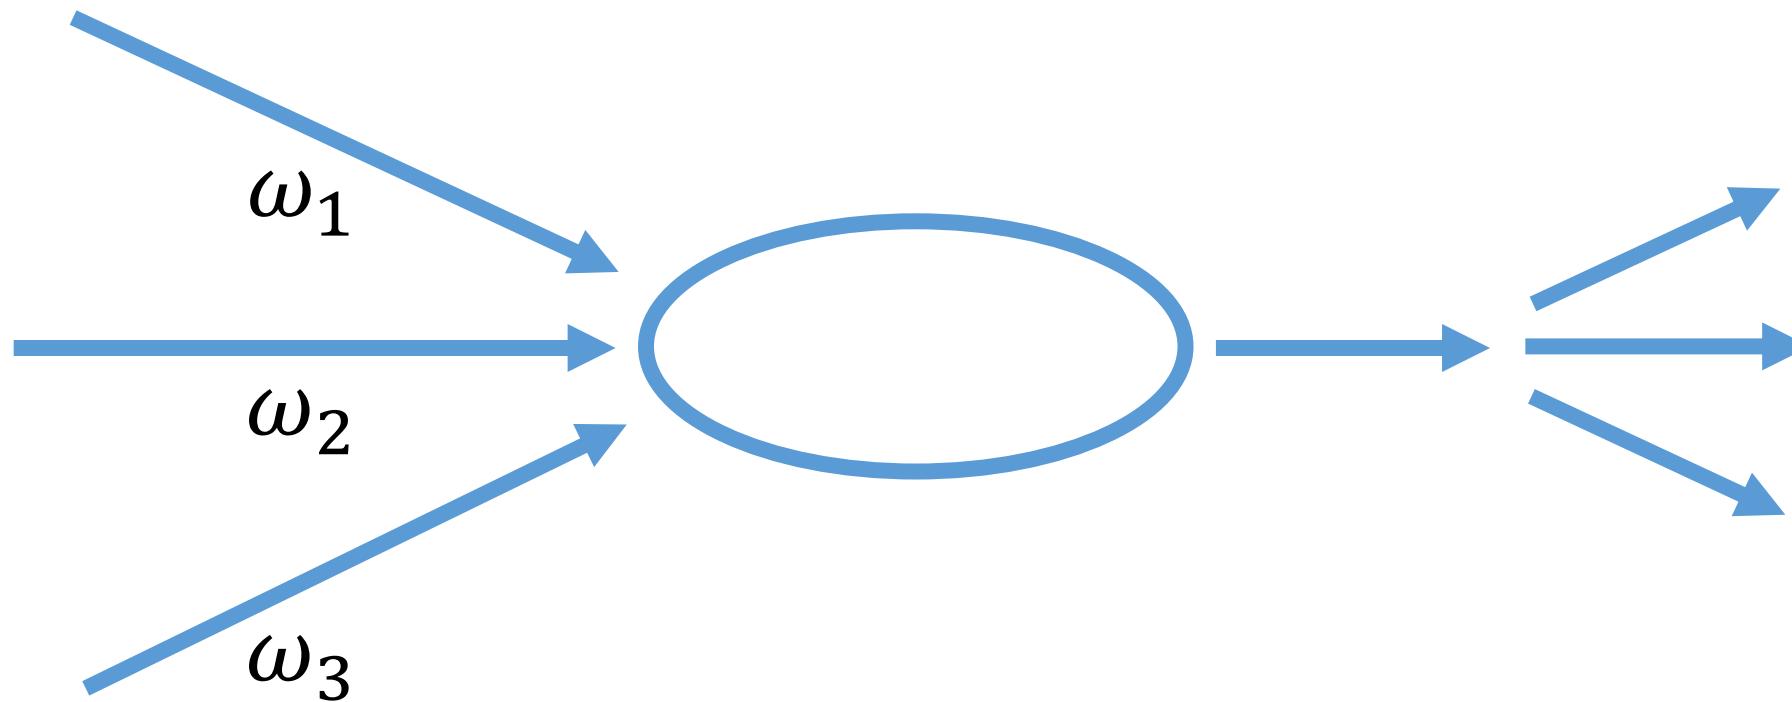
# Artificial Neuron



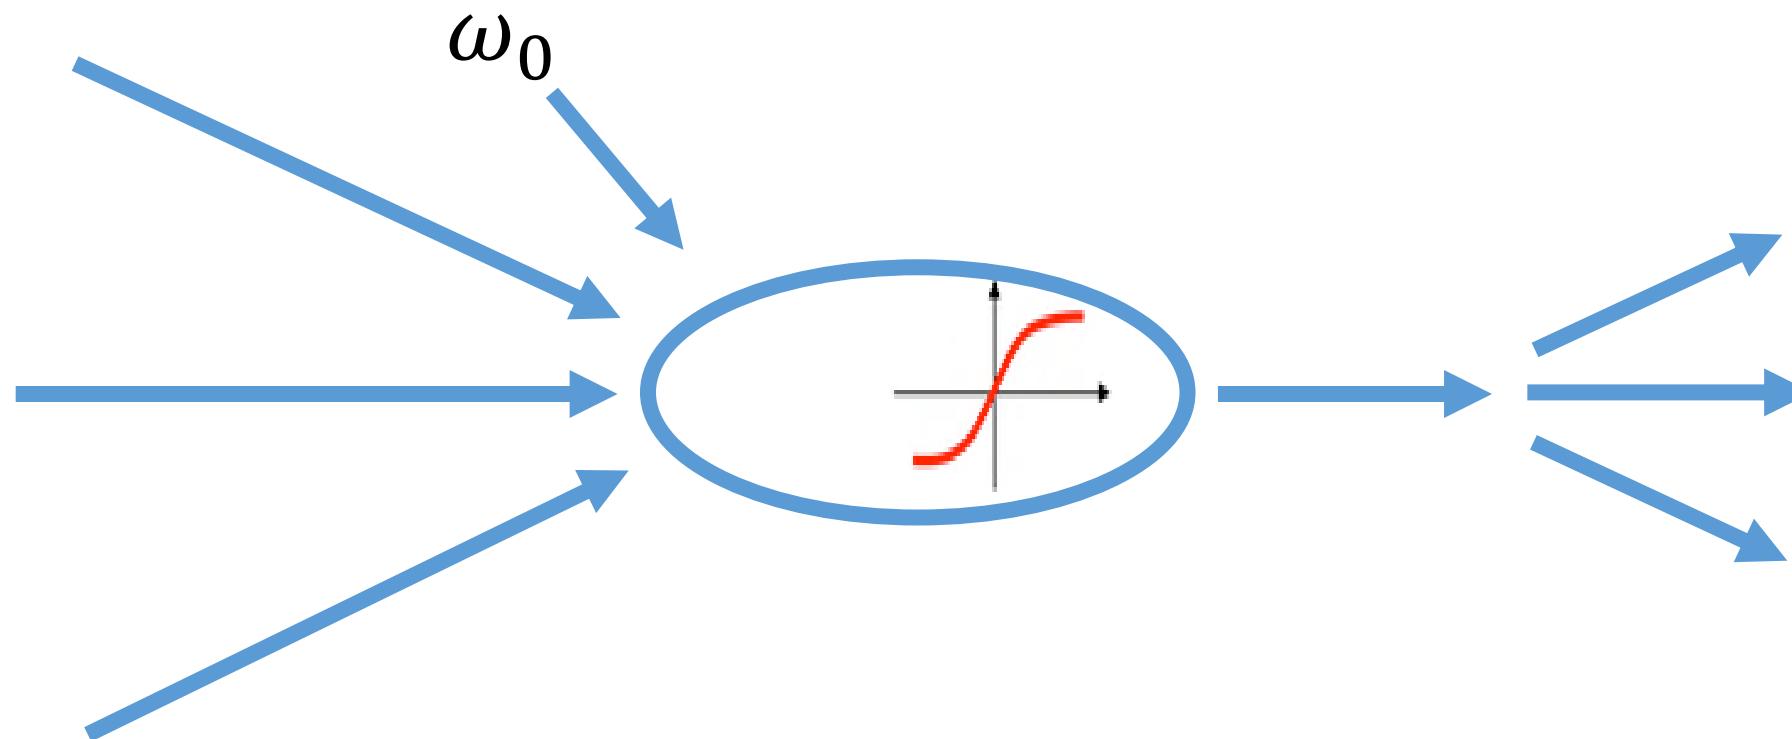
# Artificial Neuron



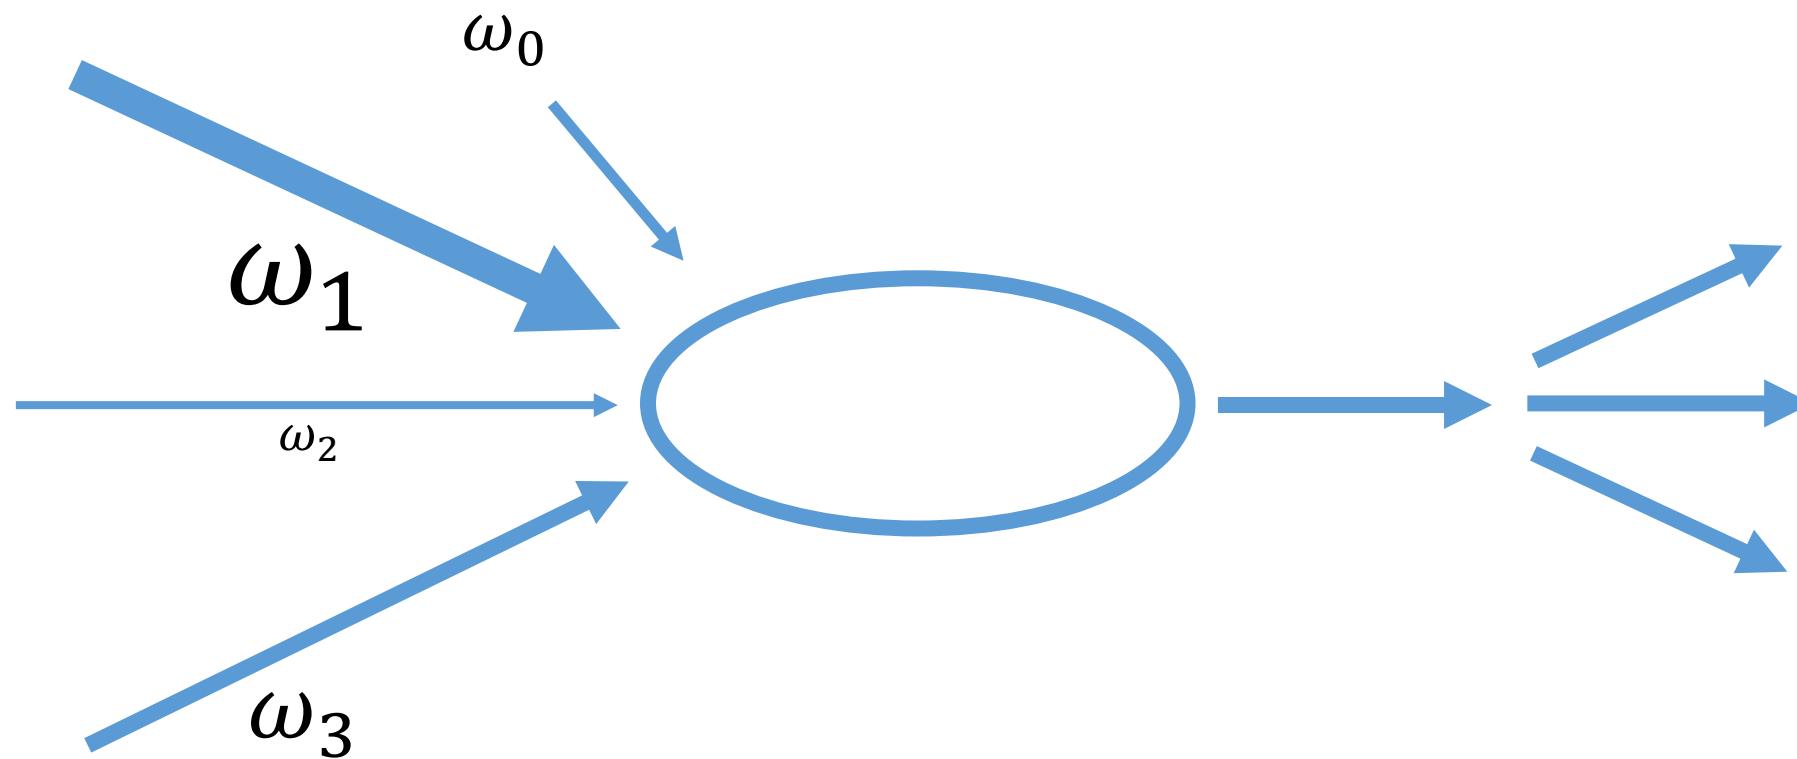
# Artificial Neuron



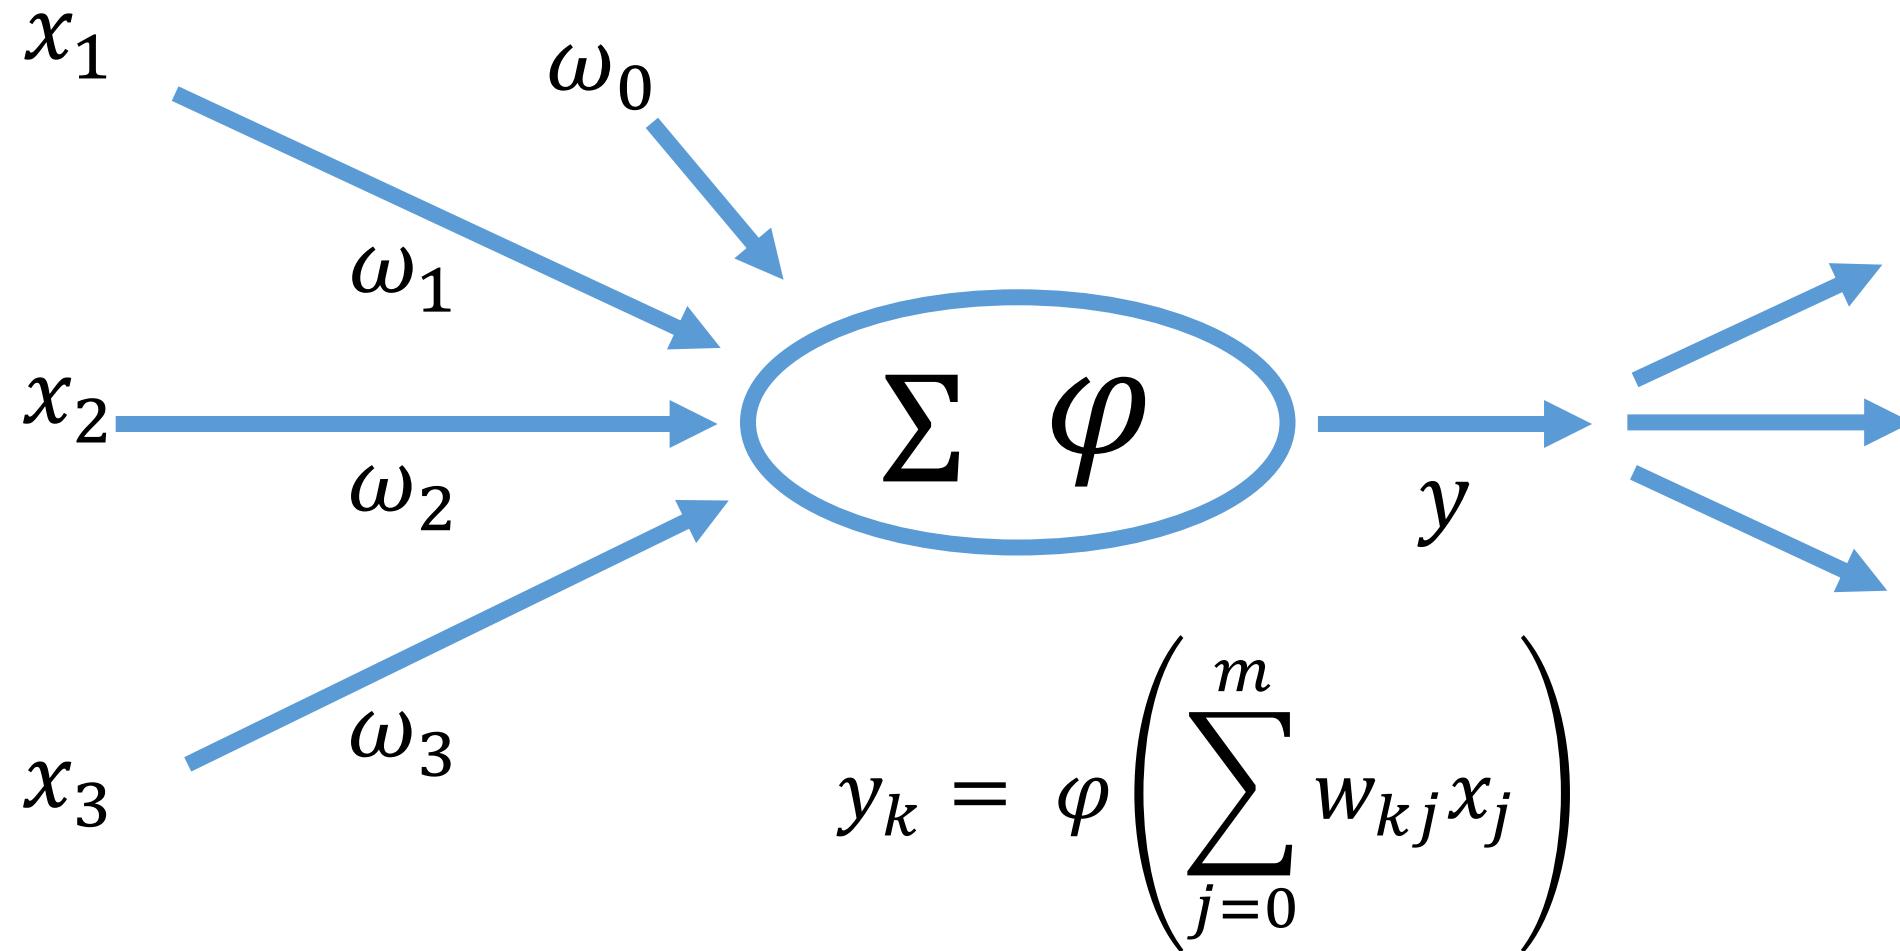
# Artificial Neuron



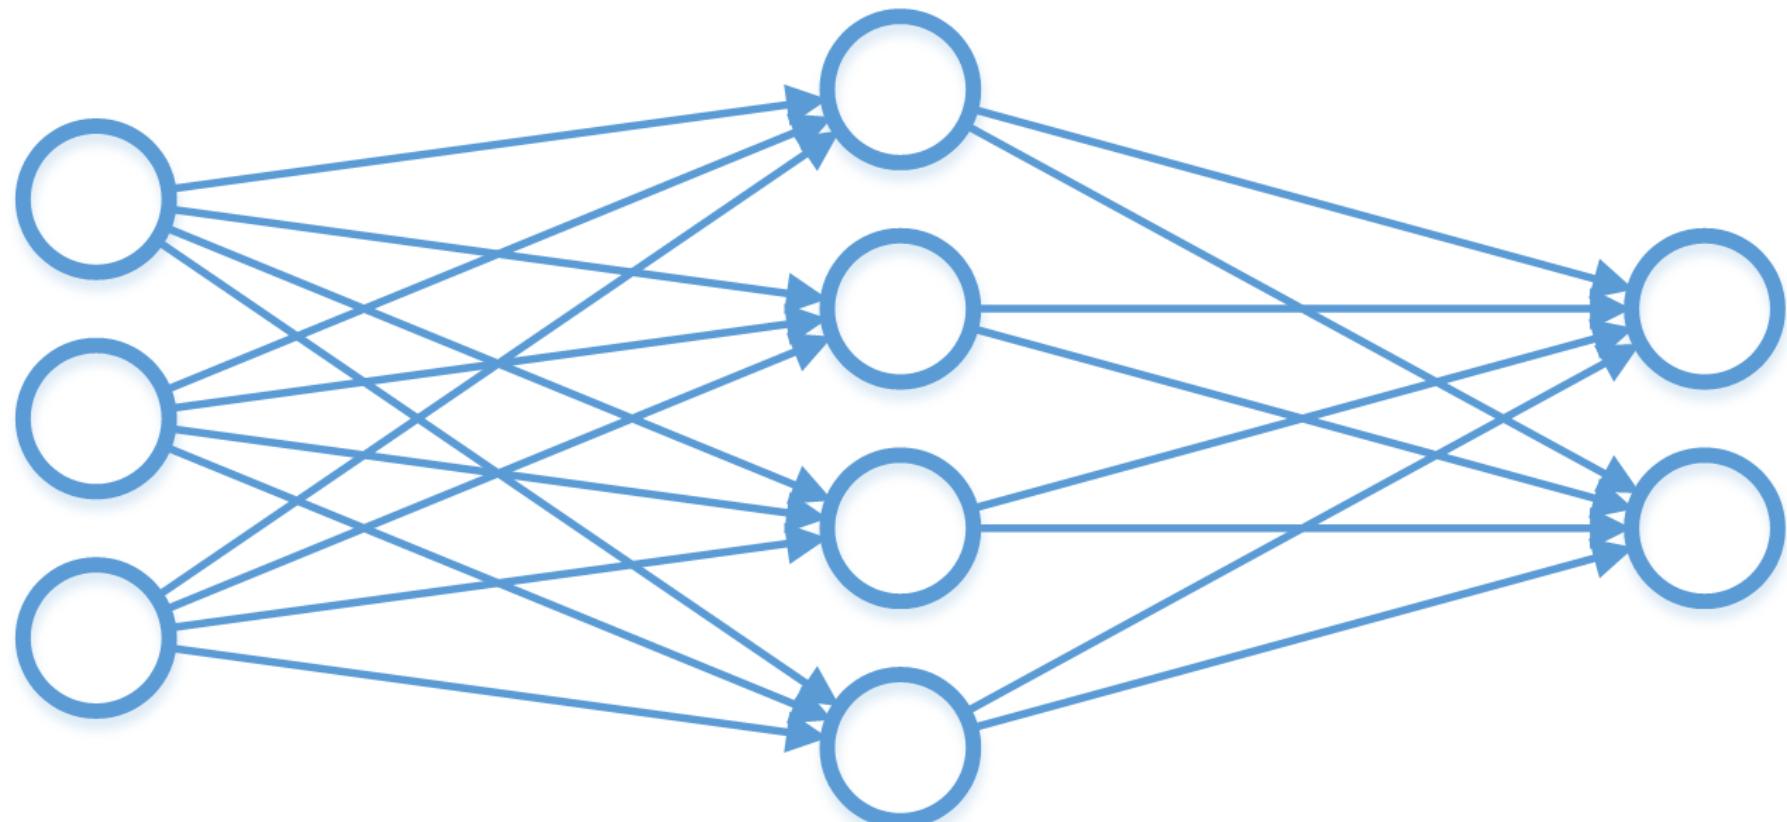
# Artificial Neuron



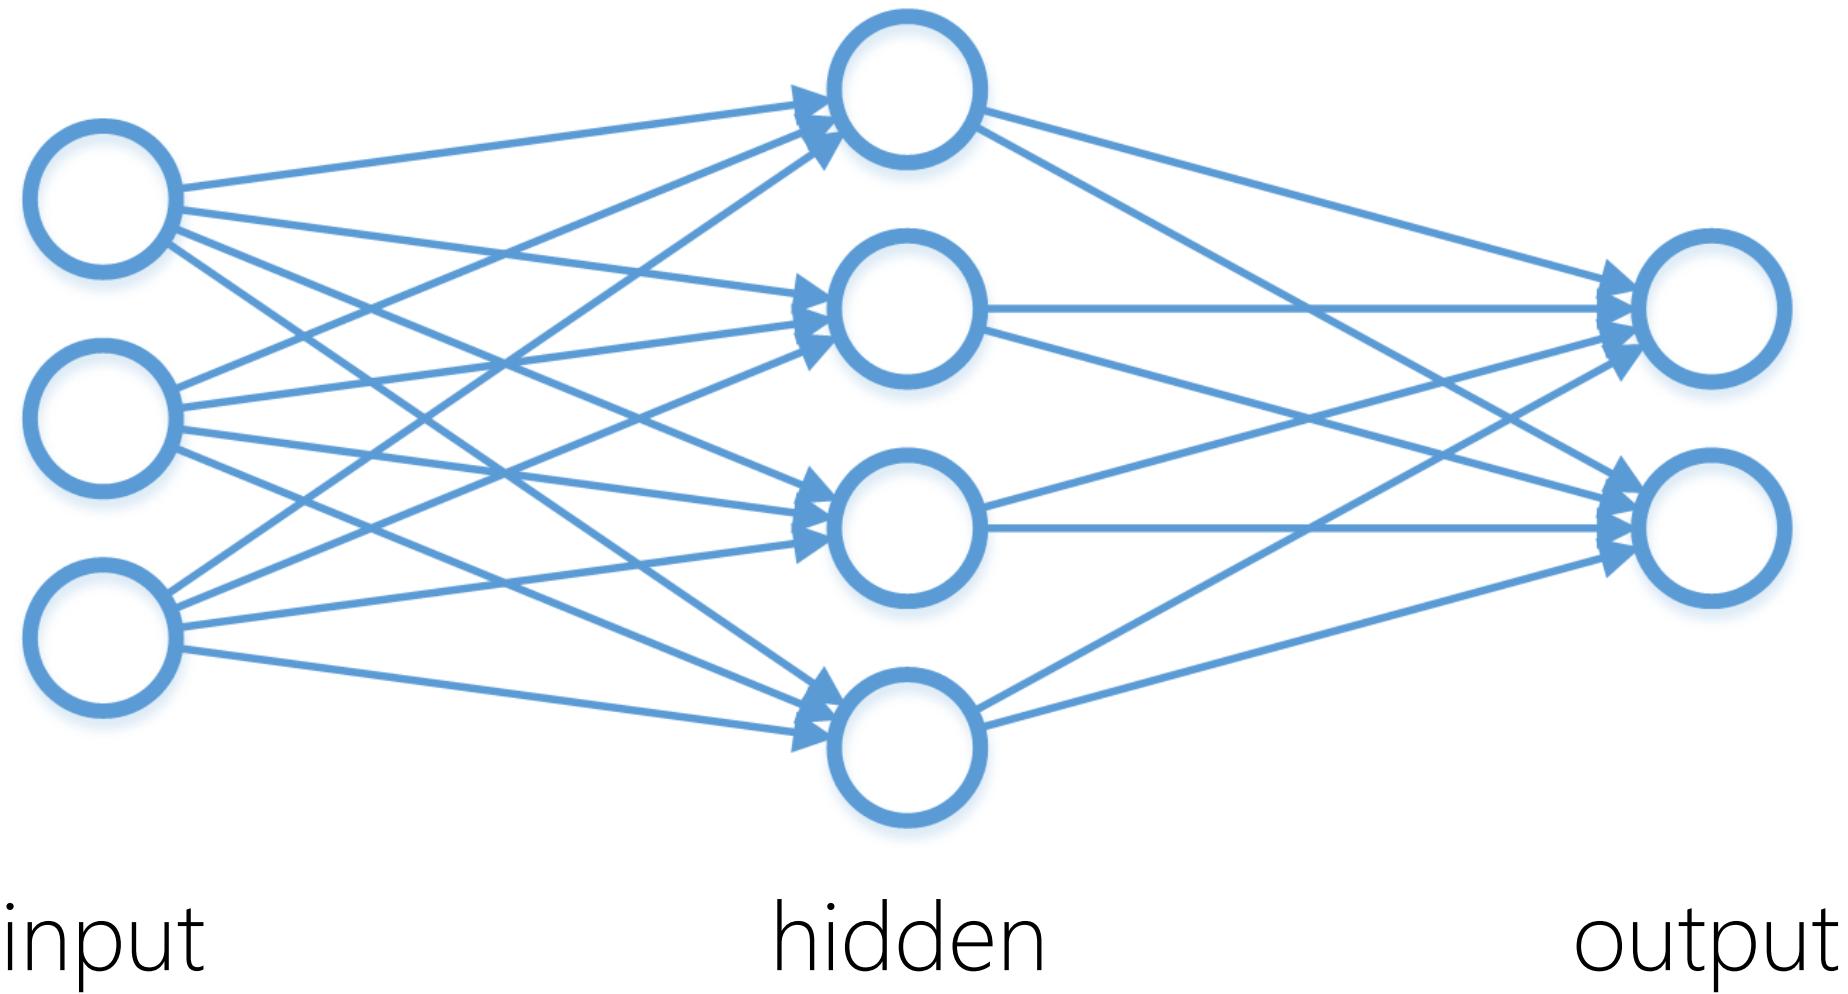
# Artificial Neuron



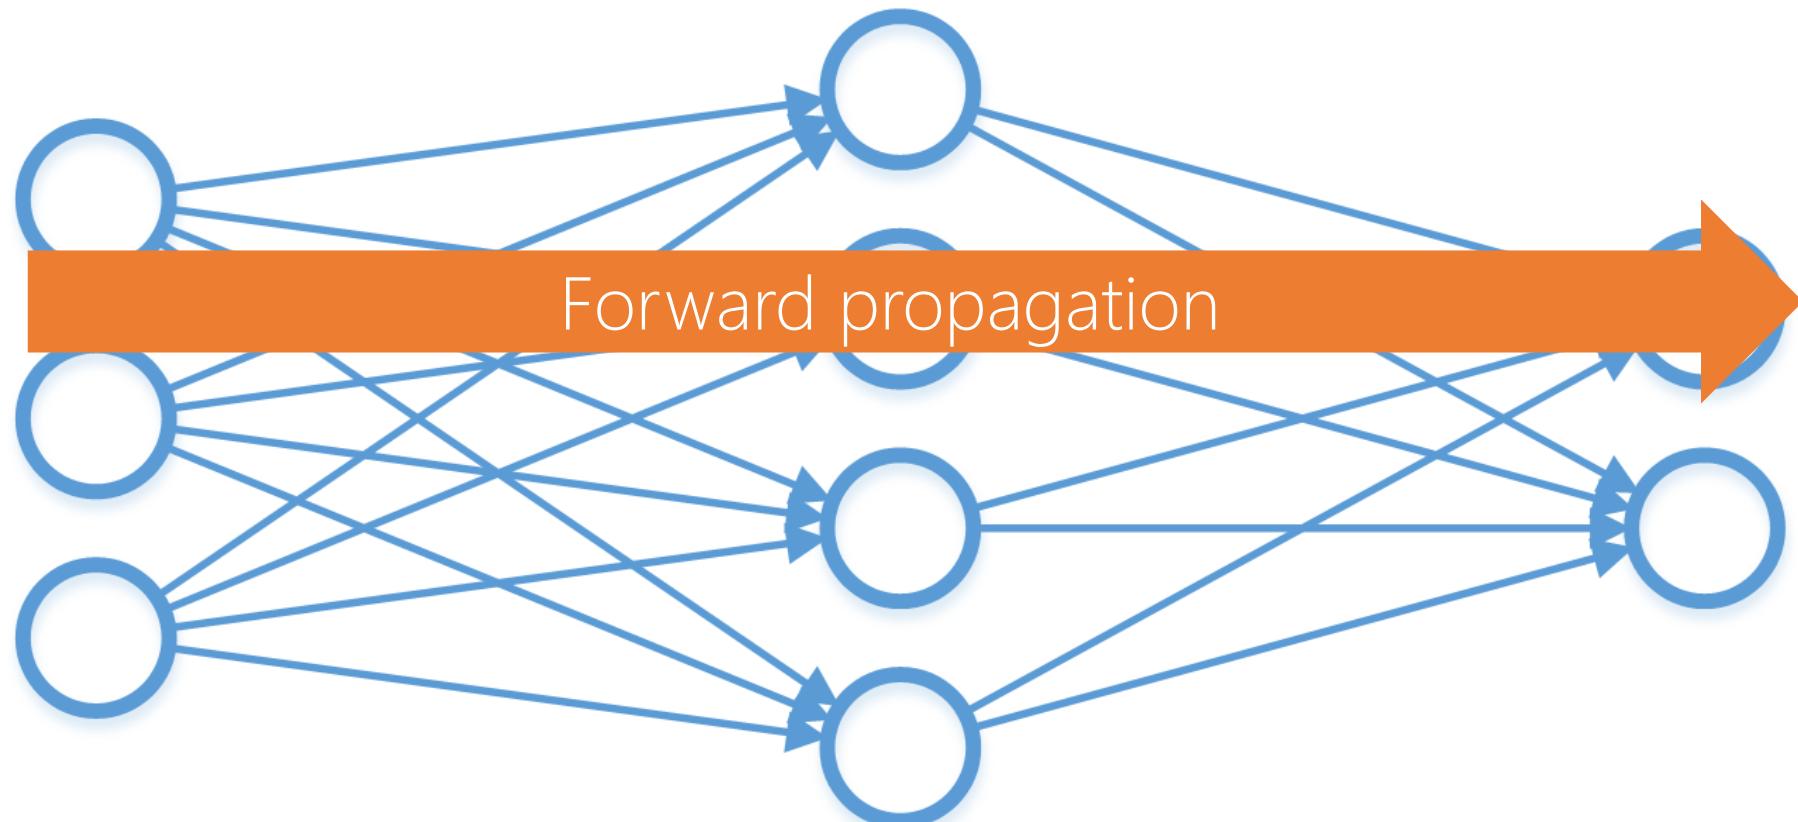
# Artificial Neural Network



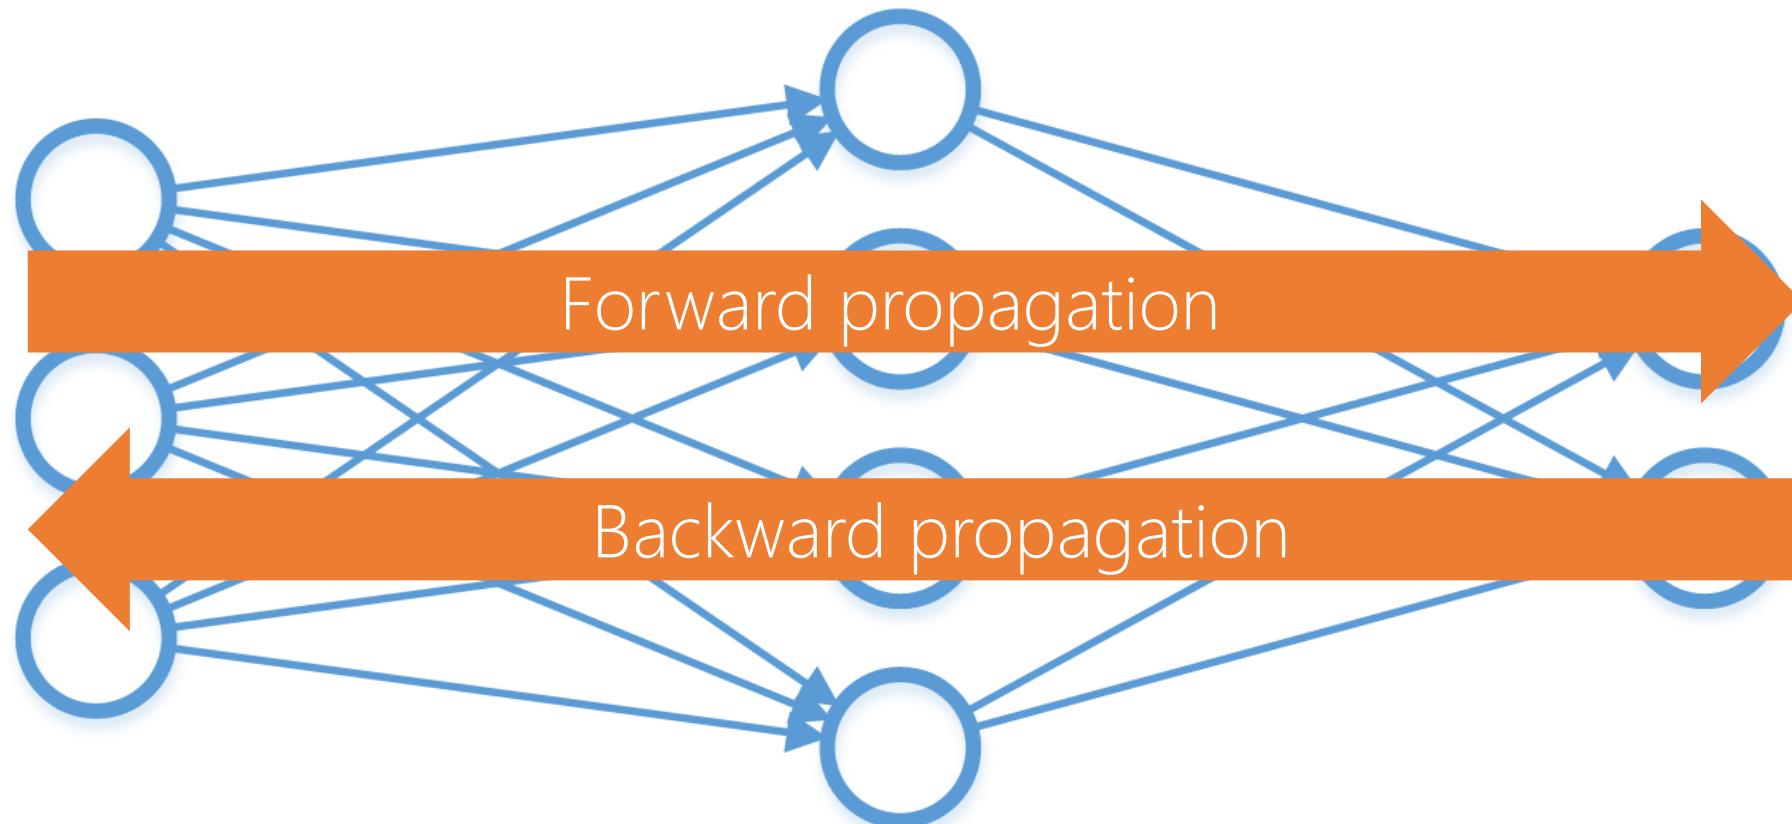
# Artificial Neural Network



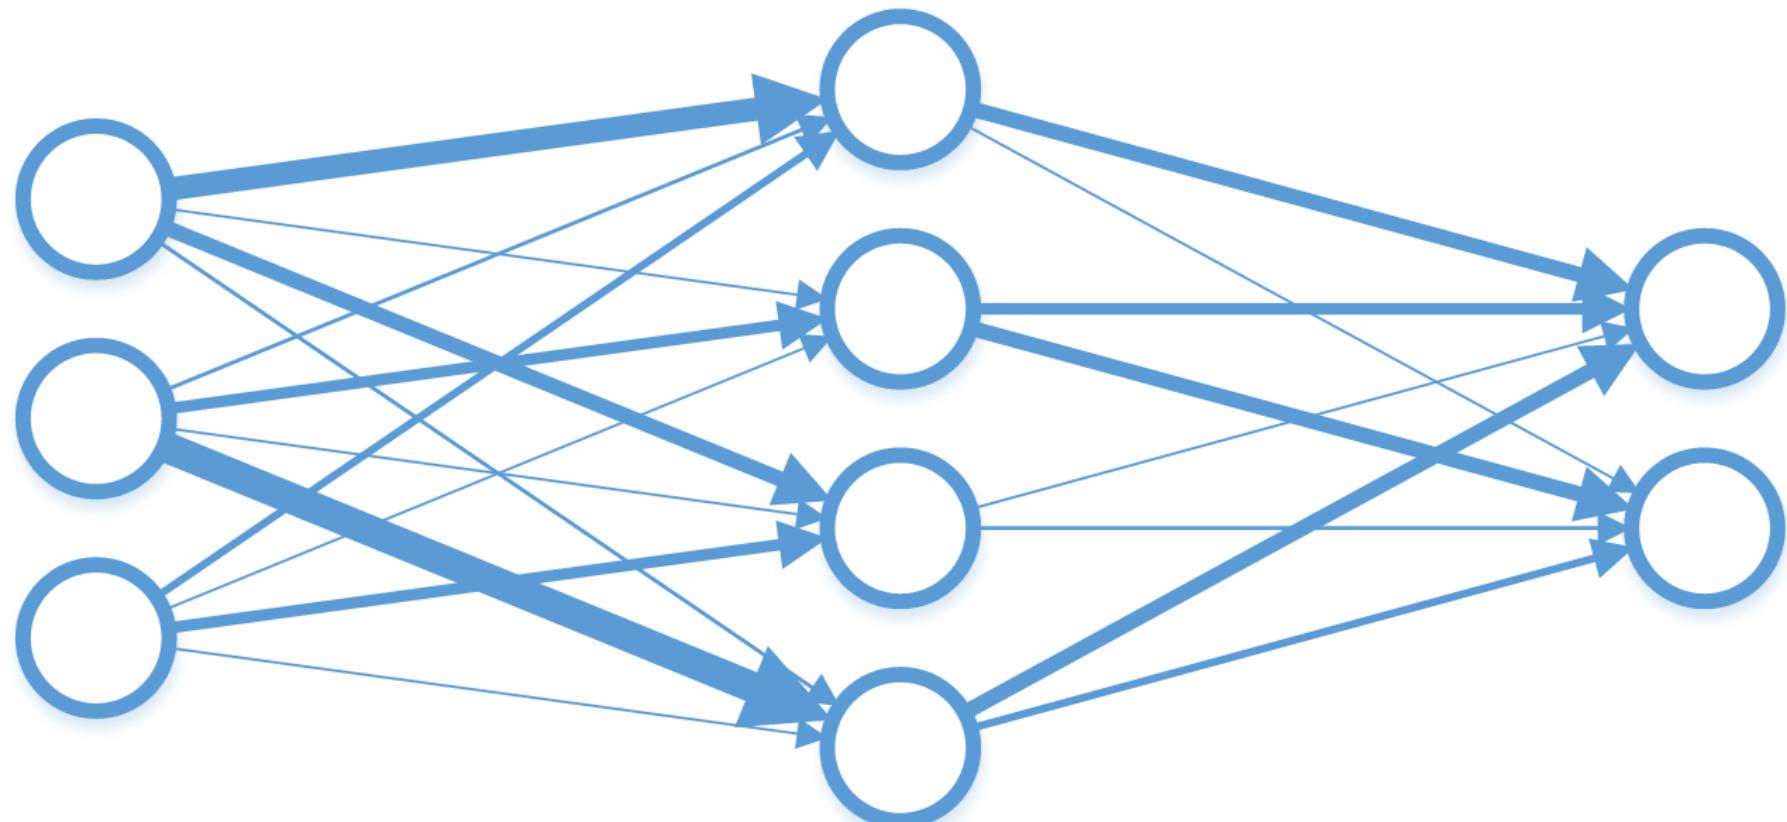
# Artificial Neural Network



# Artificial Neural Network

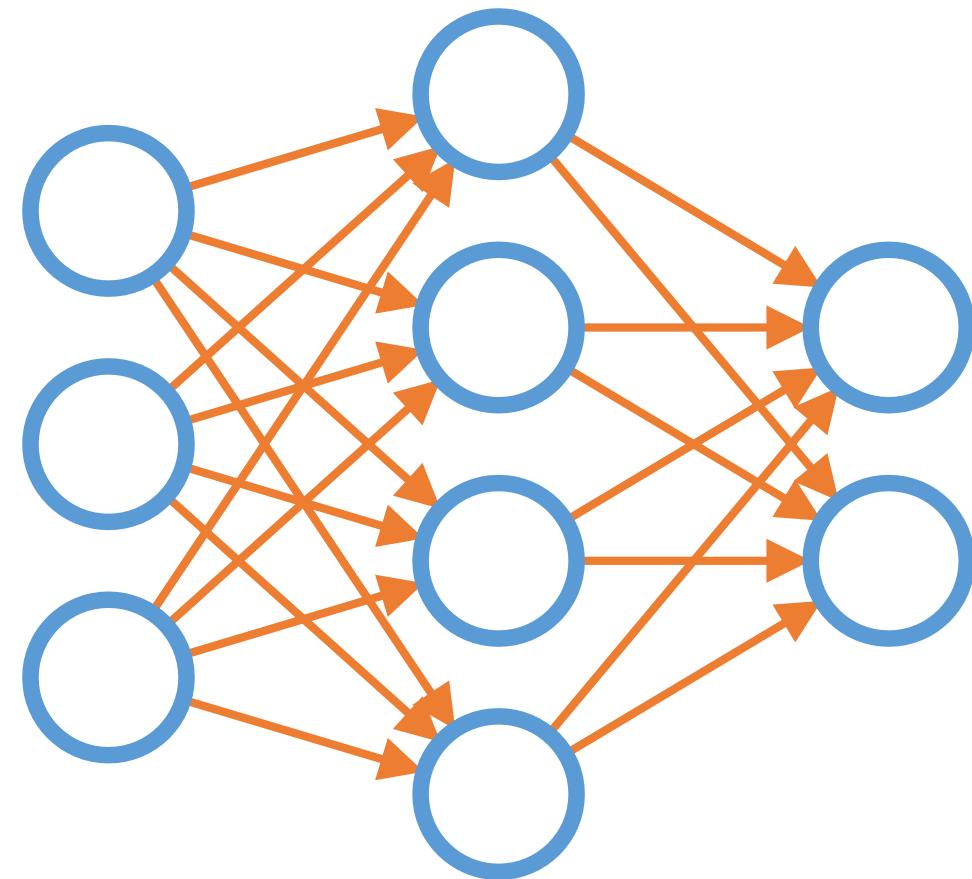


# Artificial Neural Network



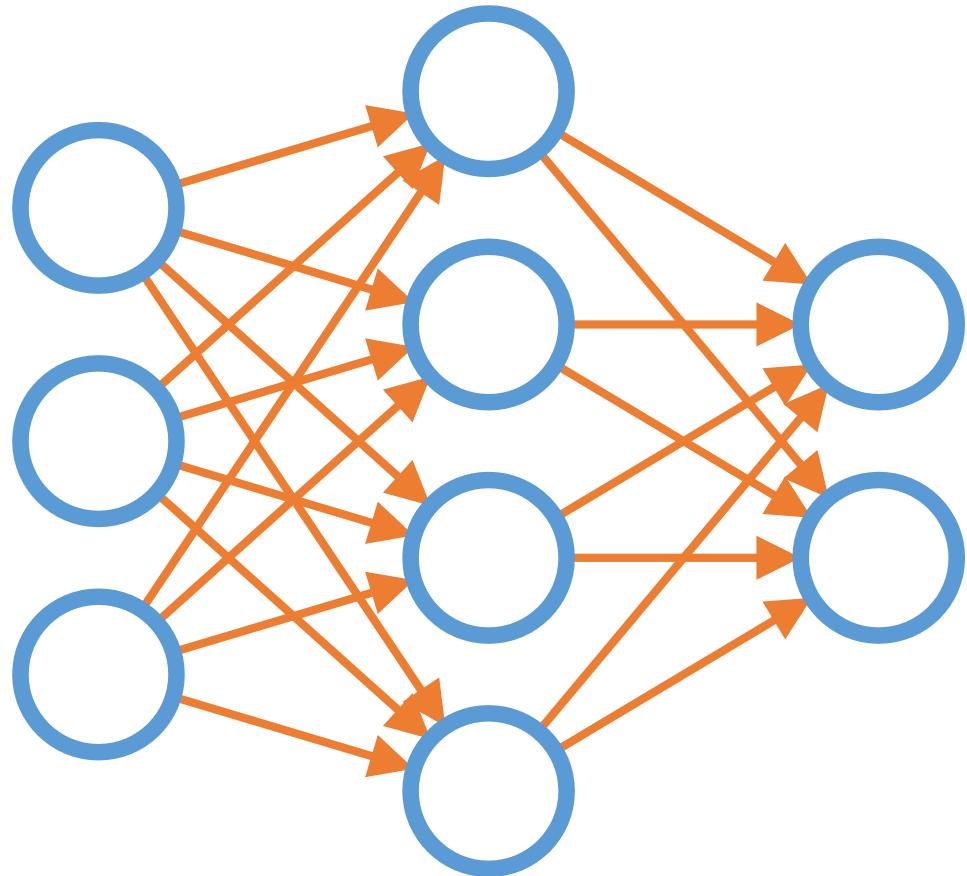
# Neural Network Classifier

Supervised learning



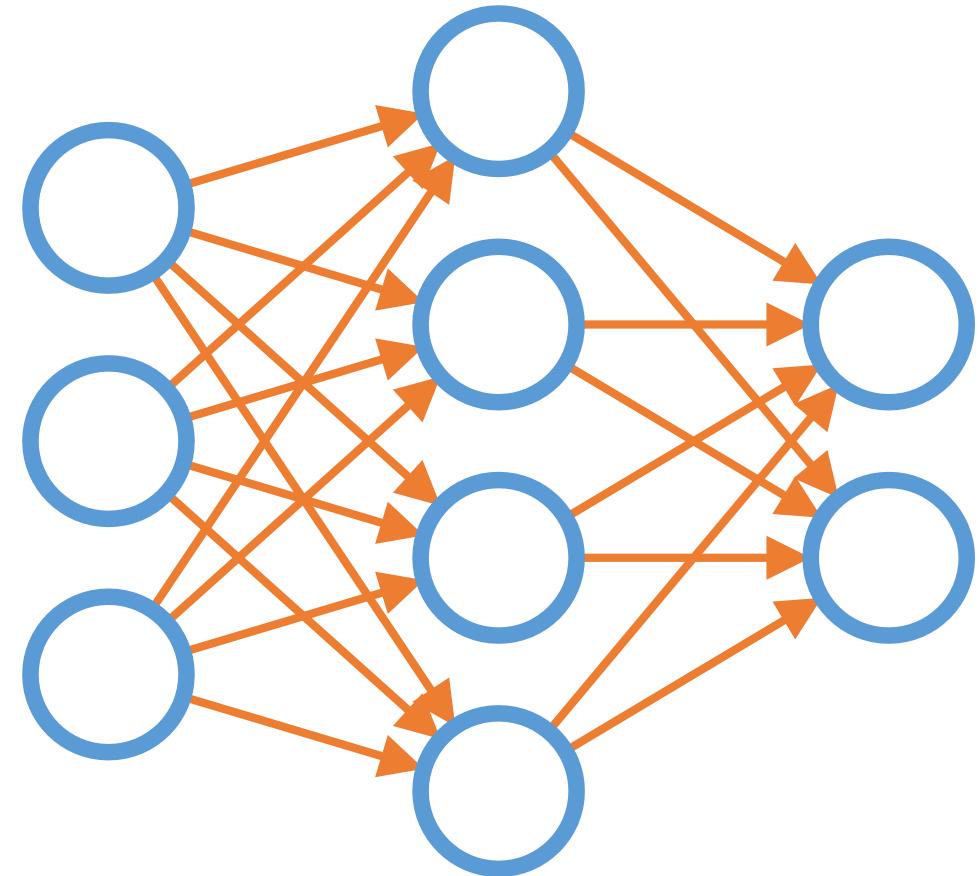
# Neural Network Classifier

Supervised learning  
Neurons in a brain



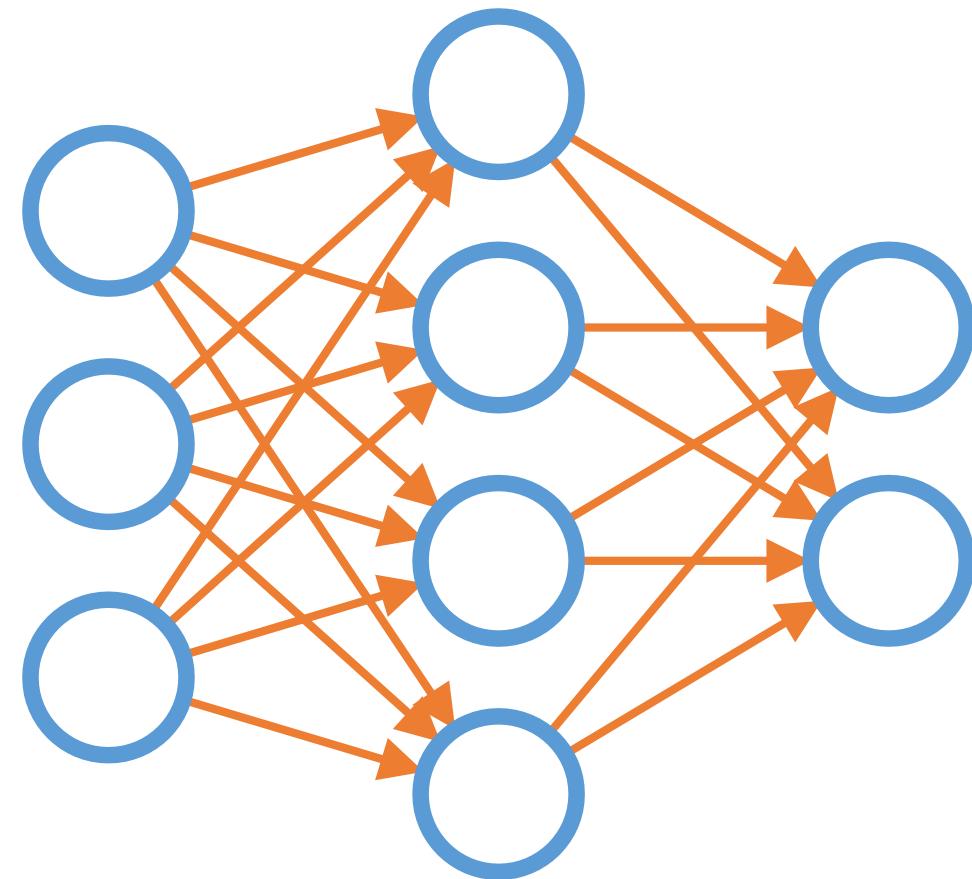
# Neural Network Classifier

Supervised learning  
Neurons in a brain  
Weighted connections



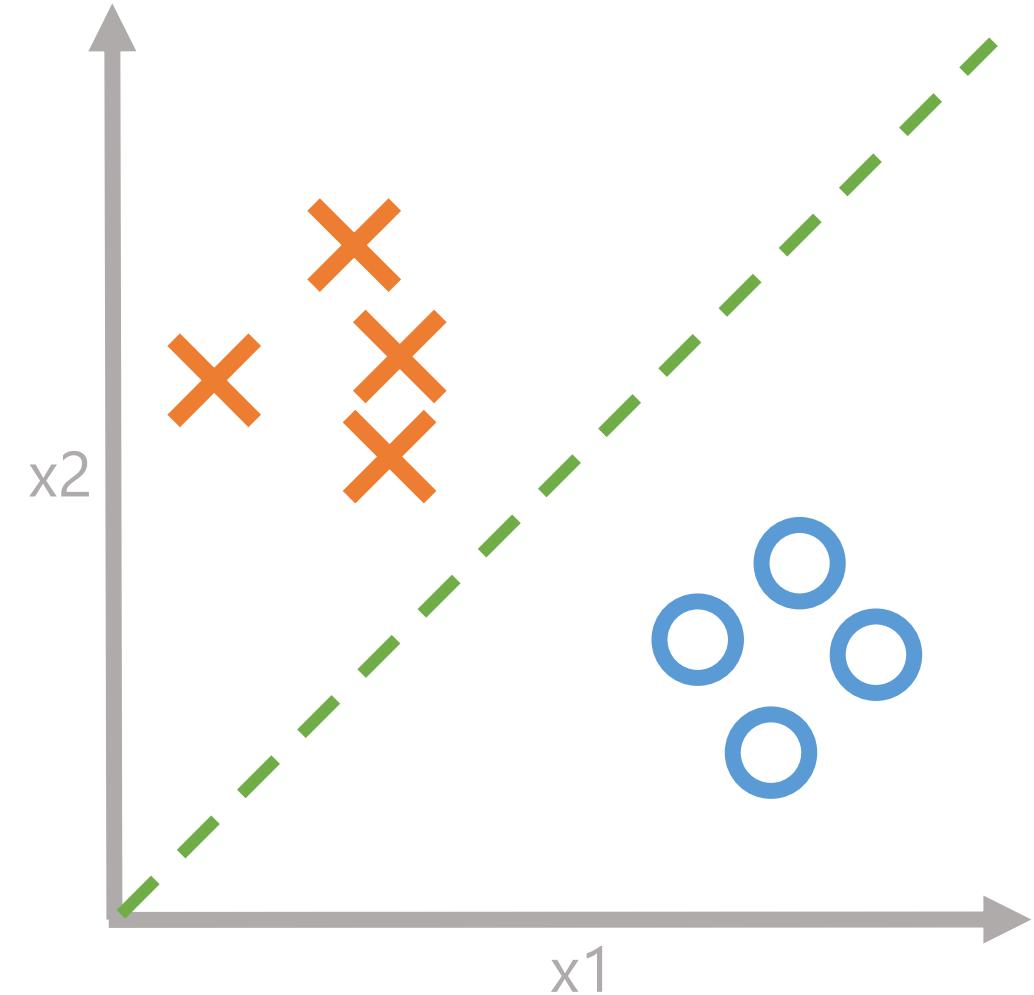
# Neural Network Classifier

Supervised learning  
Neurons in a brain  
Weighted connections  
Complex



# Real-World Examples

- Should we approve this loan?
- Will this customer buy from us?
- Should we replace this part?
- Does this person have cancer?



# Iris Data Set



Iris Setosa



Iris Versicolor



Iris Virginica

# Iris Data Set

Fisher's Iris Data				
Species	Petal Length	Petal Width	Sepal Length	Sepal Width
setosa	1.1	0.1	4.3	3
setosa	1.4	0.2	4.4	2.9
setosa	1.3	0.2	4.4	3
setosa	1.3	0.2	4.4	3.2
setosa	1.3	0.3	4.5	2.3
...	...	...	...	...

## Demo 2 - Classification

Goal: Predict species based on  
petal and sepal measurements

# Insurance Policy Risk Data Set

# Insurance Policy Rates Data Set

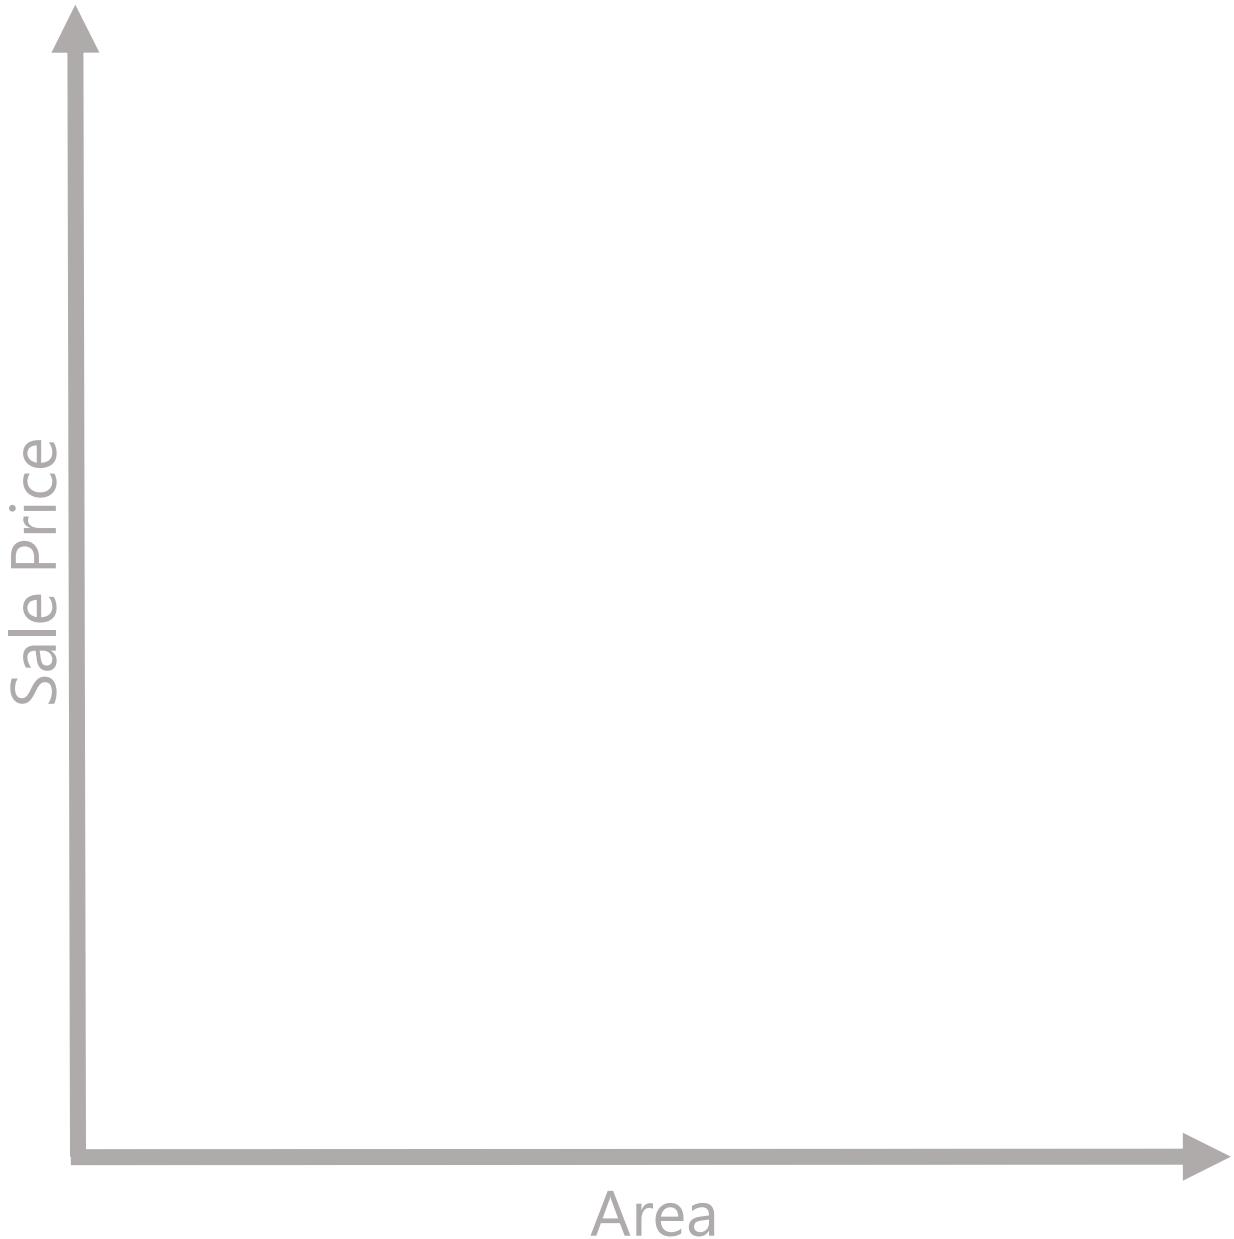
# Lab 2A – Classification (Easy)

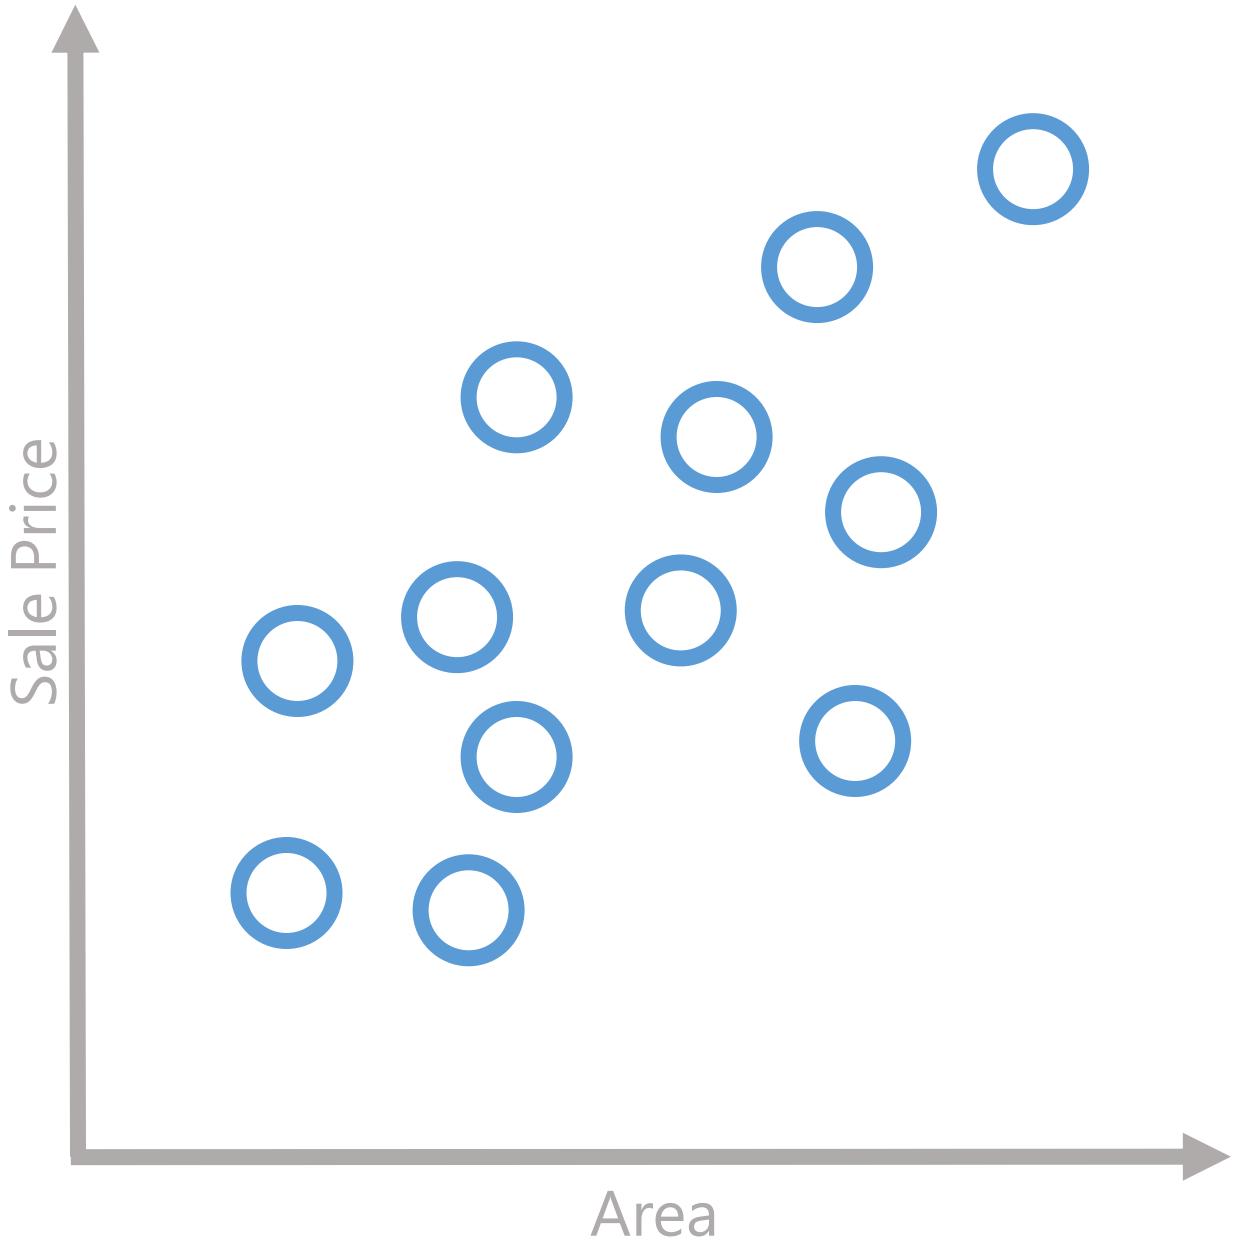
Goal: Predict species based on  
petal and sepal measurements

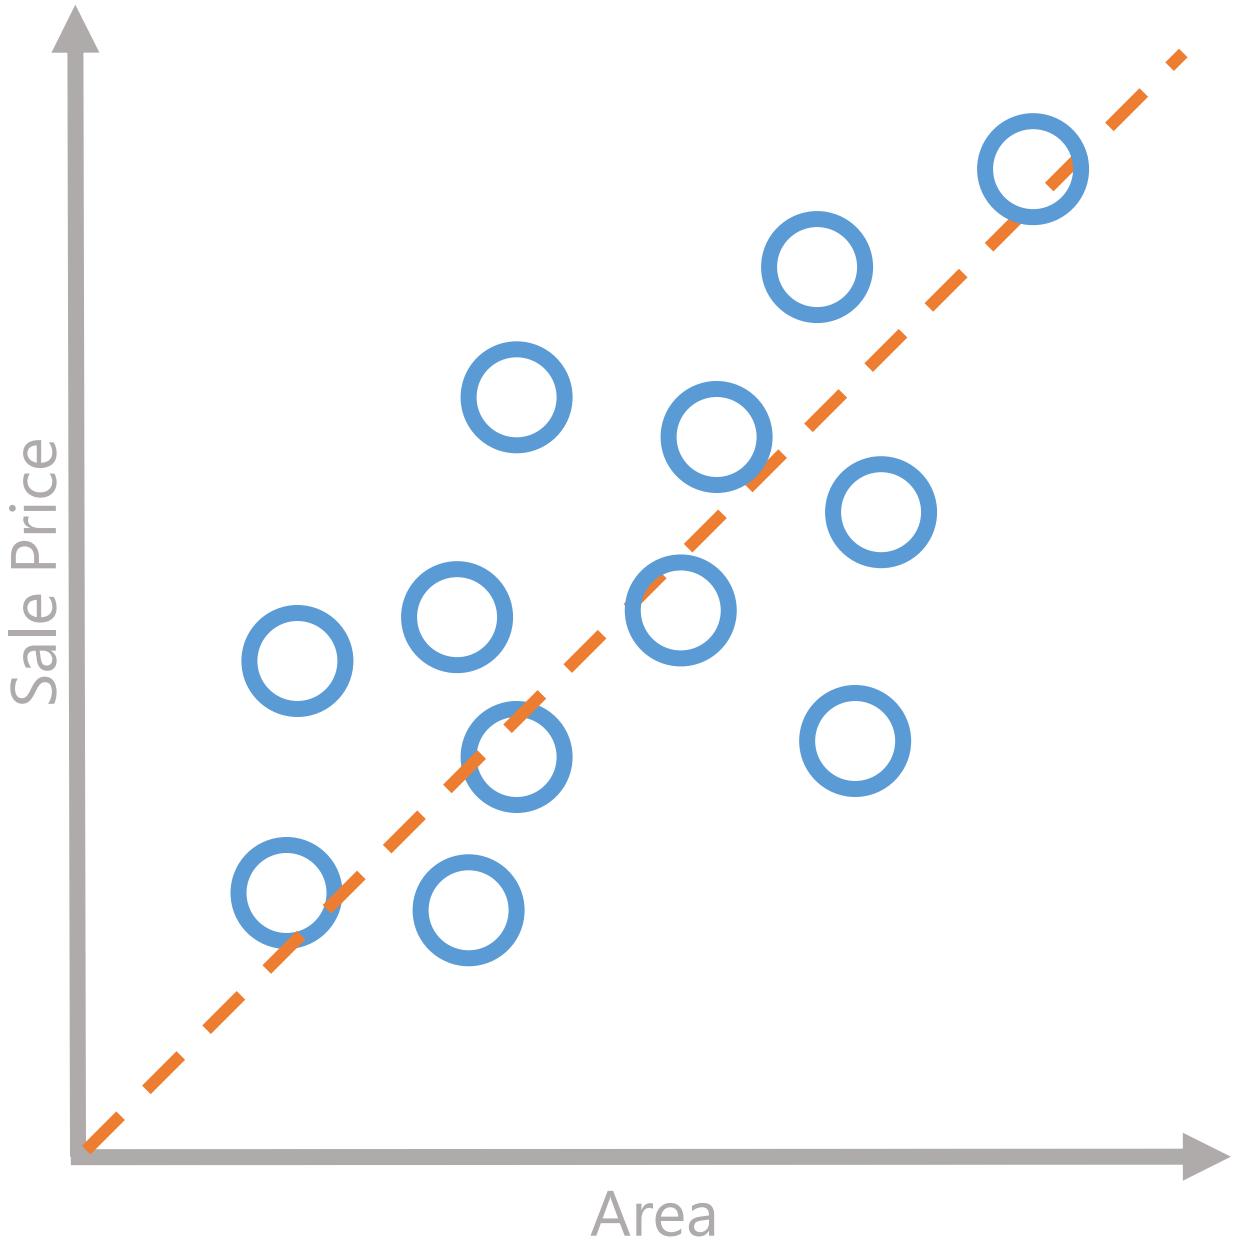
# Lab 2B – Classification (Hard)

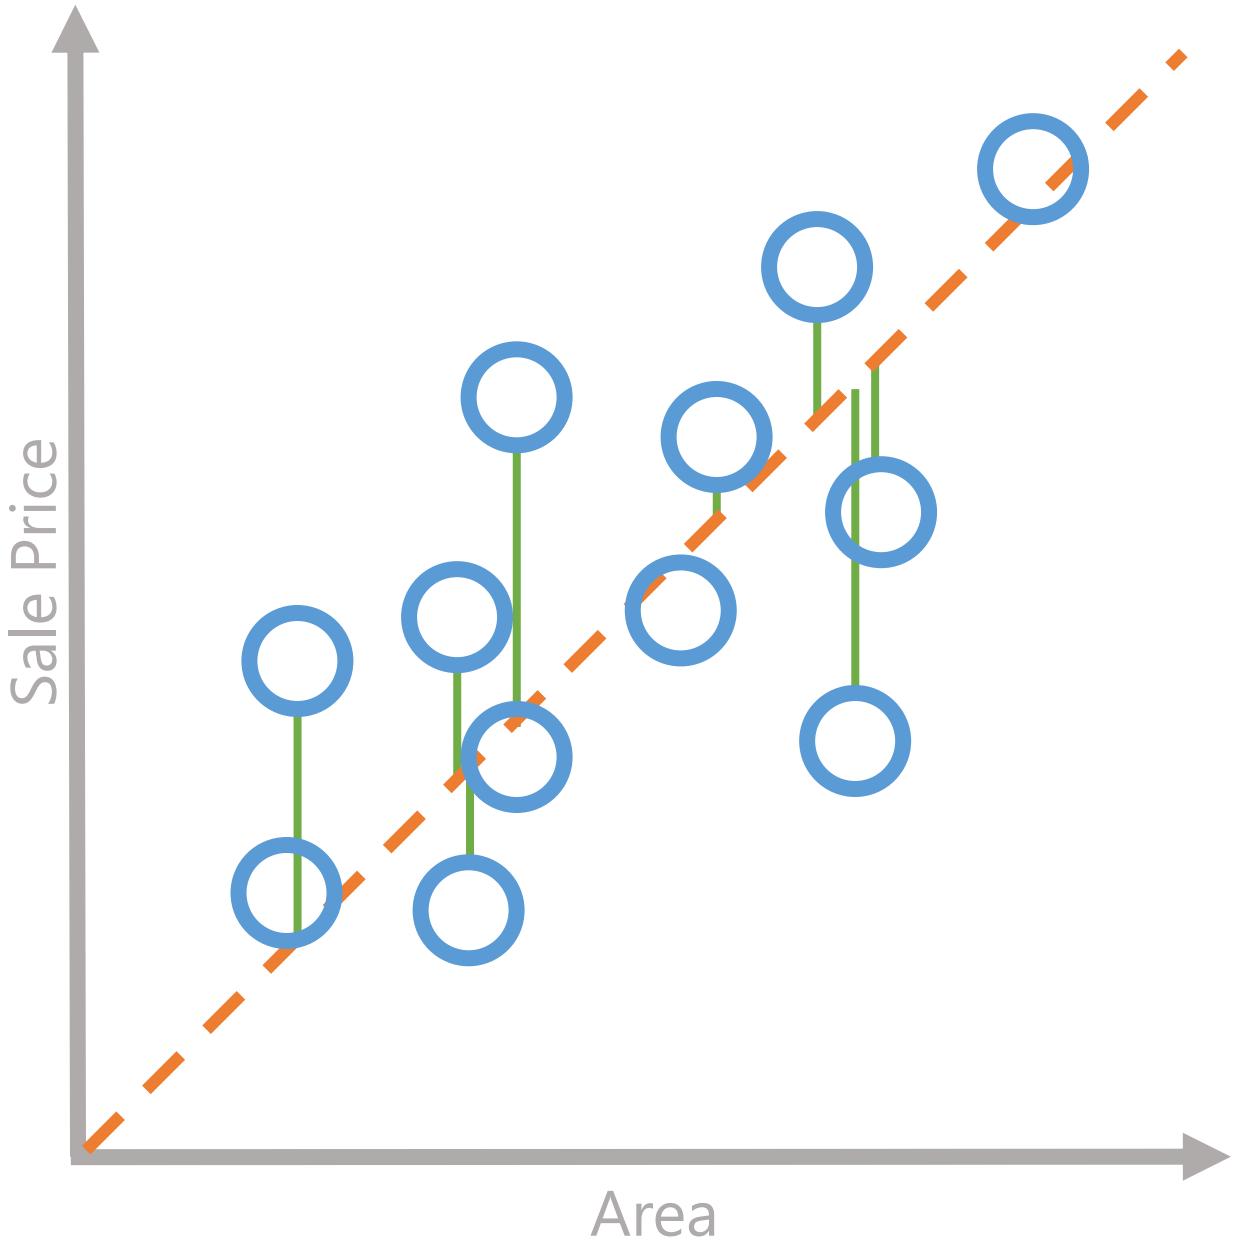
Goal: Predict the risk of  
an insurance policy

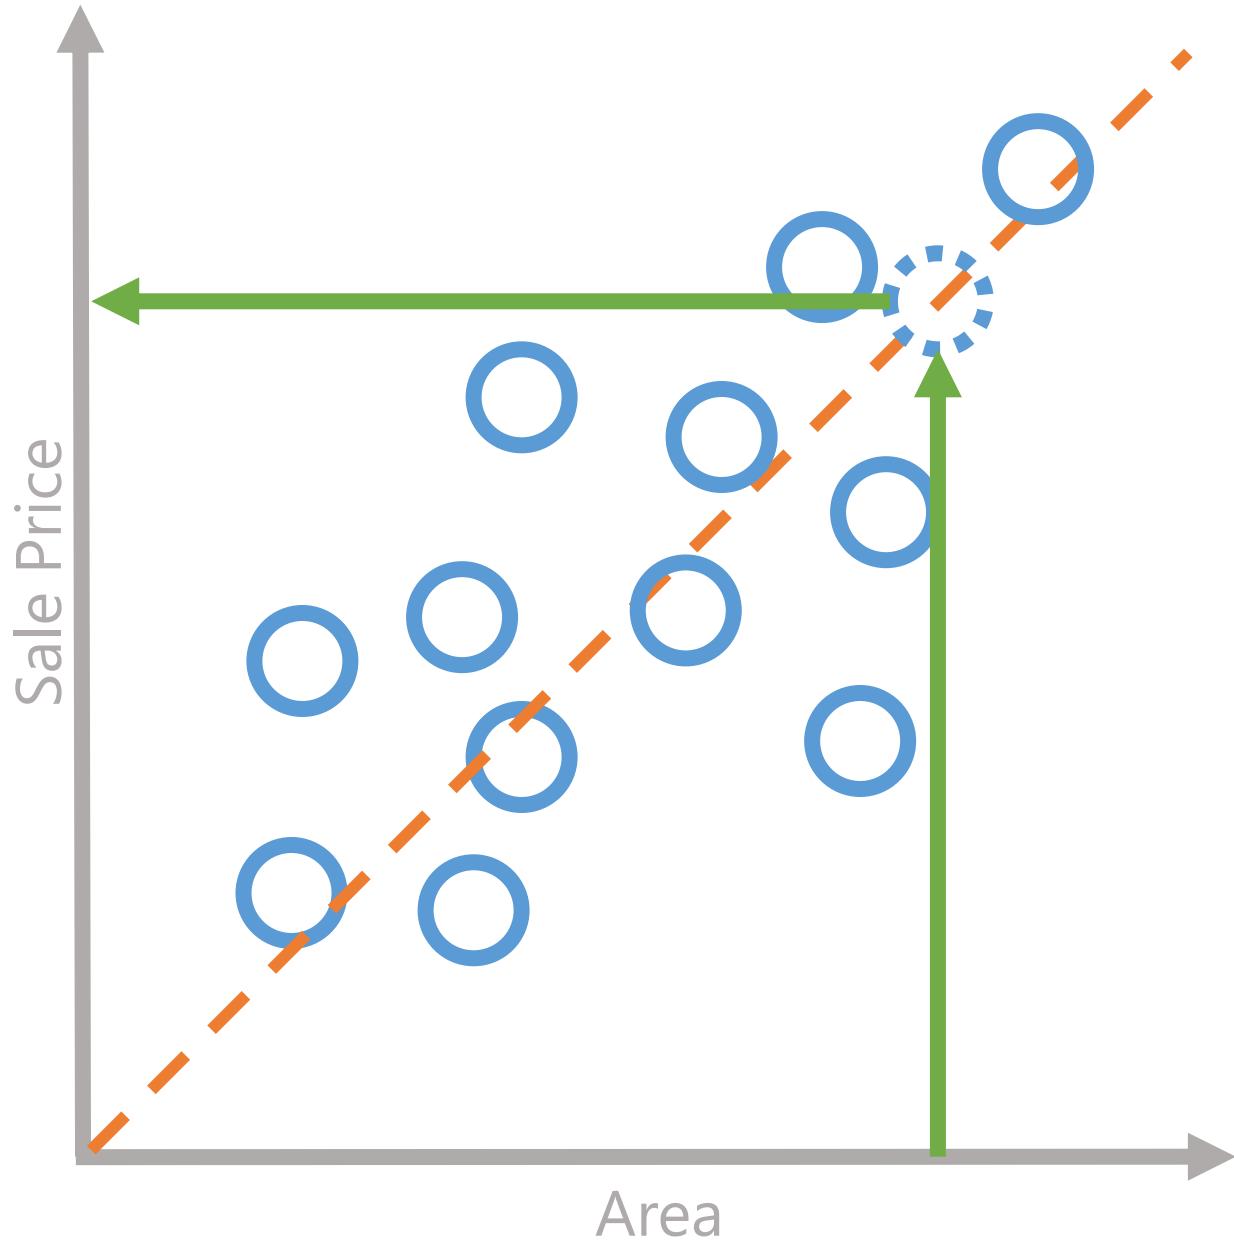
# Regression

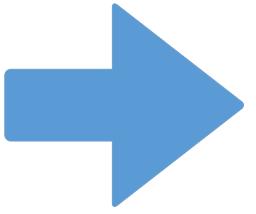
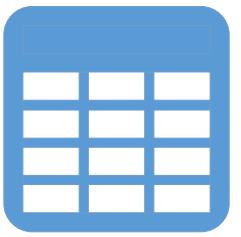
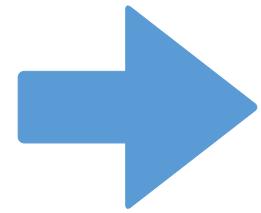










 $f(x)$ 

1.23

# Regression Algorithms

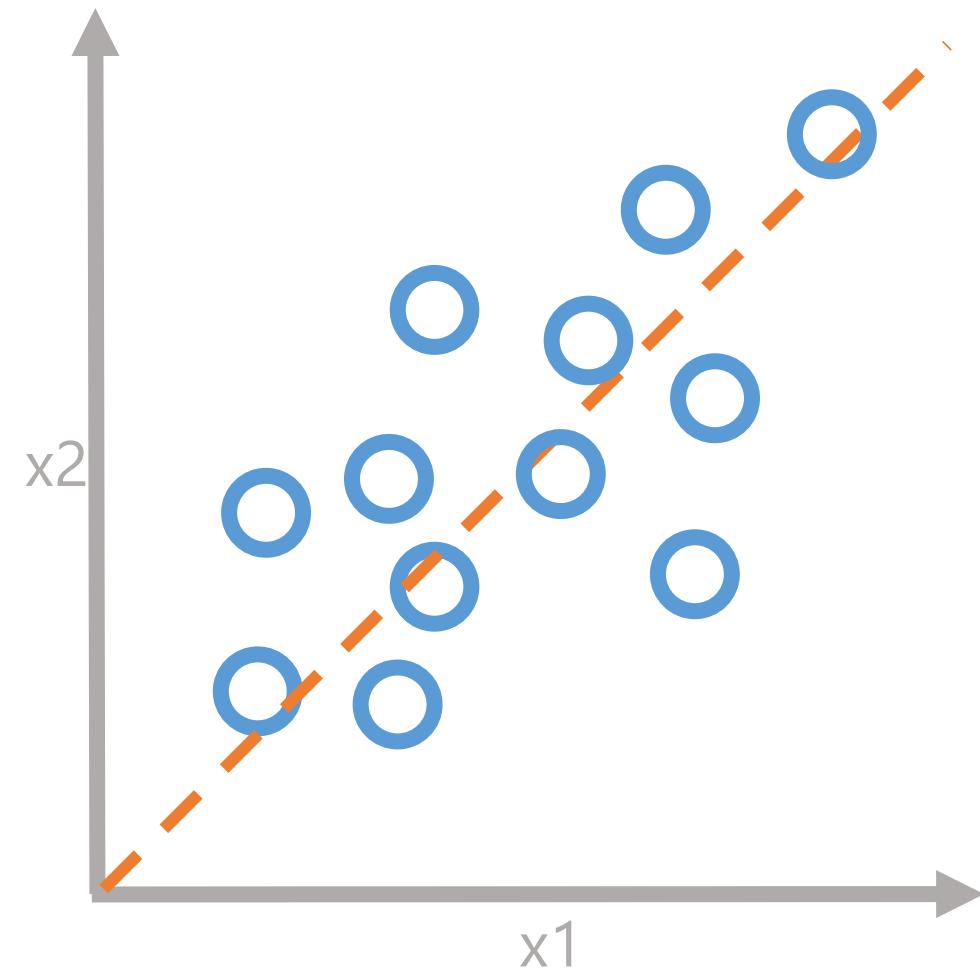
Linear Regression

Polynomial Regression

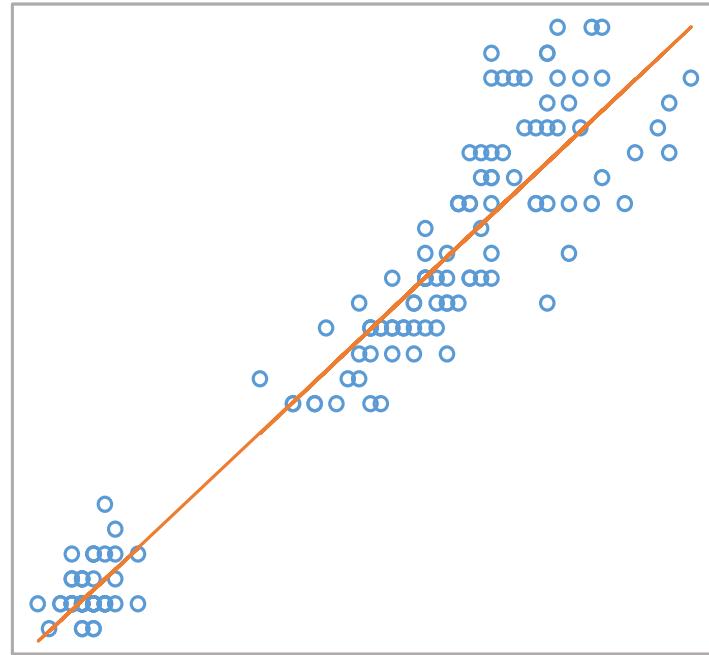
Lasso Regression

ElasticNet Regression

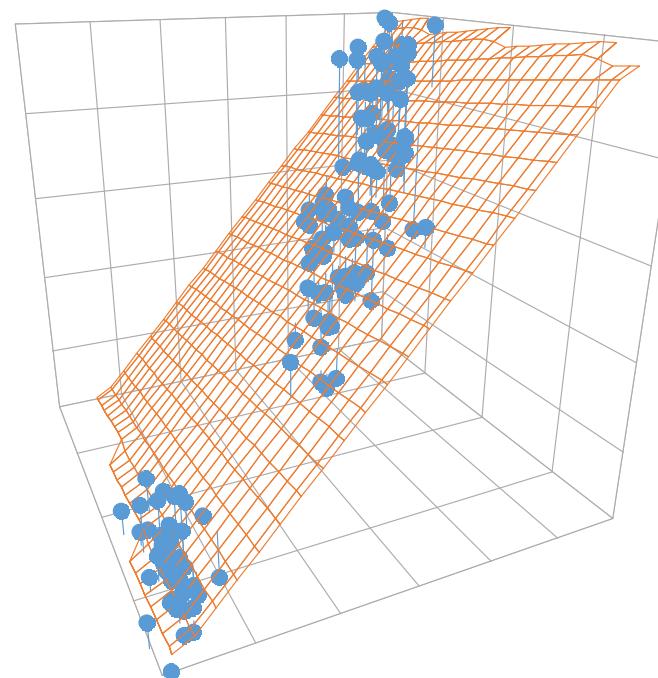
Neural Network Regression



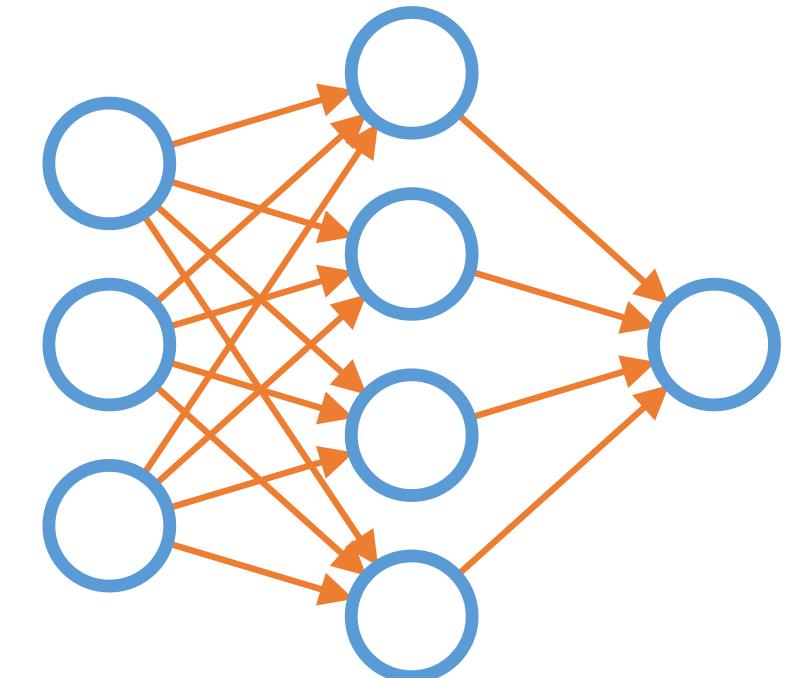
# Regression Algorithms



Simple Linear



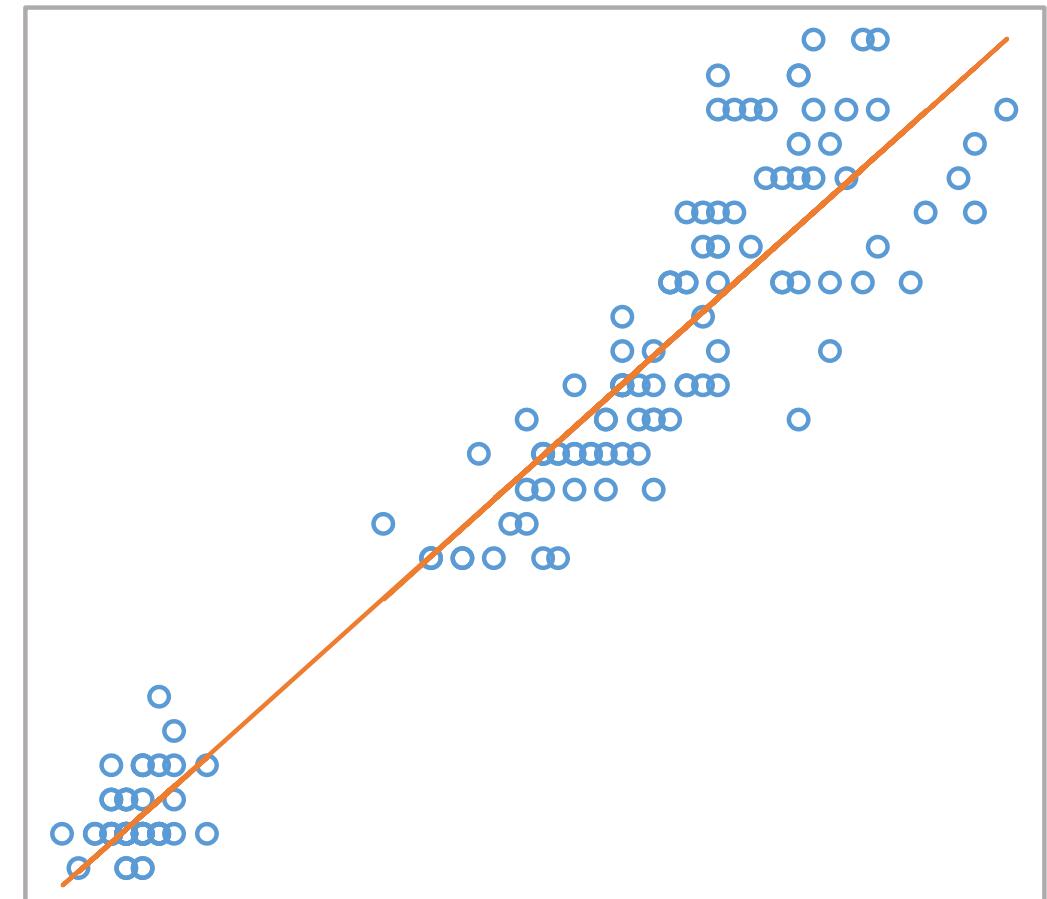
Multiple Linear



Neural Network

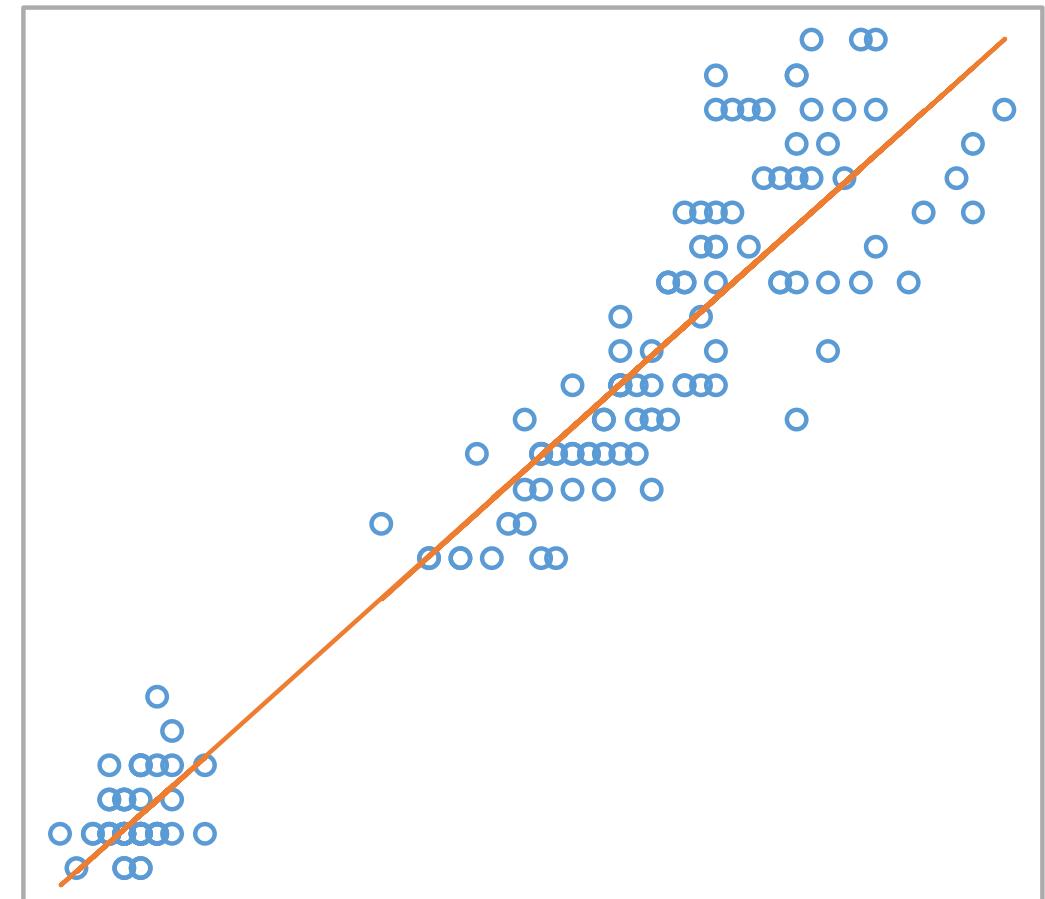
# Simple Linear Regression

Relationship



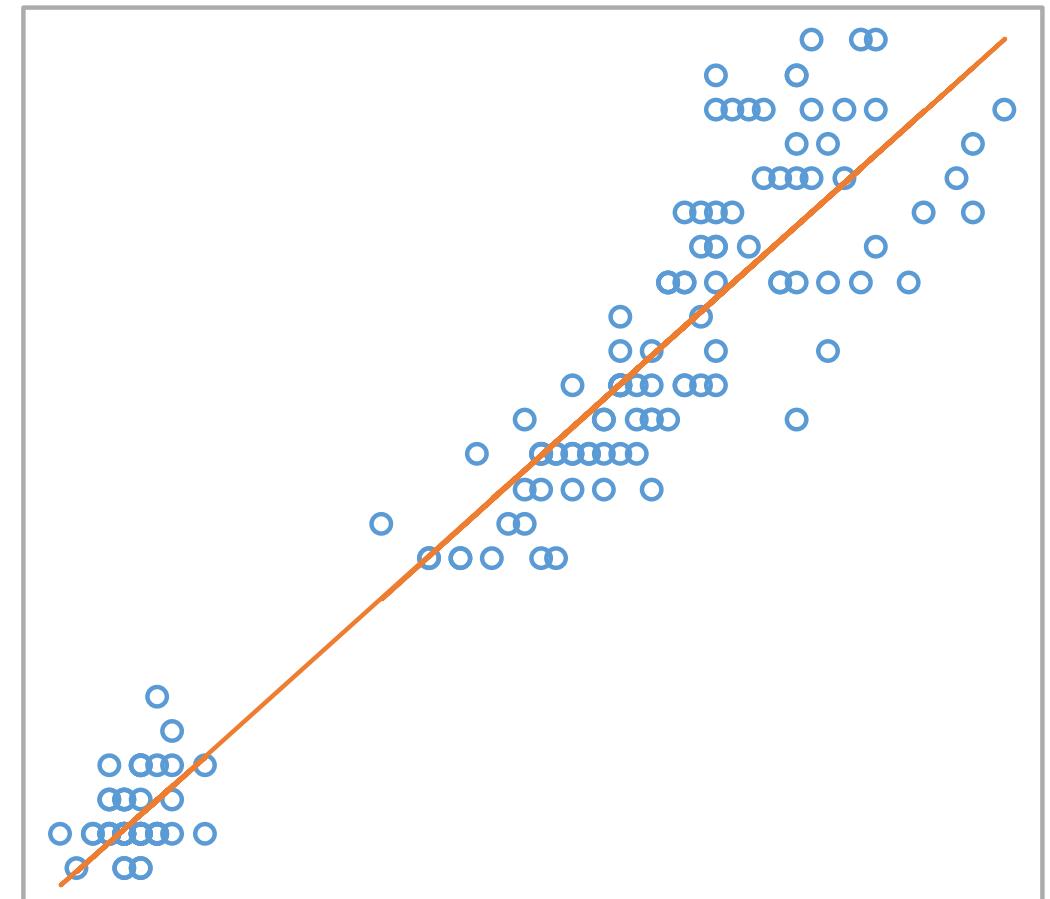
# Simple Linear Regression

Relationship  
Linear model



# Simple Linear Regression

Relationship  
Linear model  
 $y = m \cdot x + b$



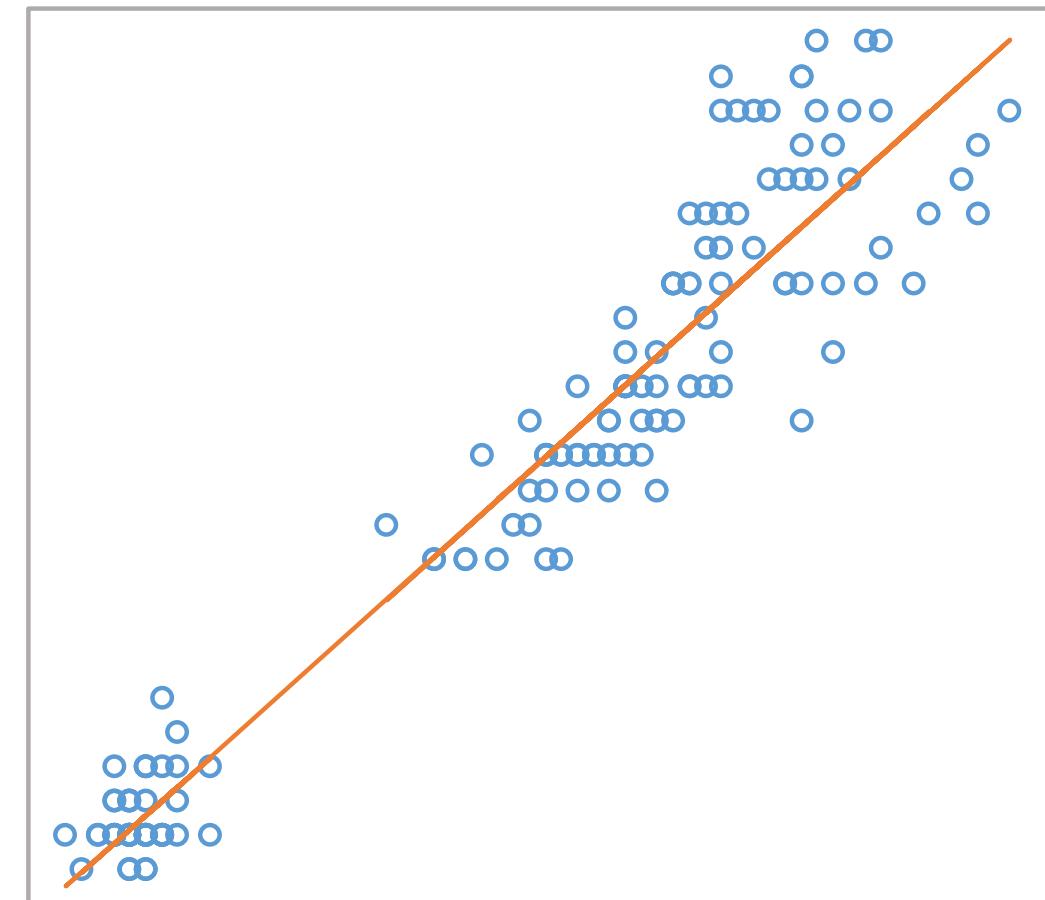
# Simple Linear Regression

Relationship

Linear model

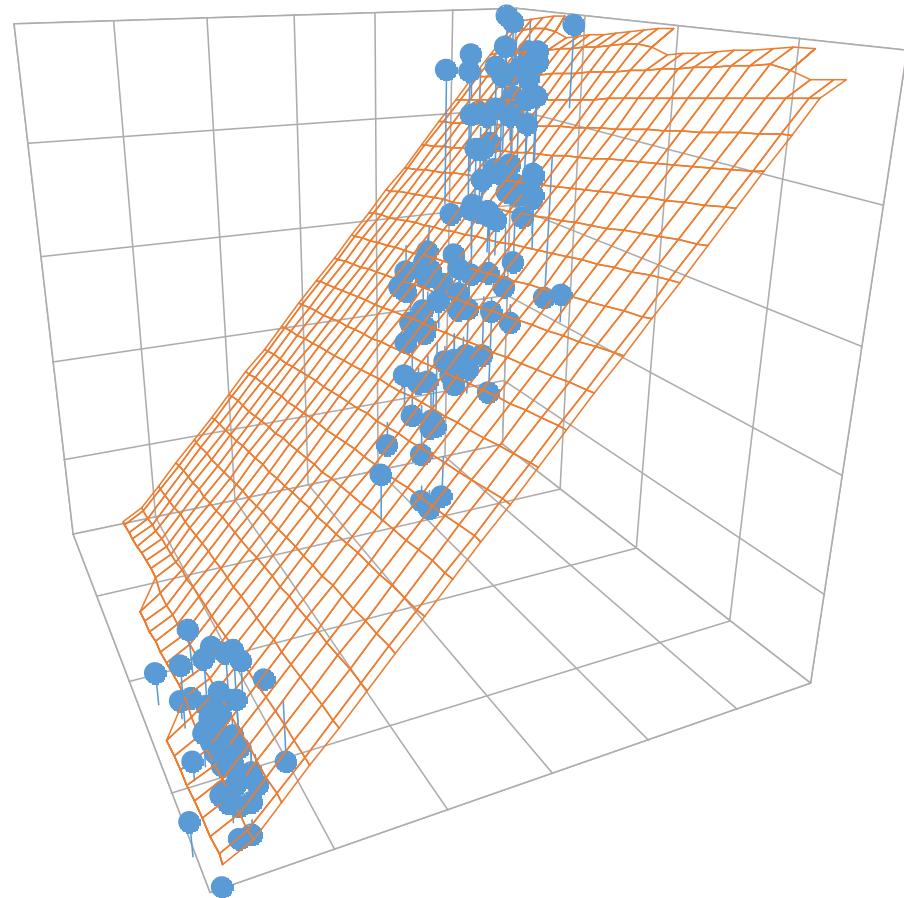
$$y = m \cdot x + b$$

Parameters estimated



# Multiple Linear Regression

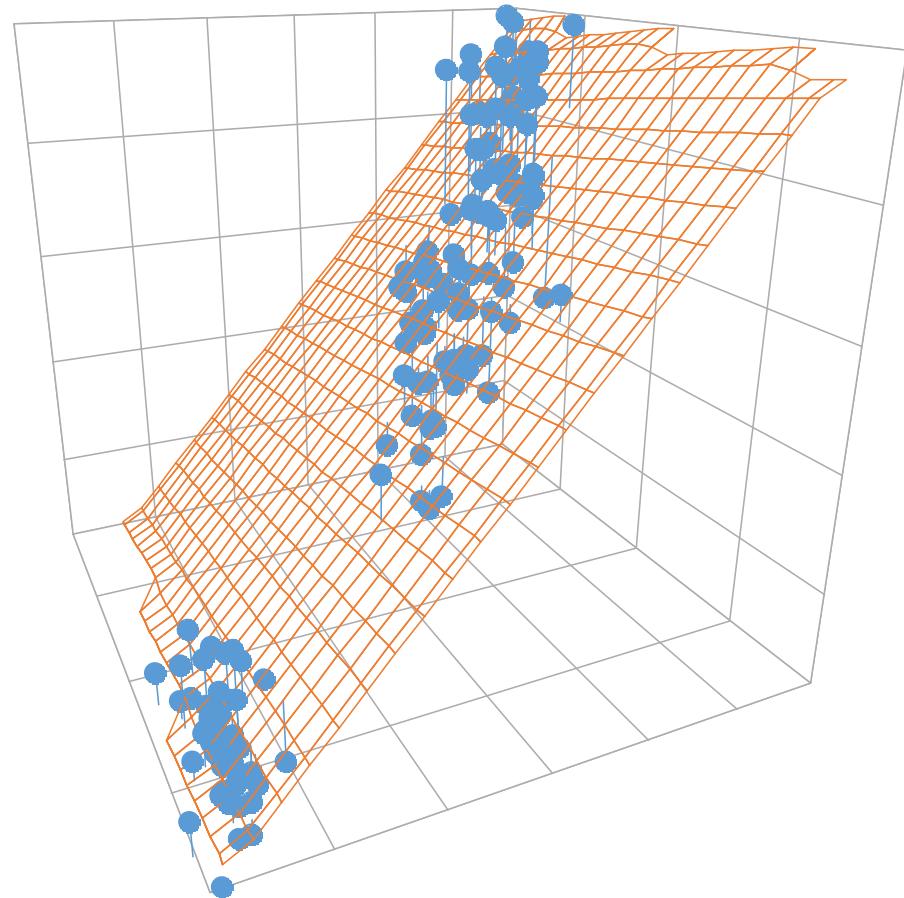
Similar to SLR



# Multiple Linear Regression

Similar to SLR

Multiple variables

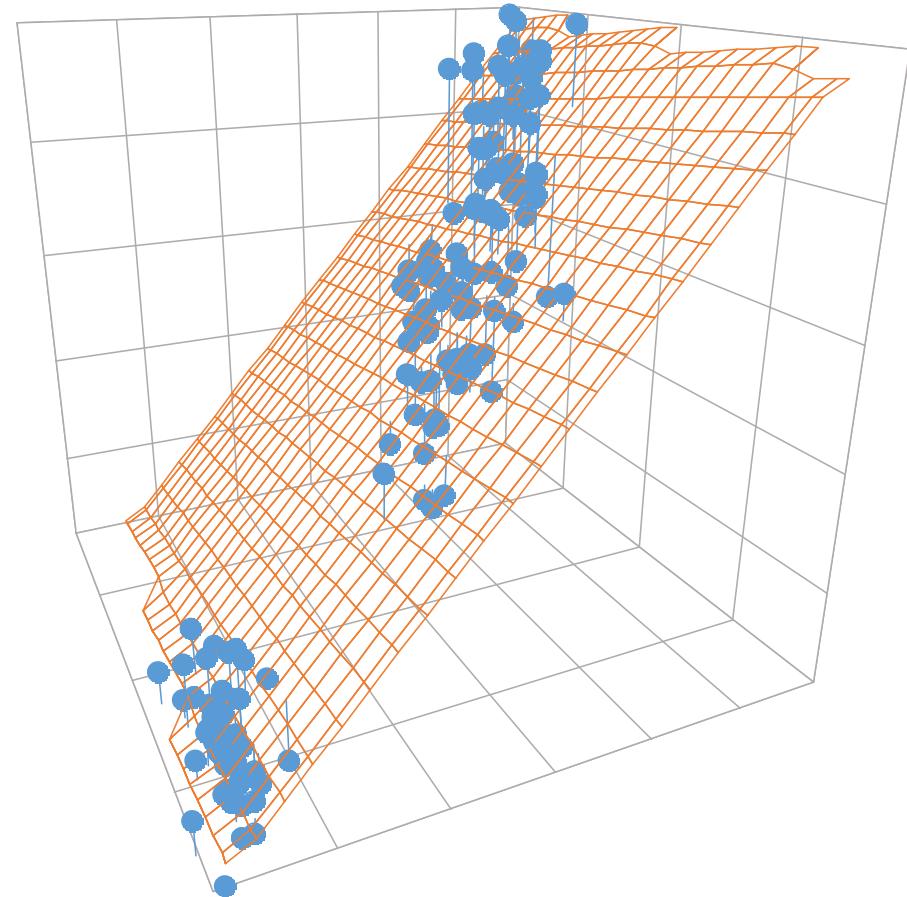


# Multiple Linear Regression

Similar to SLR

Multiple variables

Multiple slopes



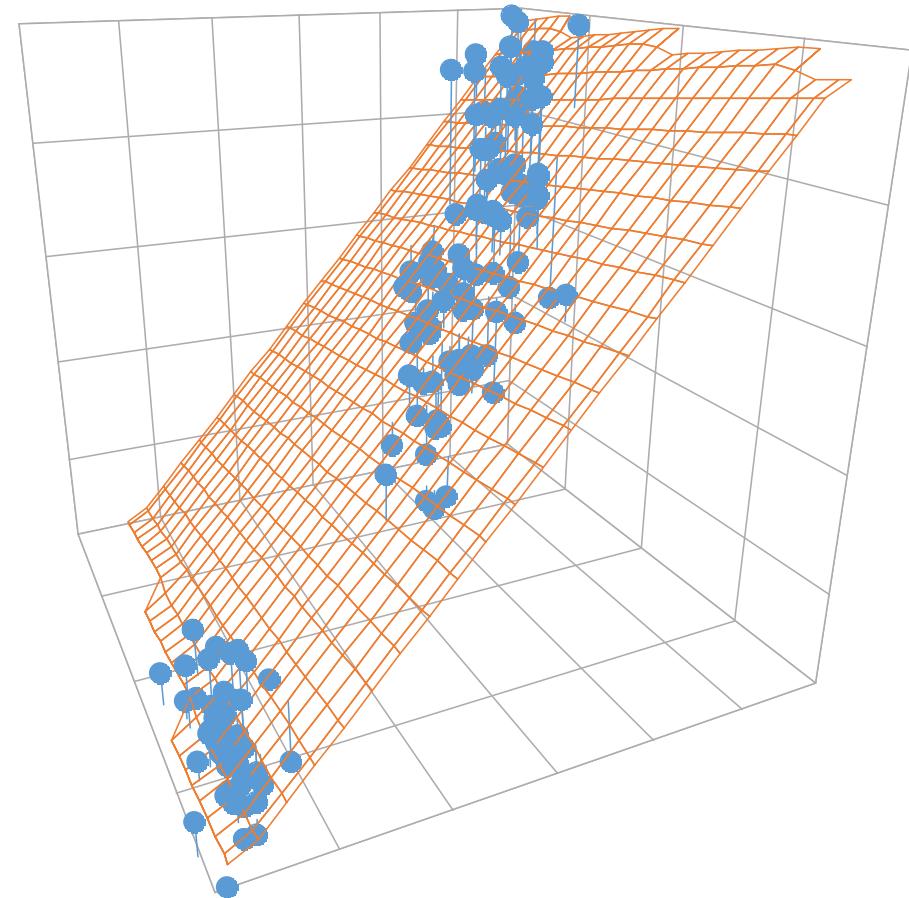
# Multiple Linear Regression

Similar to SLR

Multiple variables

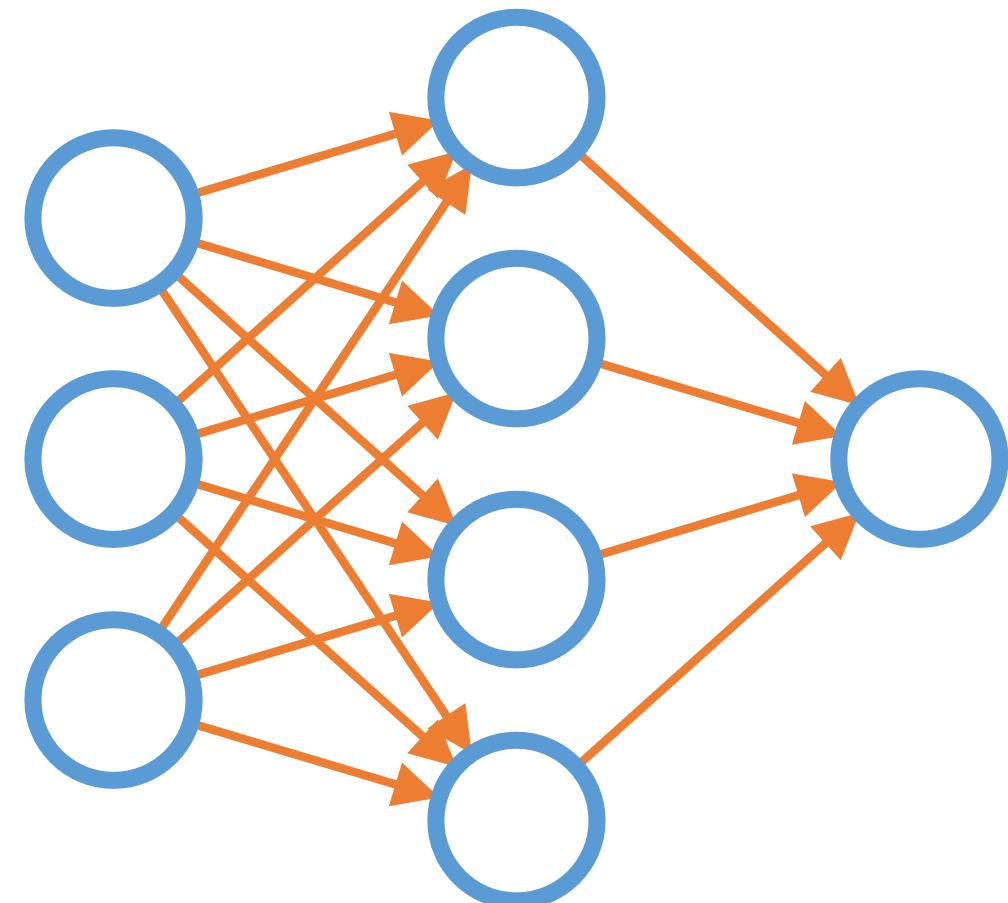
Multiple slopes

Categorical variables



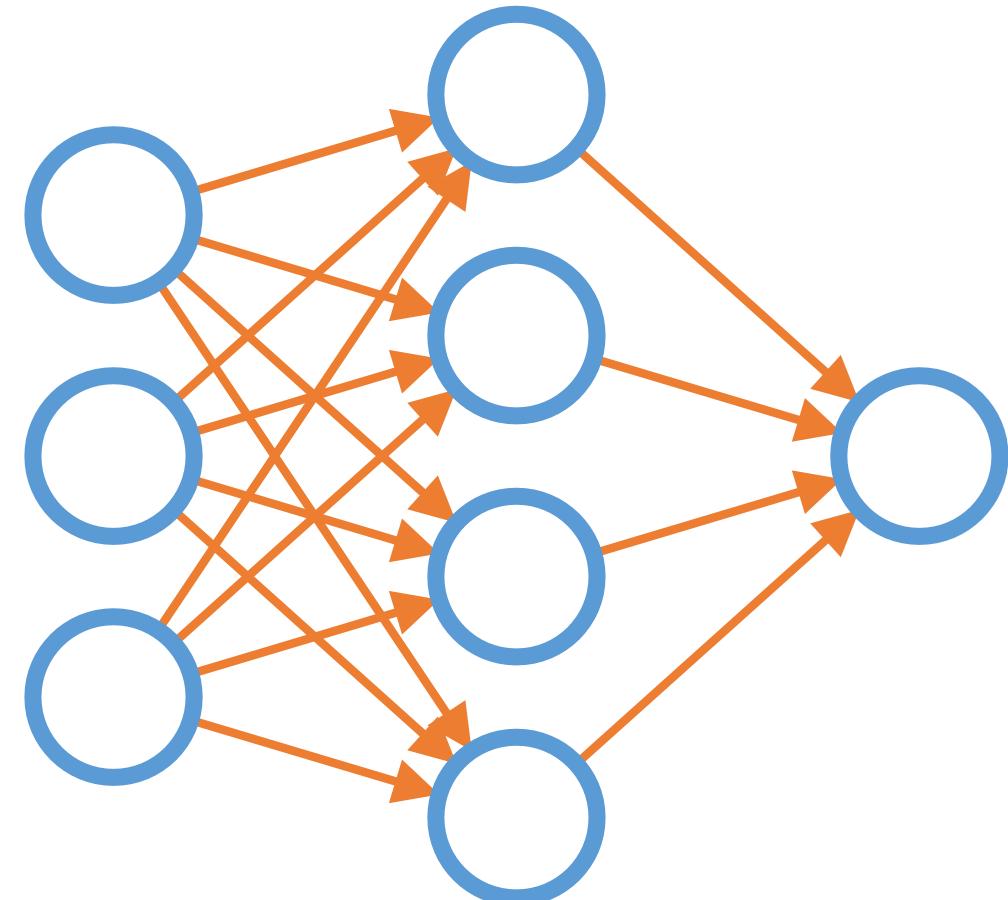
# Neural Network Regression

Similar to NN classifier



# Neural Network Regression

Similar to NN classifier  
Numeric output



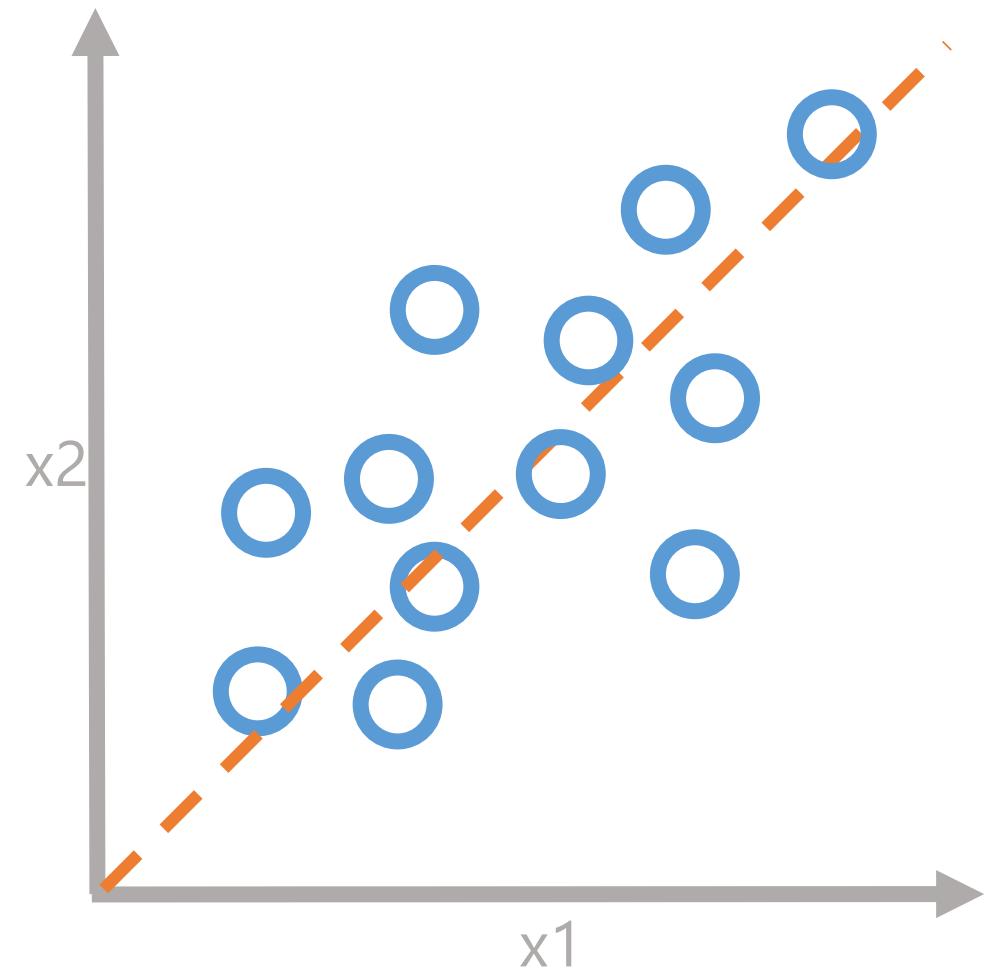
# Real-World Examples

How much profit will we make?

What will the price be tomorrow?

How many units will they buy?

How long until this part fails?



# Demo 3 - Regression

Goal: Predict petal width  
based on petal length

# Lab 3A – Regression (Easy)

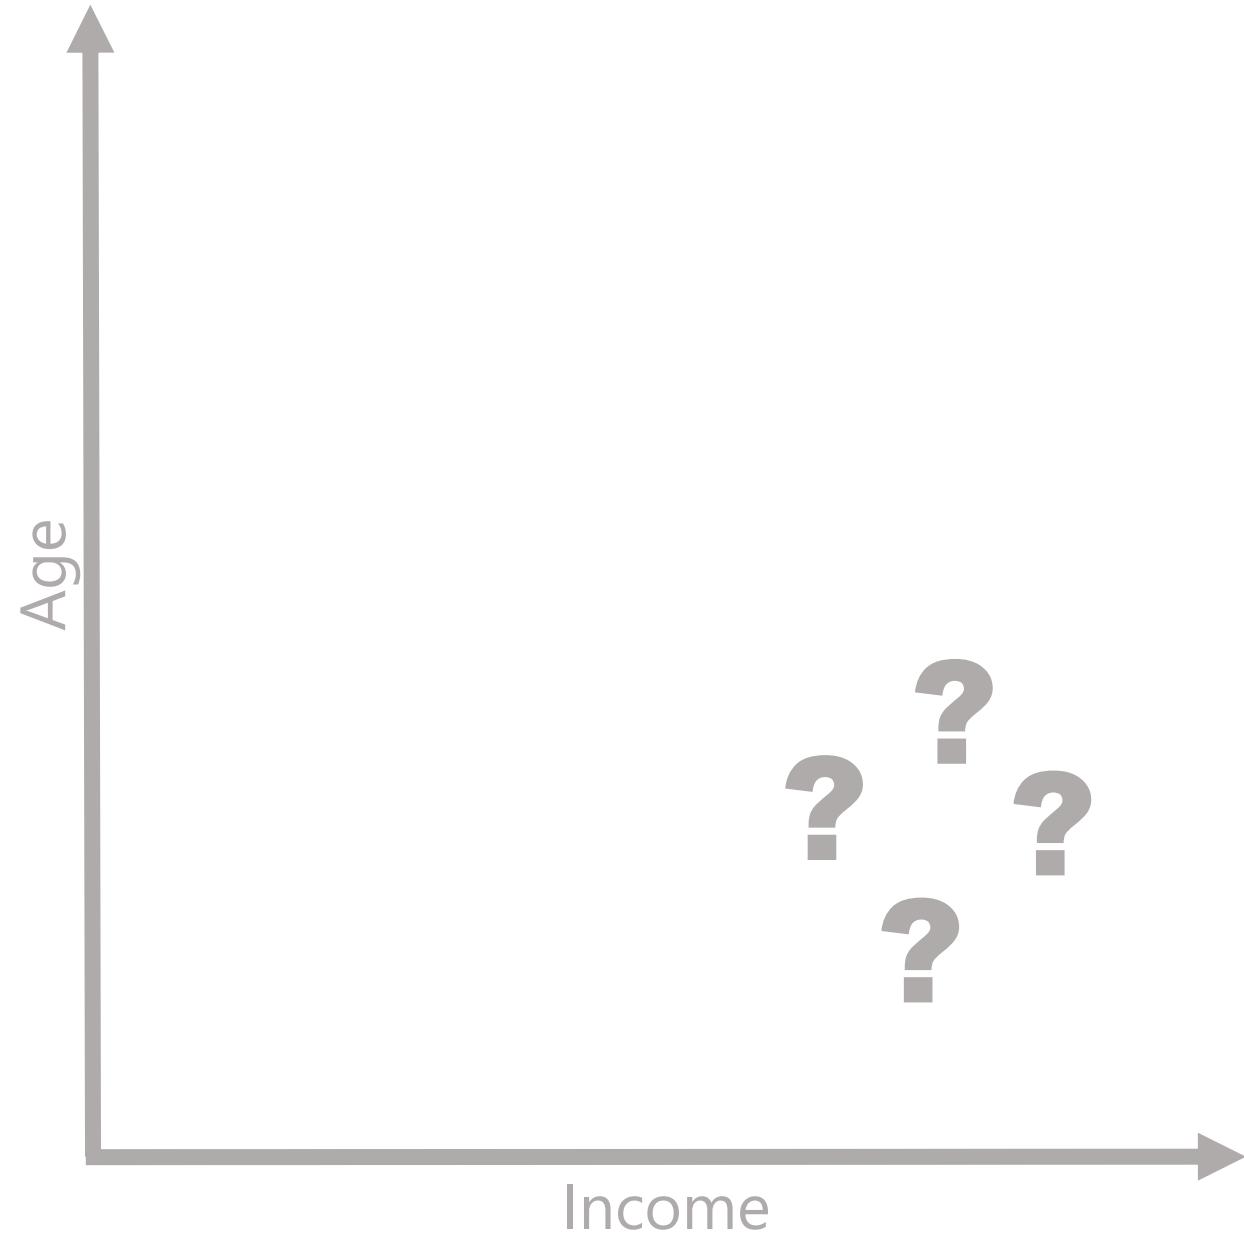
Goal: Predict petal width

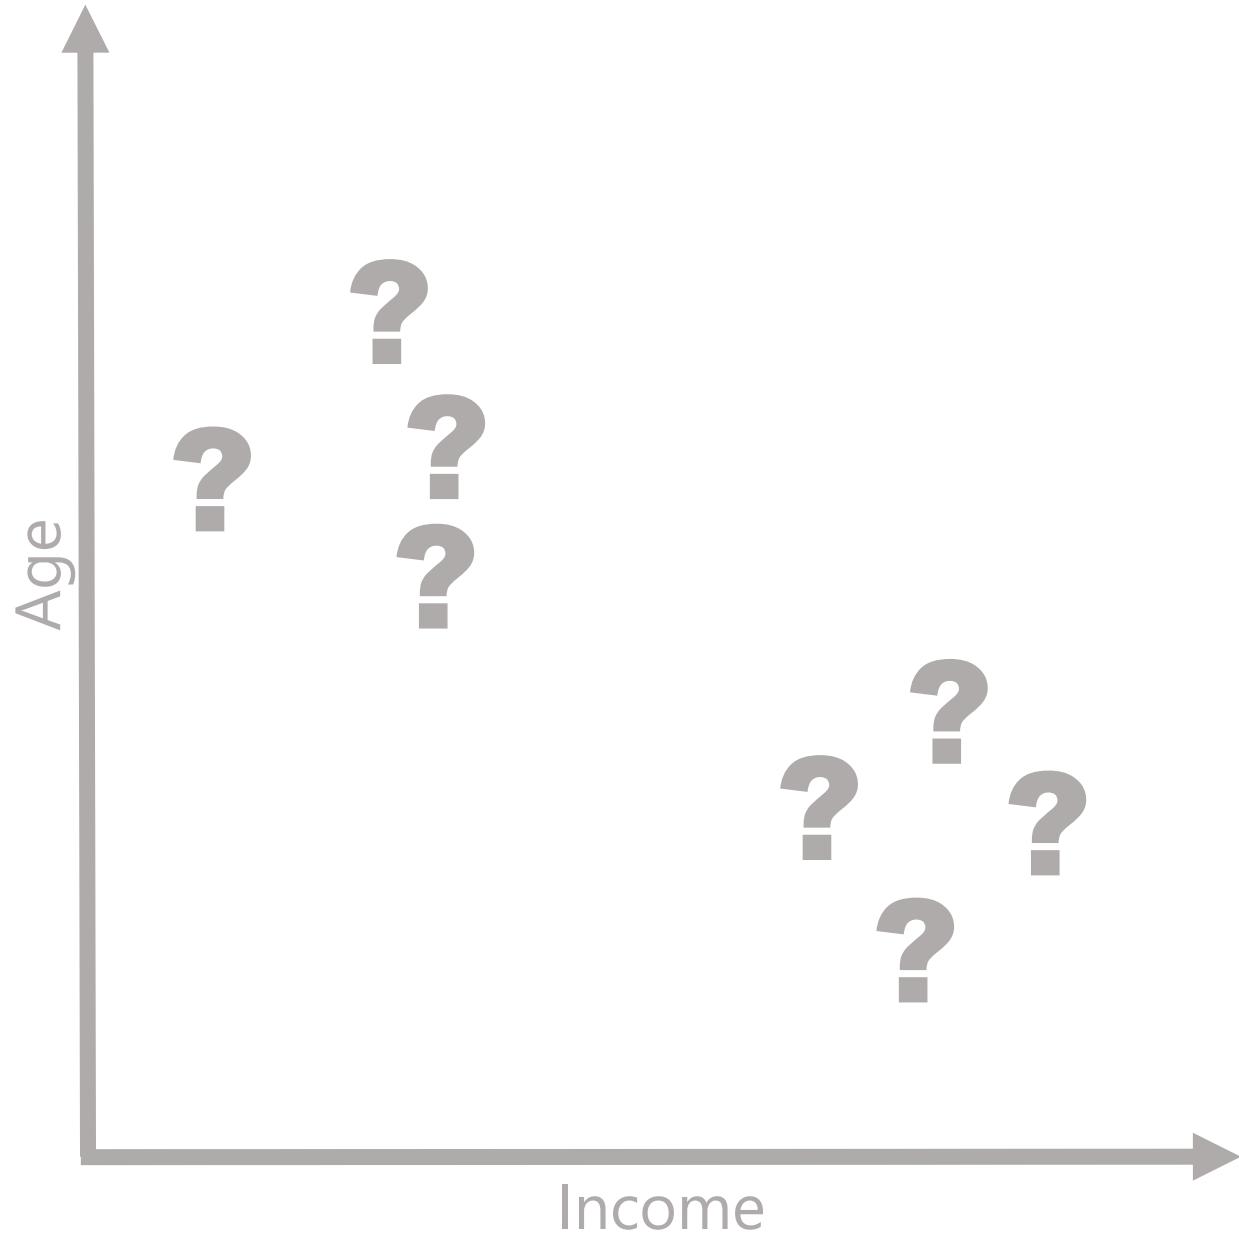
# Lab 3B – Regression (Hard)

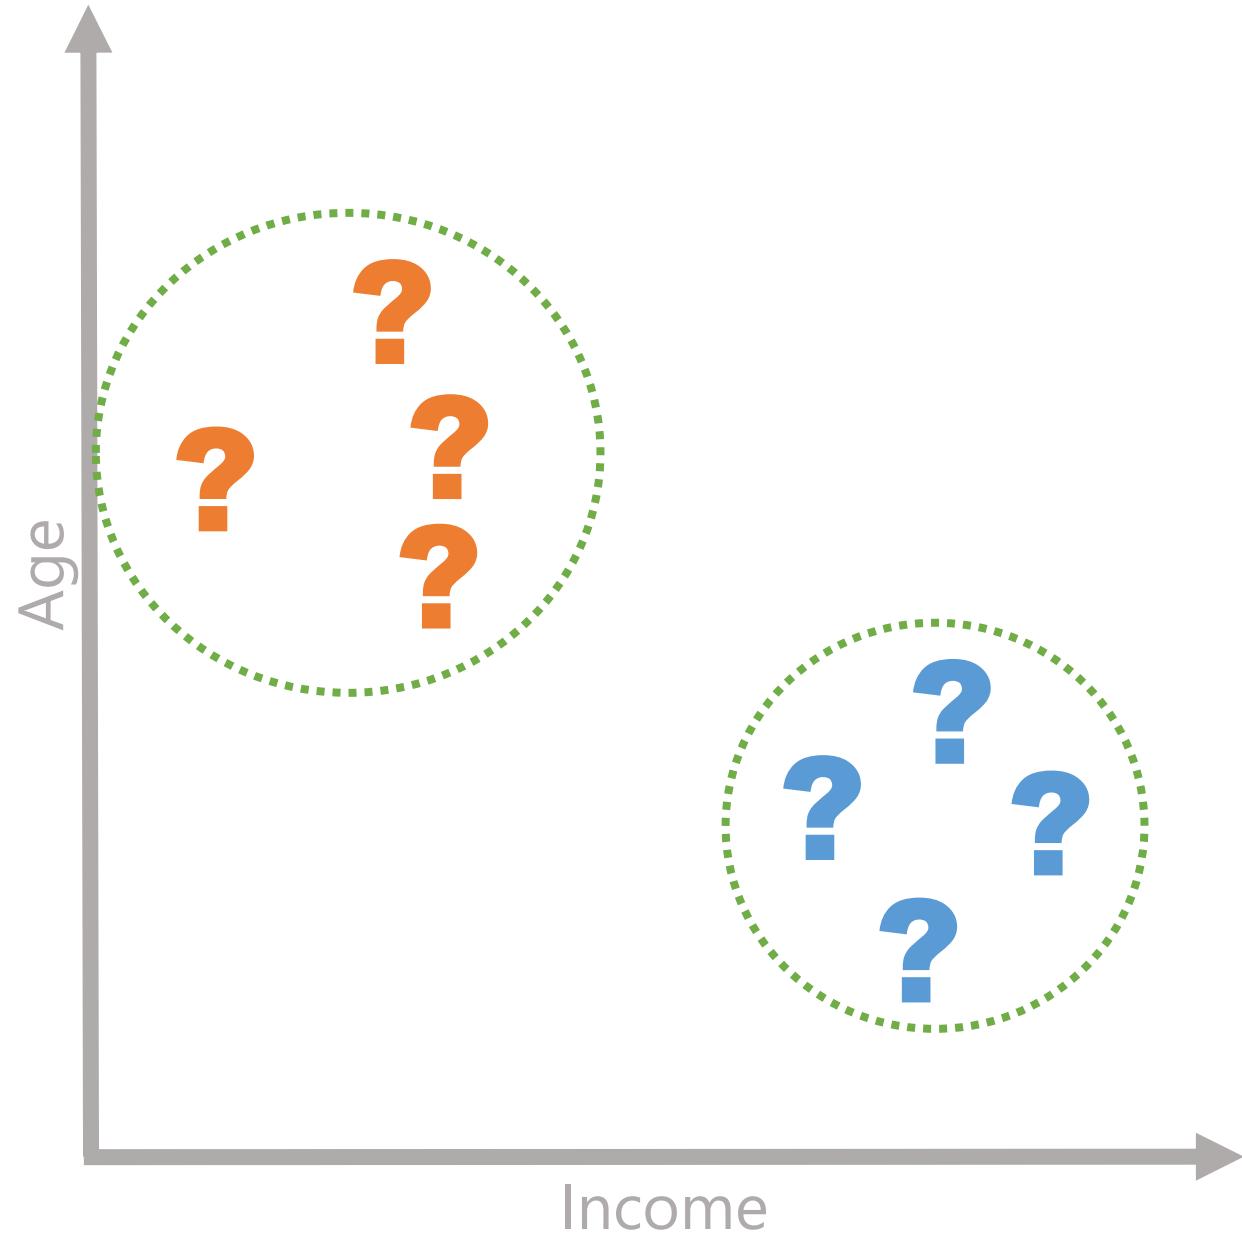
Goal: Predict mortality rate

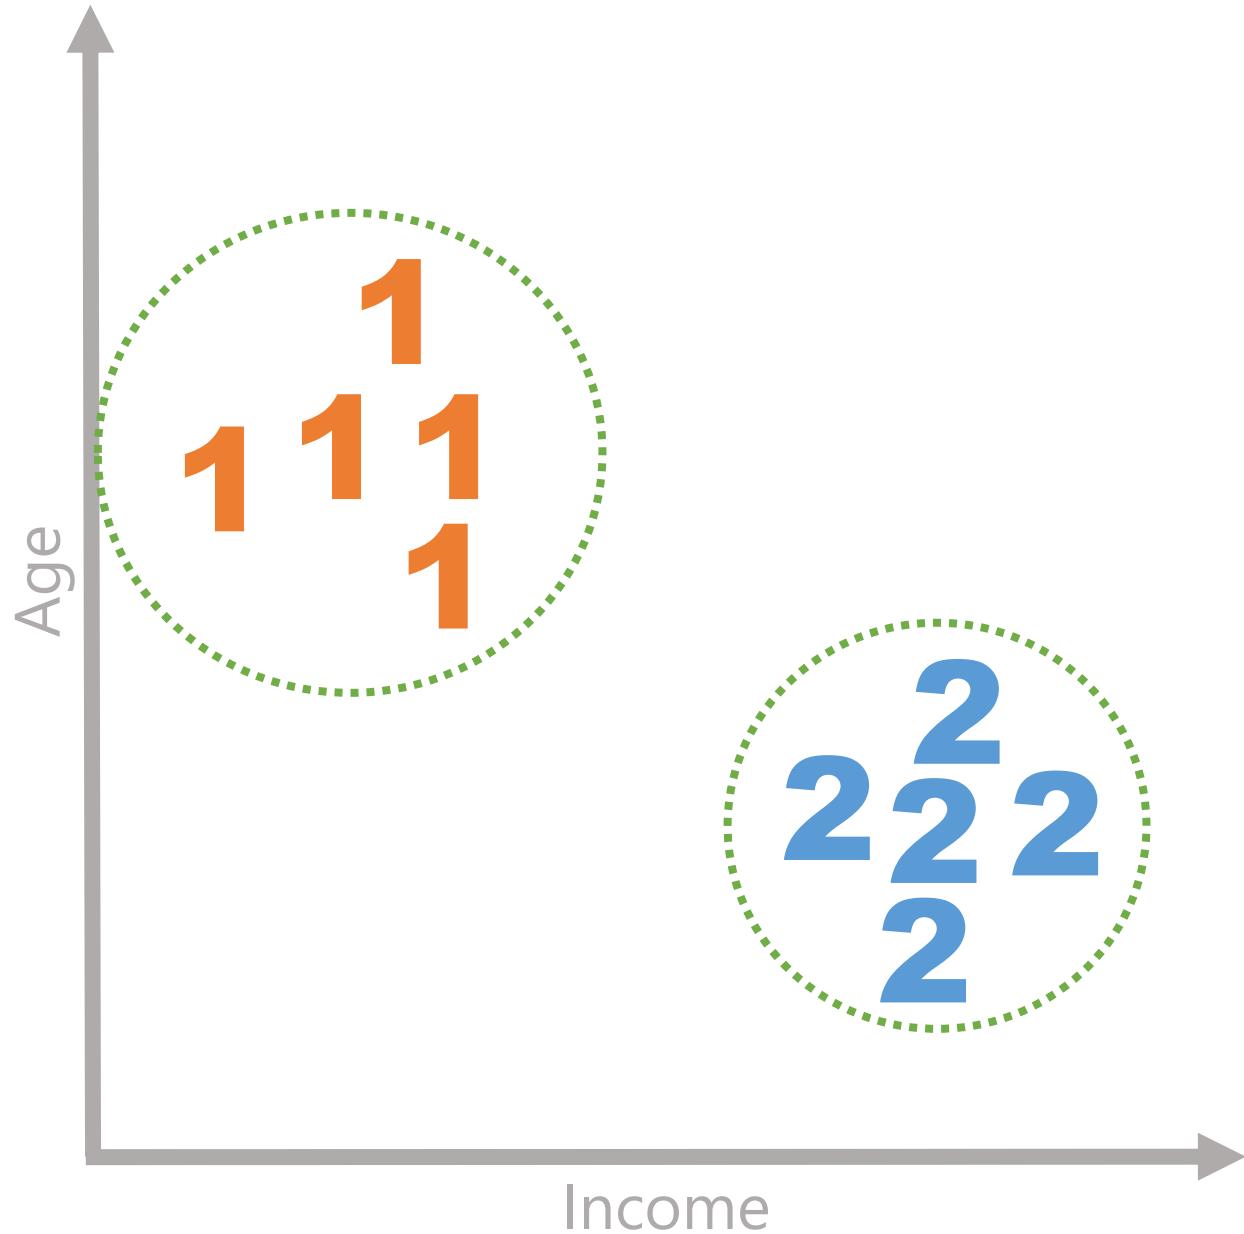
# Clustering

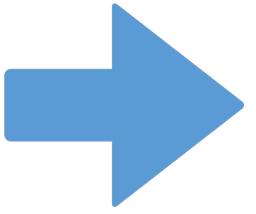
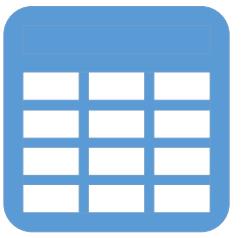




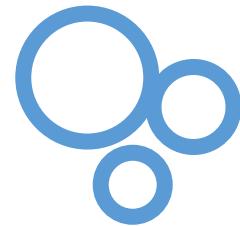
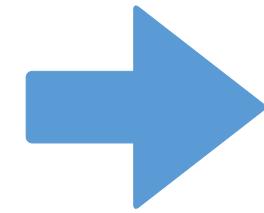








$f(x)$

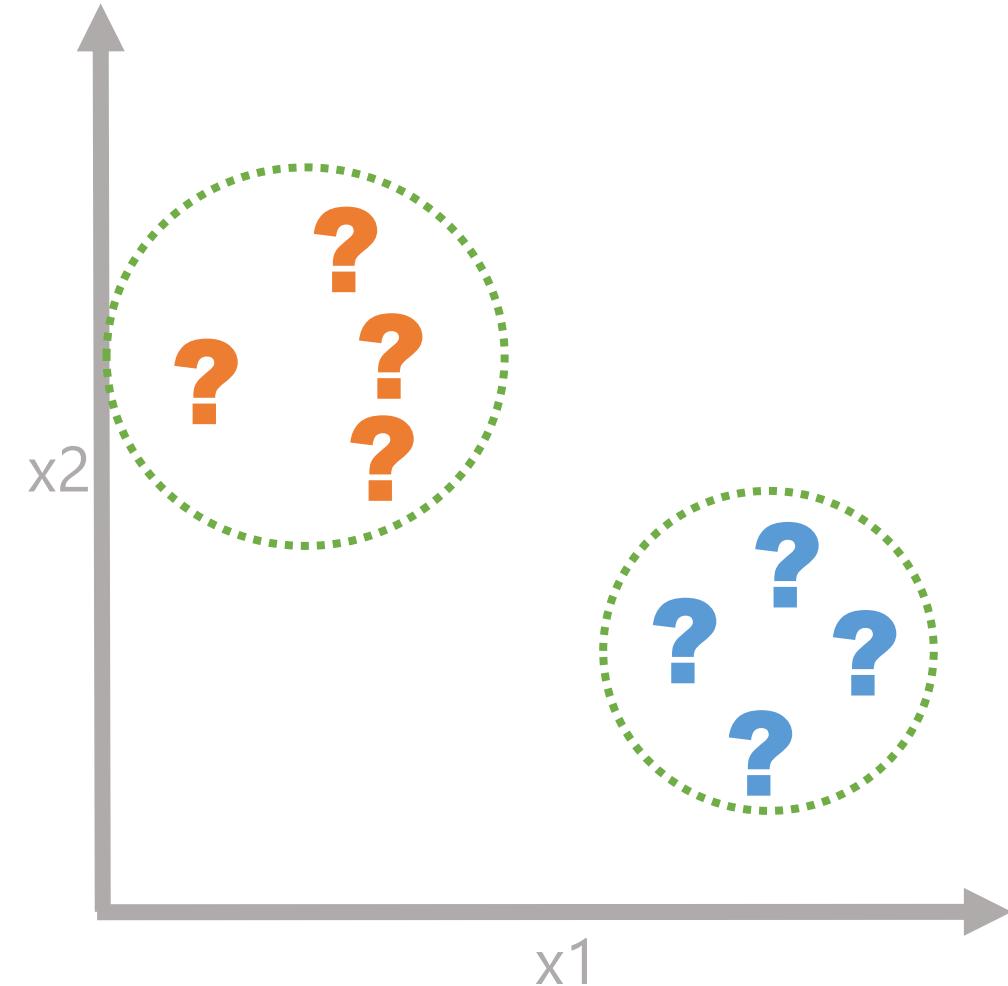


# Clustering Algorithms

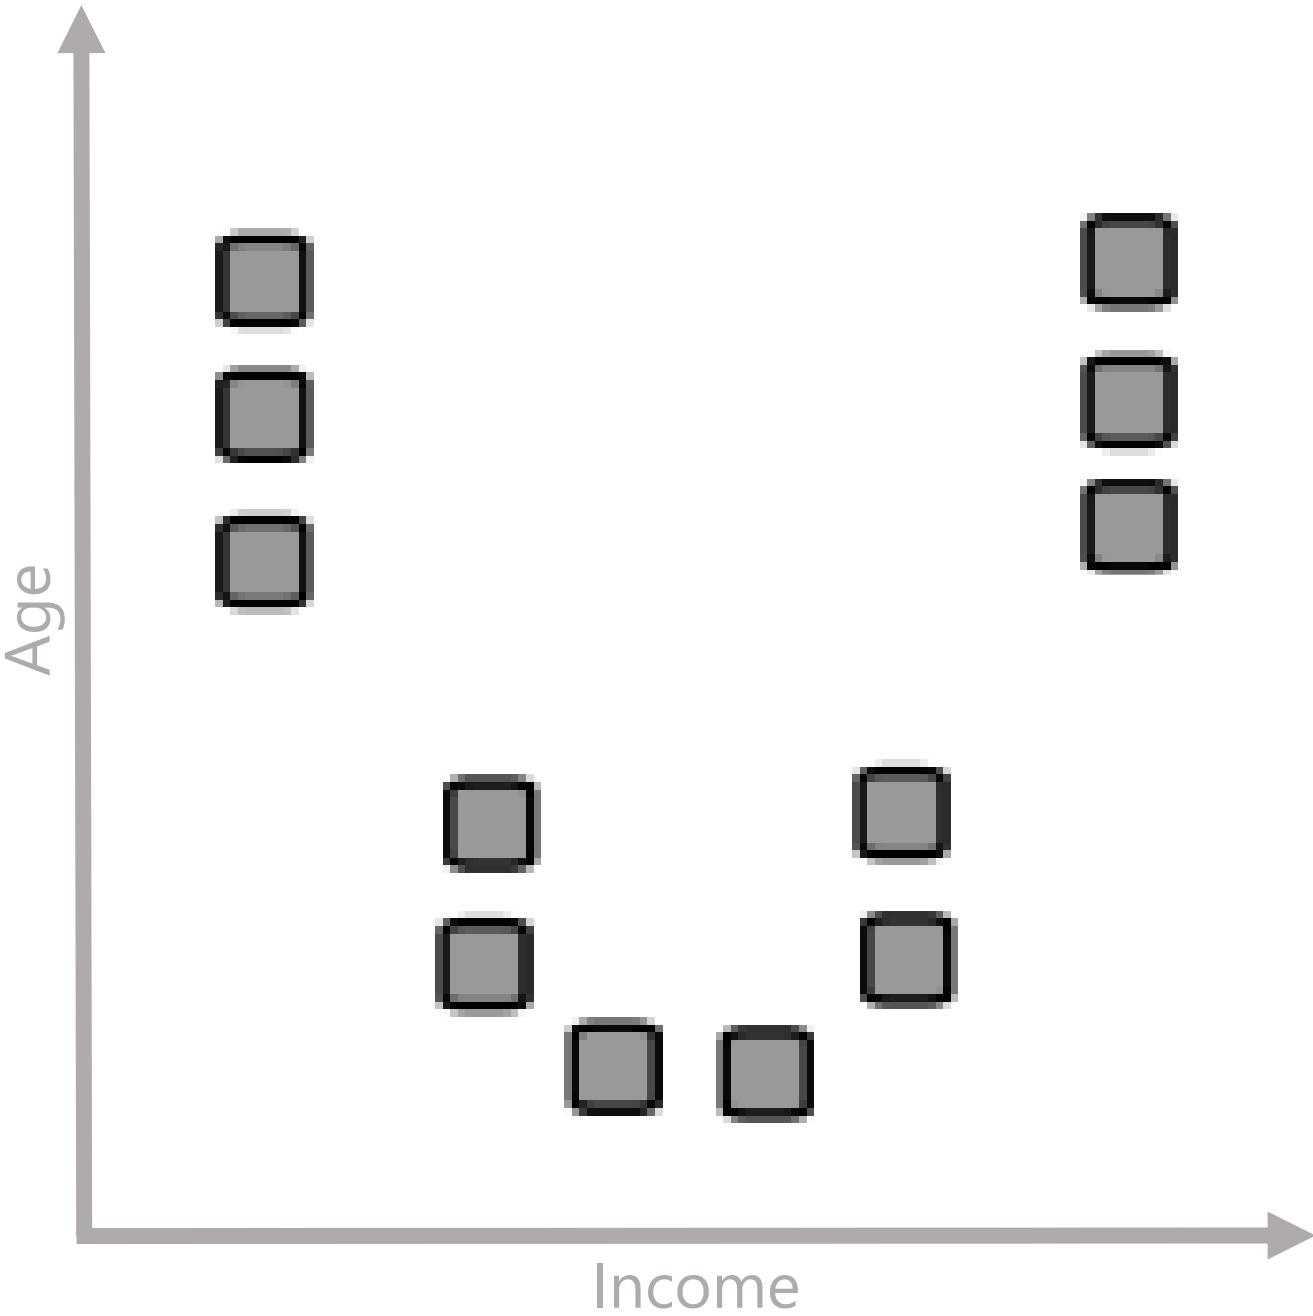
K-means

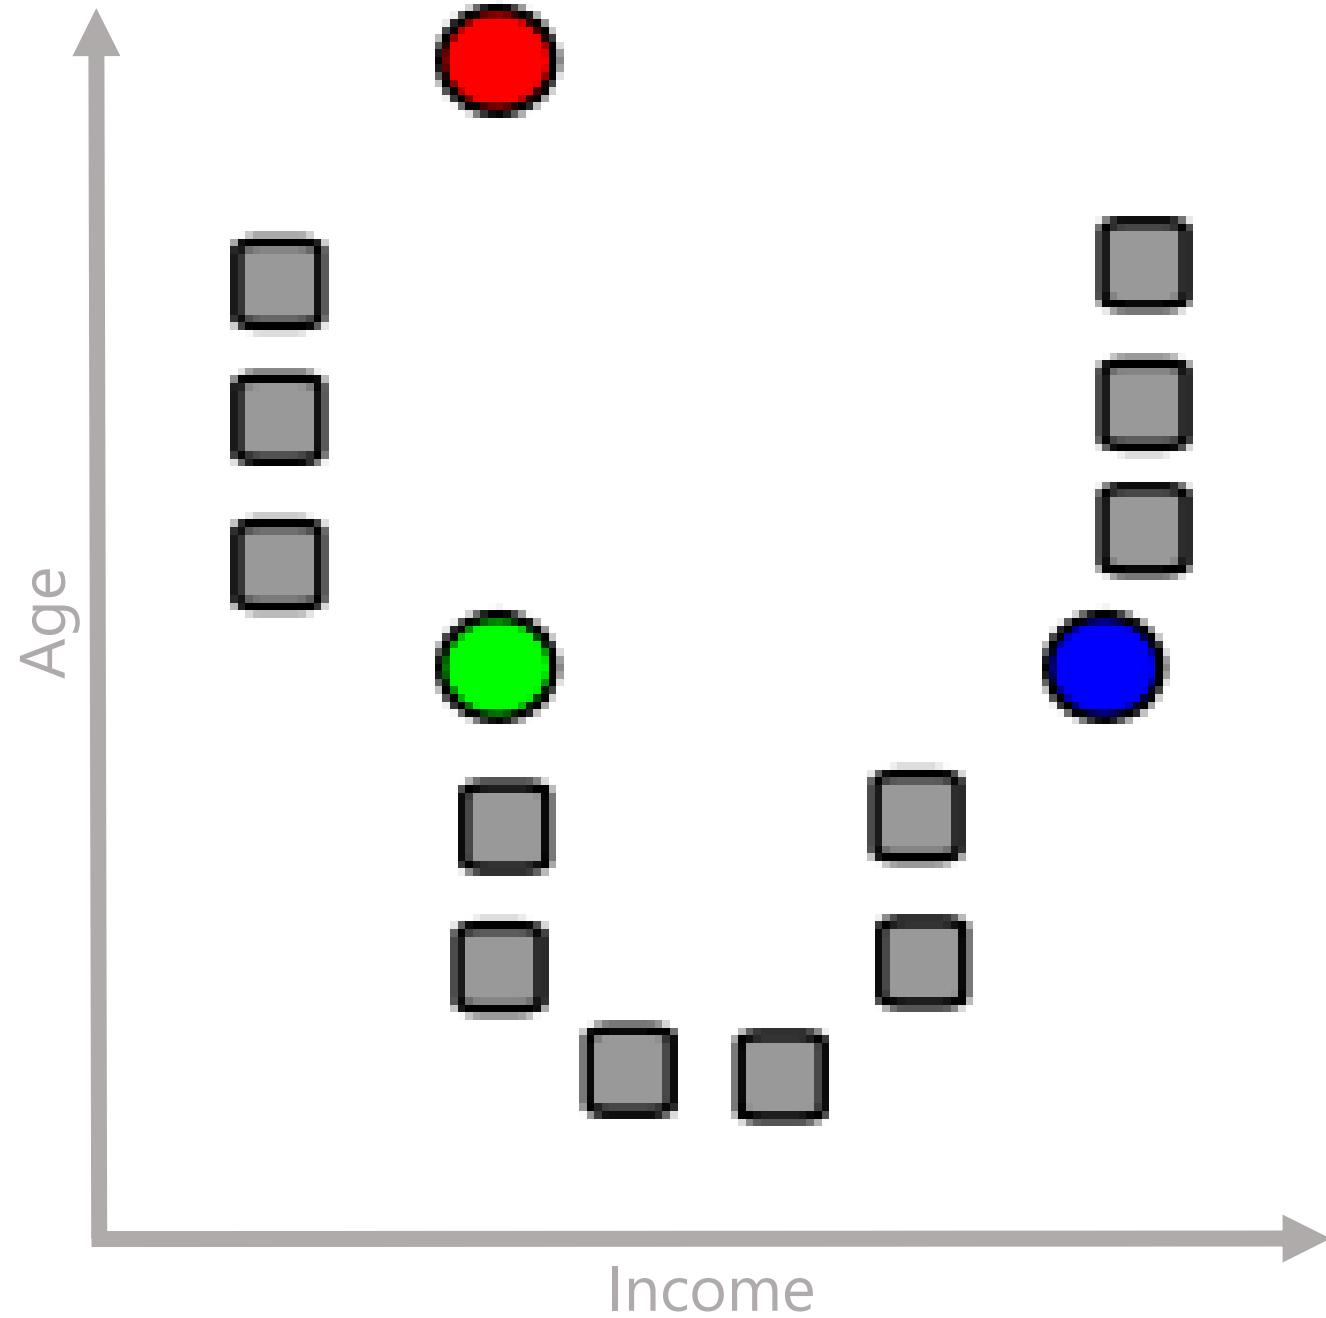
Hierarchical clustering

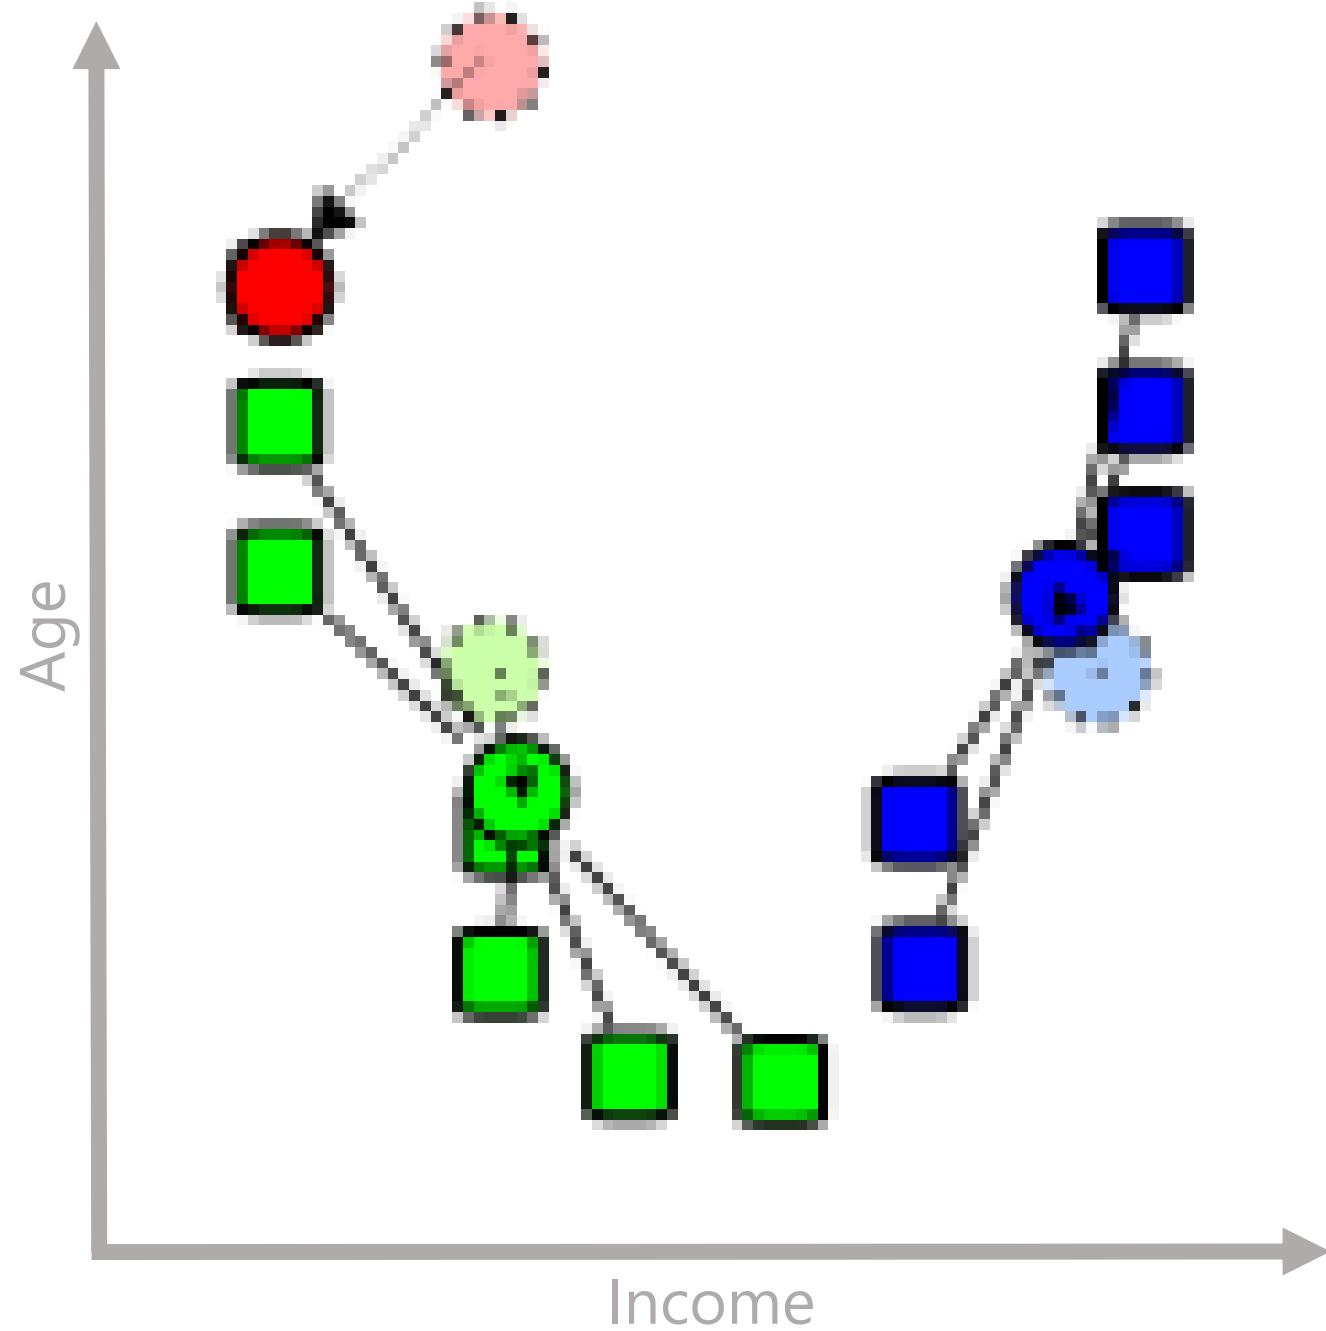
Expectation maximization

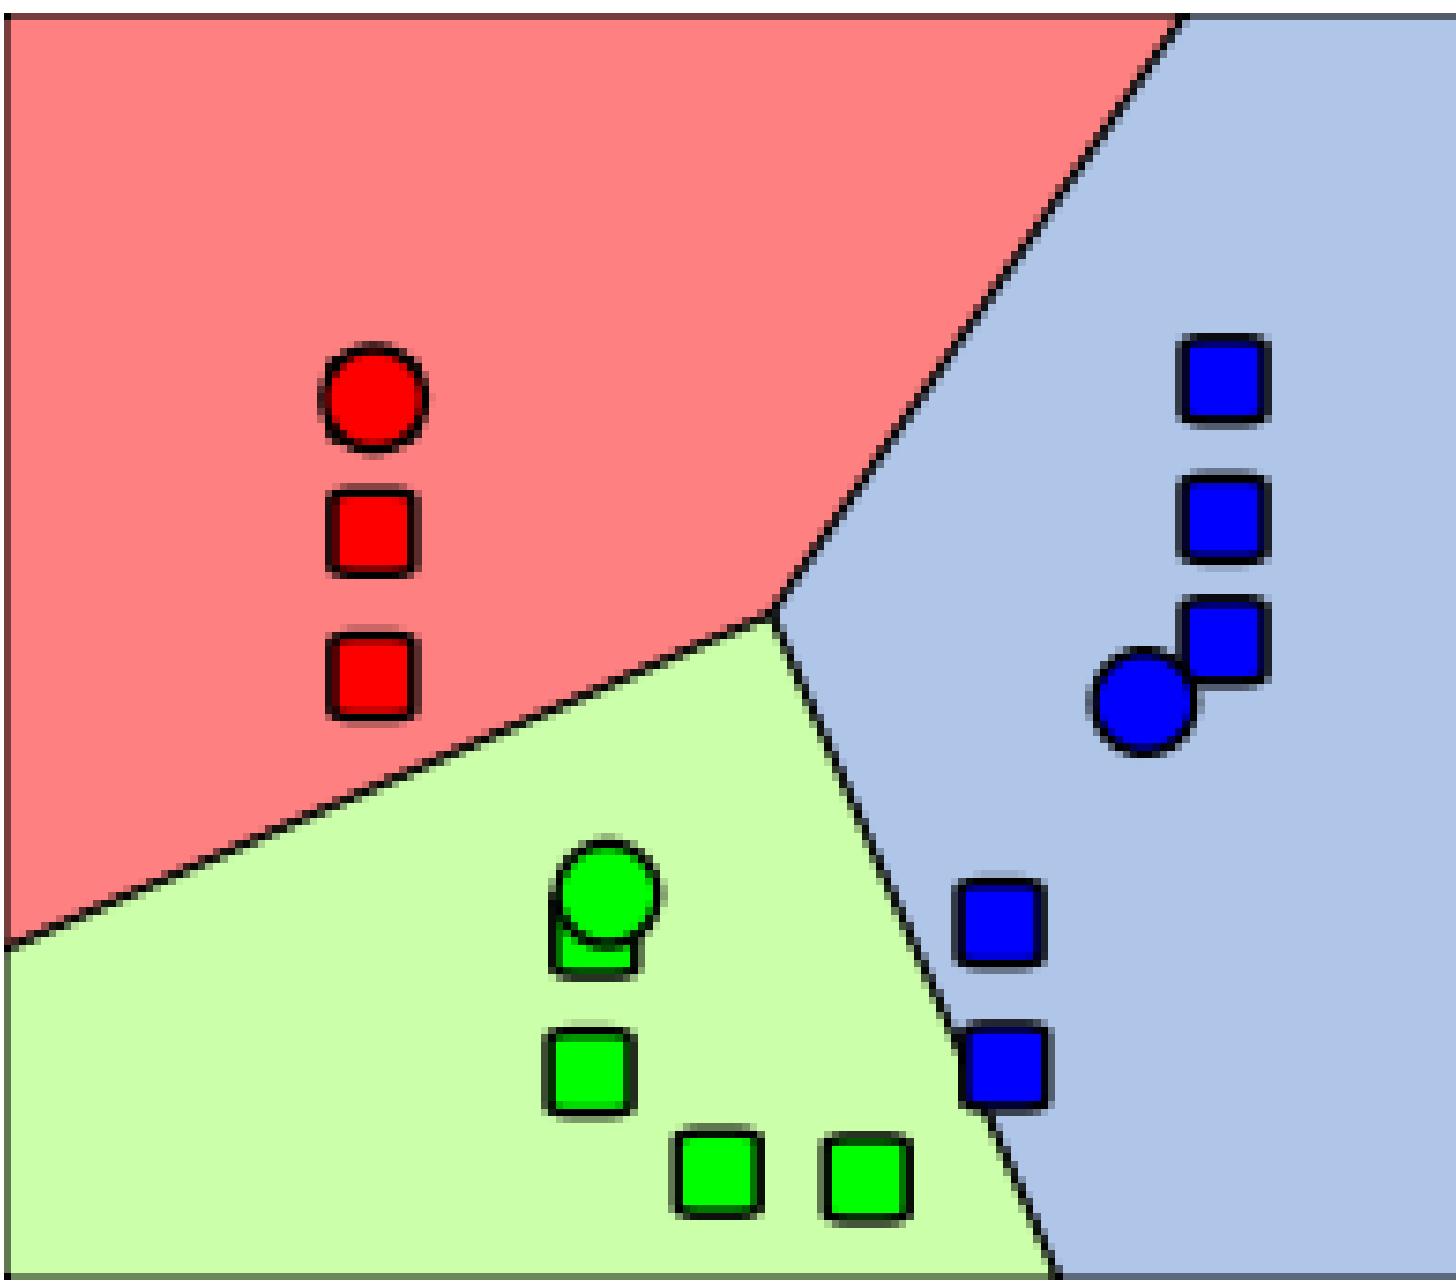


# k-Means Clustering



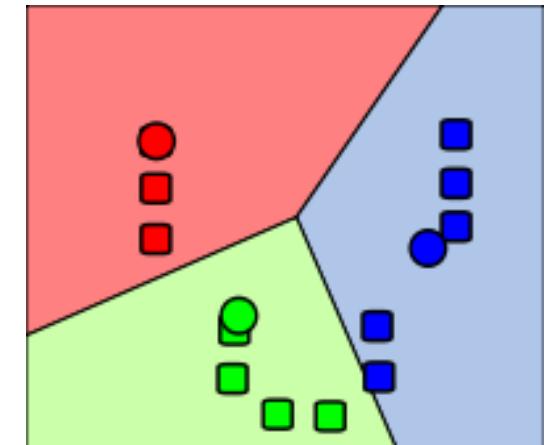
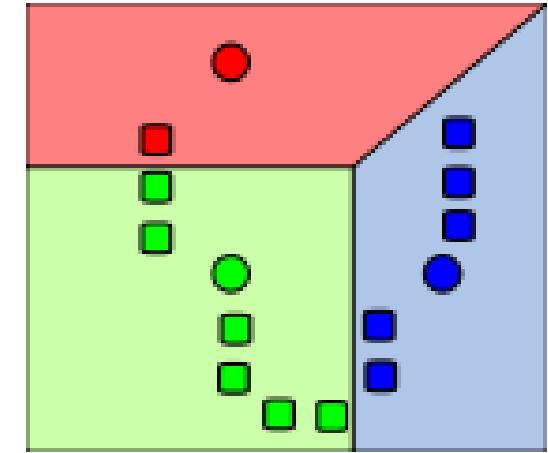
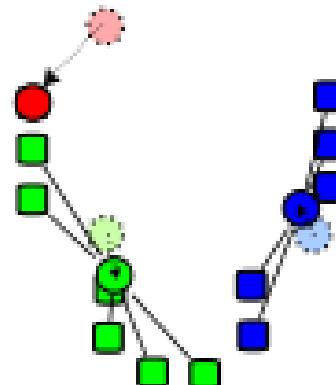
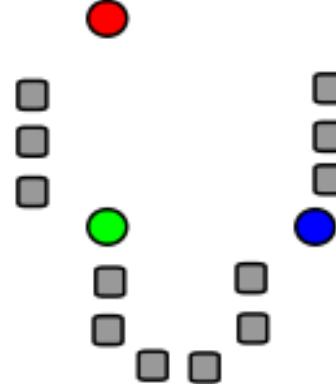






# k-Means Clustering

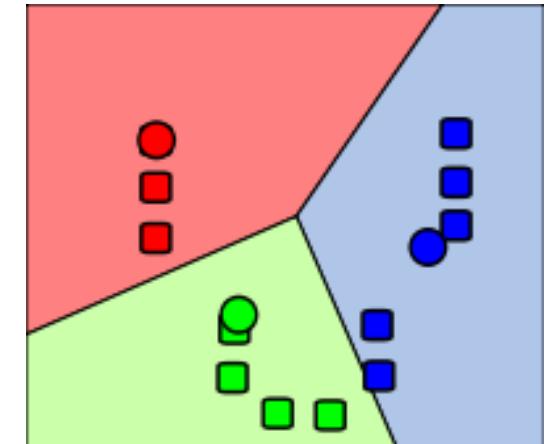
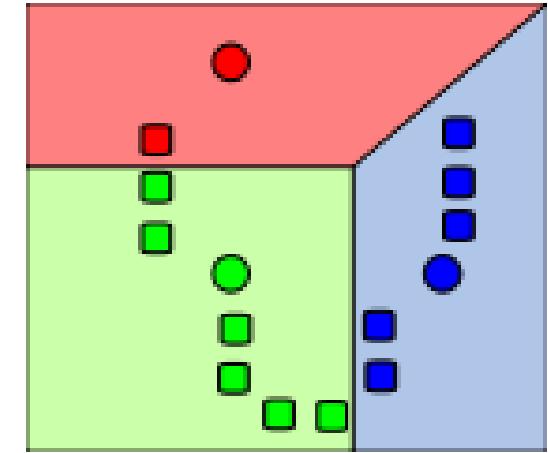
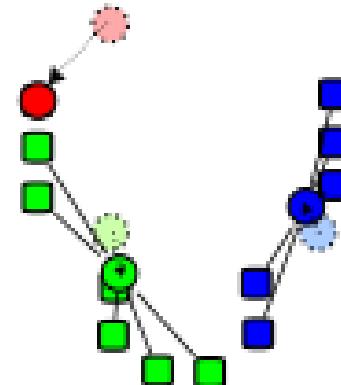
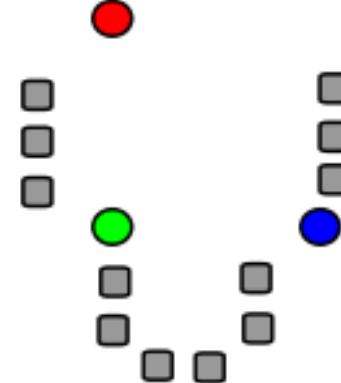
Unsupervised learning



Source: Wikipedia

# k-Means Clustering

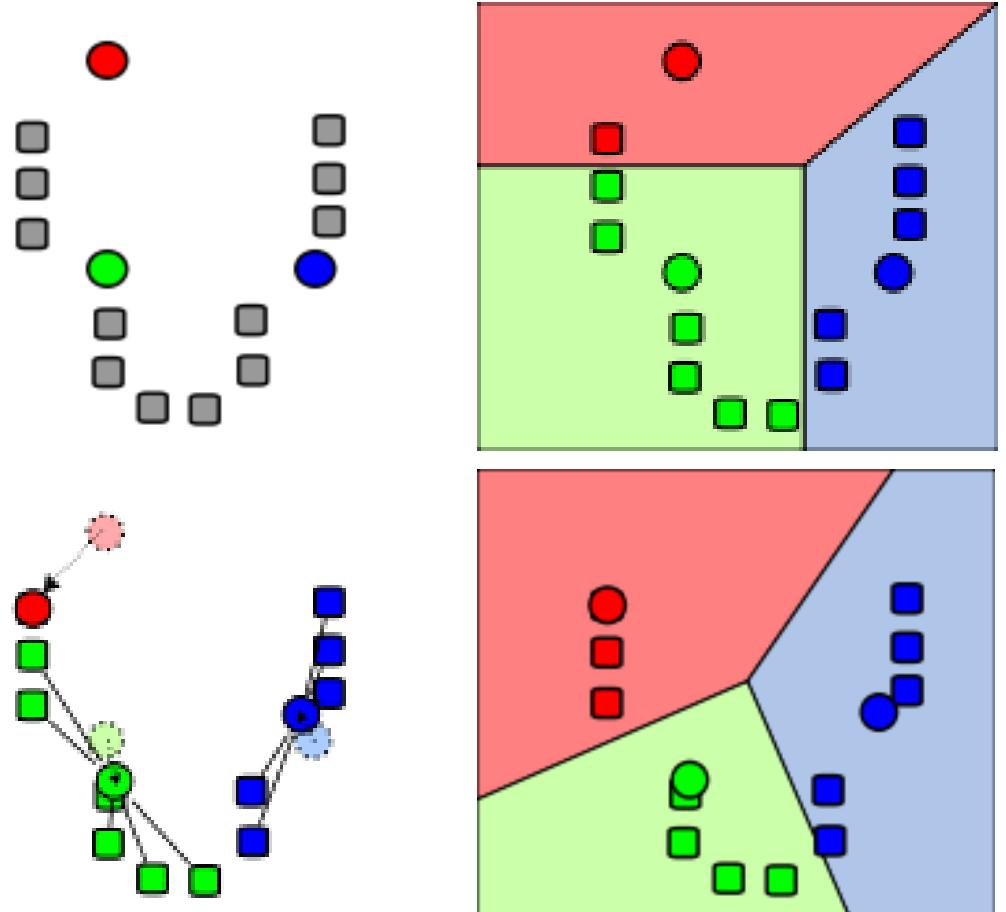
Unsupervised learning  
Specify k (# of clusters)



Source: Wikipedia

# k-Means Clustering

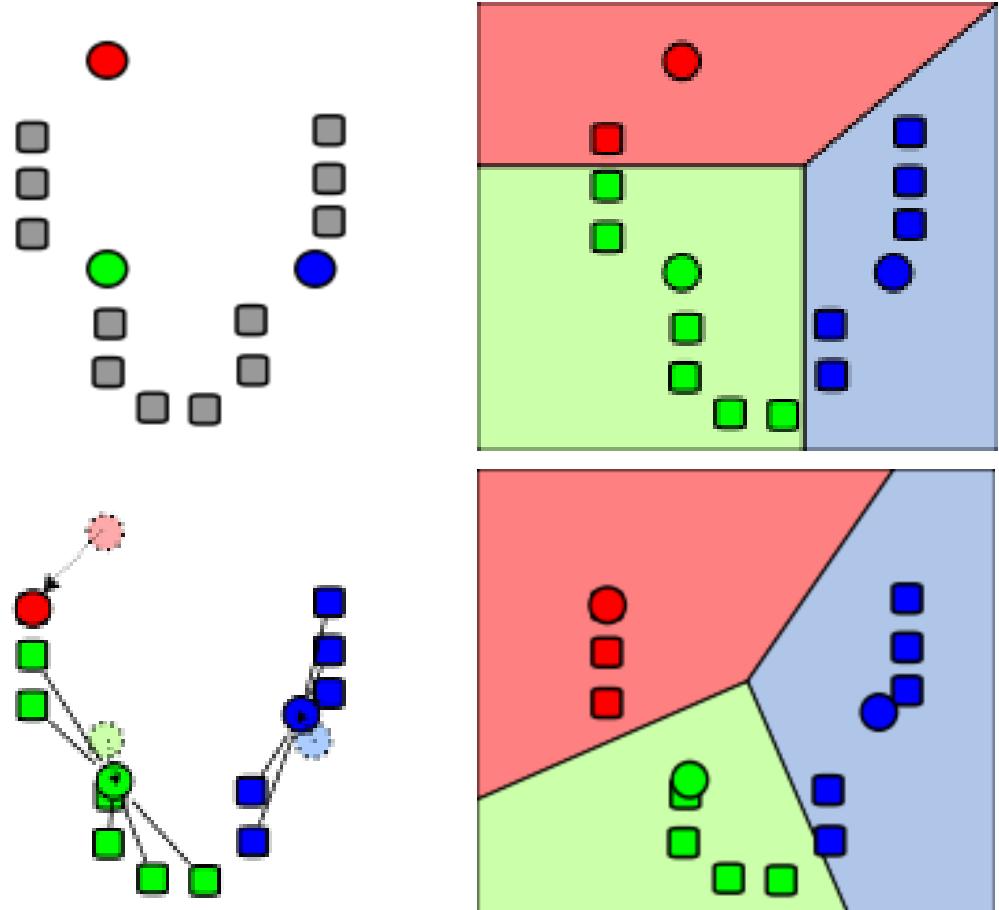
Unsupervised learning  
Specify k (# of clusters)  
Algorithm finds centers



Source: Wikipedia

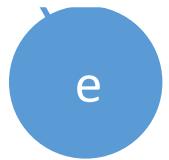
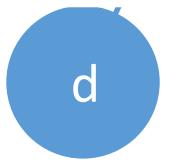
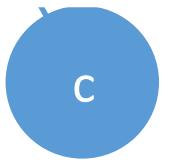
# k-Means Clustering

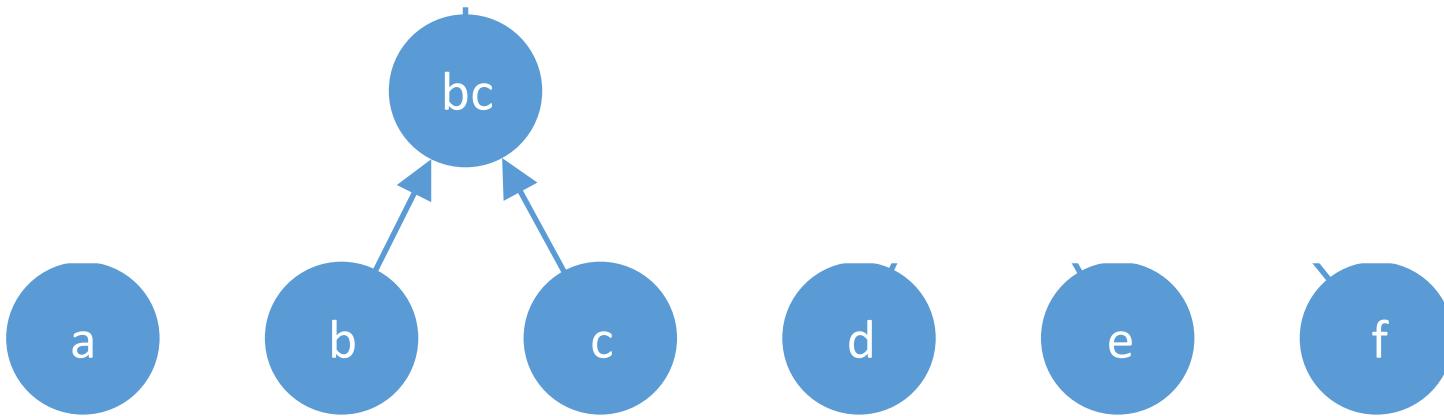
Unsupervised learning  
Specify k (# of clusters)  
Algorithm finds centers  
Random restarts

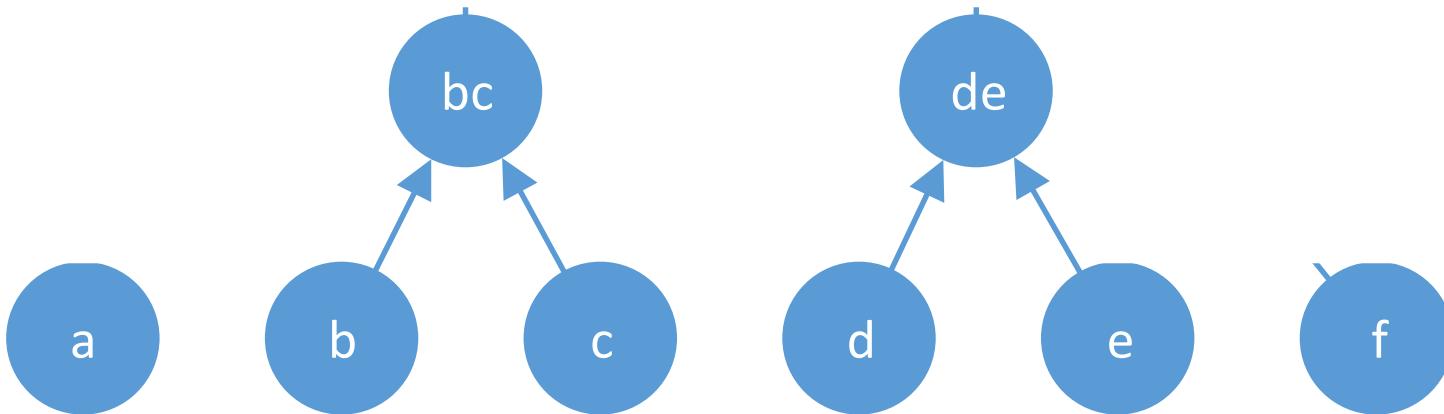


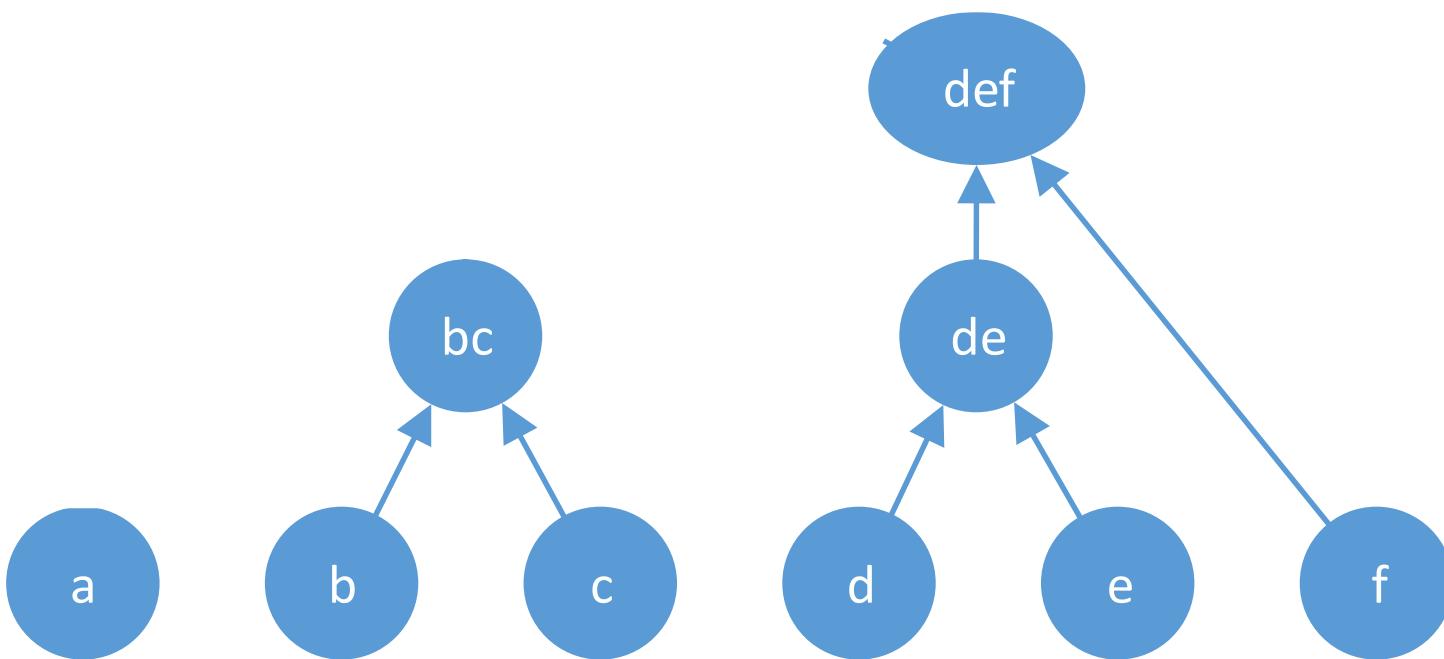
Source: Wikipedia

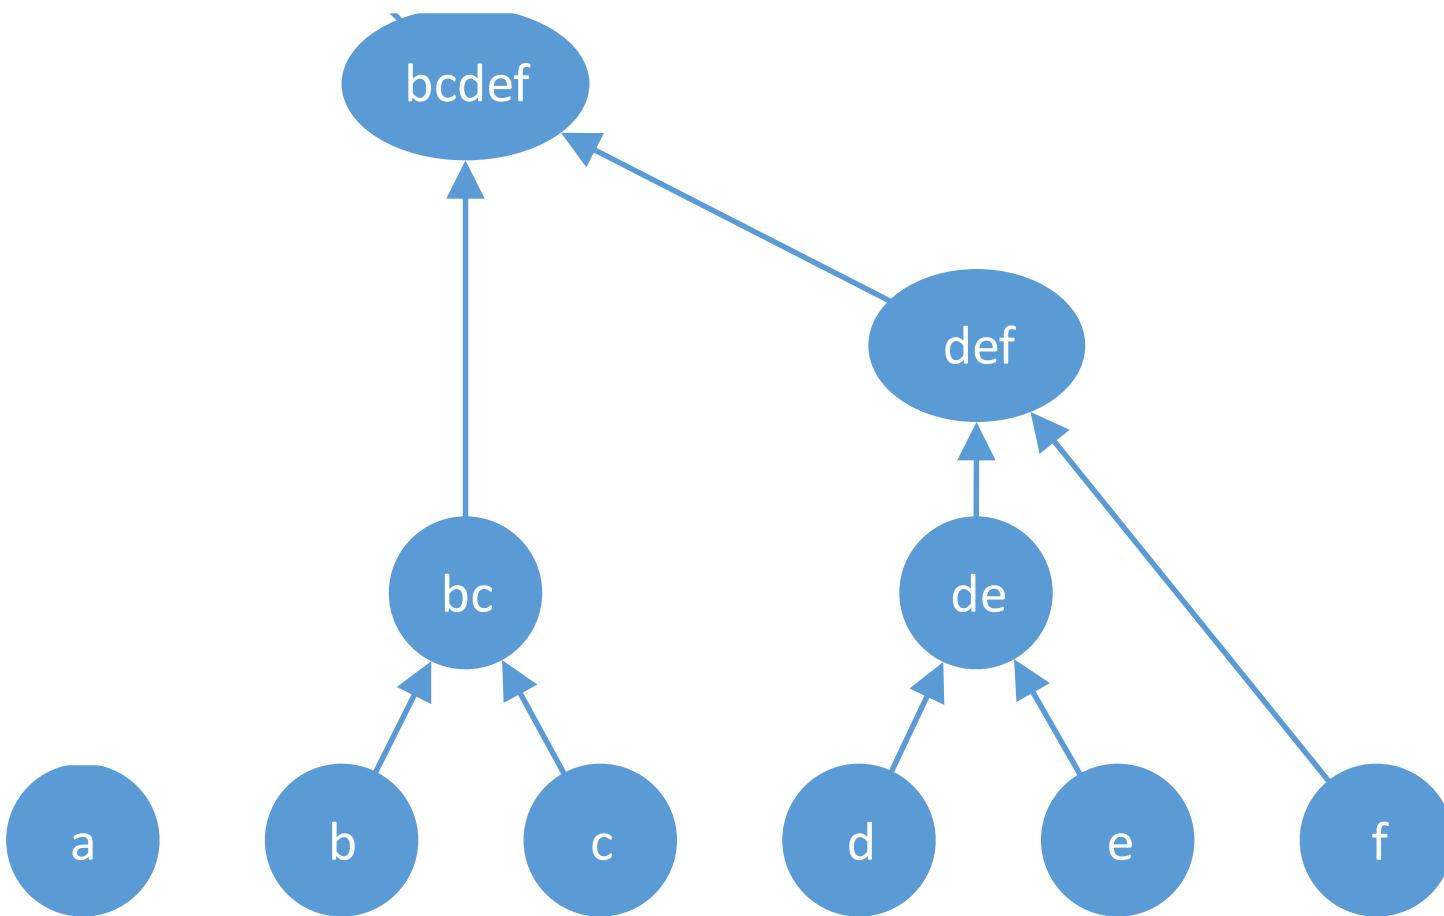
# Hierarchical Clustering

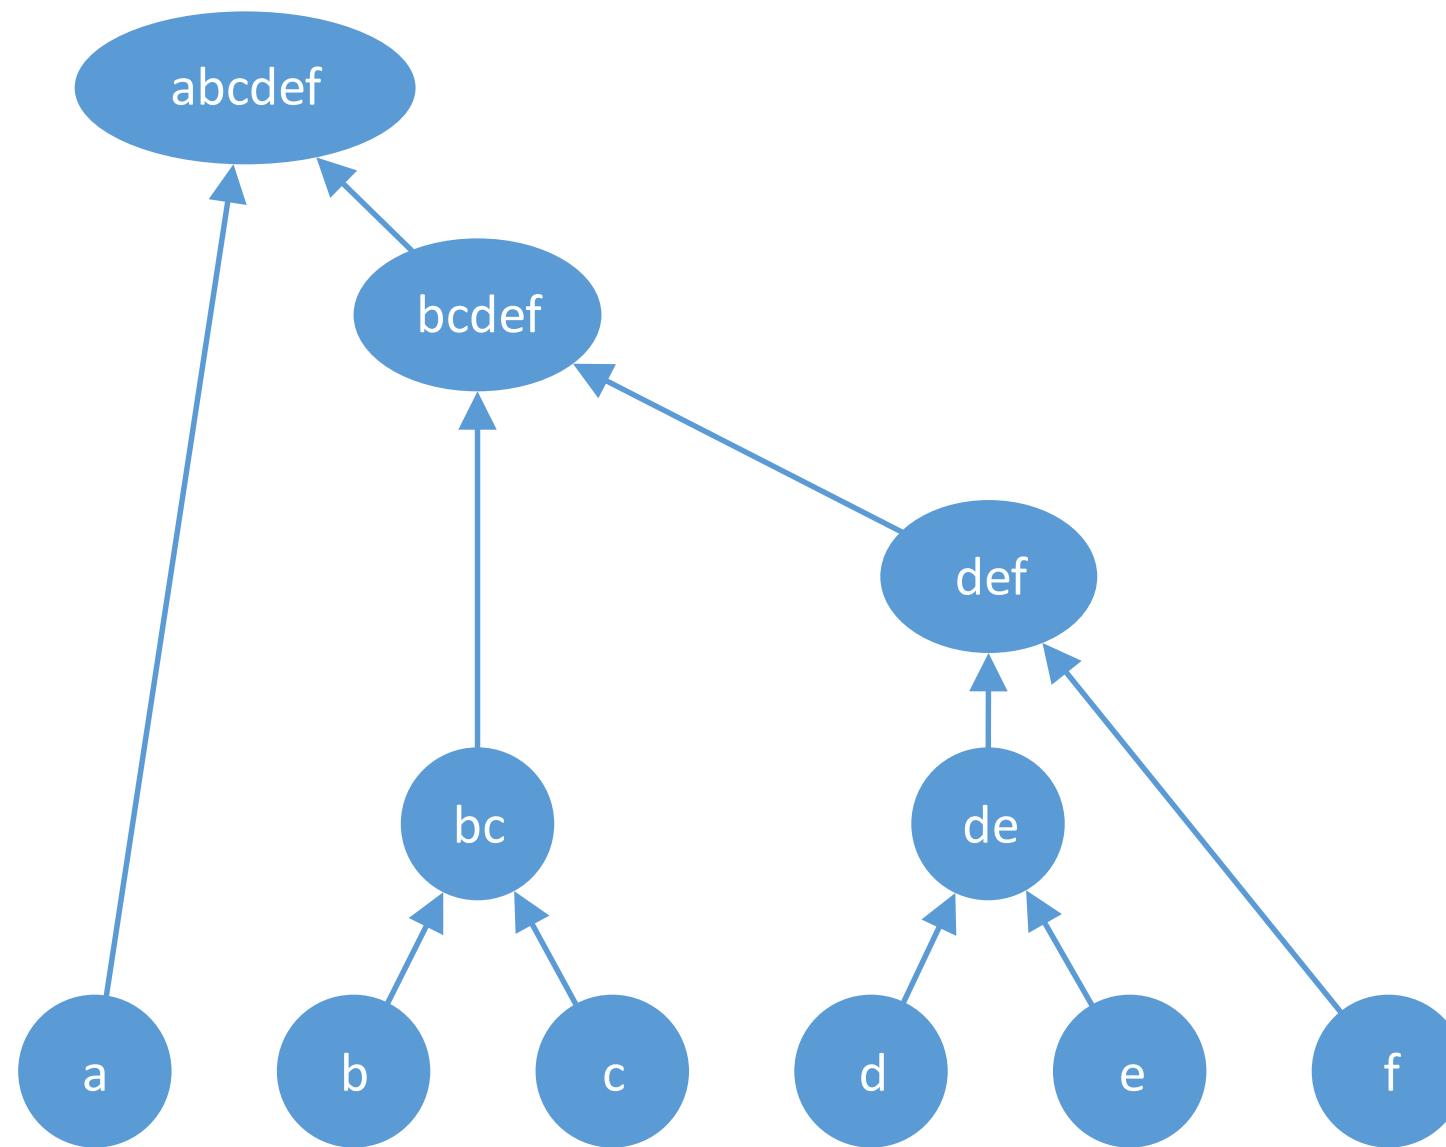


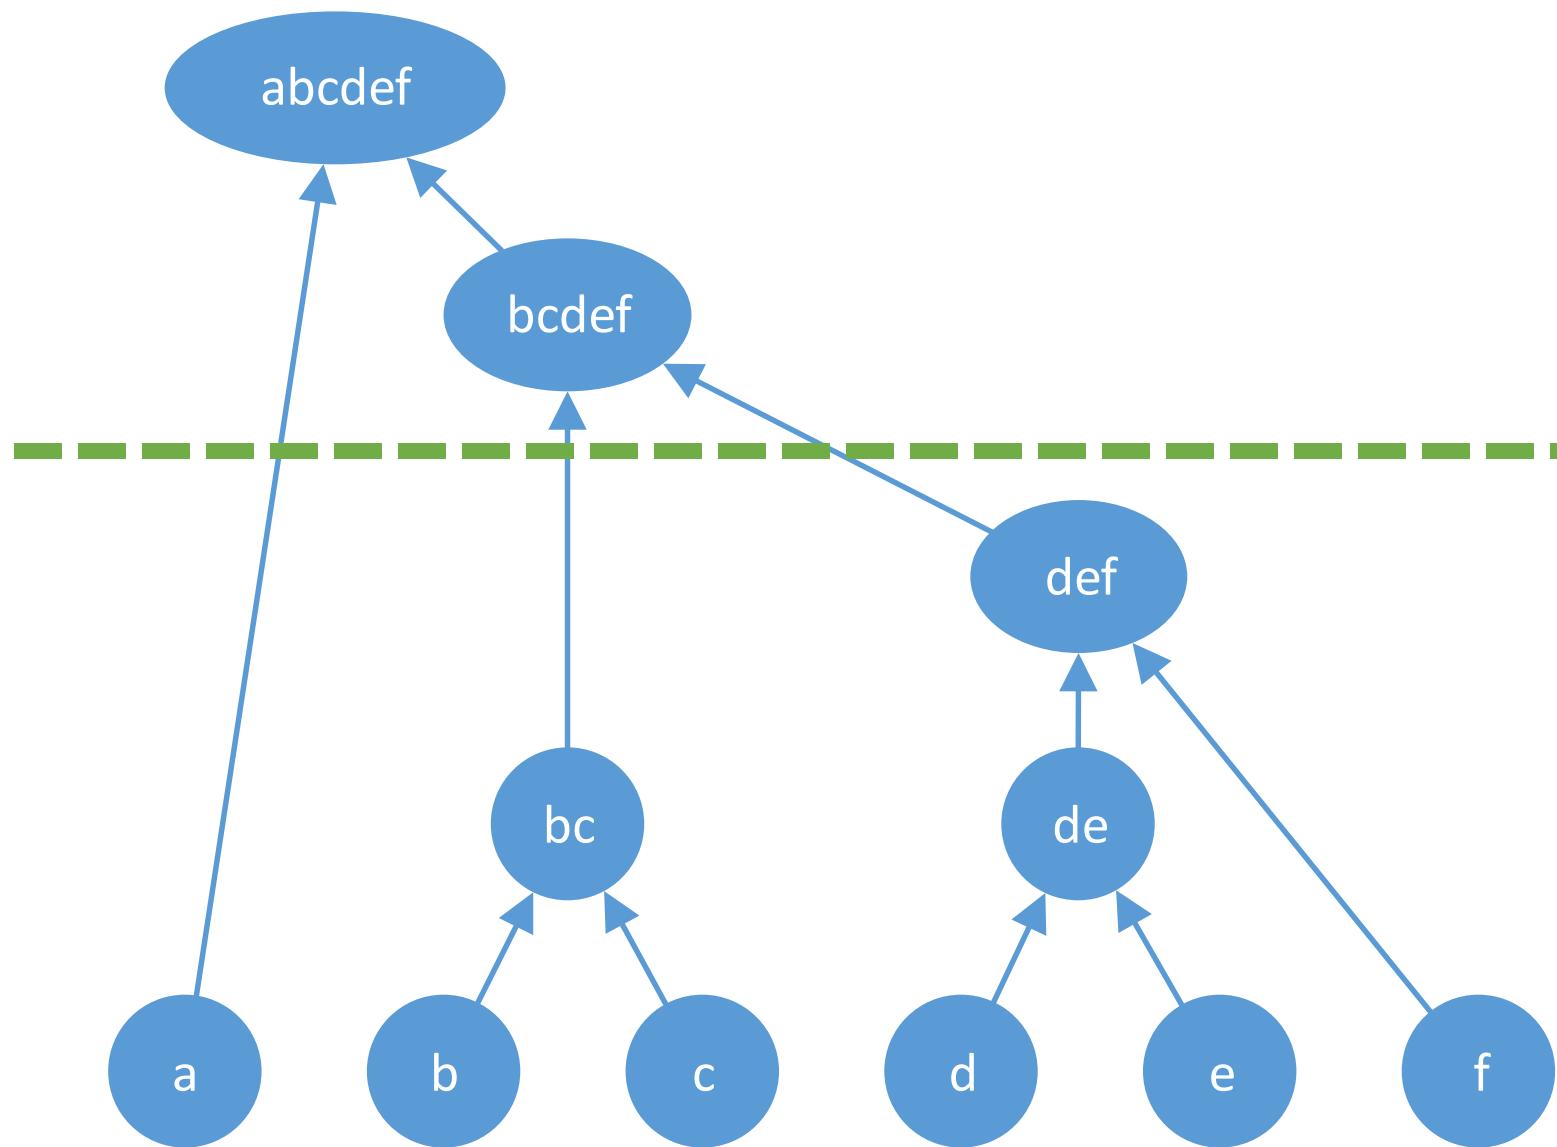






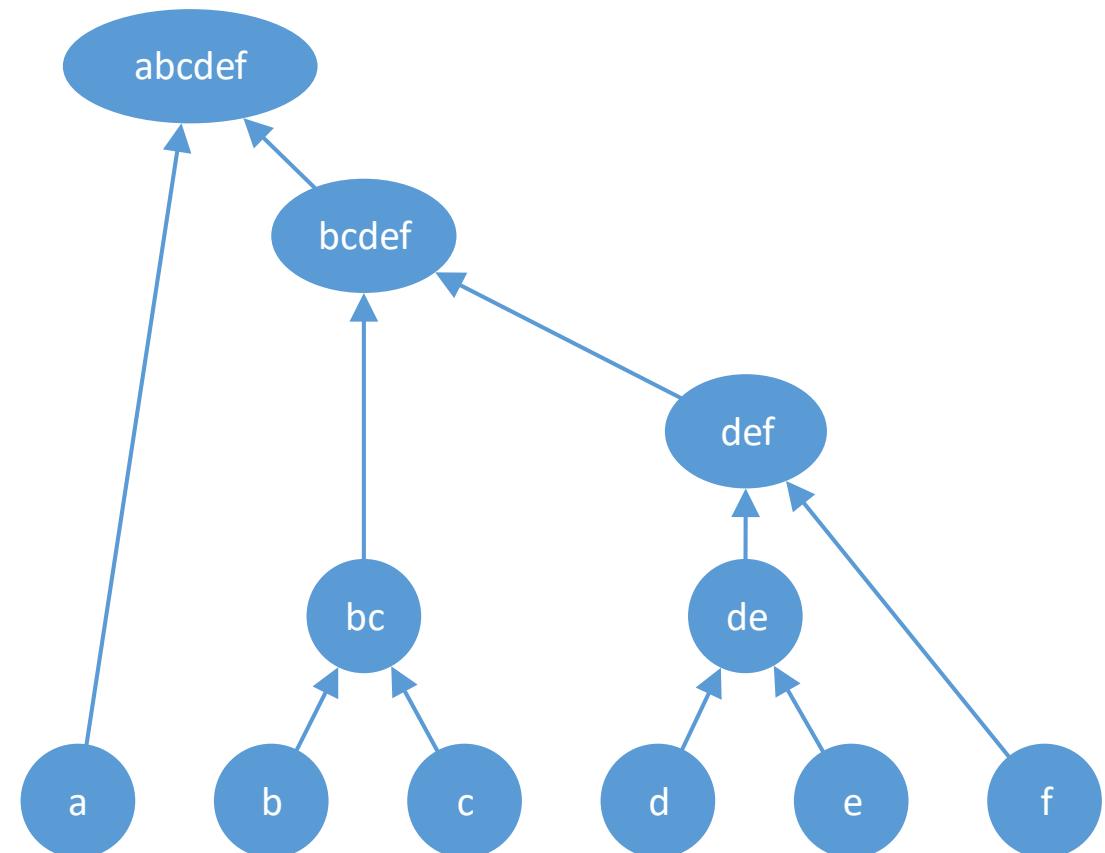






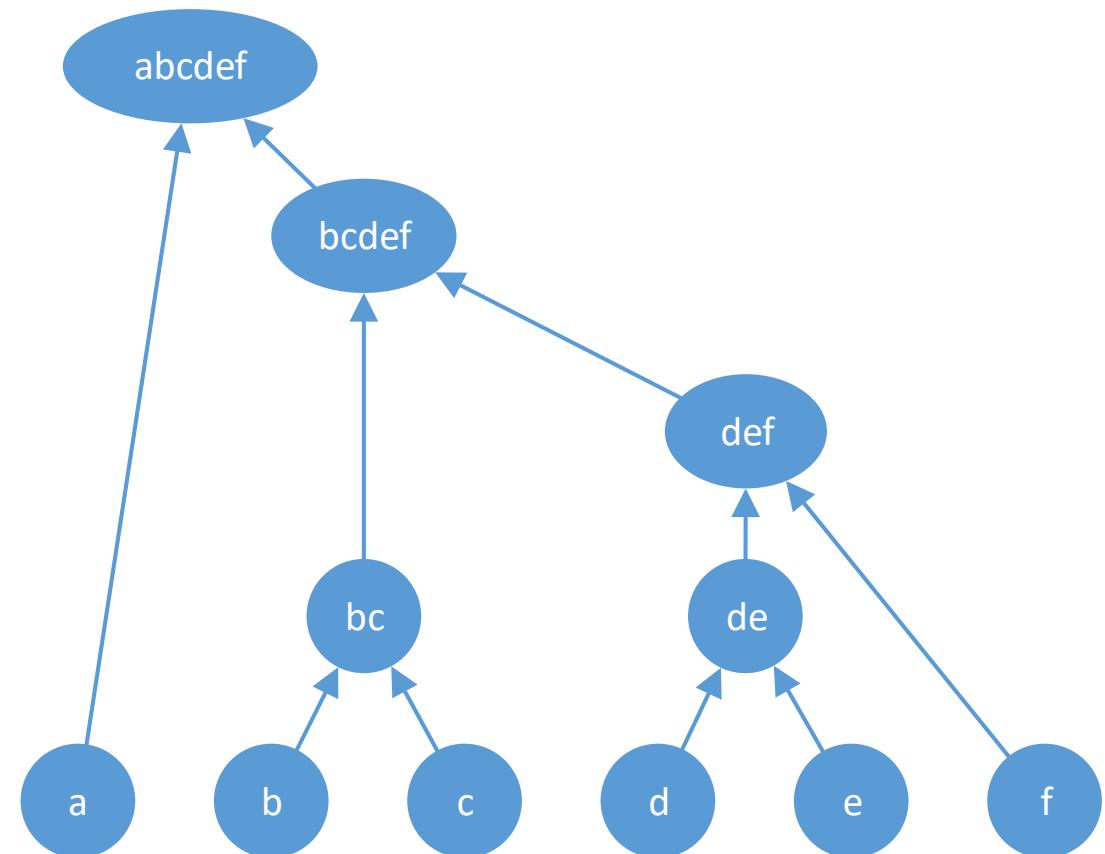
# Hierarchical Clustering

Unsupervised learning



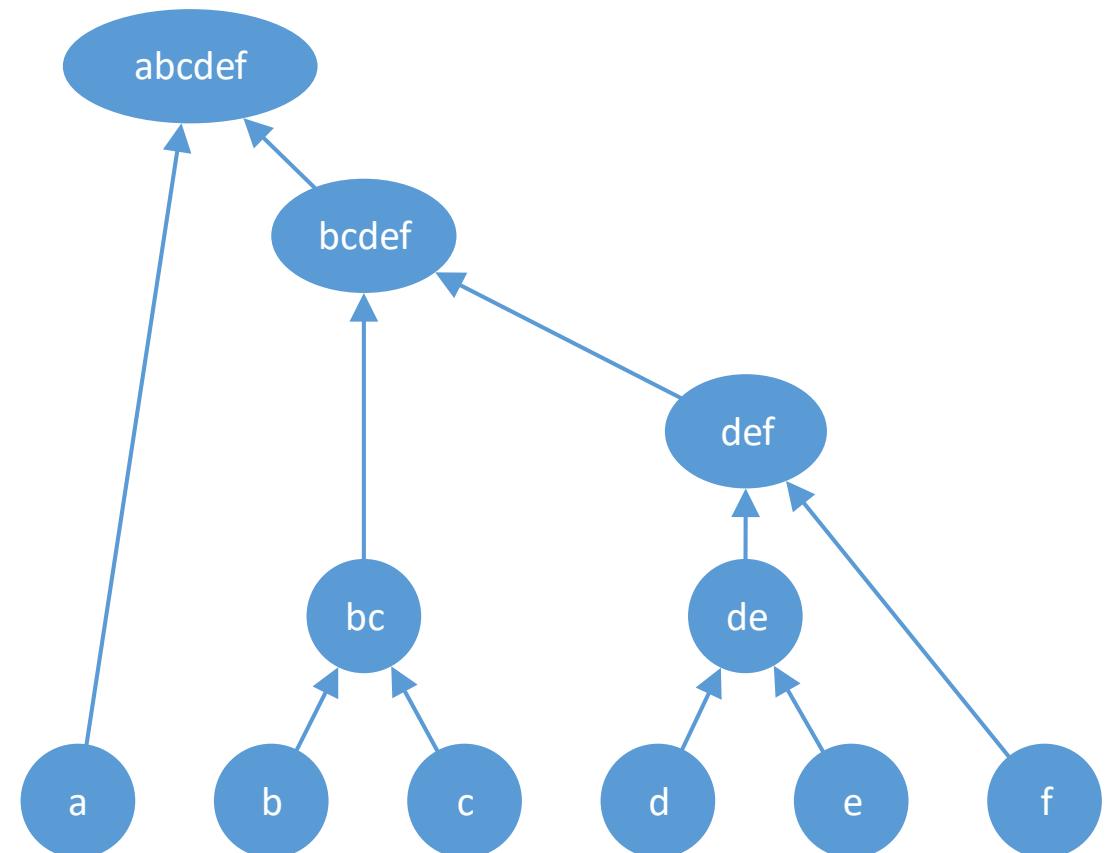
# Hierarchical Clustering

Unsupervised learning  
Tree of connectedness



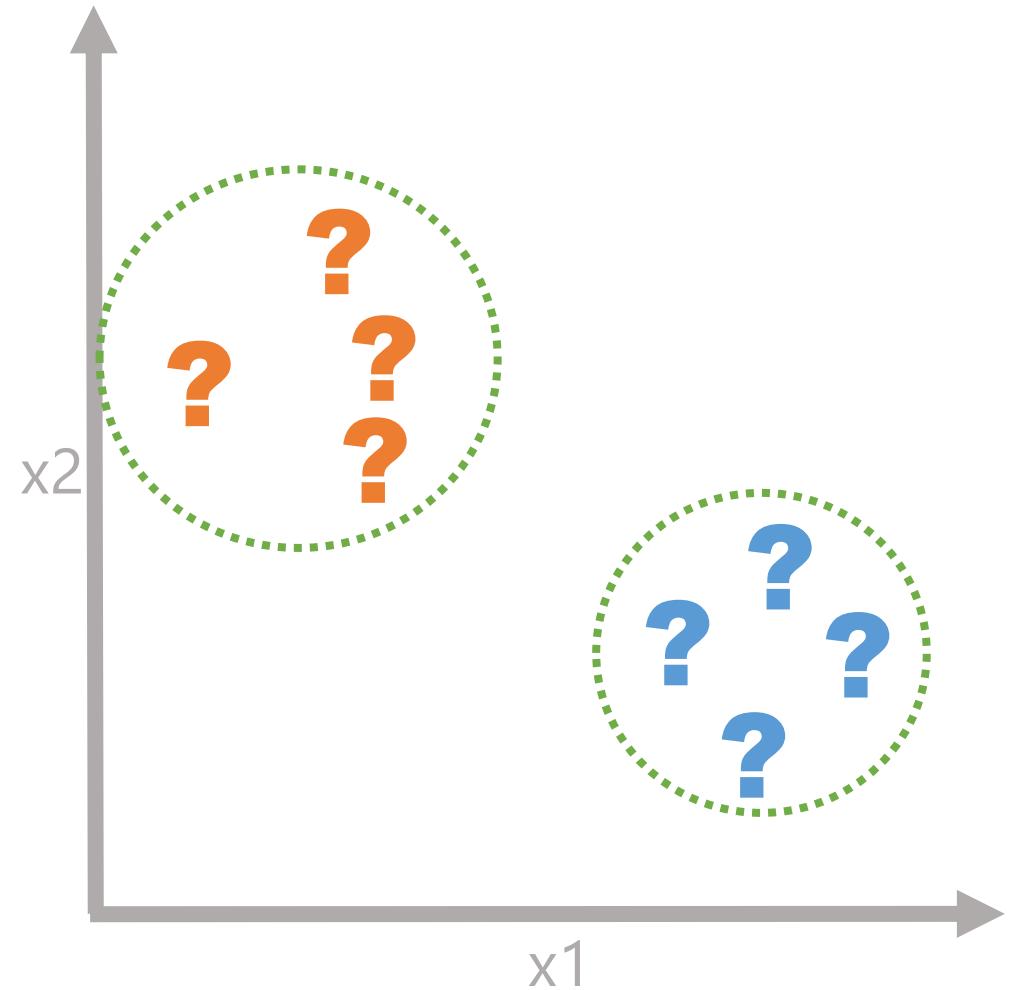
# Hierarchical Clustering

Unsupervised learning  
Tree of connectedness  
Cuts create clusters



# Real-world Examples

What are our market segments?  
How to group our documents?  
Which products to recommend?



# Demo 4 - Clustering

Goal: Group flowers by similarity

# Lab 4A – Clustering (Easy)

Goal: Group flowers by similarity

# Lab 4B – Clustering (Hard)

Goal: Group insurance policies

# Ensemble Learning





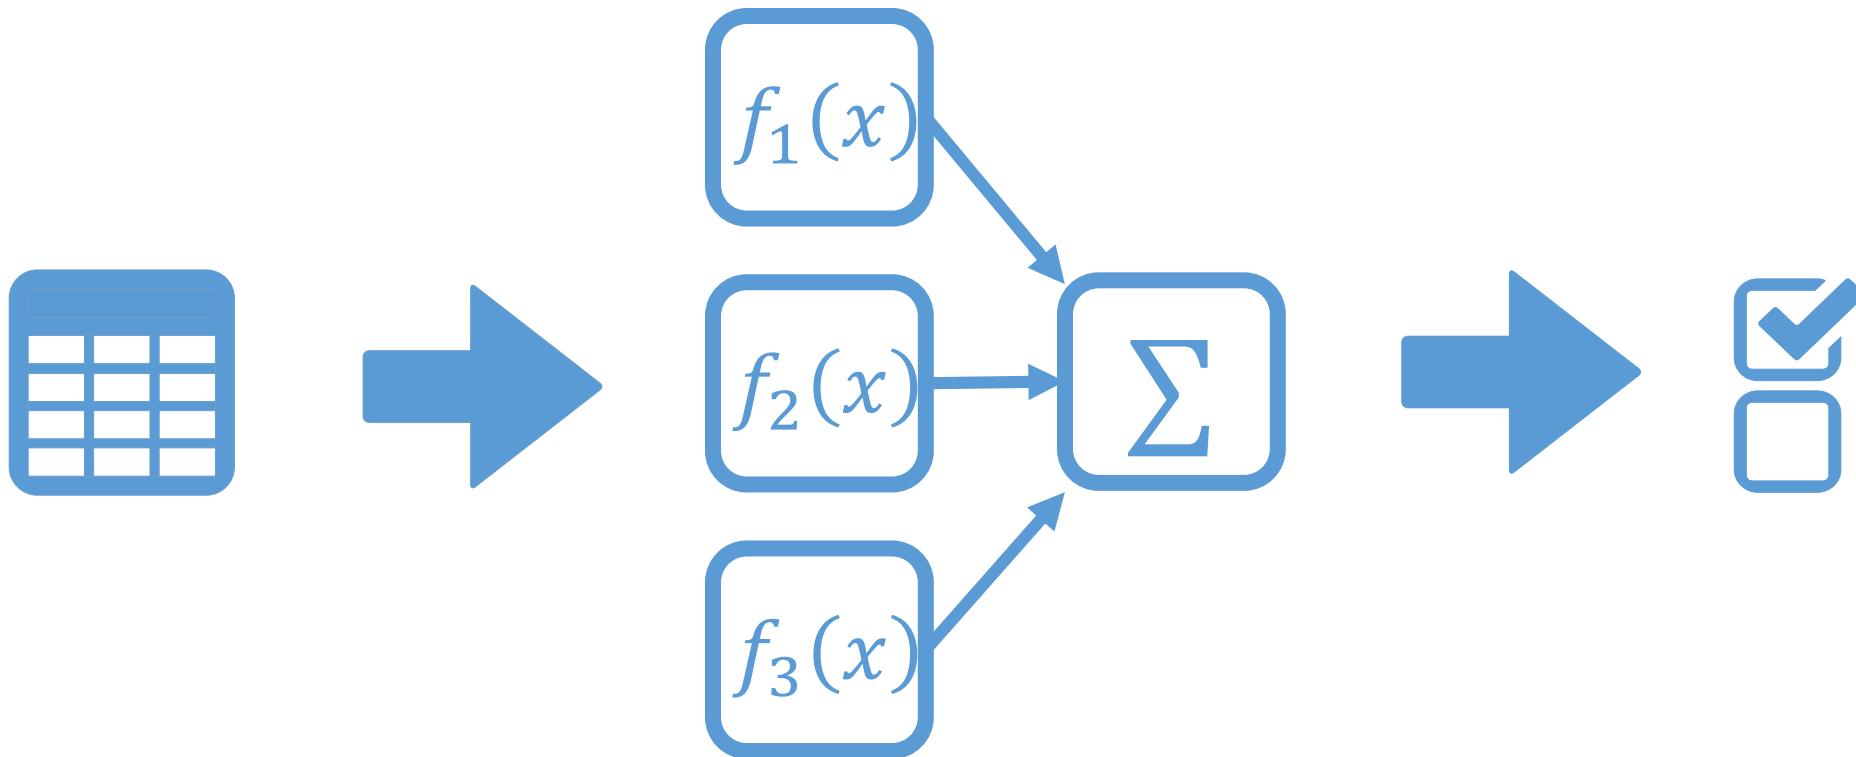
# Wisdom of the Crowds

# HOORAY! OUR SPORTS TEAM WON!

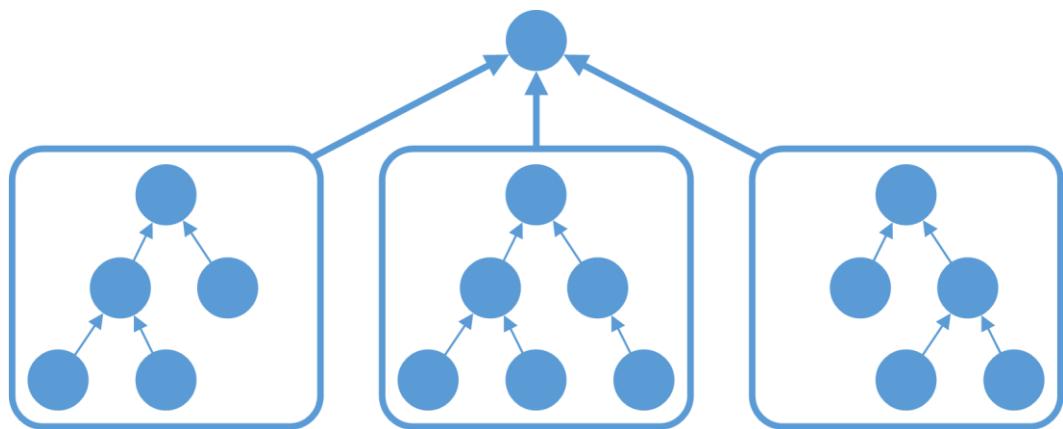


LET'S CELEBRATE  
BY DESTROYING OUR OWN CITY!

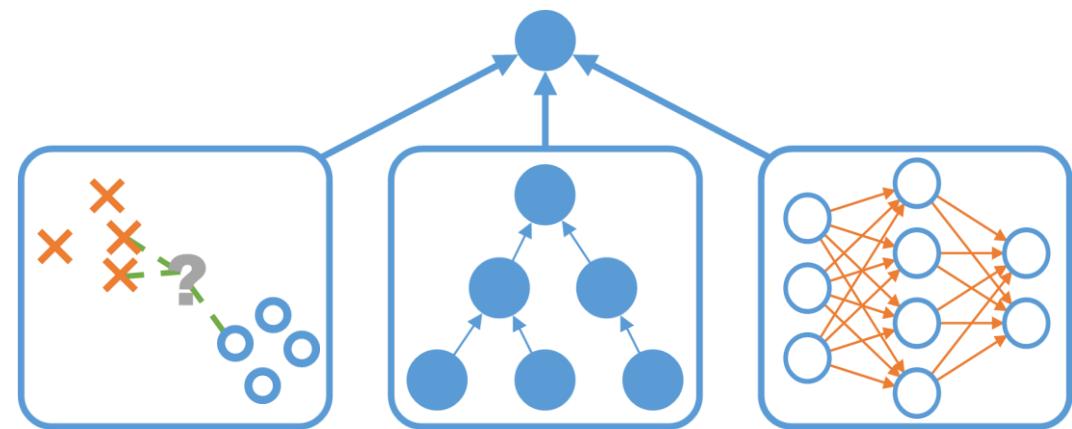
# Ensemble Learning



# Types of Ensembles



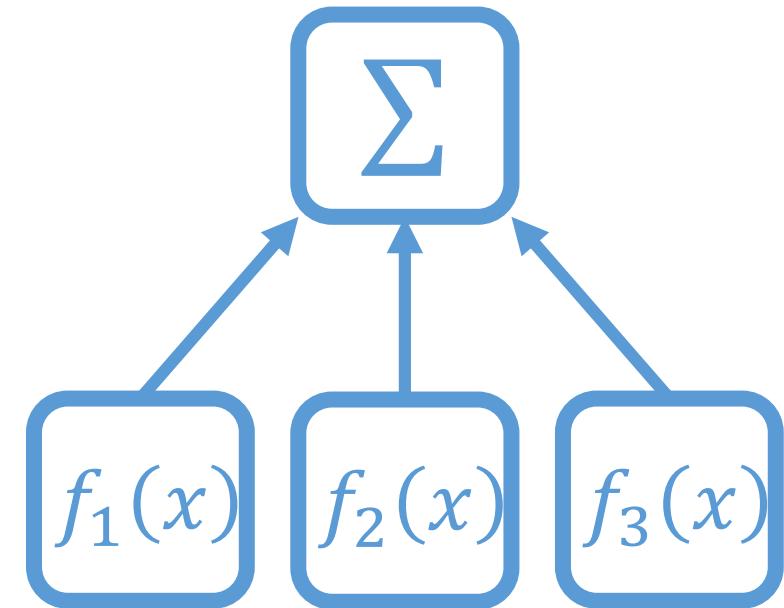
Same Type of Model



Different Types of Models

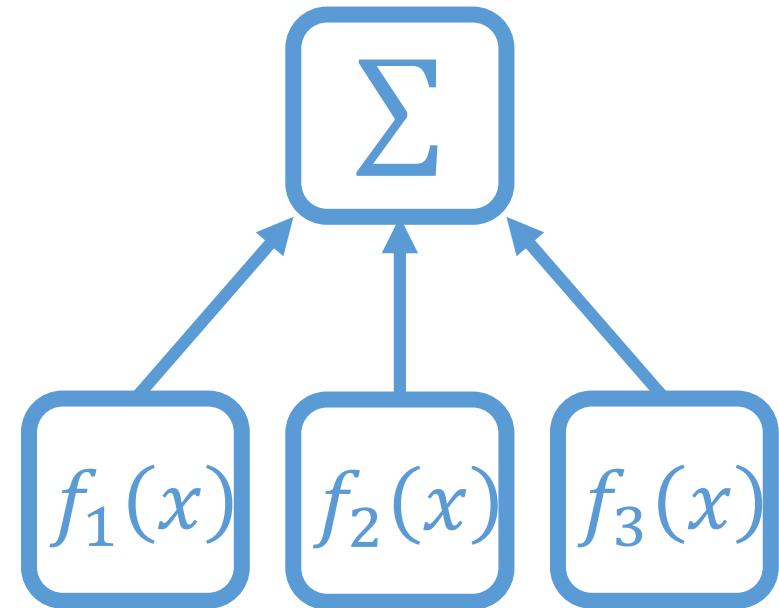
# Ensemble Creation Techniques

Bagging  
Boosting  
Stacking

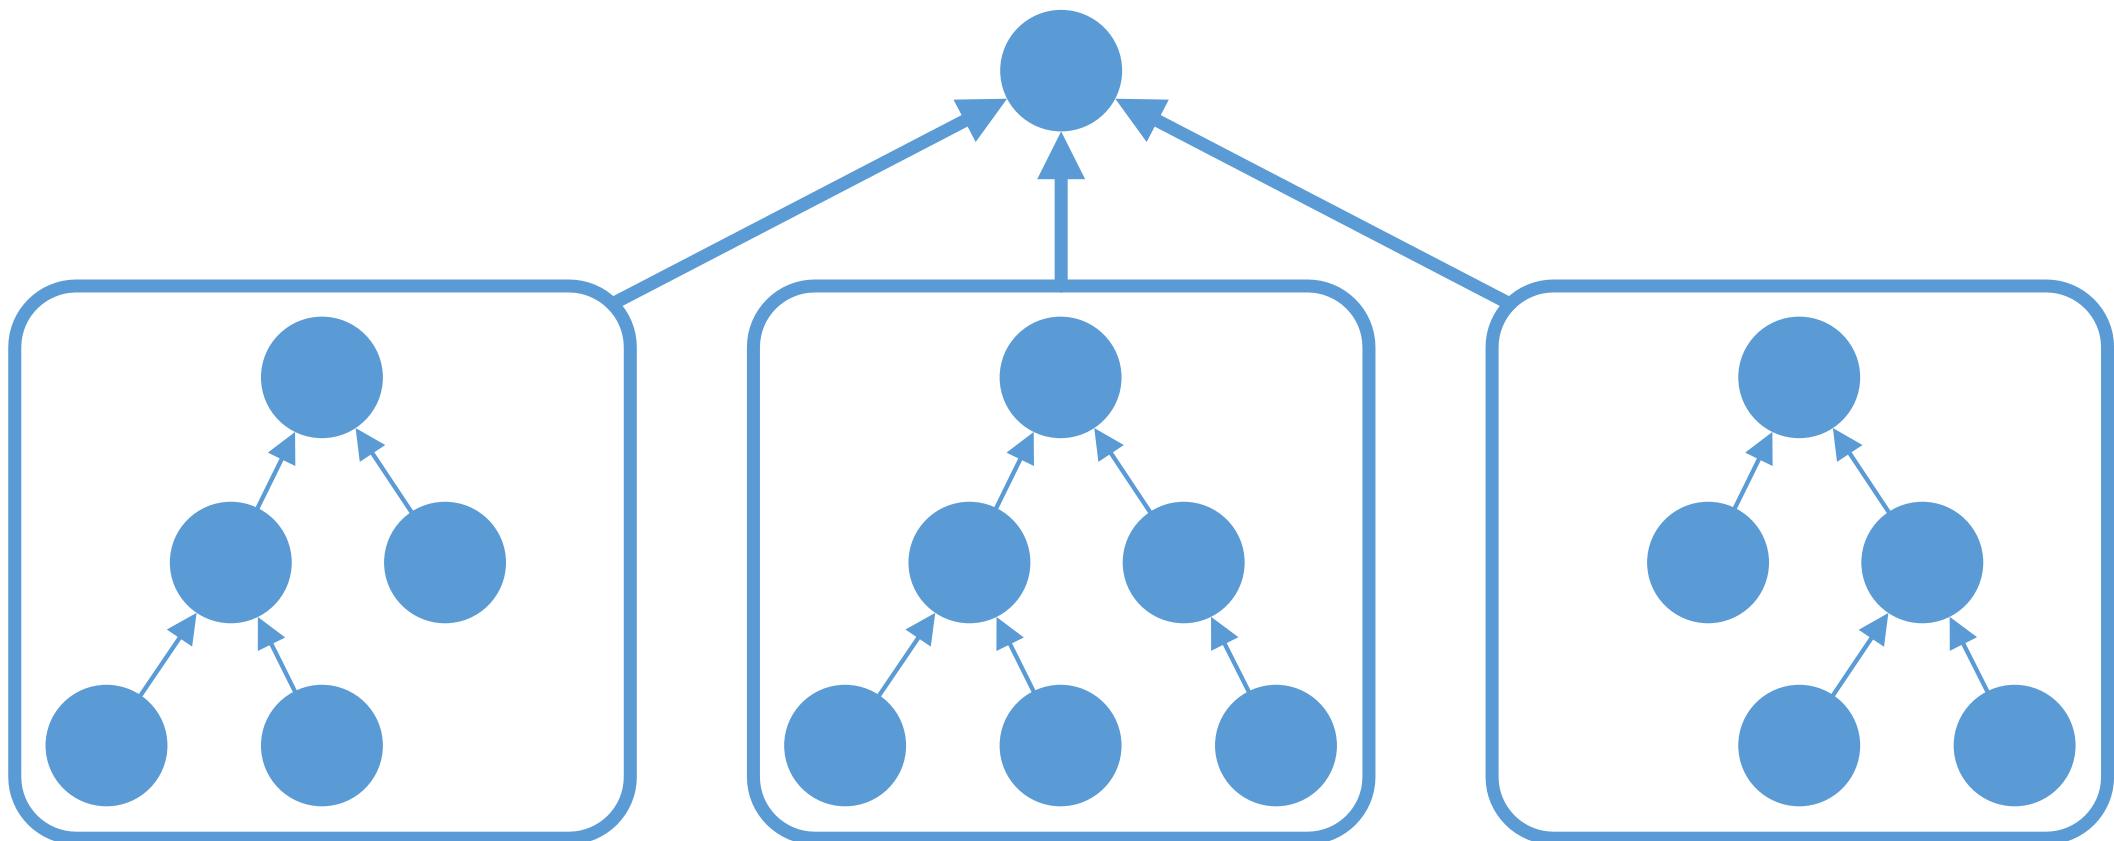


# Ensemble Aggregation Techniques

- Averaging
- Majority Vote
- Weighted Average
- Weighted Majority Vote

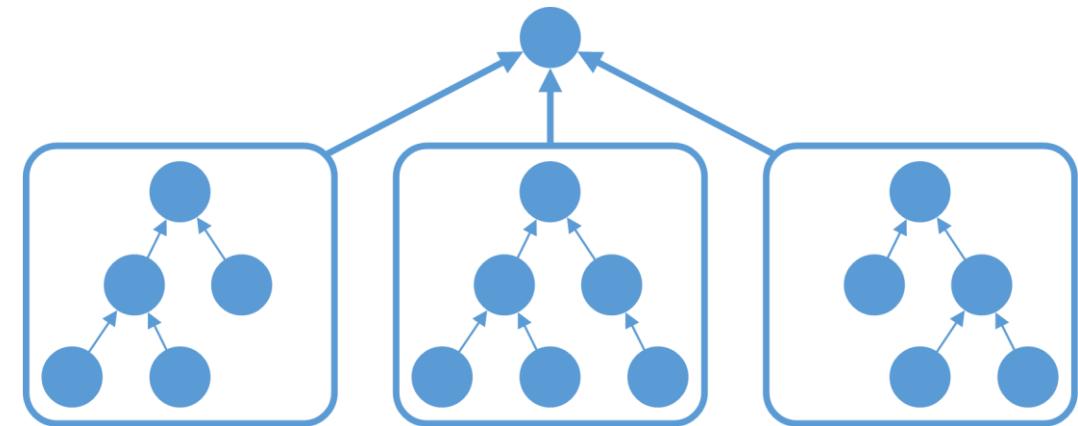


# Random Forest Classifier



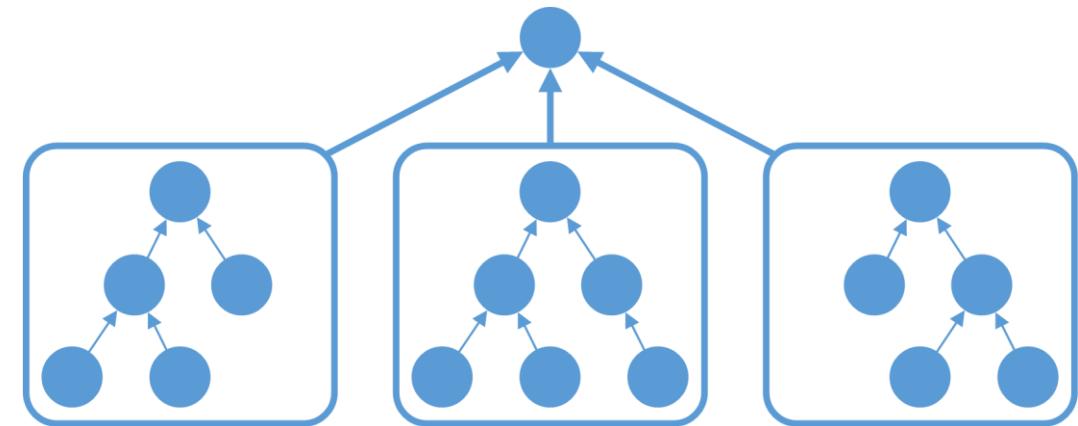
# Random Forest Classifier

Multiple trees



# Random Forest Classifier

Multiple trees  
Created by bagging

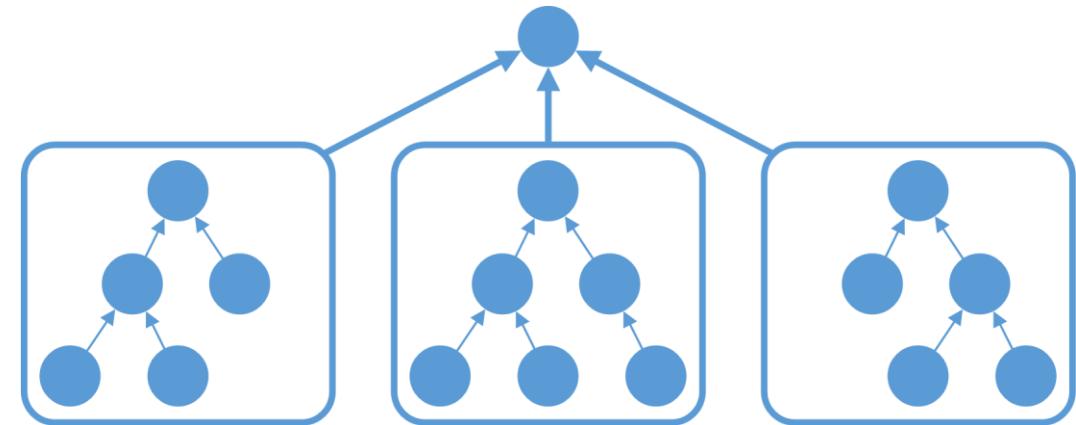


# Random Forest Classifier

Multiple trees

Created by bagging

Majority vote



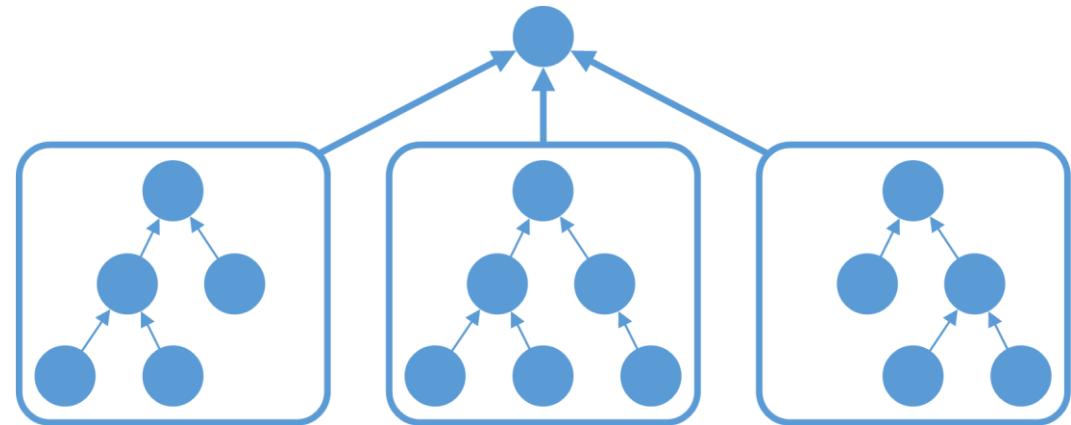
# Random Forest Classifier

Multiple trees

Created by bagging

Majority vote

More robust



# Why Use Ensemble Learning?

## **Pros**

More accurate

More robust

More stable

# Why Use Ensemble Learning?

## **Pros**

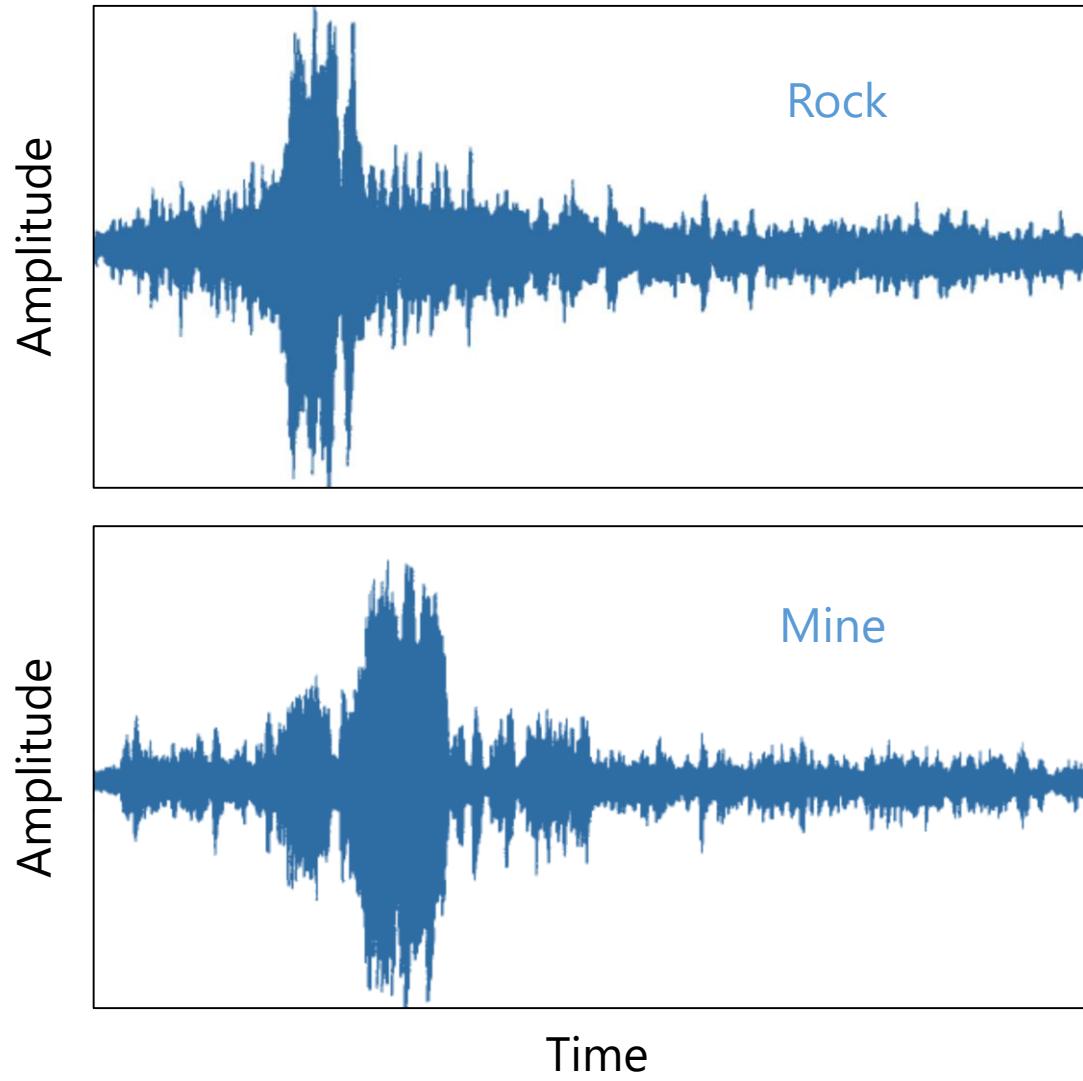
- More accurate
- More robust
- More stable

## **Cons**

- More complex
- More CPU time
- More art than science

# Ensemble Learning Demo





# Sonar



# Demo 5 – ML in Practice

Goal: Predict rock or mine

# Lab 5A – ML in Practice (Easy)

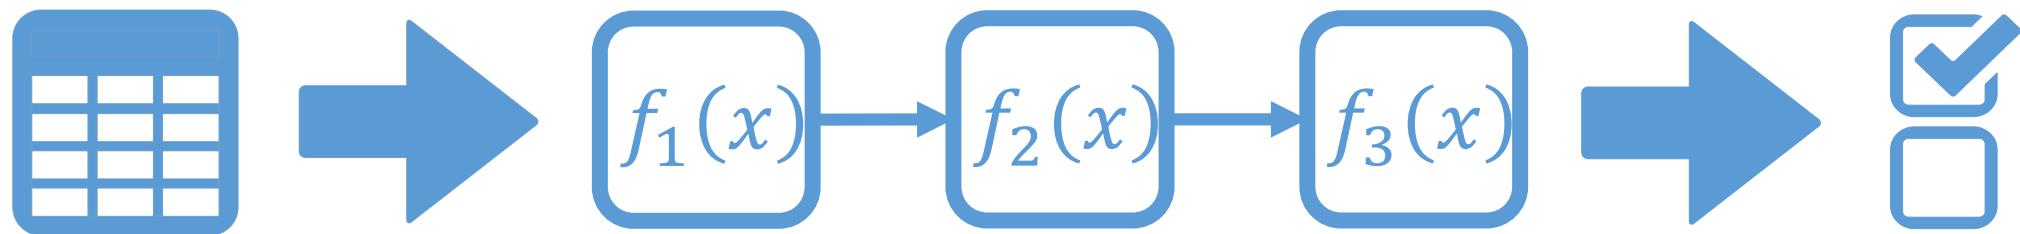
Goal: Predict rock or mine

# Lab 5B – ML in Practice (Hard)

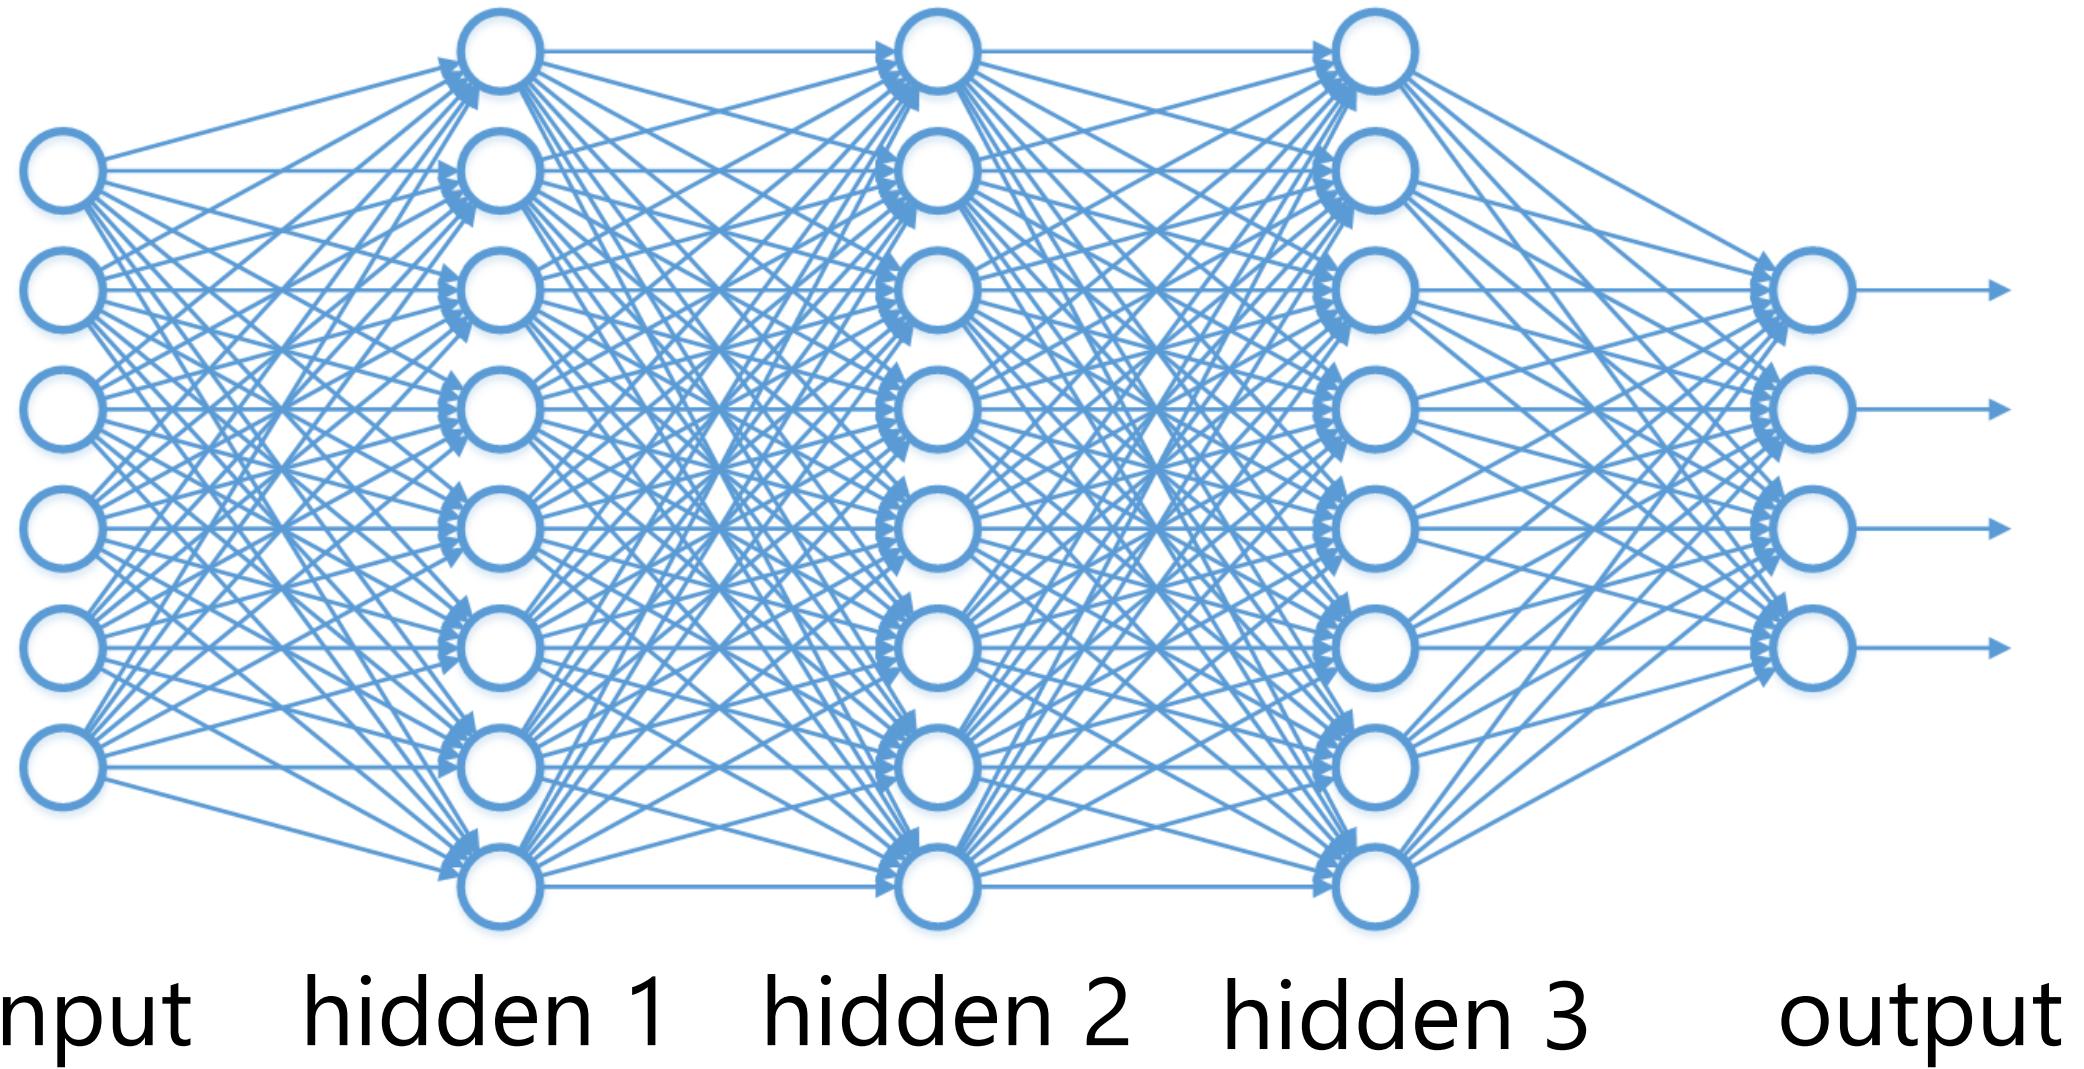
Goal: Predict risk class

# Deep Learning

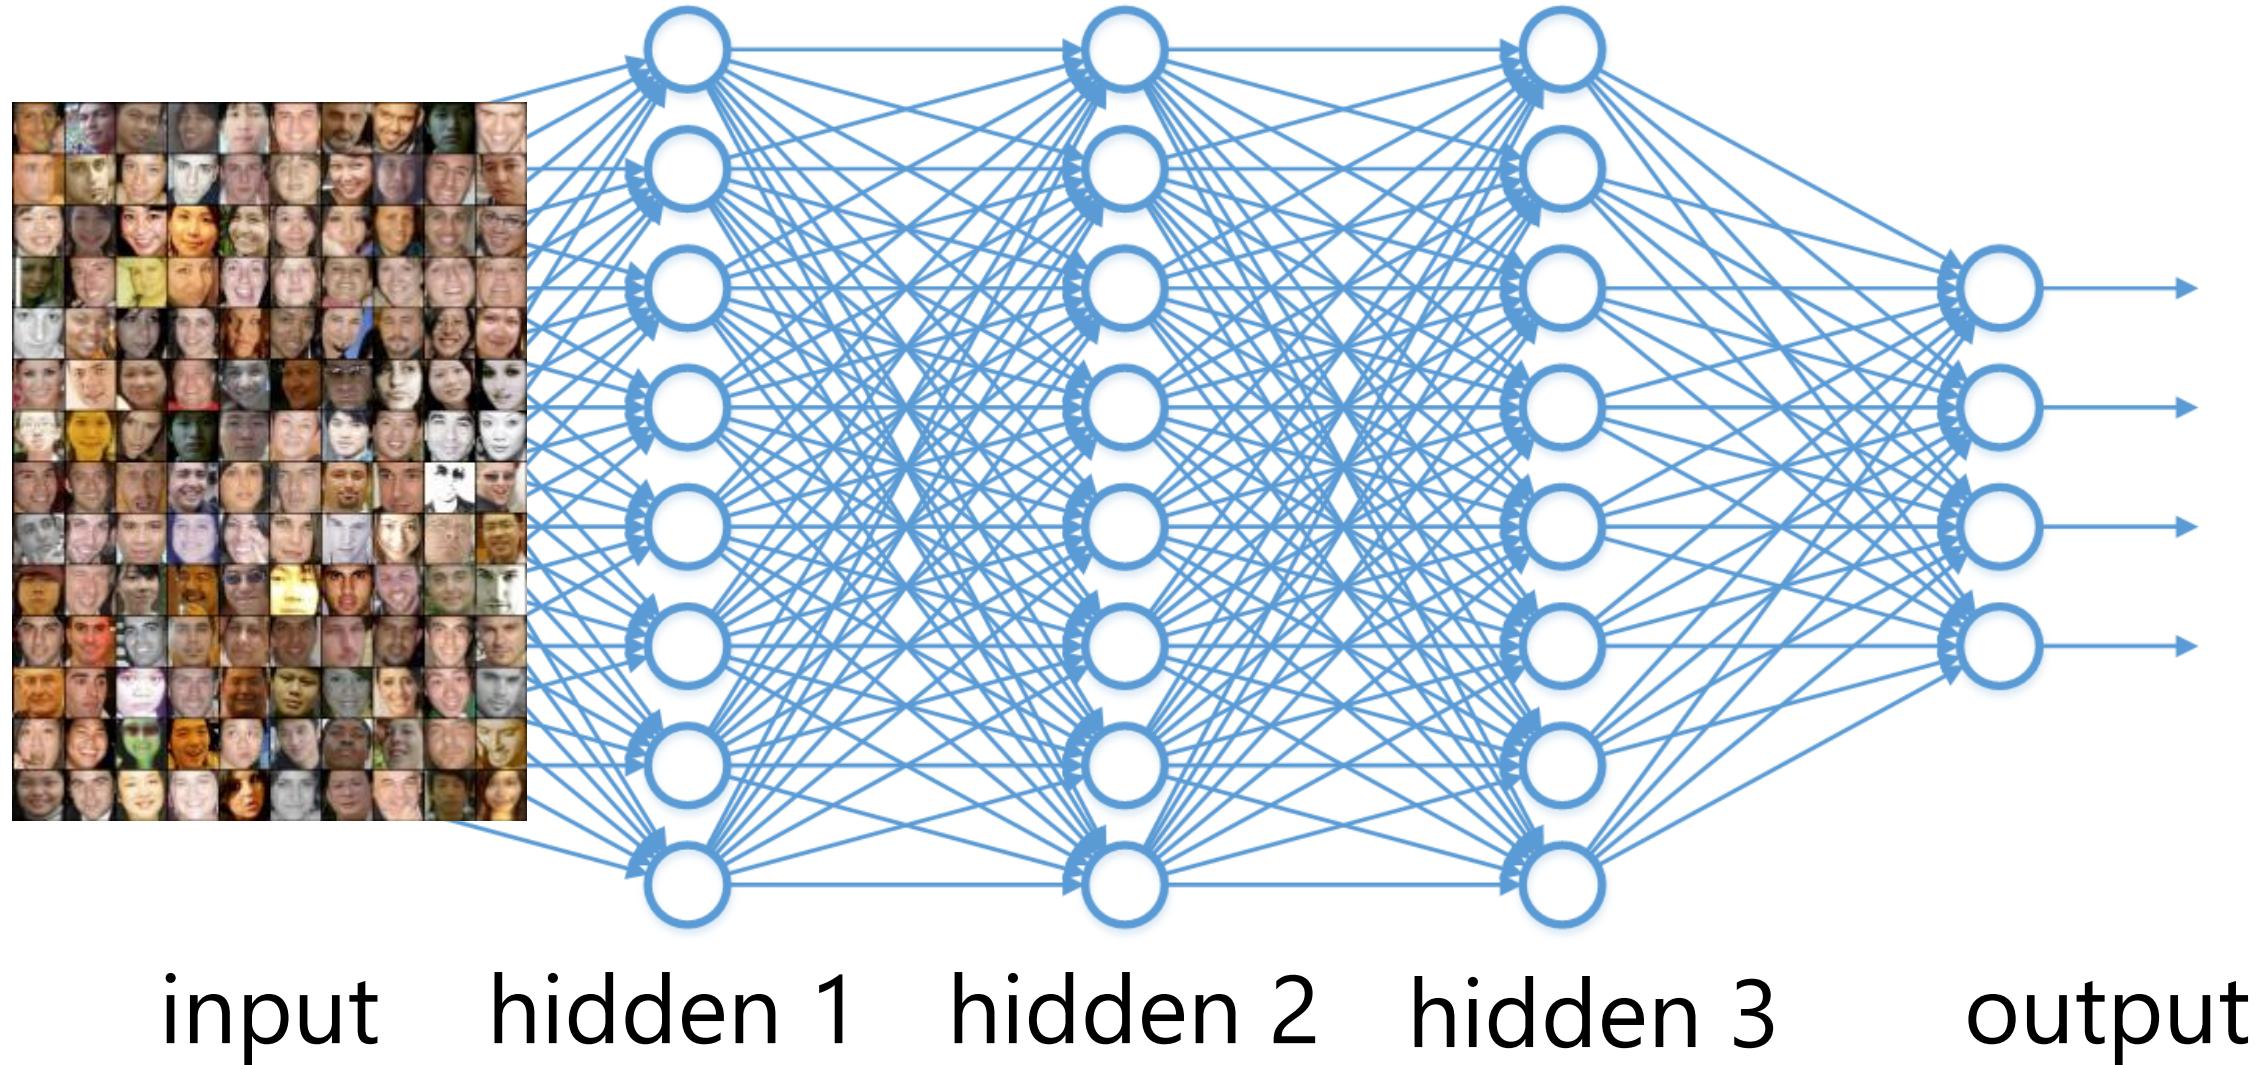
# Deep Learning



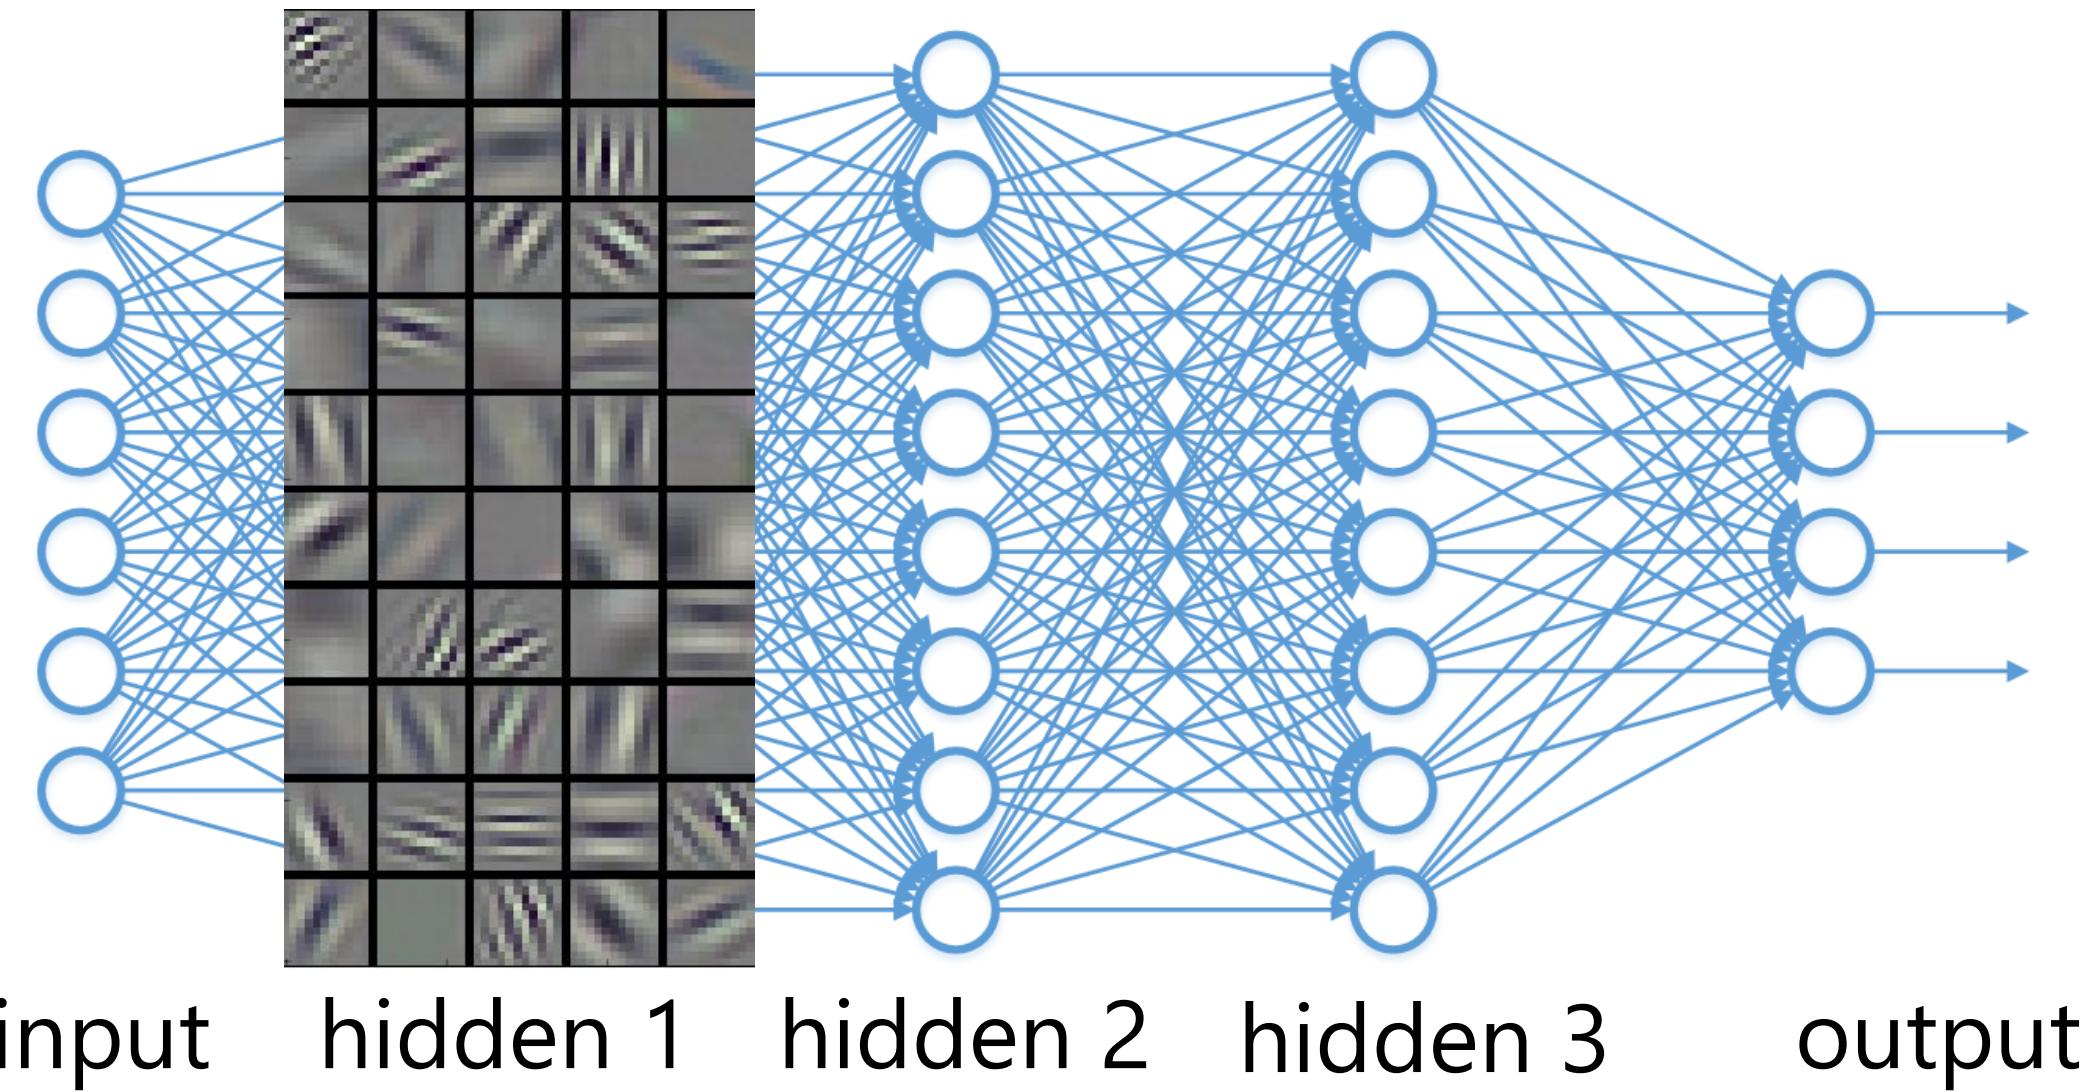
# Deep Neural Network



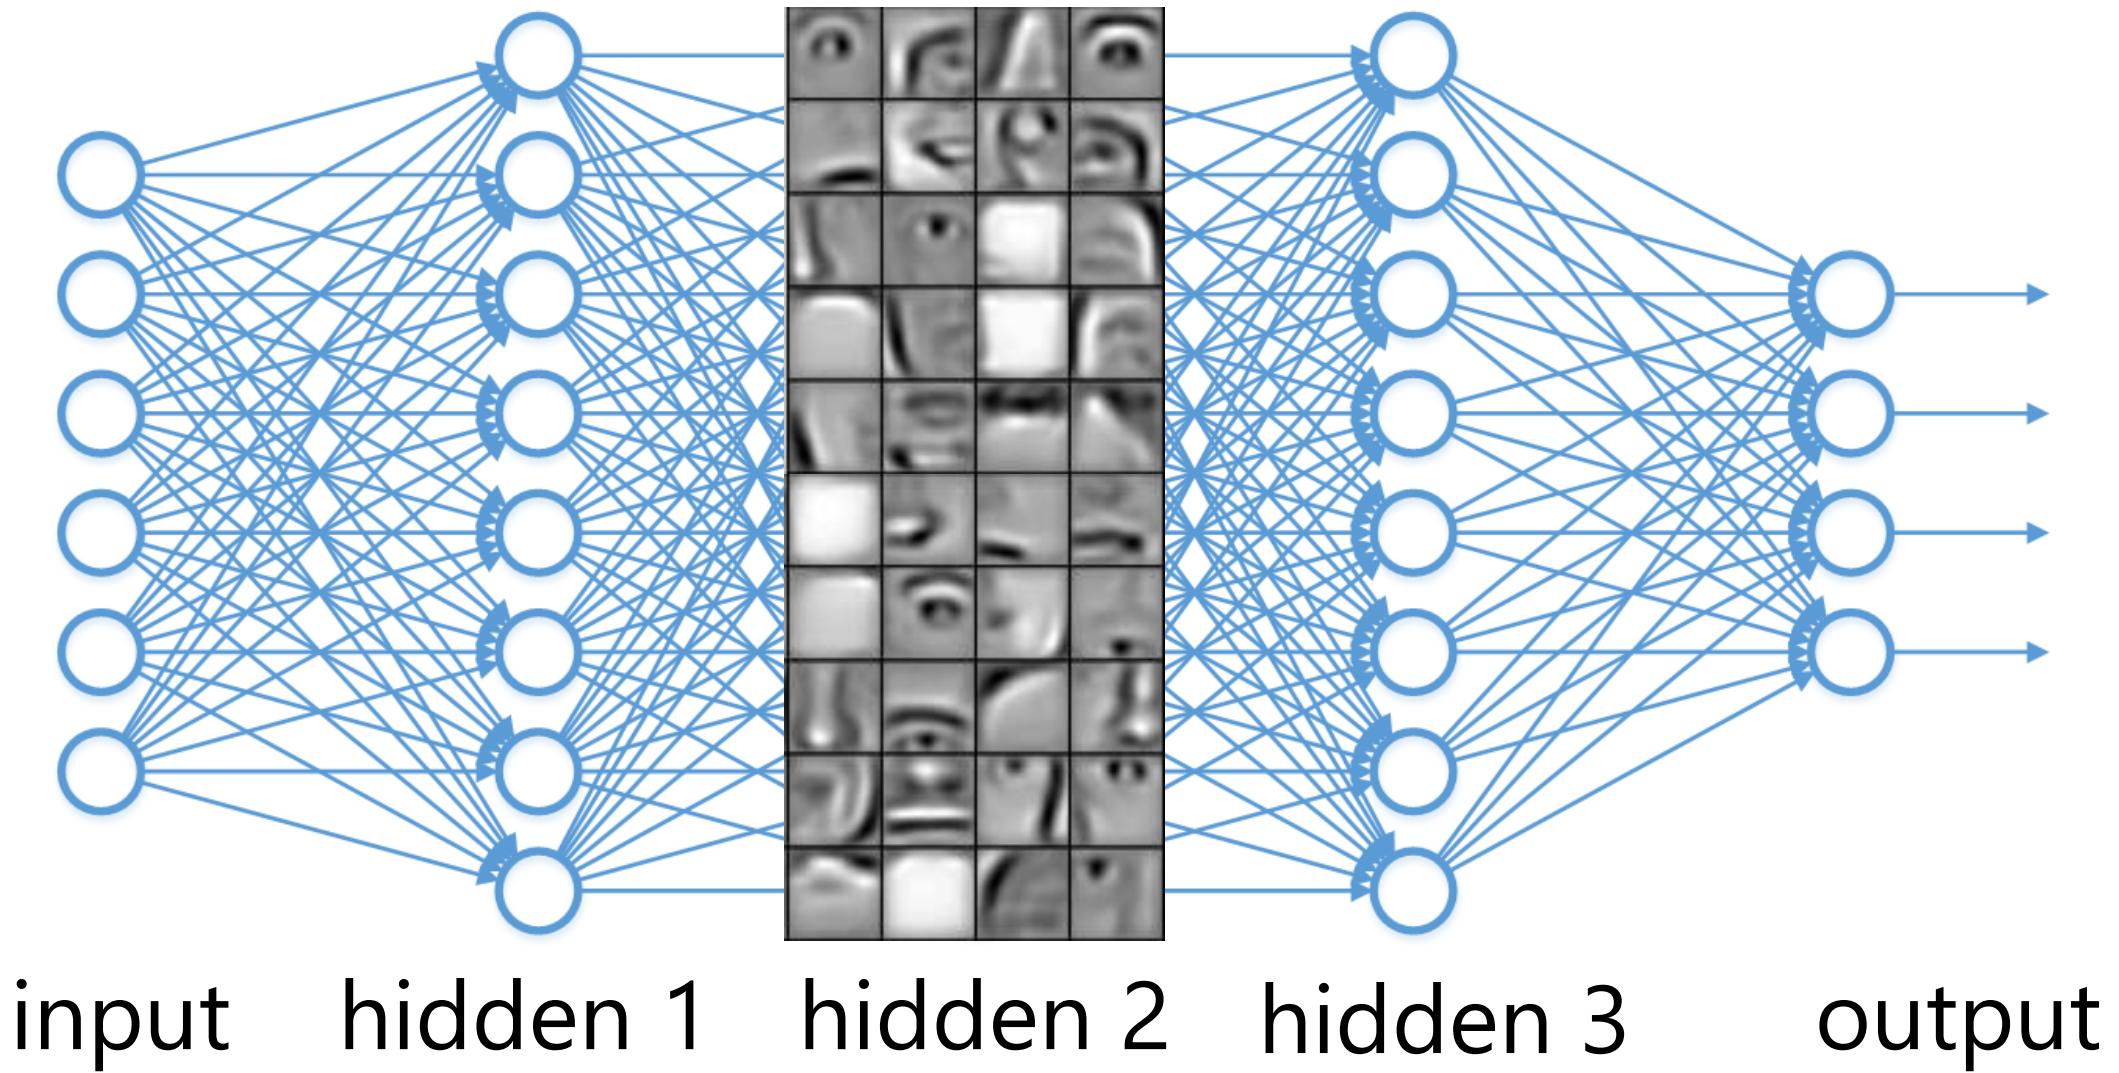
# Deep Neural Network



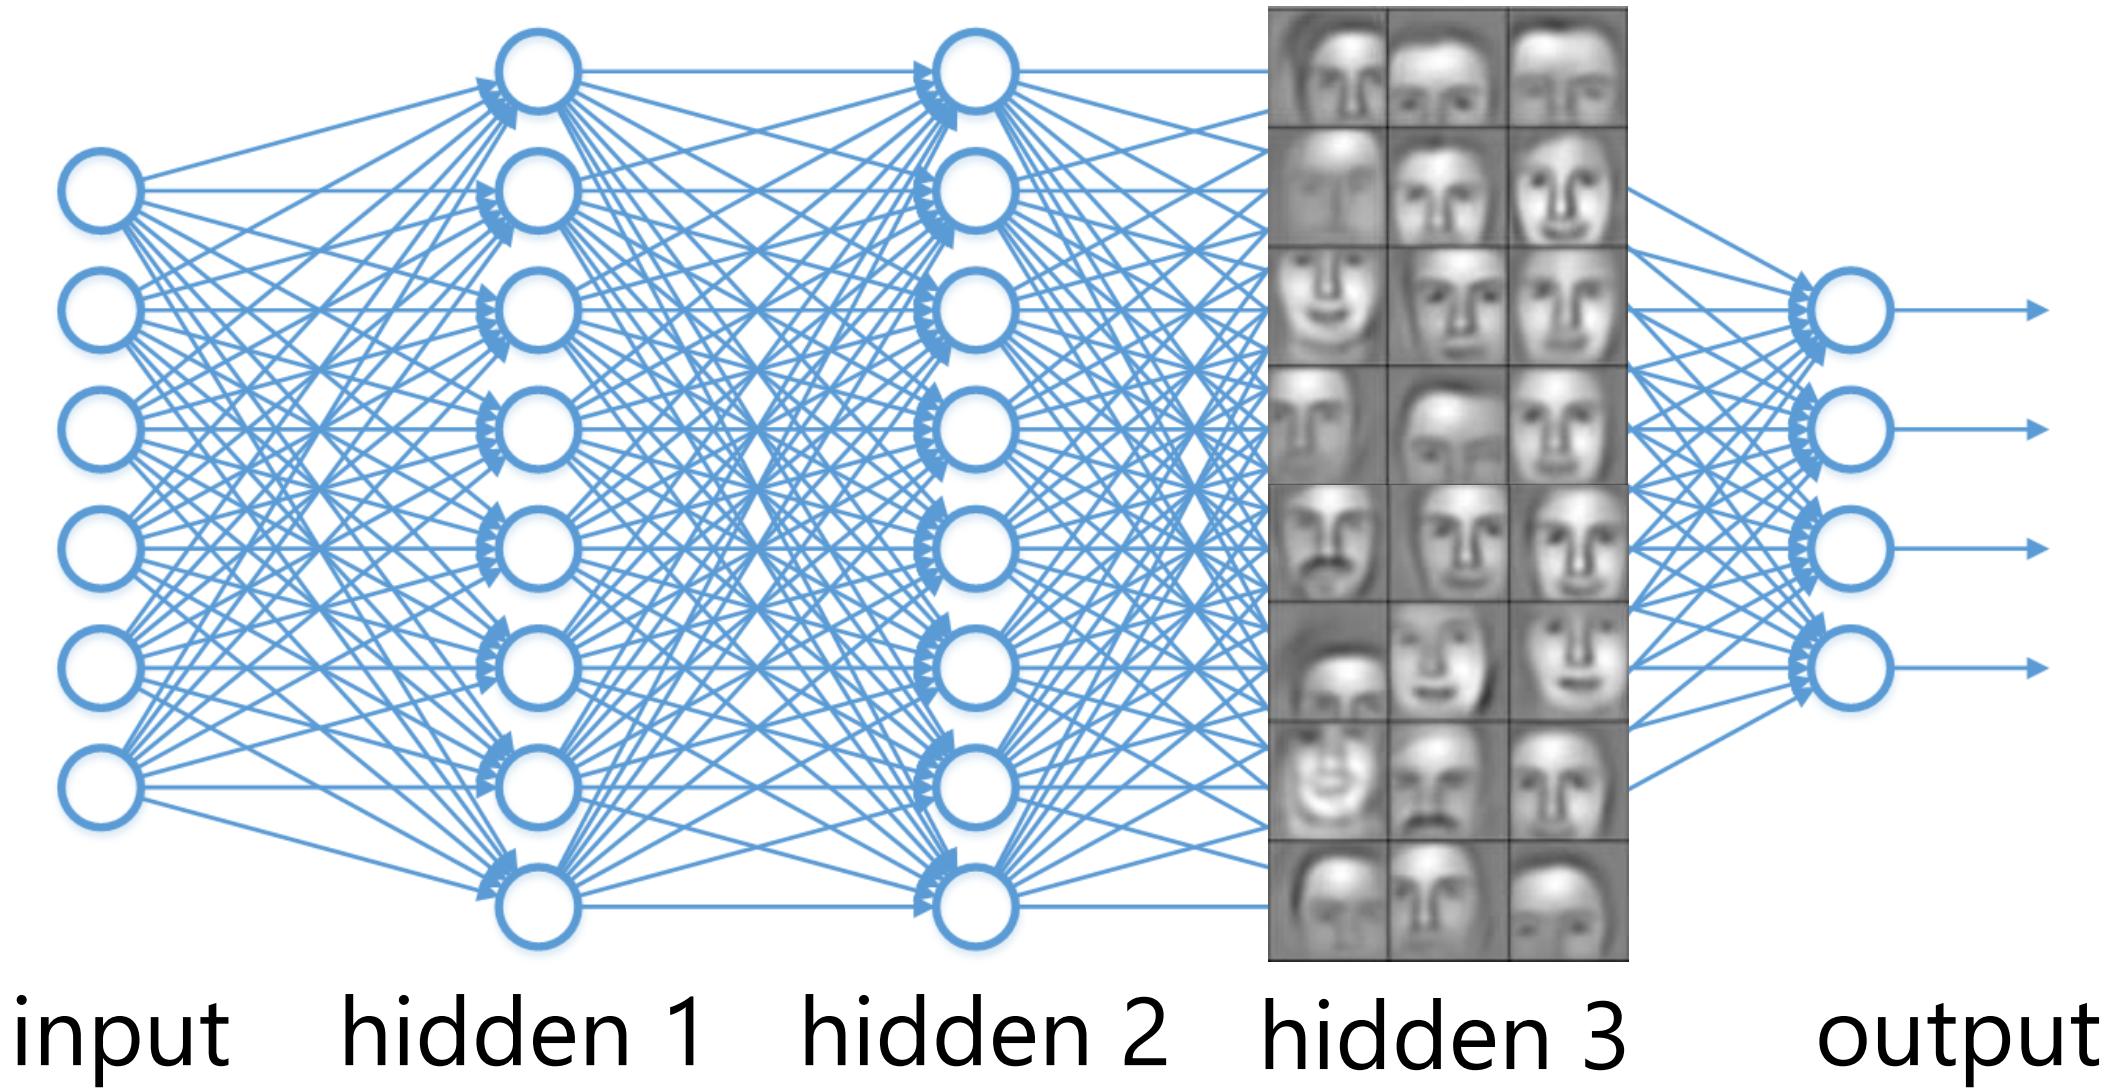
# Deep Neural Network



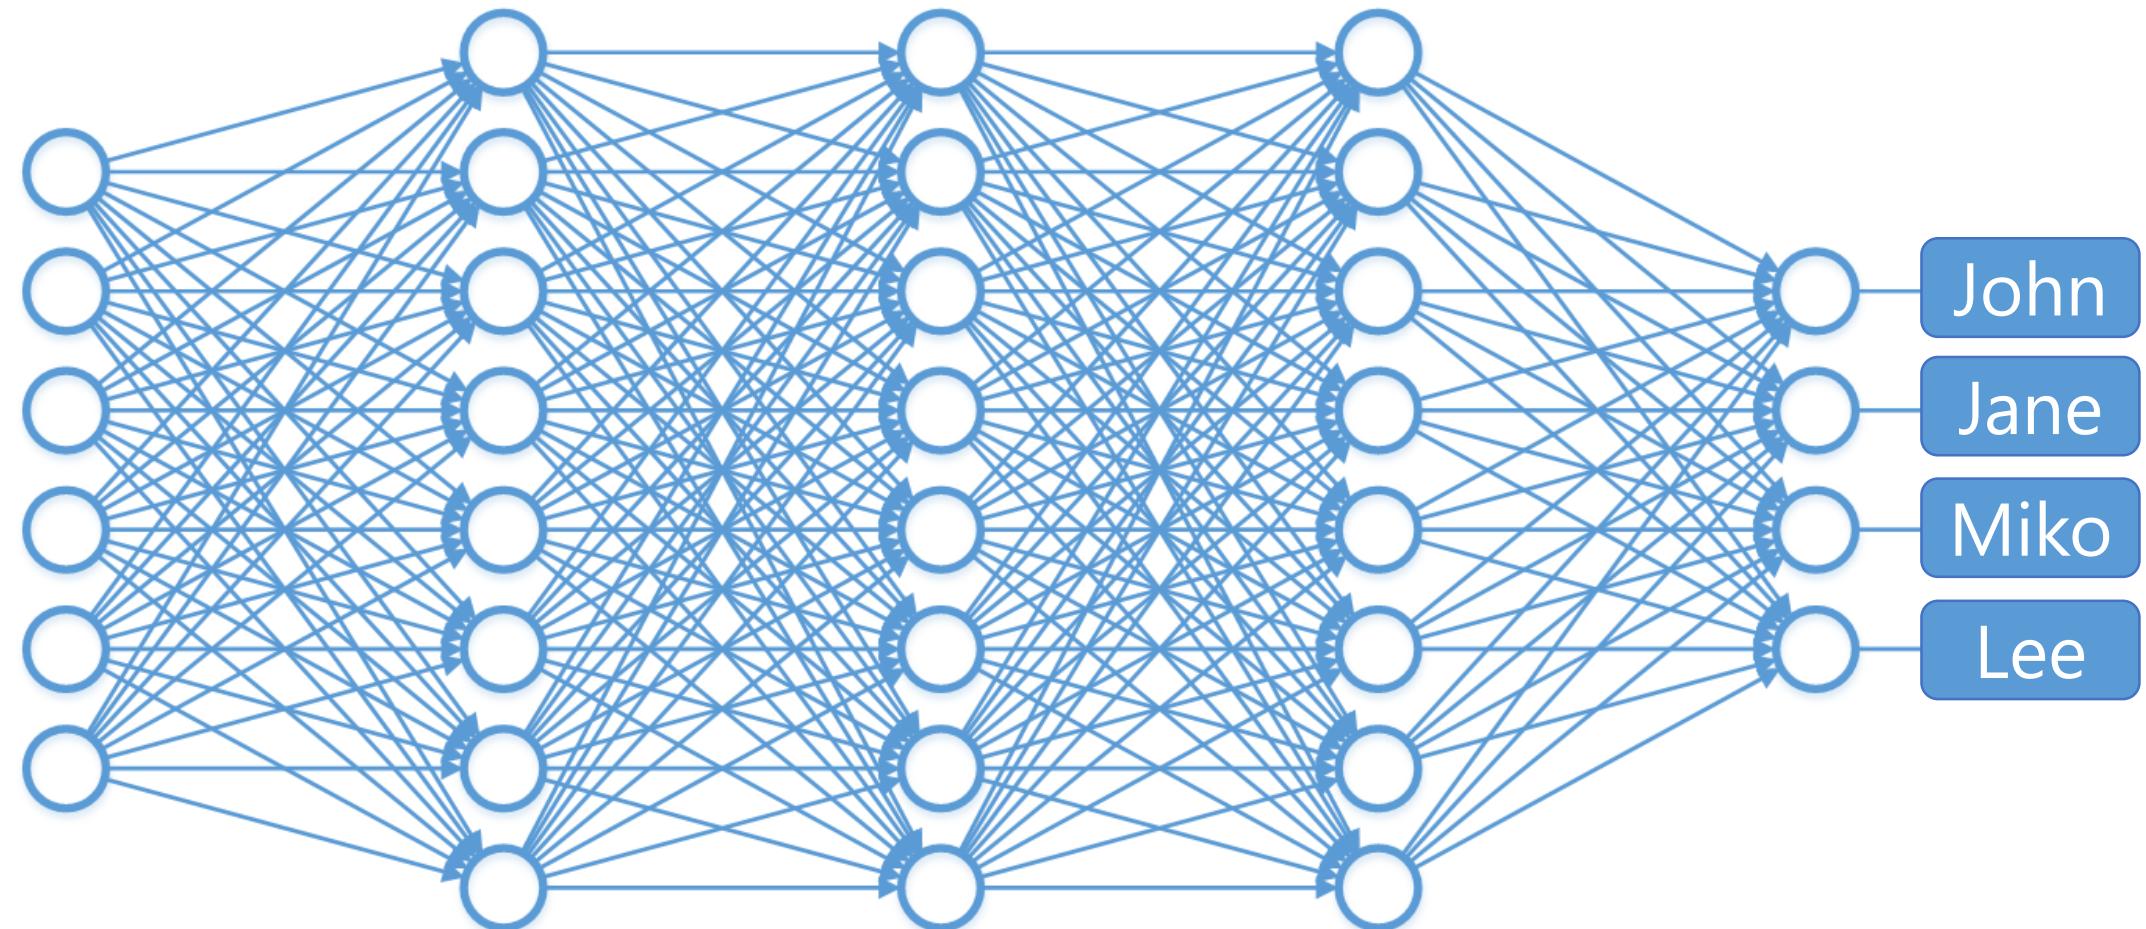
# Deep Neural Network



# Deep Neural Network



# Deep Neural Network



input

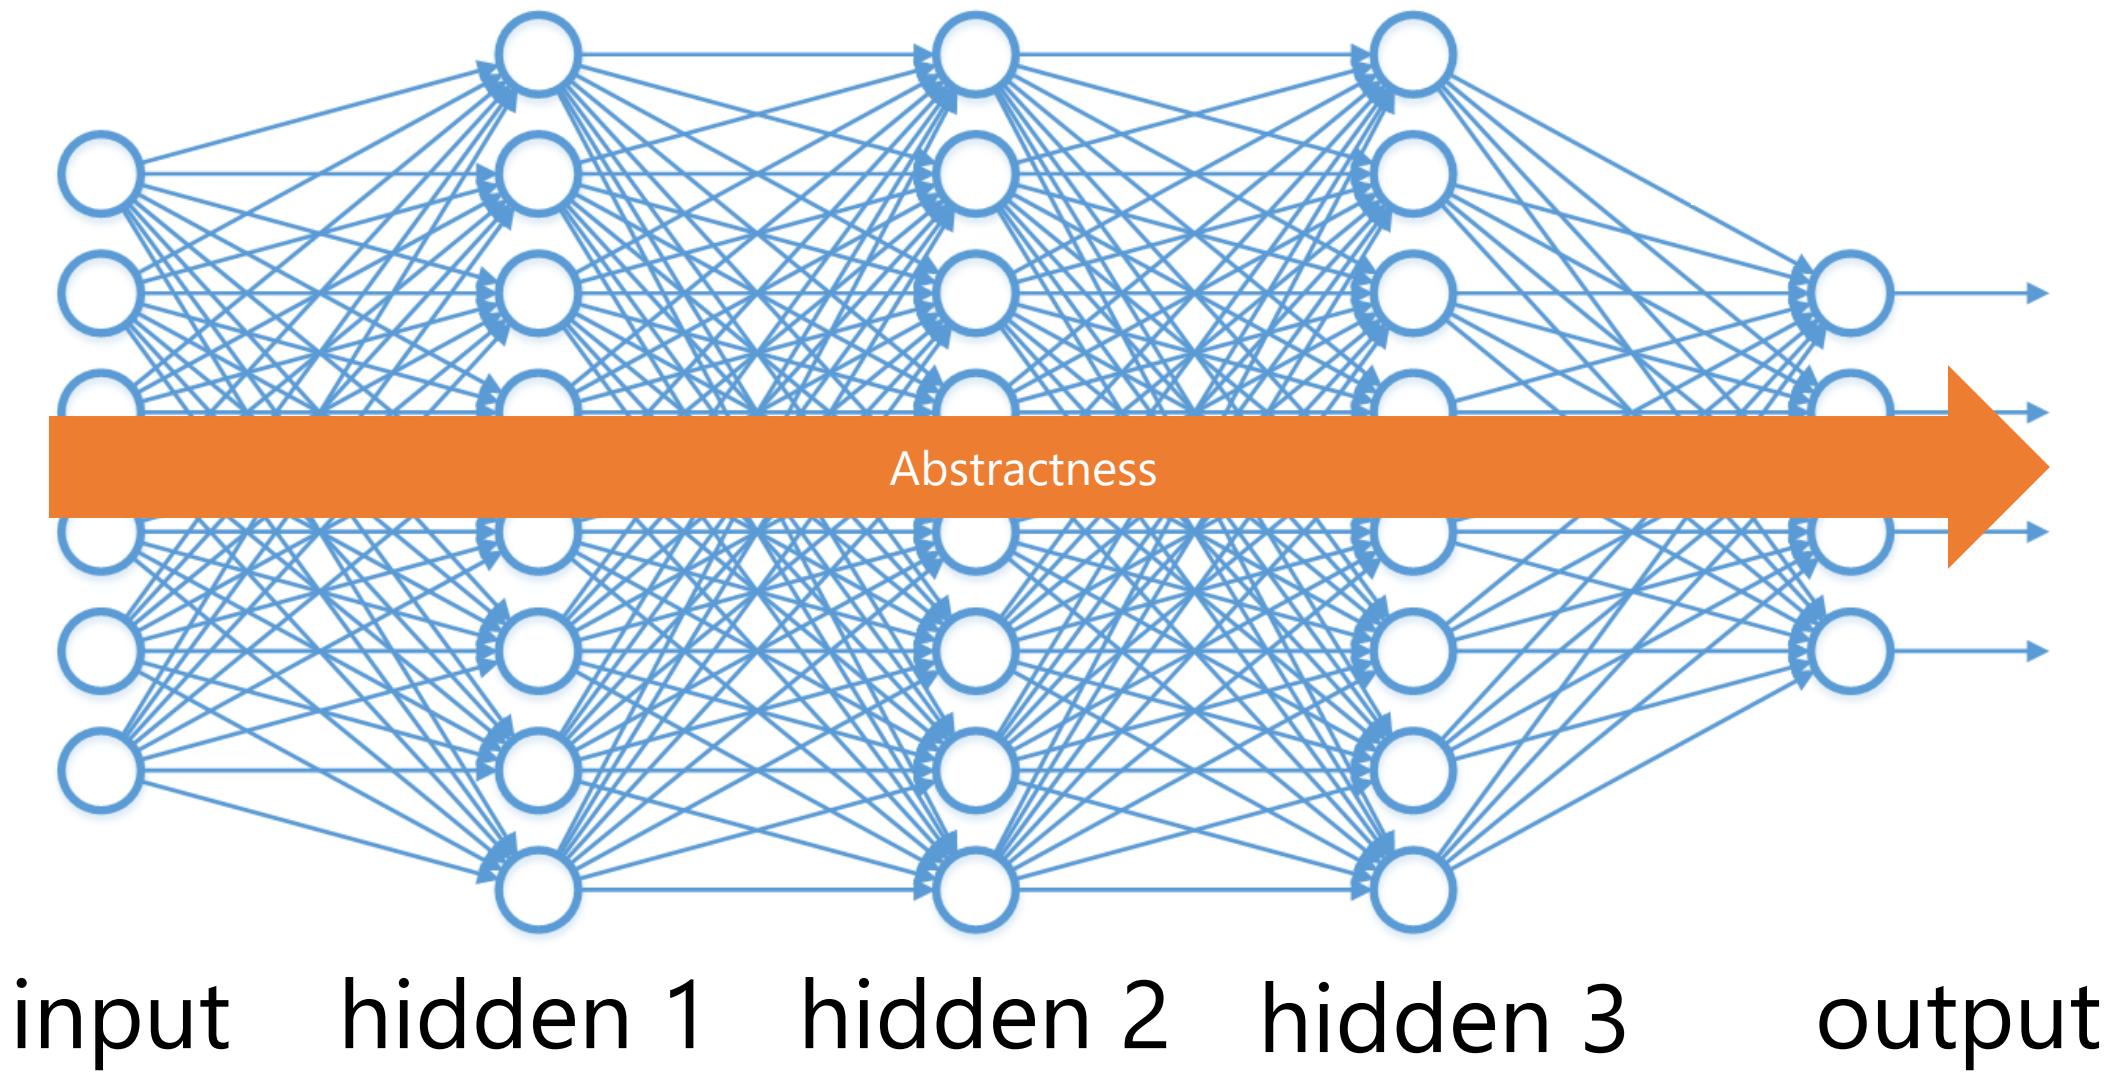
hidden 1

hidden 2

hidden 3

output

# Deep Neural Network



# Deep Learning Techniques

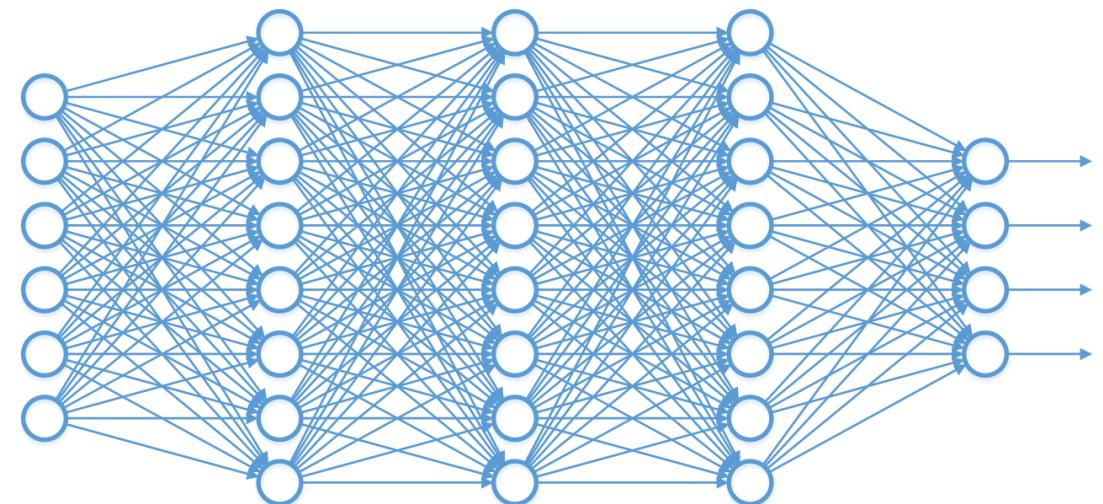
Fully connected (DNN)

Convolutional (CNN)

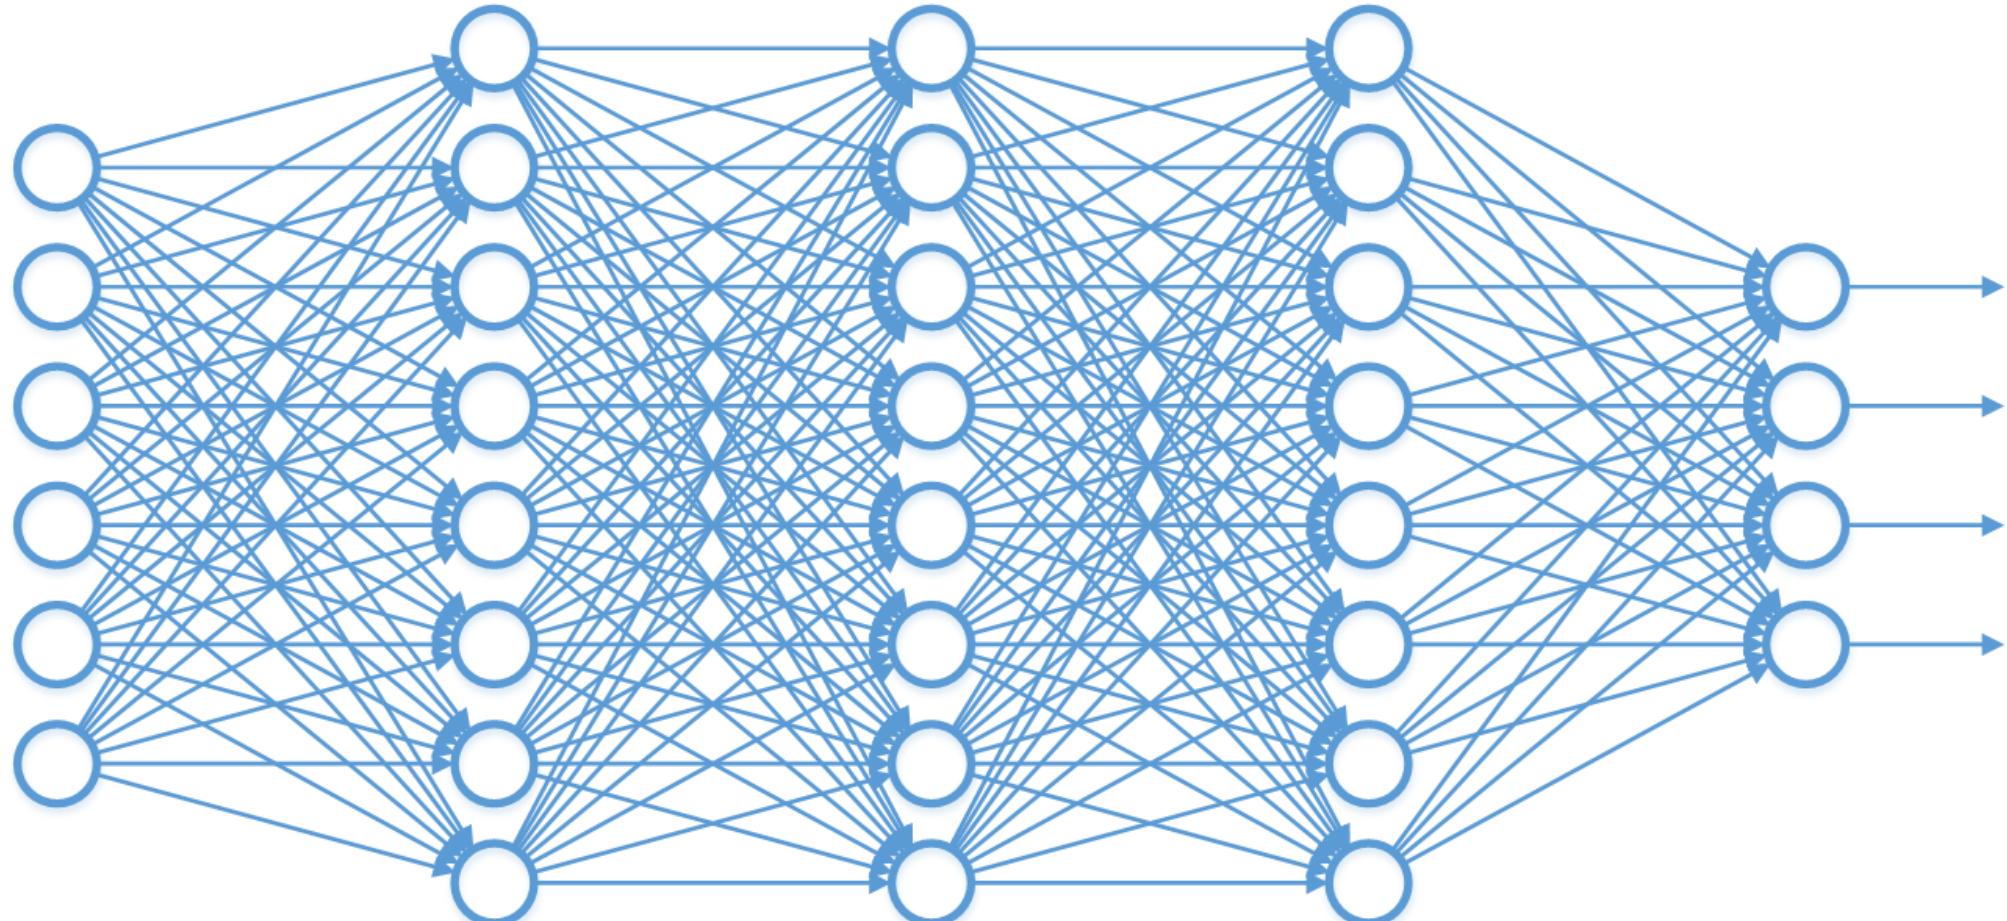
Recurrent (RNN)

Generative Adversarial (GAN)

Deep Q Learning (DQN)

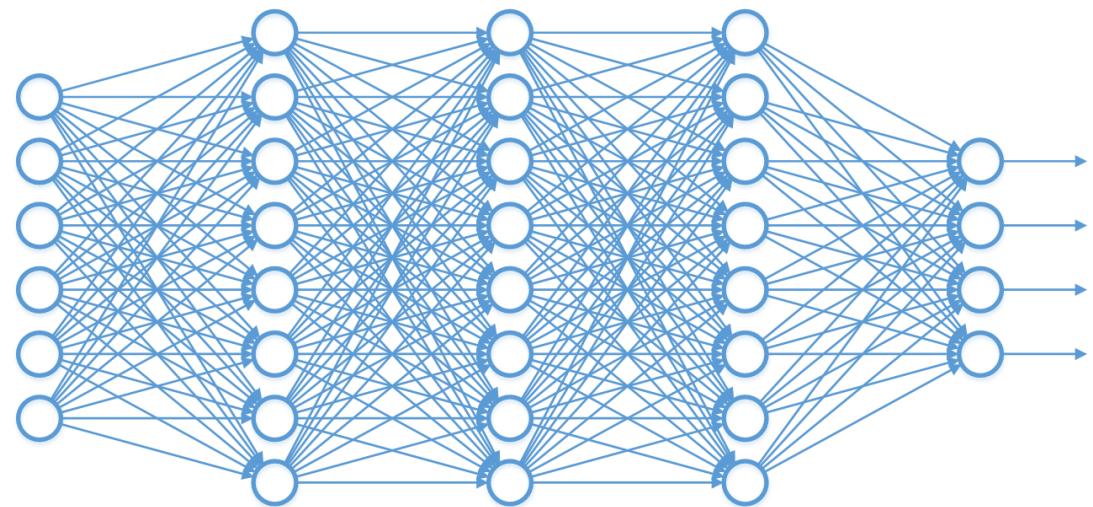


# Deep Neural Network



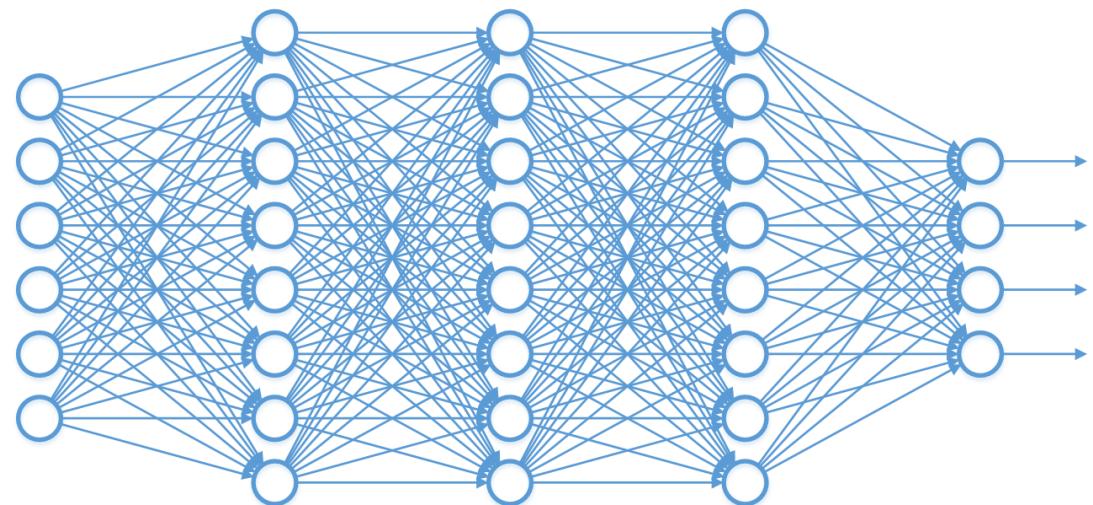
# Deep Neural Network

Neural network



# Deep Neural Network

Neural network  
Multiple hidden layers

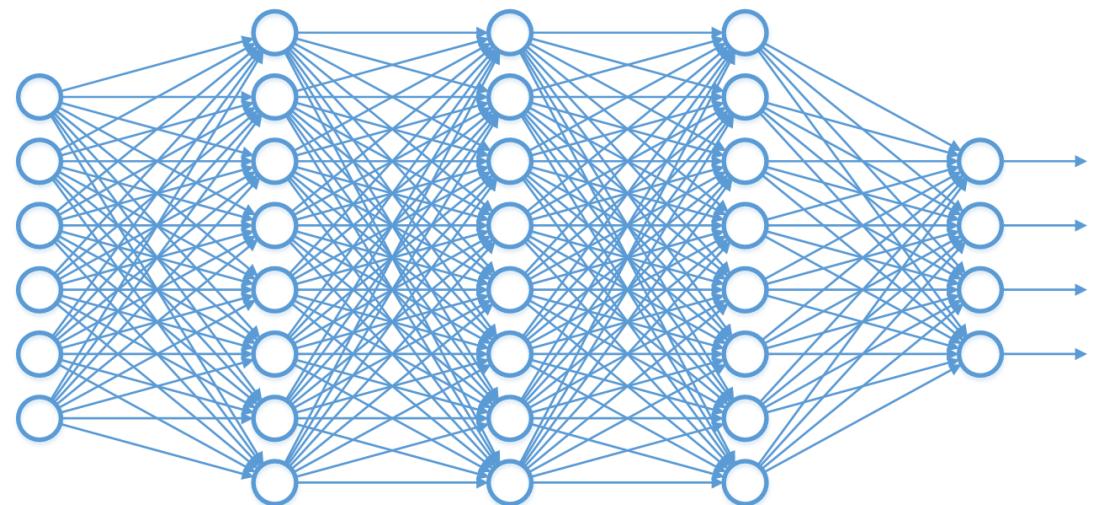


# Deep Neural Network

Neural network

Multiple hidden layers

Non-linear activation



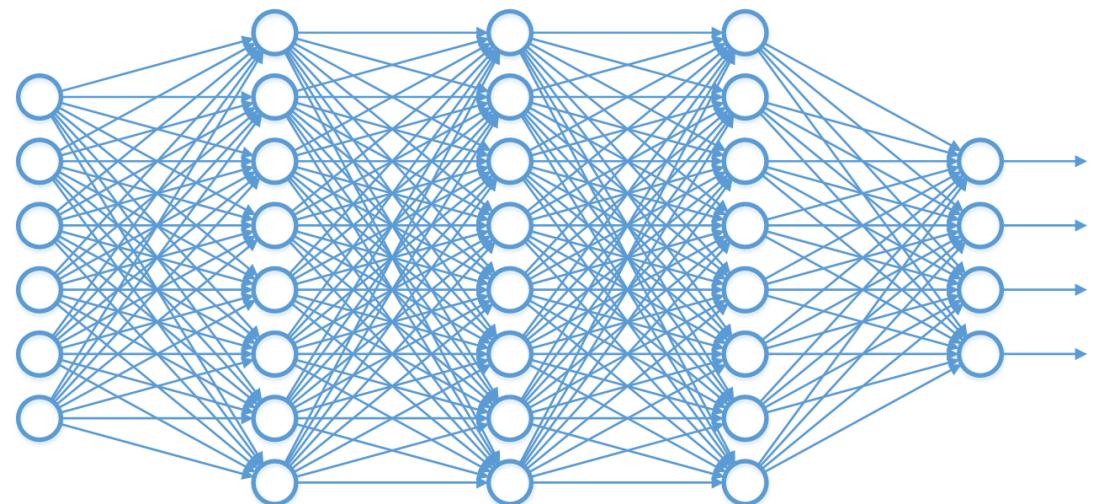
# Deep Neural Network

Neural network

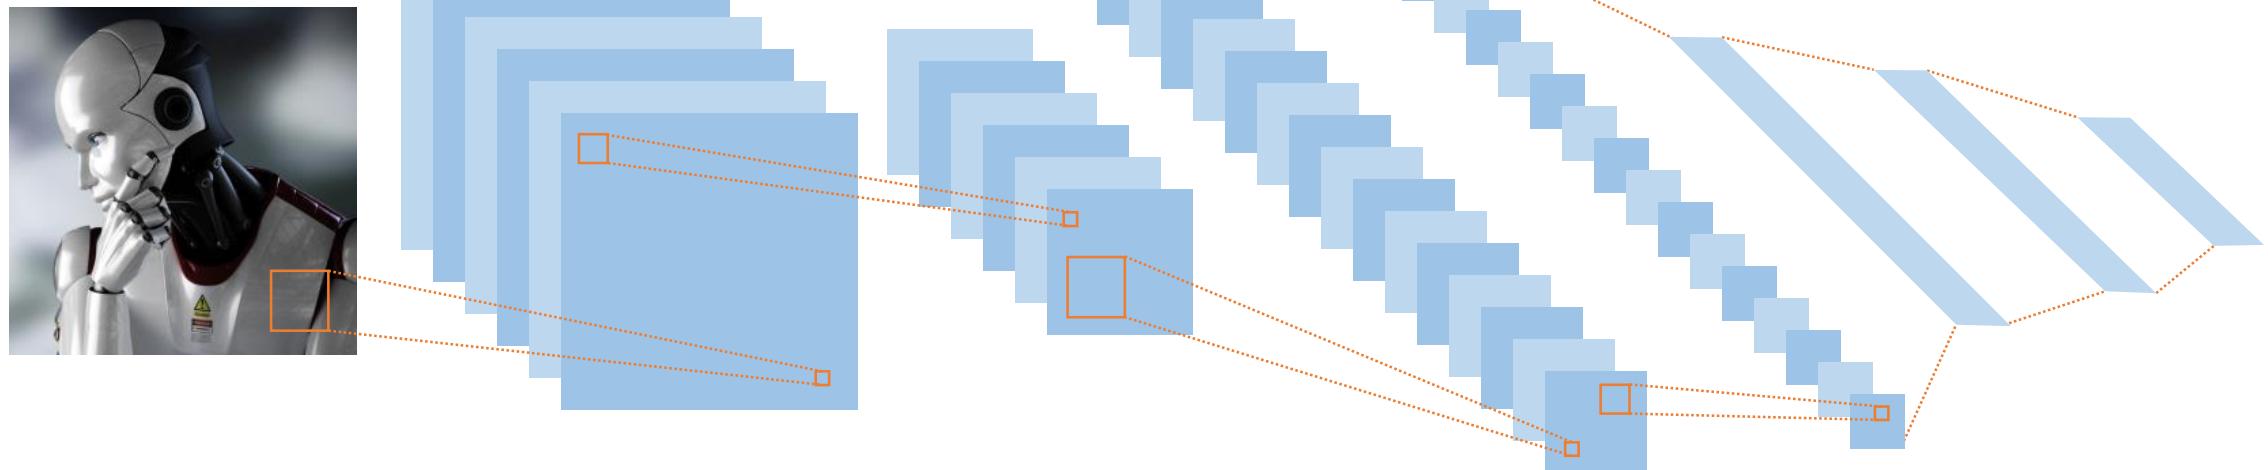
Multiple hidden layers

Non-linear activation

Fully connected

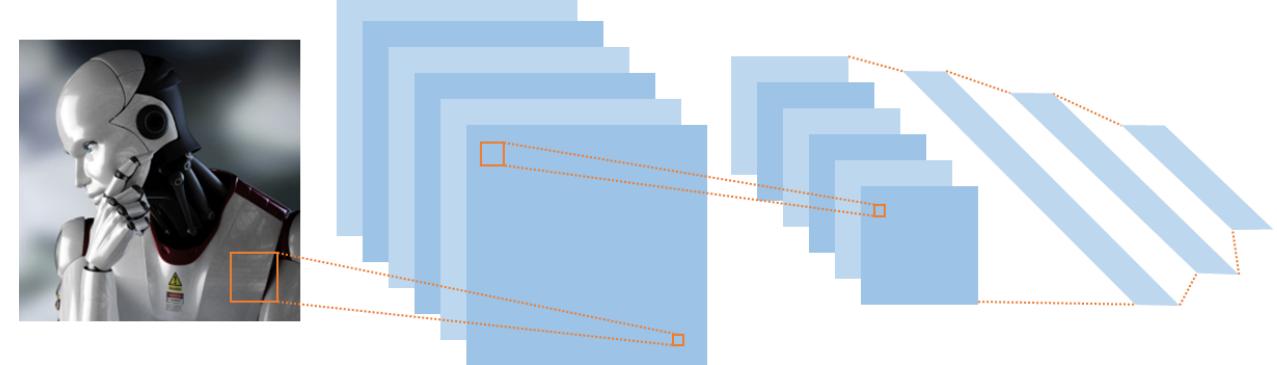


# Convolutional Neural Networks (CNN)



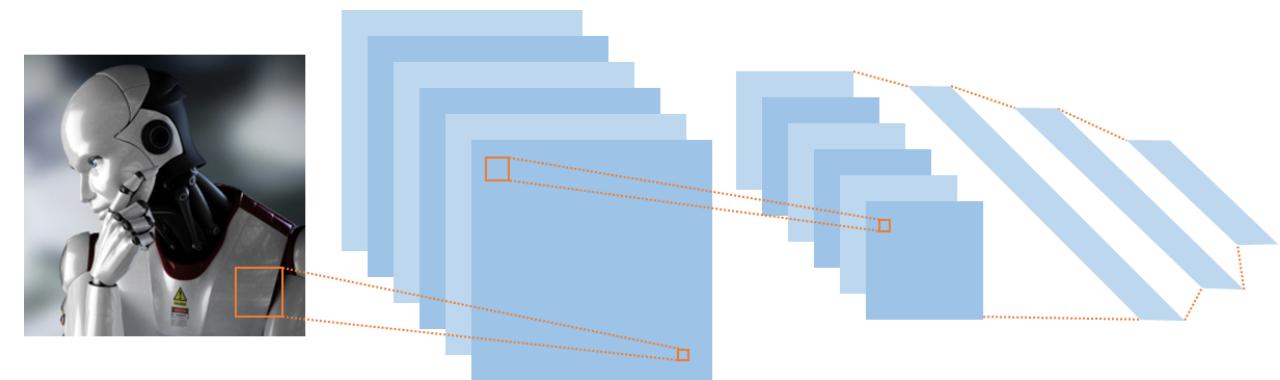
# Convolutional Neural Network (CNN)

Sparse



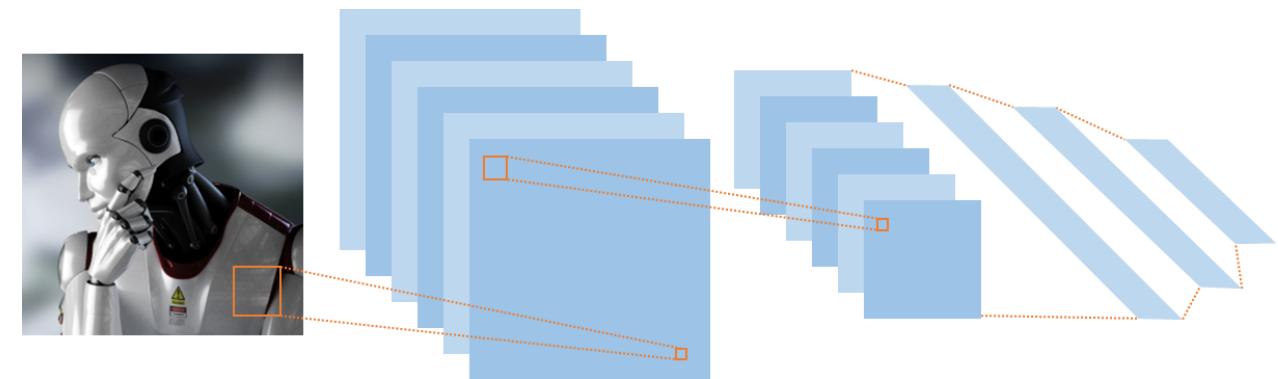
# Convolutional Neural Network (CNN)

Sparse  
Convolutions



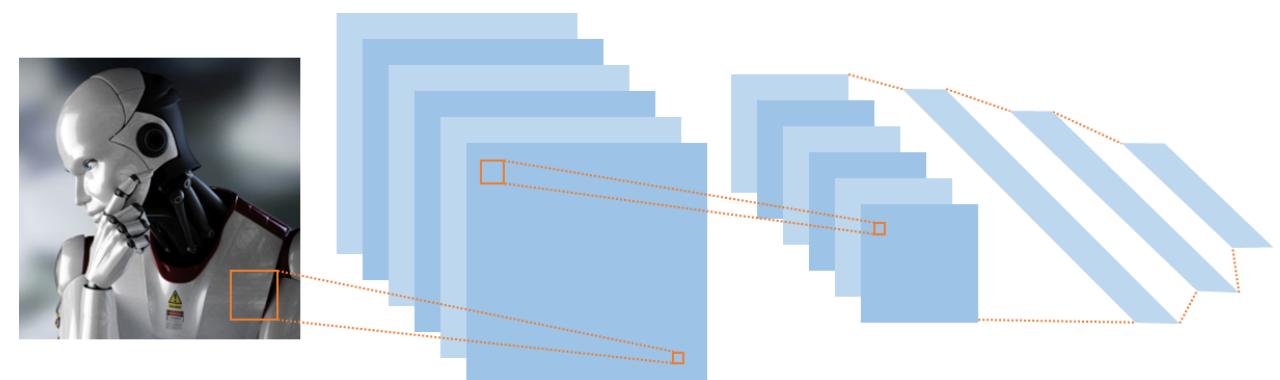
# Convolutional Neural Network (CNN)

Sparse  
Convolutions  
Filters

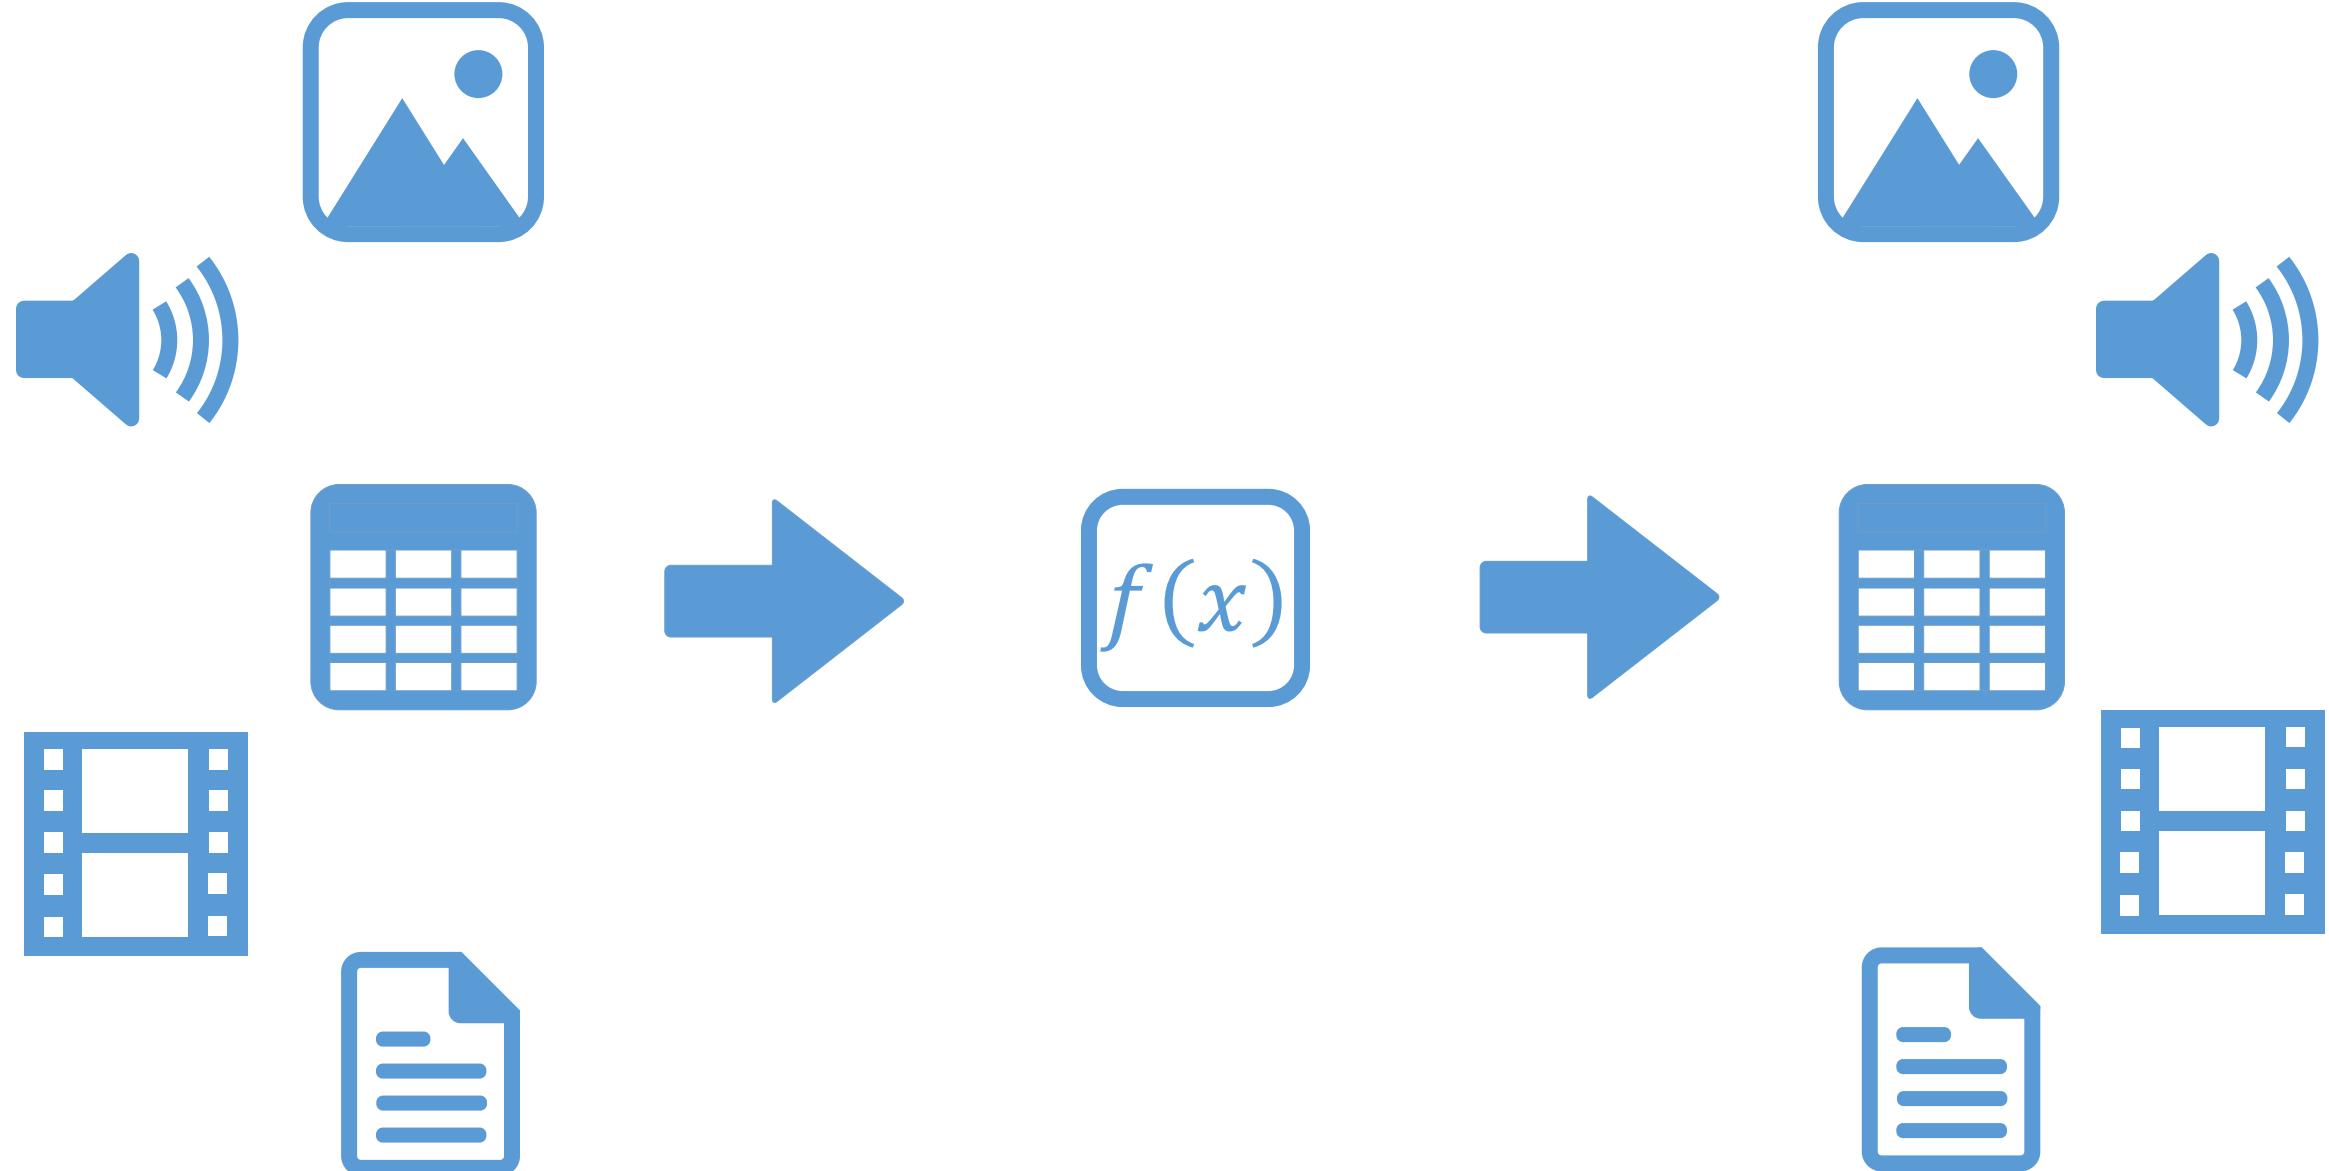


# Convolutional Neural Network (CNN)

Sparse  
Convolutions  
Filters  
Pooling







# Why Use Deep Learning?

## **Pros**

More powerful

More accurate

Data synthesis

# Why Use Deep Learning?

## **Pros**

- More powerful
- More accurate
- Data synthesis

## **Cons**

- More complex
- More training
- Less transparent

# Deep Learning Demo

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1

2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2

3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3

4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4

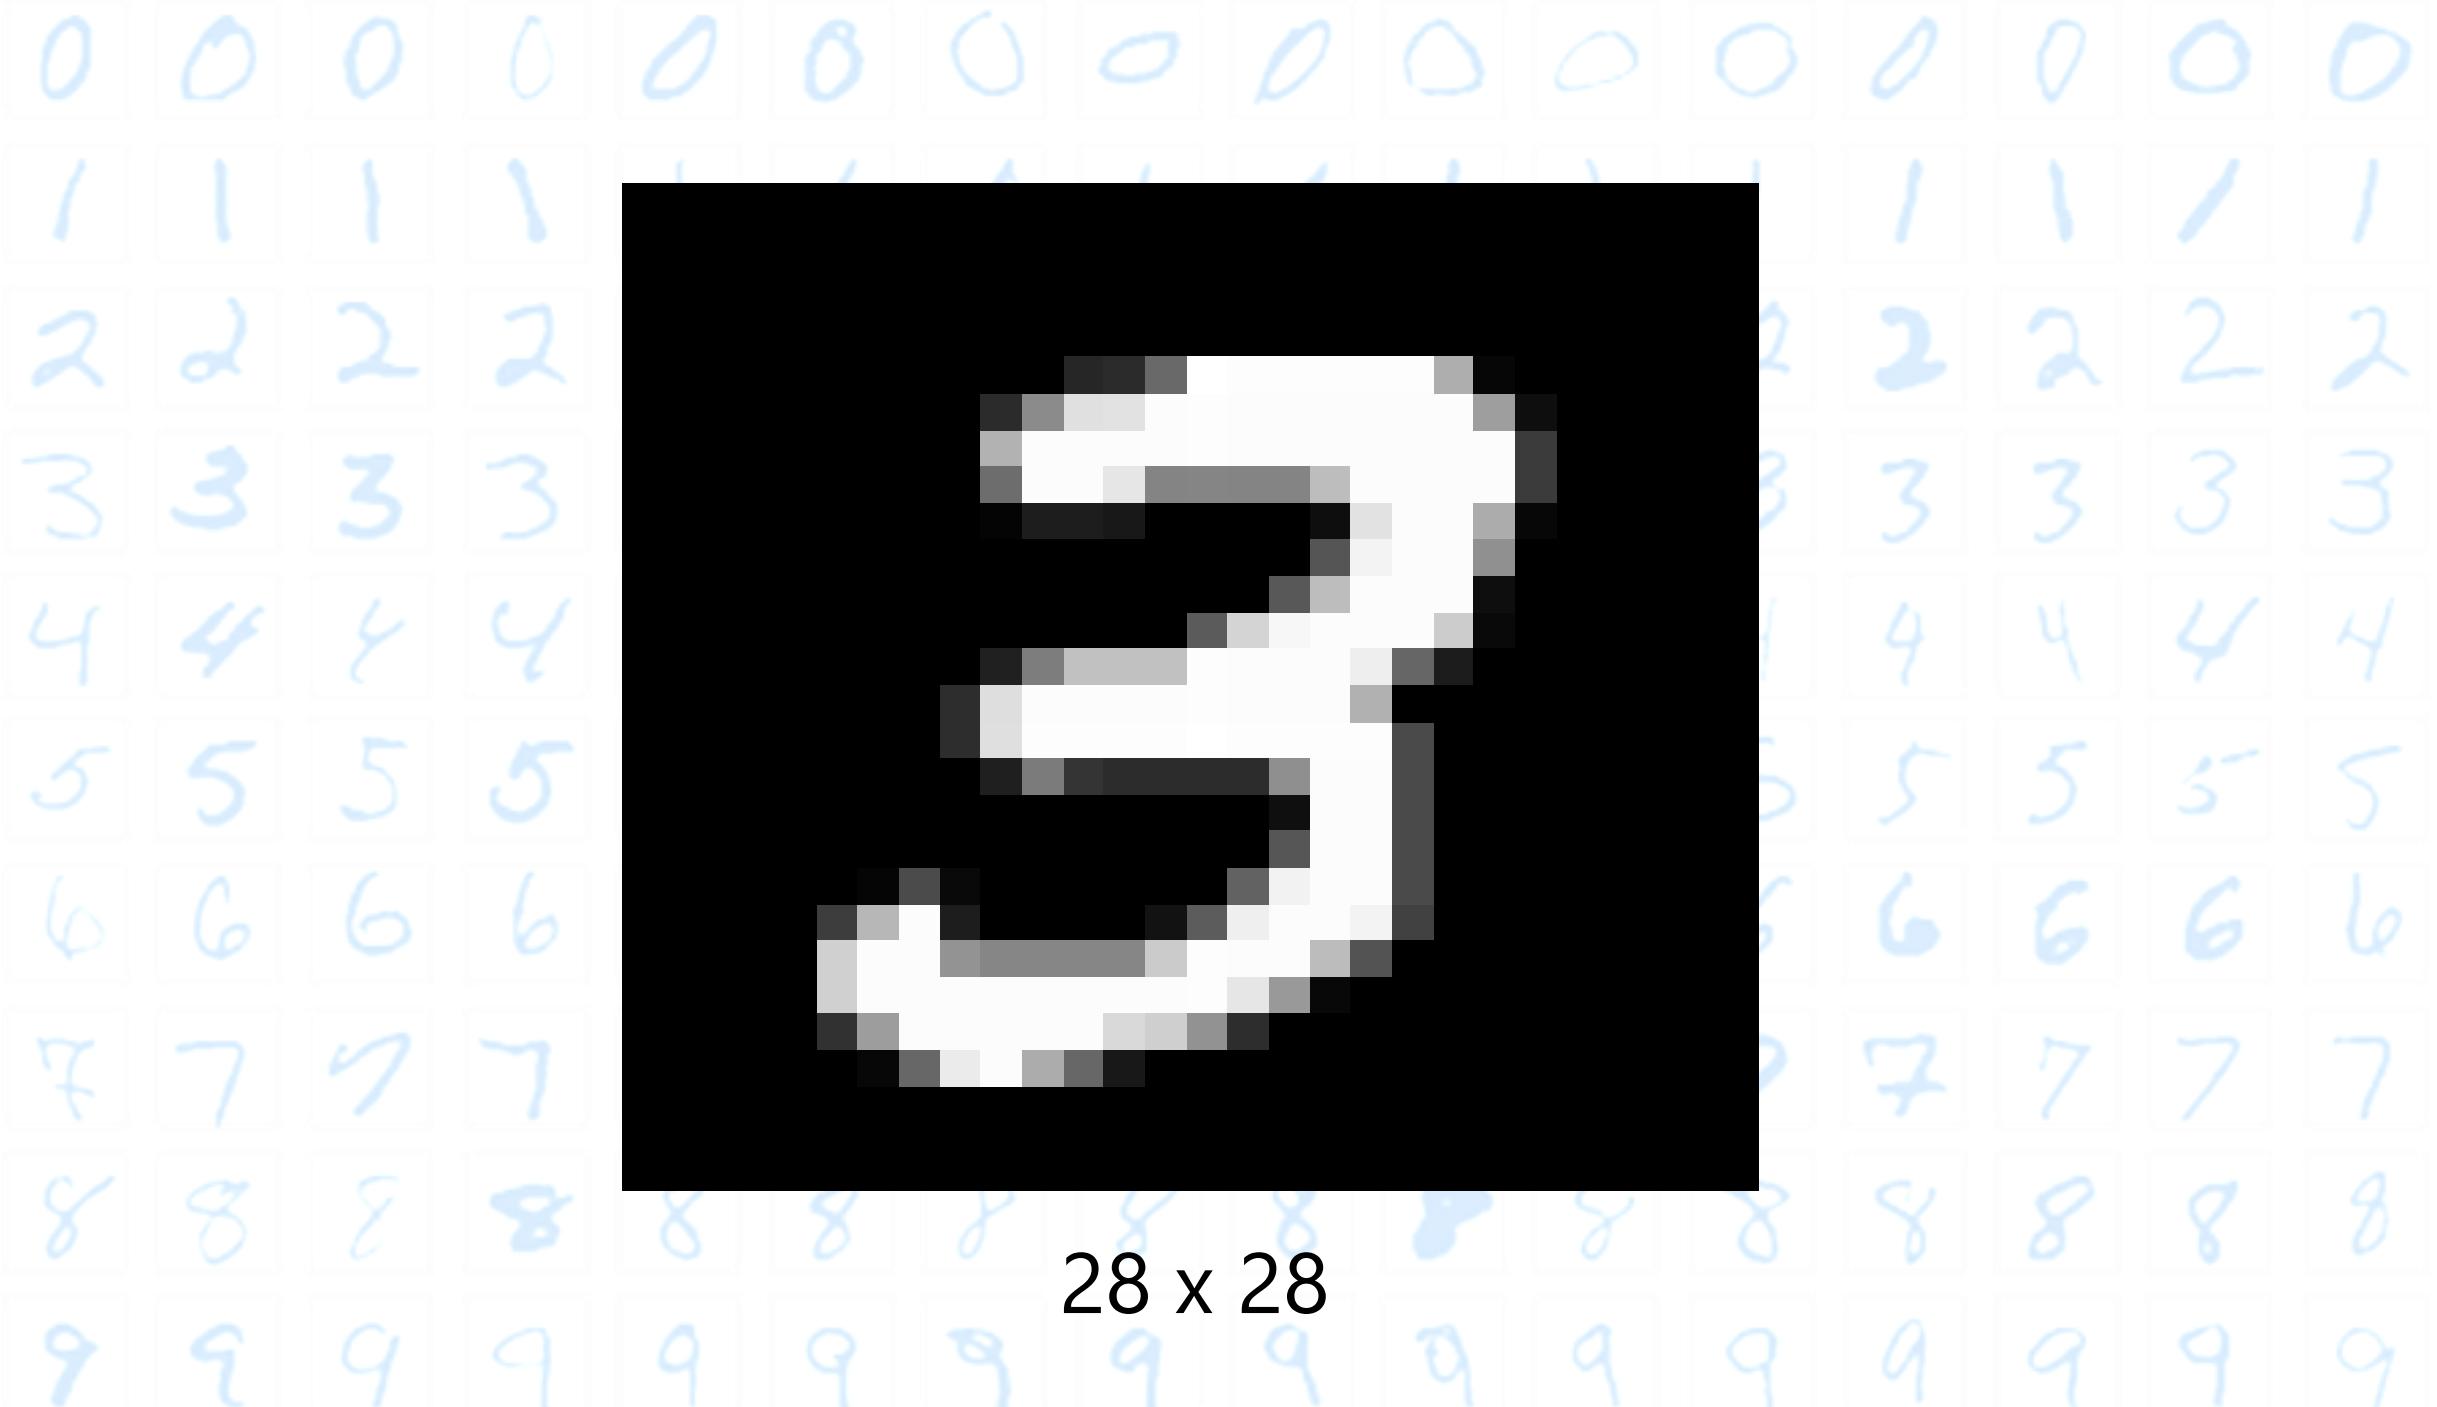
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5

6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6

7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7

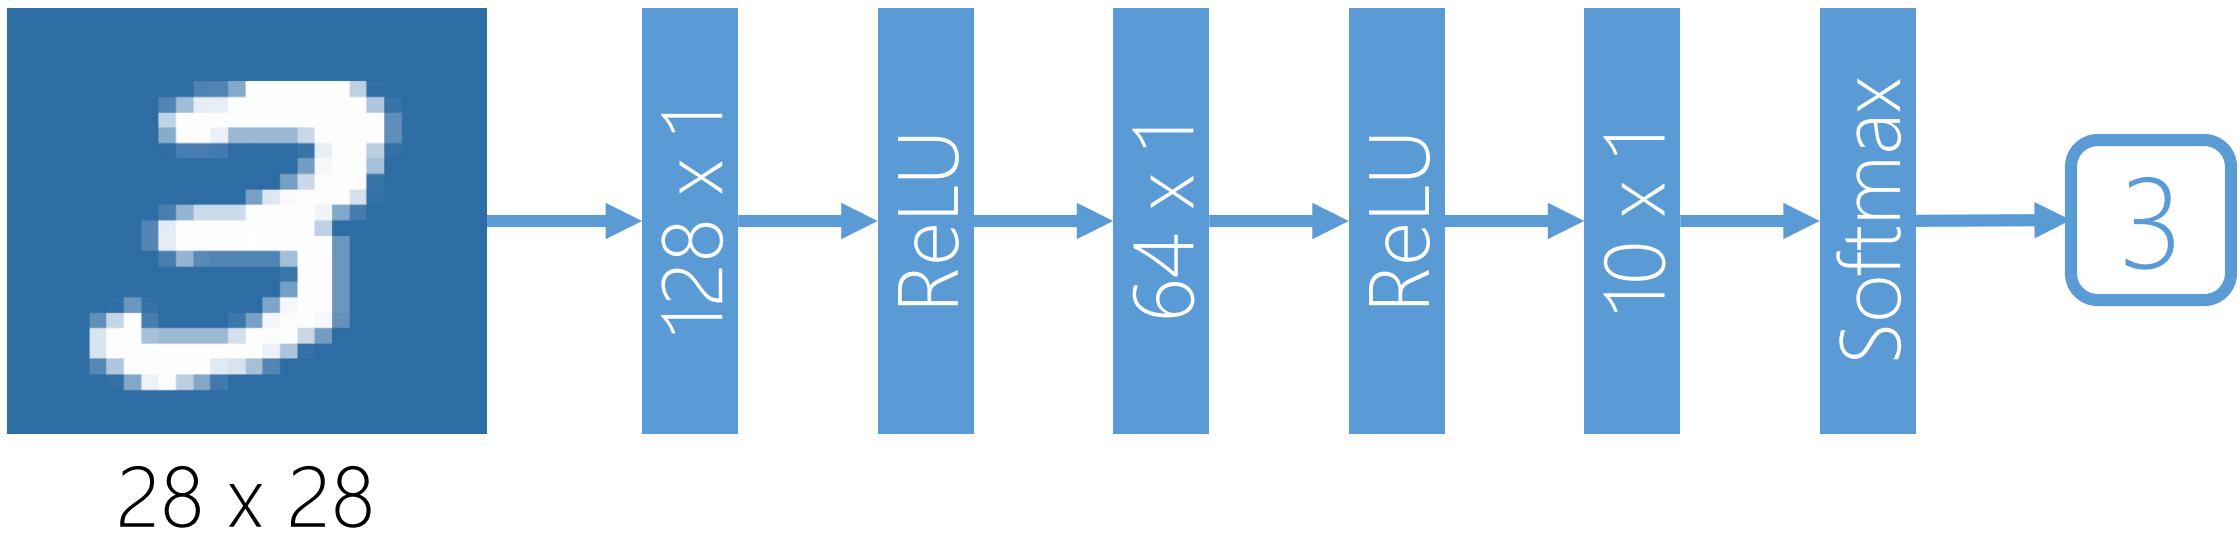
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8

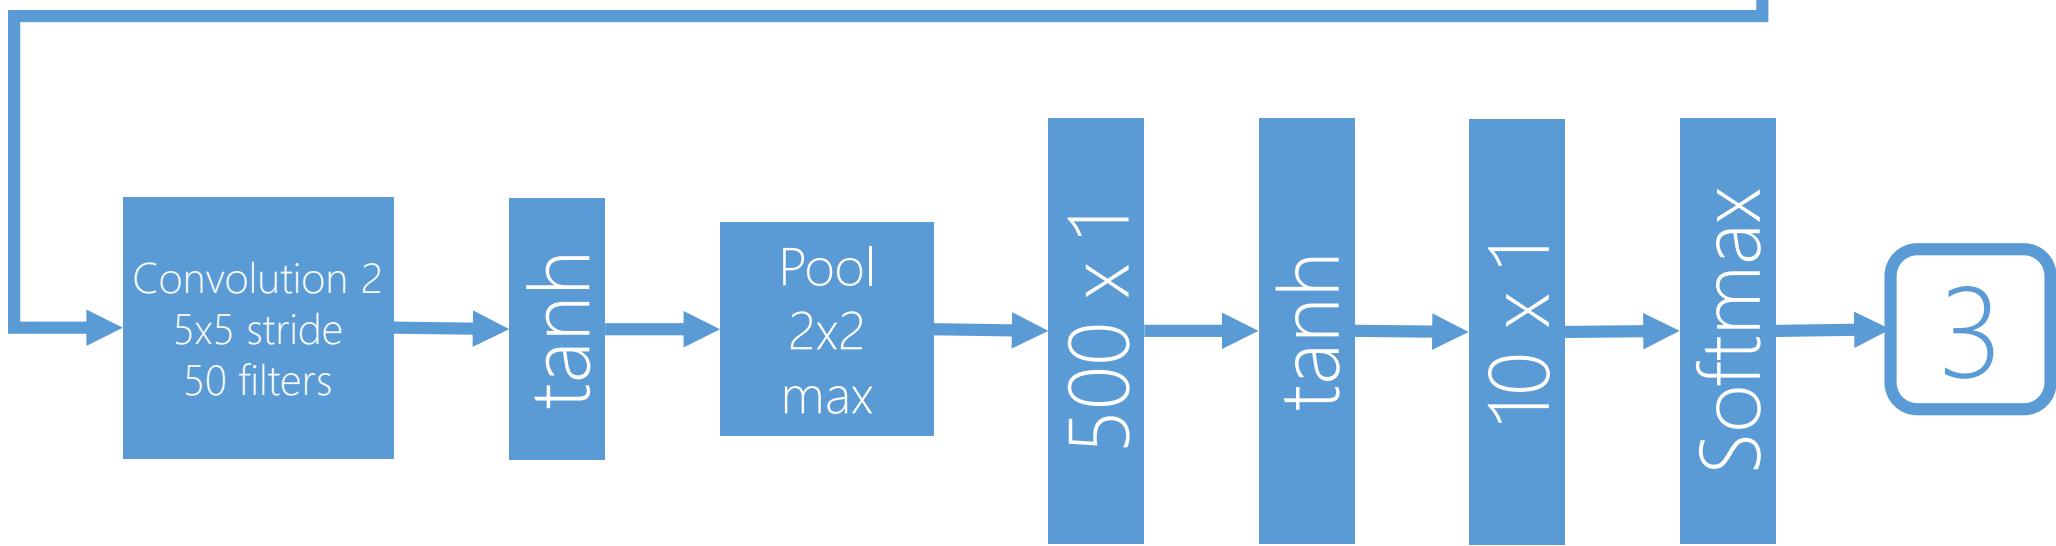
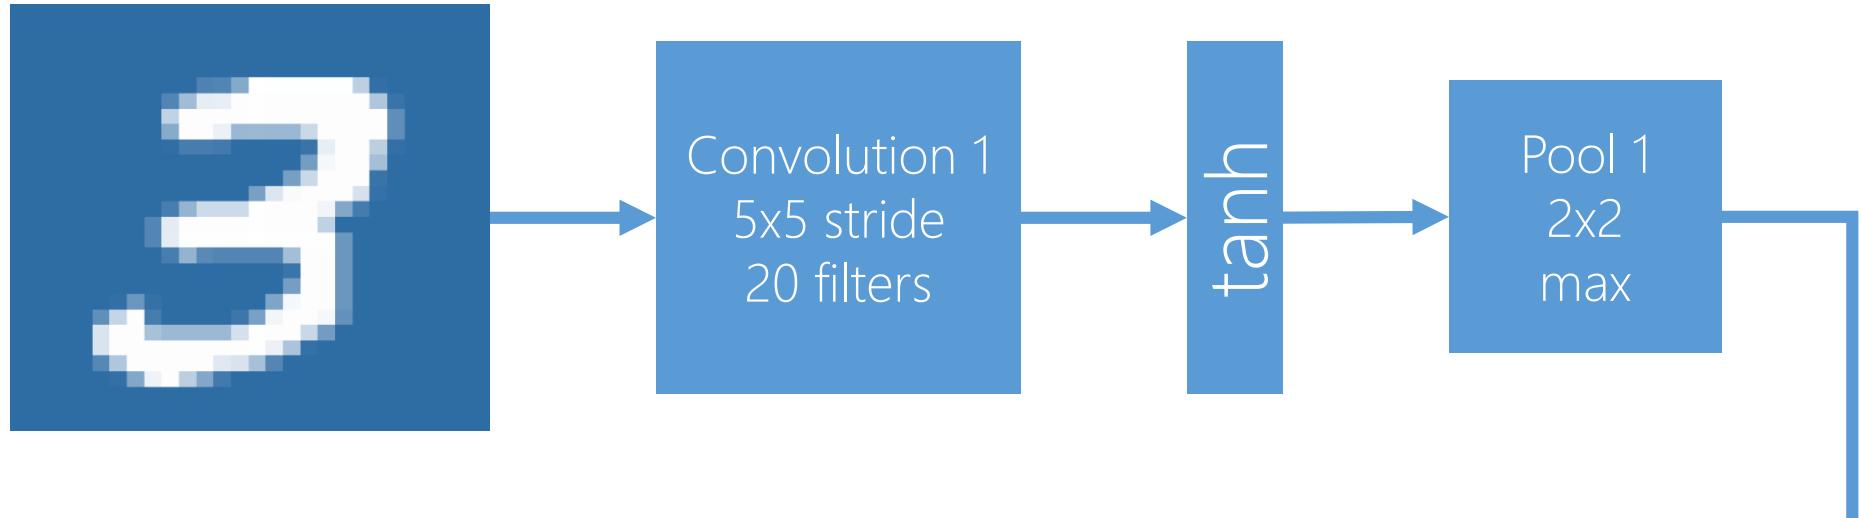
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9



28 x 28

## MNIST





## Demo 6 – Deep Learning

Goal: Predict handwritten digits  
with a deep neural network

Lab 6A – Deep Learning (Easy)

Goal: Predict handwritten digits  
with a deep neural network

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9

## Lab 6B – ML in Practice (Hard)

Goal: Predict handwritten digits  
with CNN (LeNet)

Lab 6B – ML in Practice (Hard)

Goal: Predict handwritten digits  
with CNN (LeNet)

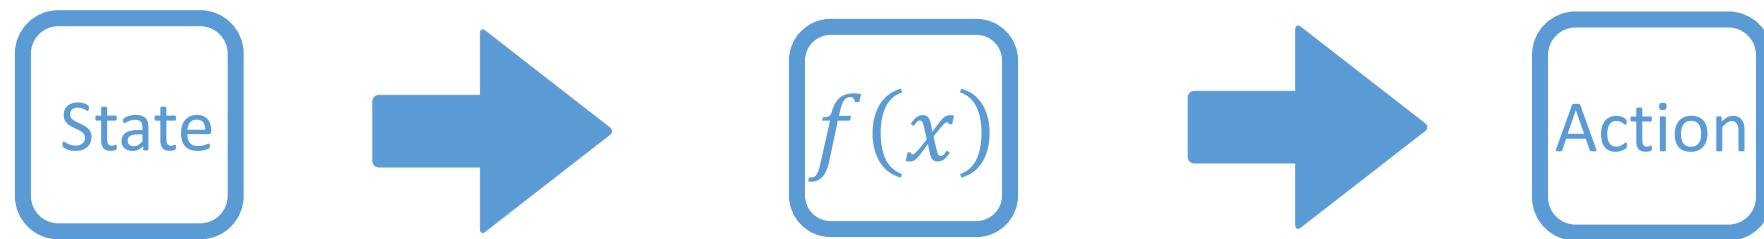
Lab 6B – ML in Practice (Hard)

Goal: Predict handwritten digits  
with CNN (LeNet)

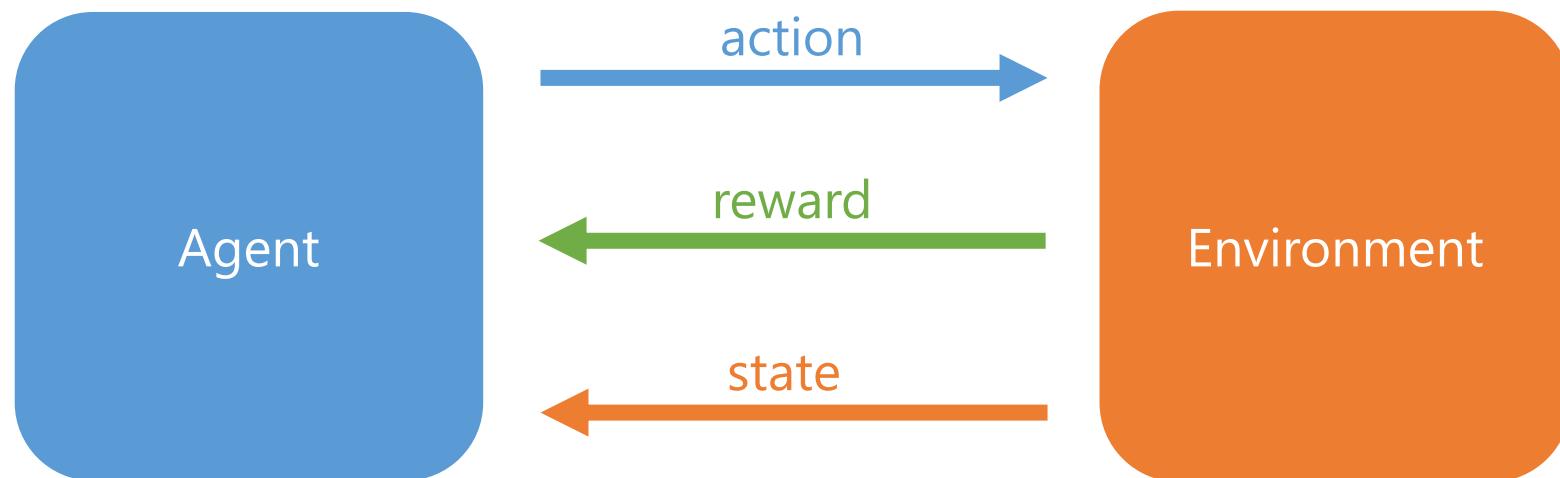
# Reinforcement Learning

NOTE: Add video of RL playing video game

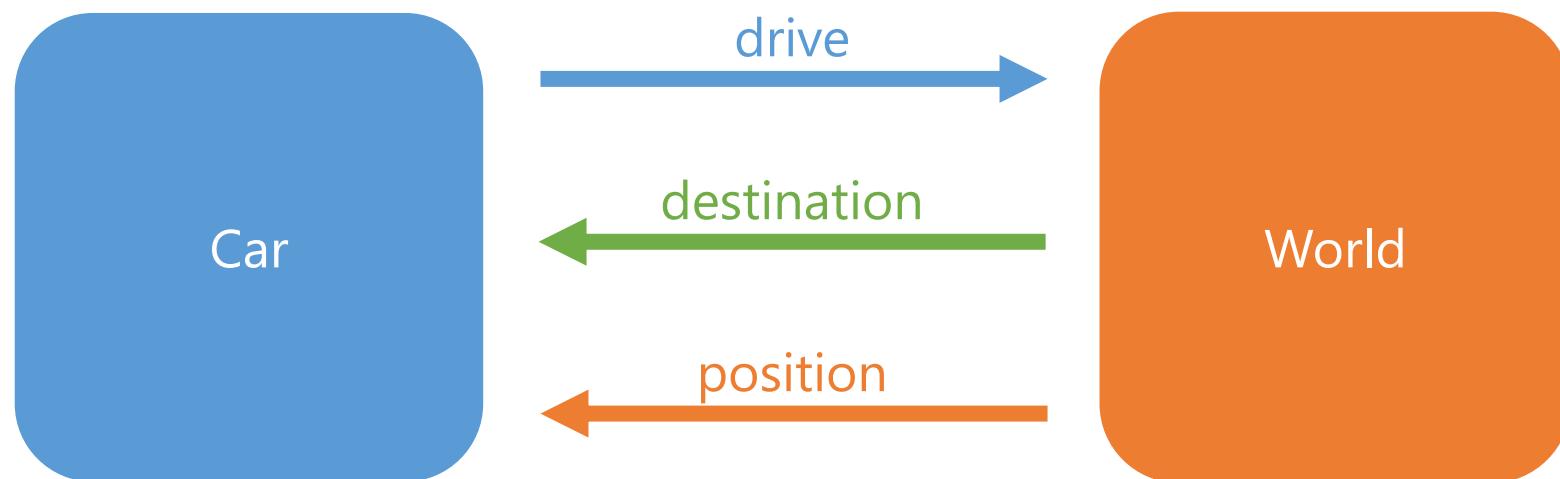
# Reinforcement Learning



# Reinforcement Learning

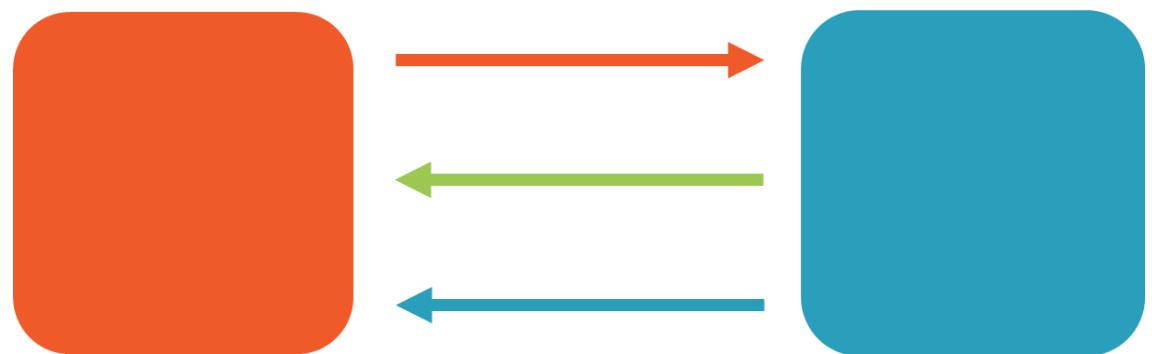


# Reinforcement Learning



# Reinforcement Learning

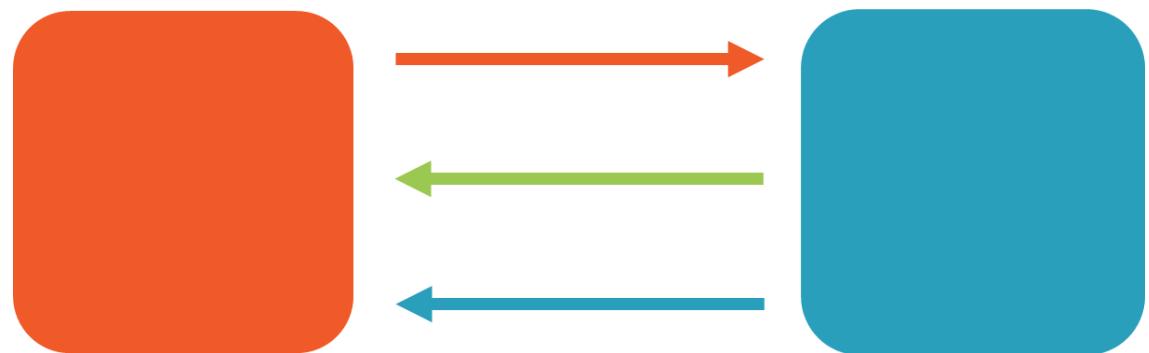
Action replay



# Reinforcement Learning

Action replay

Optimal policy

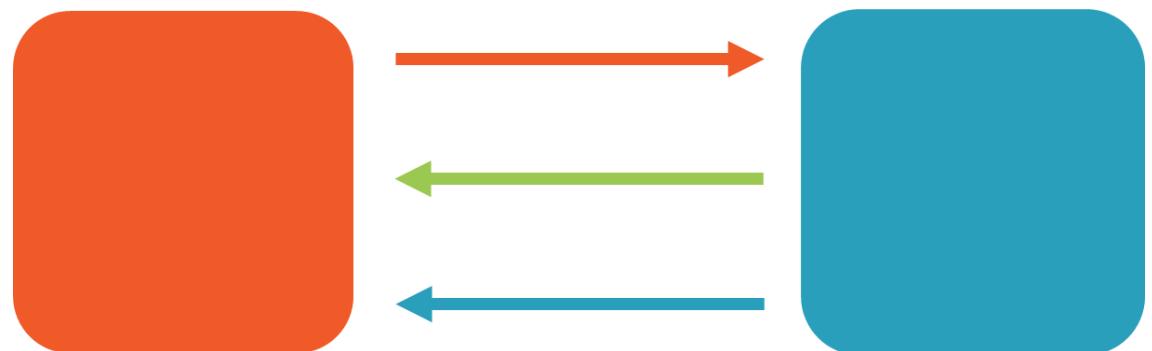


# Reinforcement Learning

Action replay

Optimal policy

Discounted reward



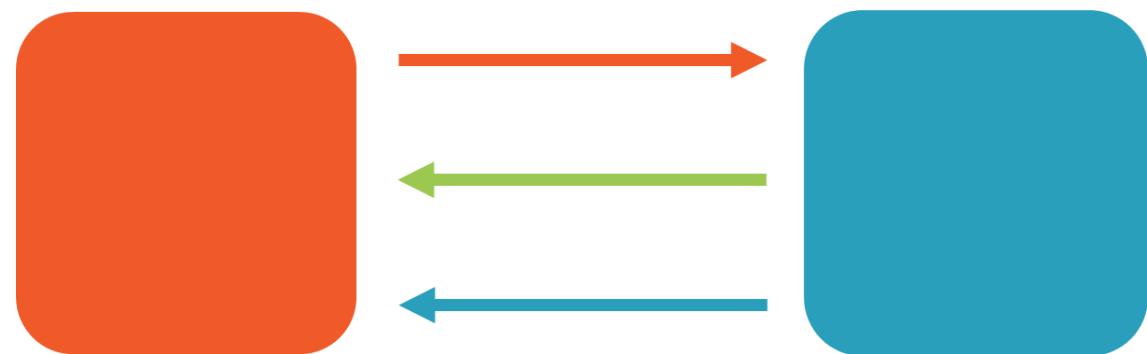
# Reinforcement Learning

Action replay

Optimal policy

Discounted reward

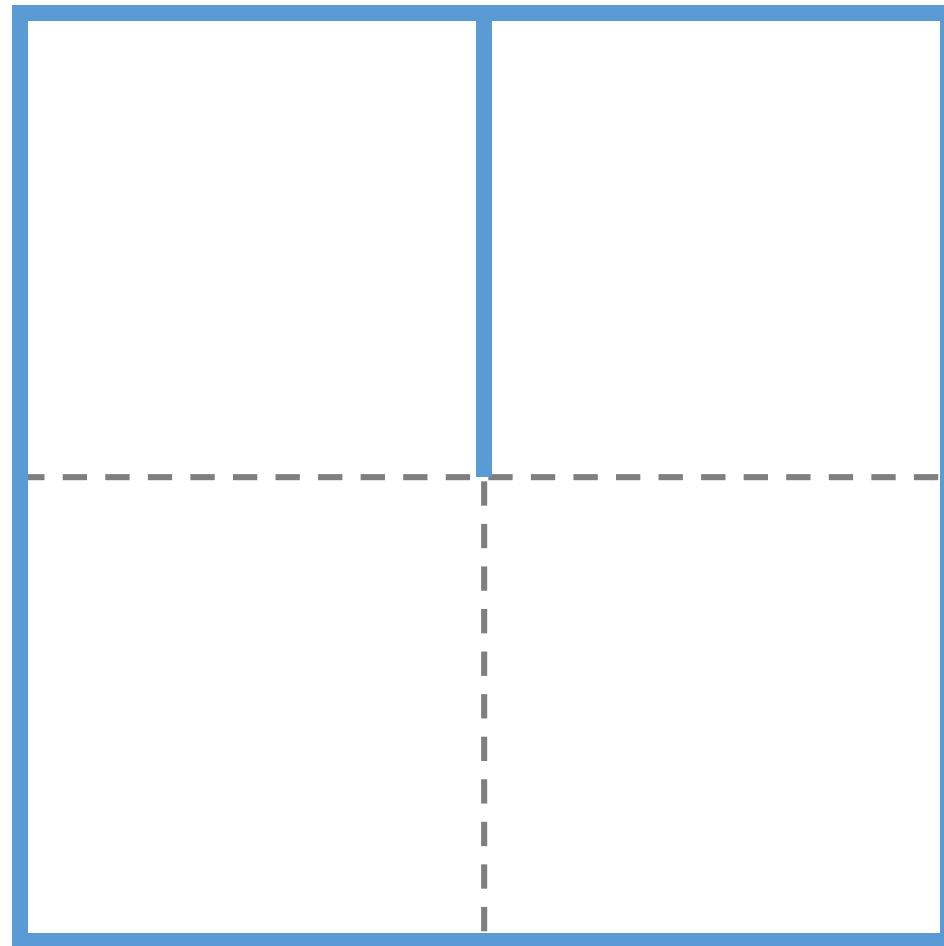
Markov decision process



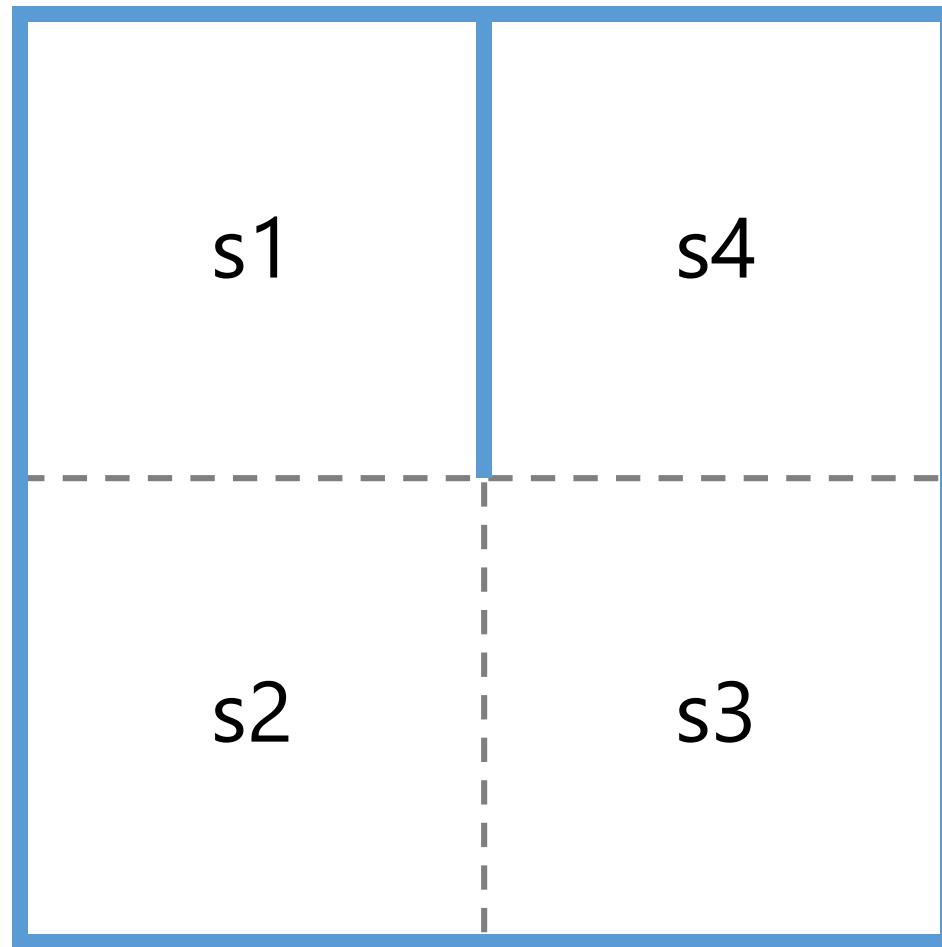
# Reinforcement Learning Demo



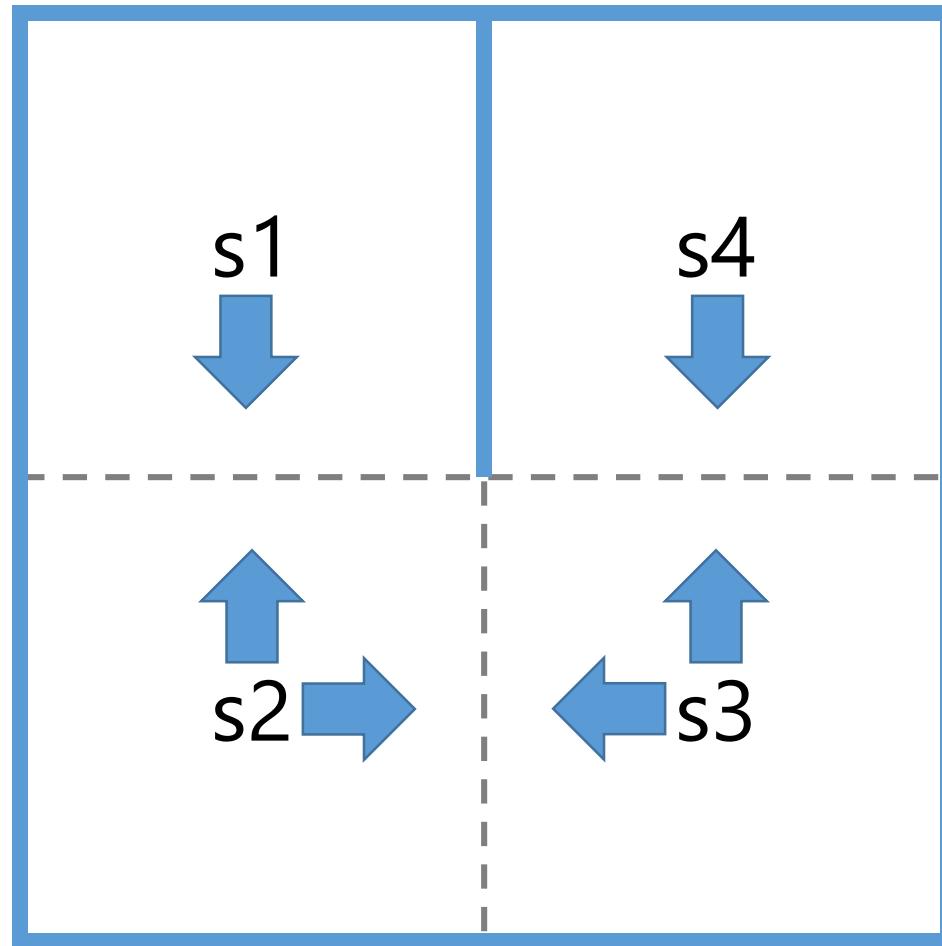
# Grid World



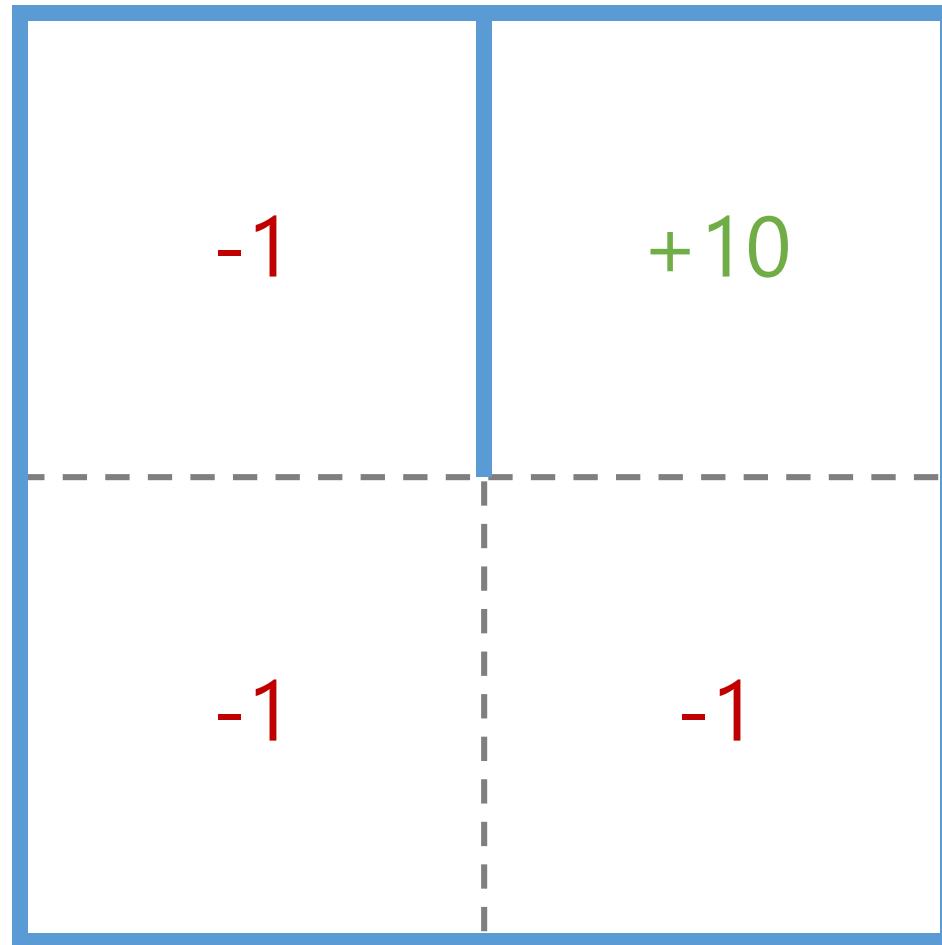
# States



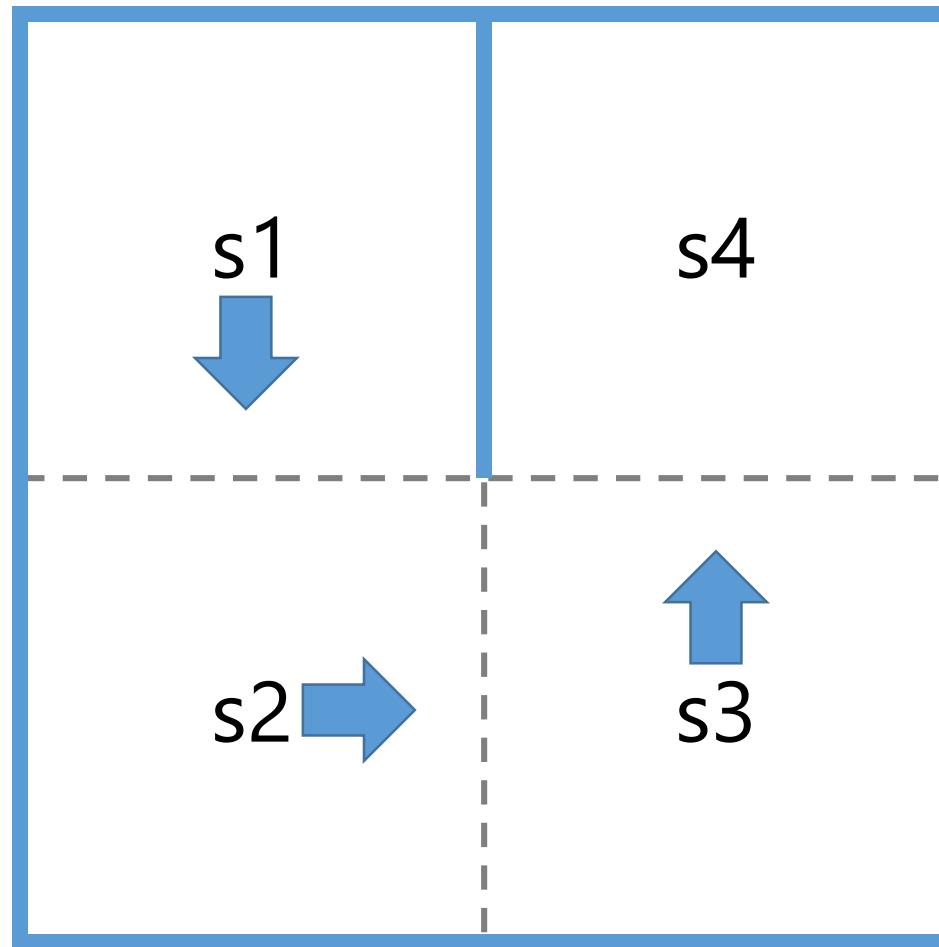
# Actions



# Rewards



# Optimal Policy



# Recap

States:  $s_1, s_2, s_3, s_4$

Actions: up, down, left, right

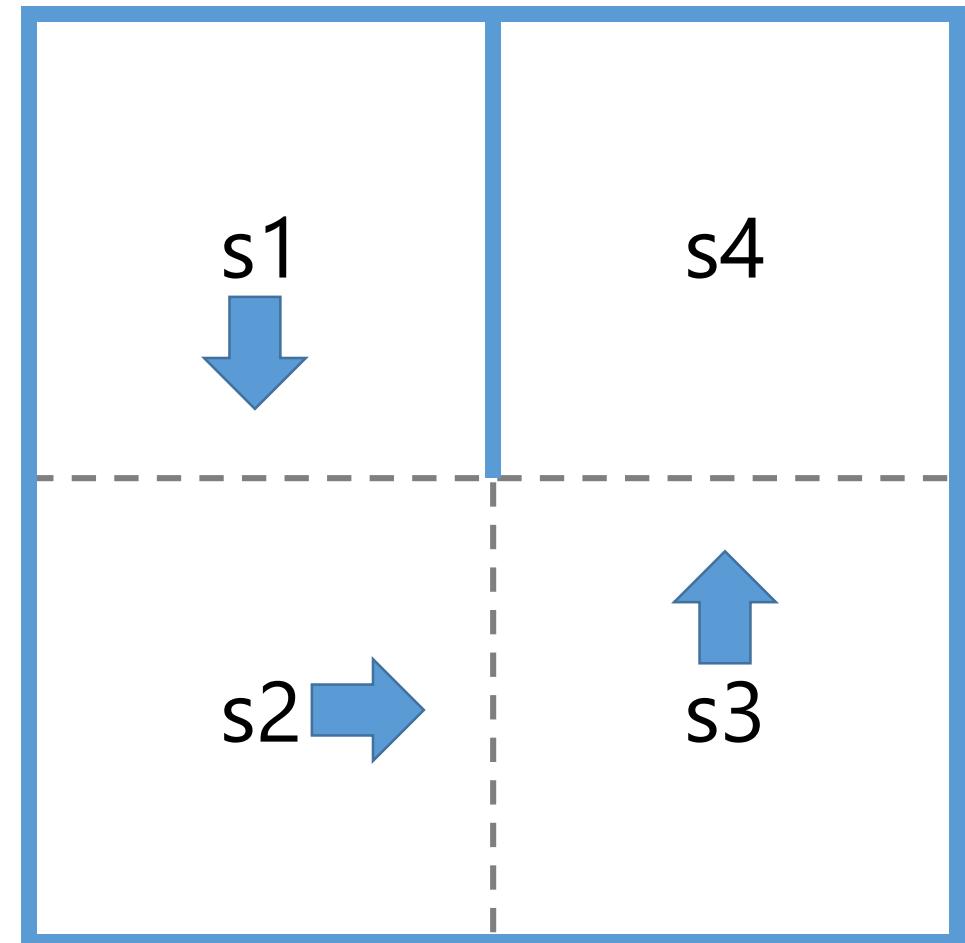
Rewards:  $s_1, s_3 = -1;$

$s_4 = 10$

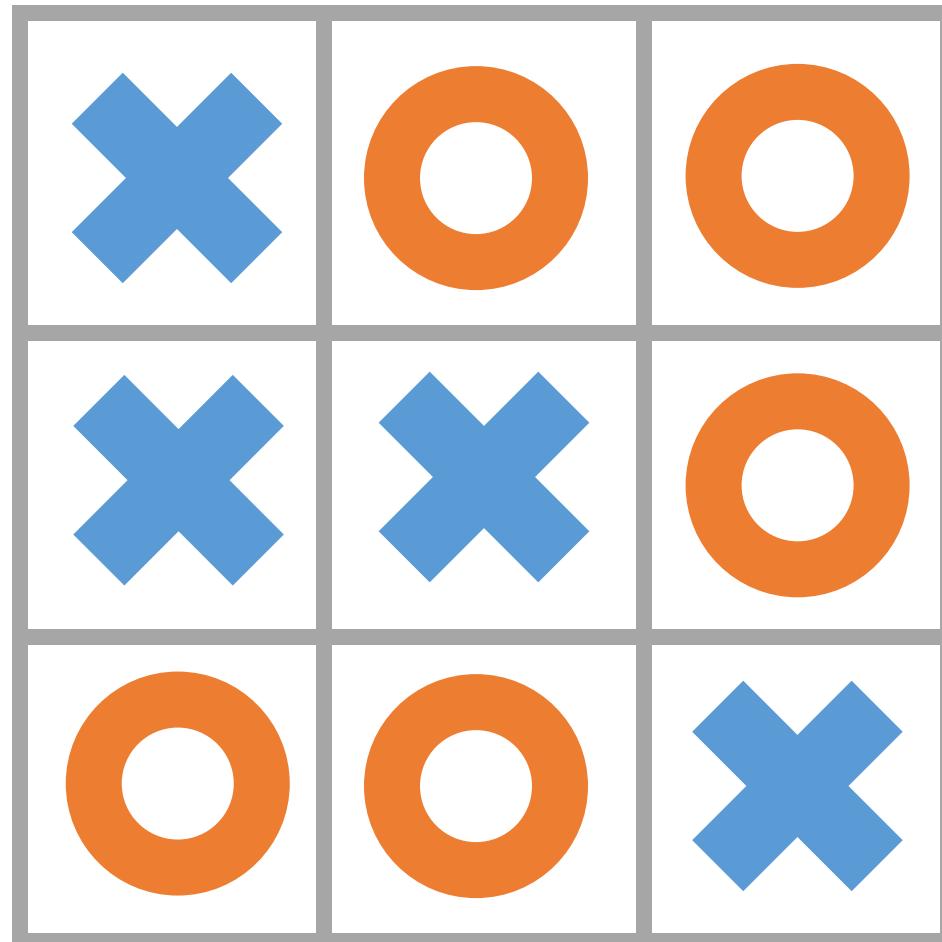
Policy:  $s_1 = \text{down}$

$s_2 = \text{right}$

$s_3 = \text{up}$



# Tic-Tac-Toe



# ML in Practice

What is the machine learning process?

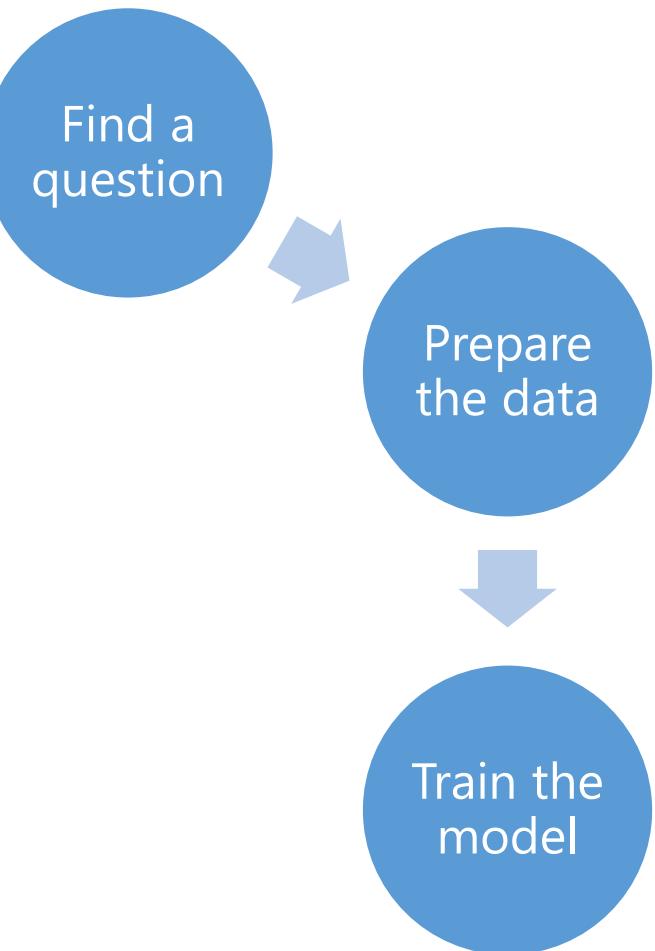


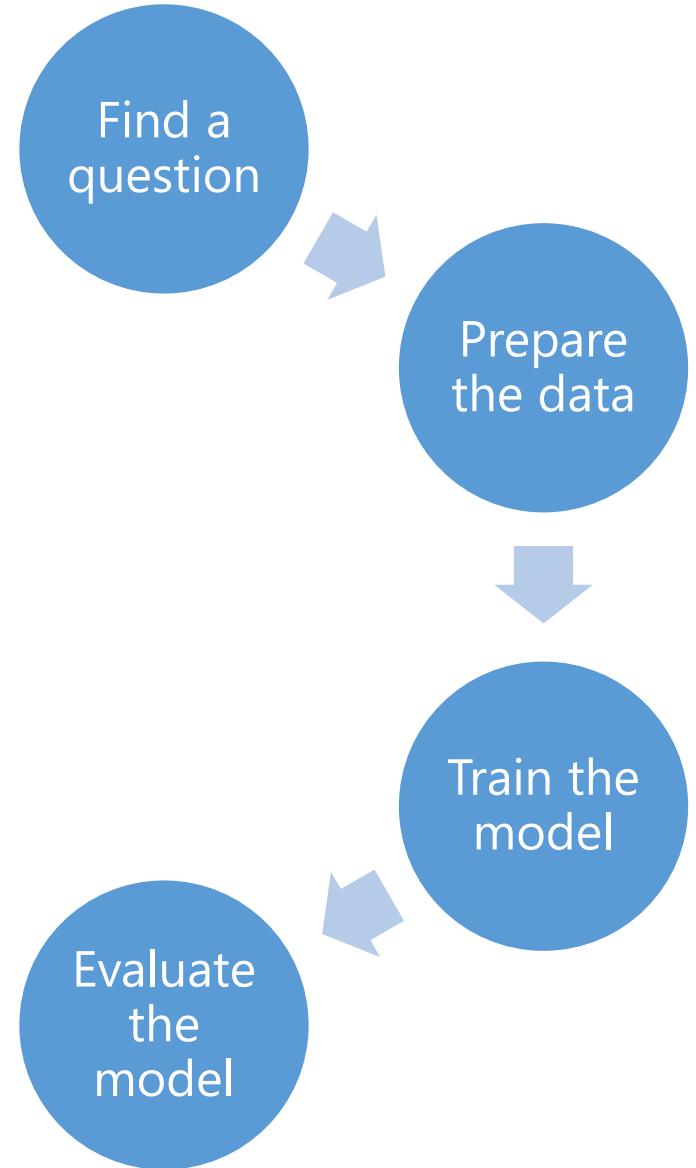
Find a  
question

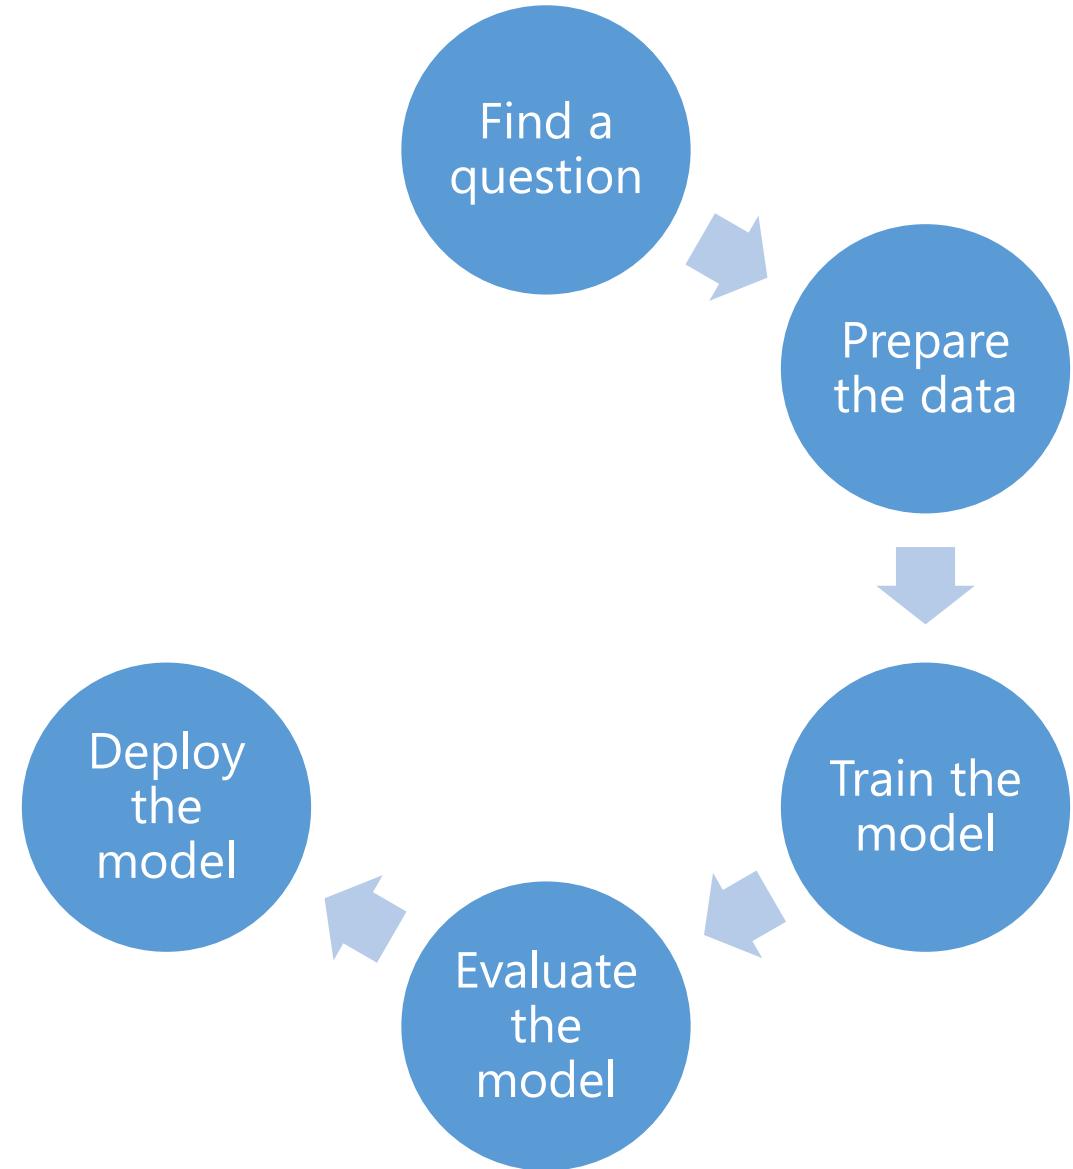
```
graph TD; A((Find a question)) --> B((Prepare the data))
```

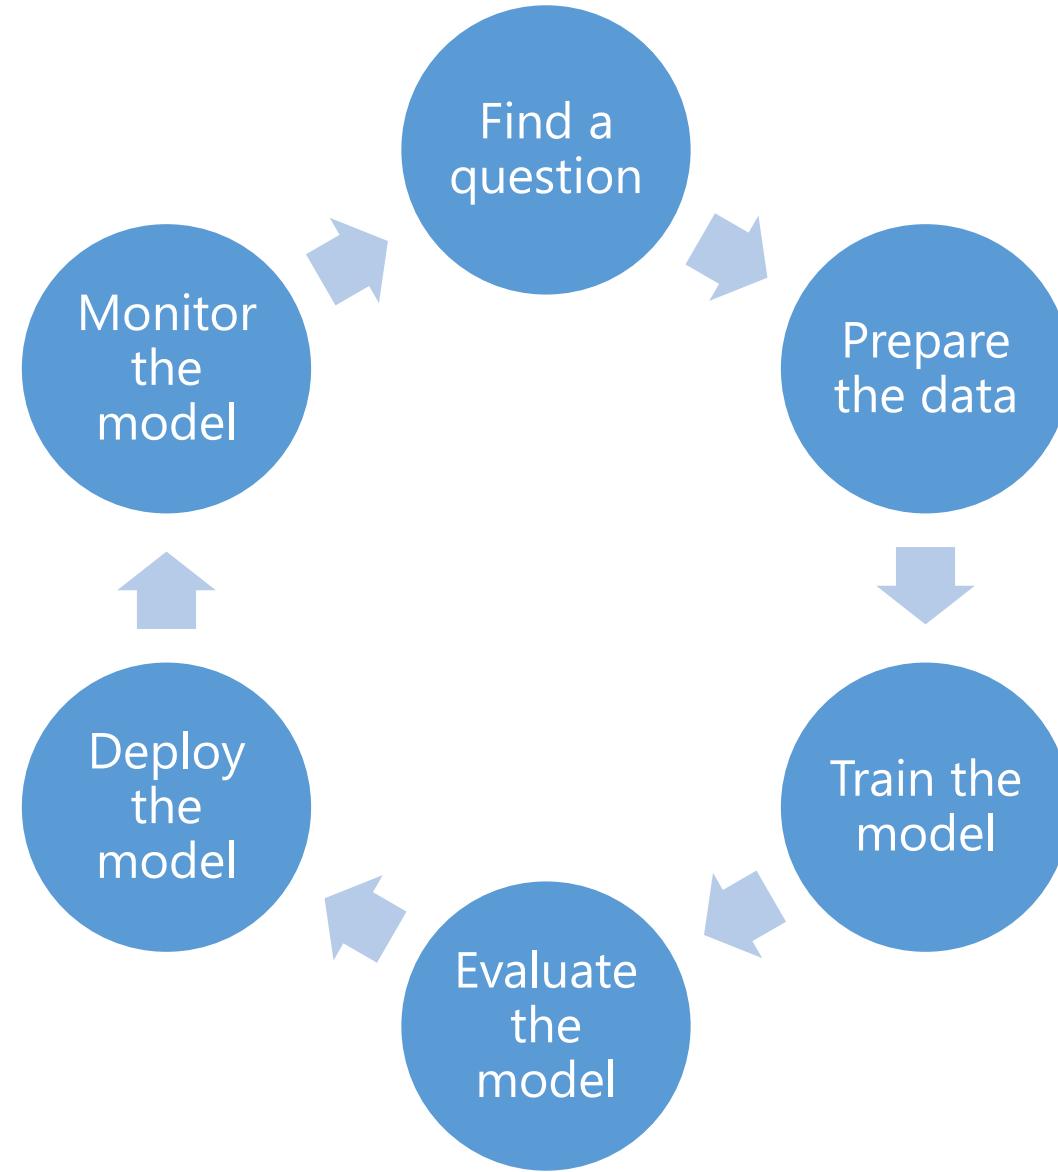
Find a  
question

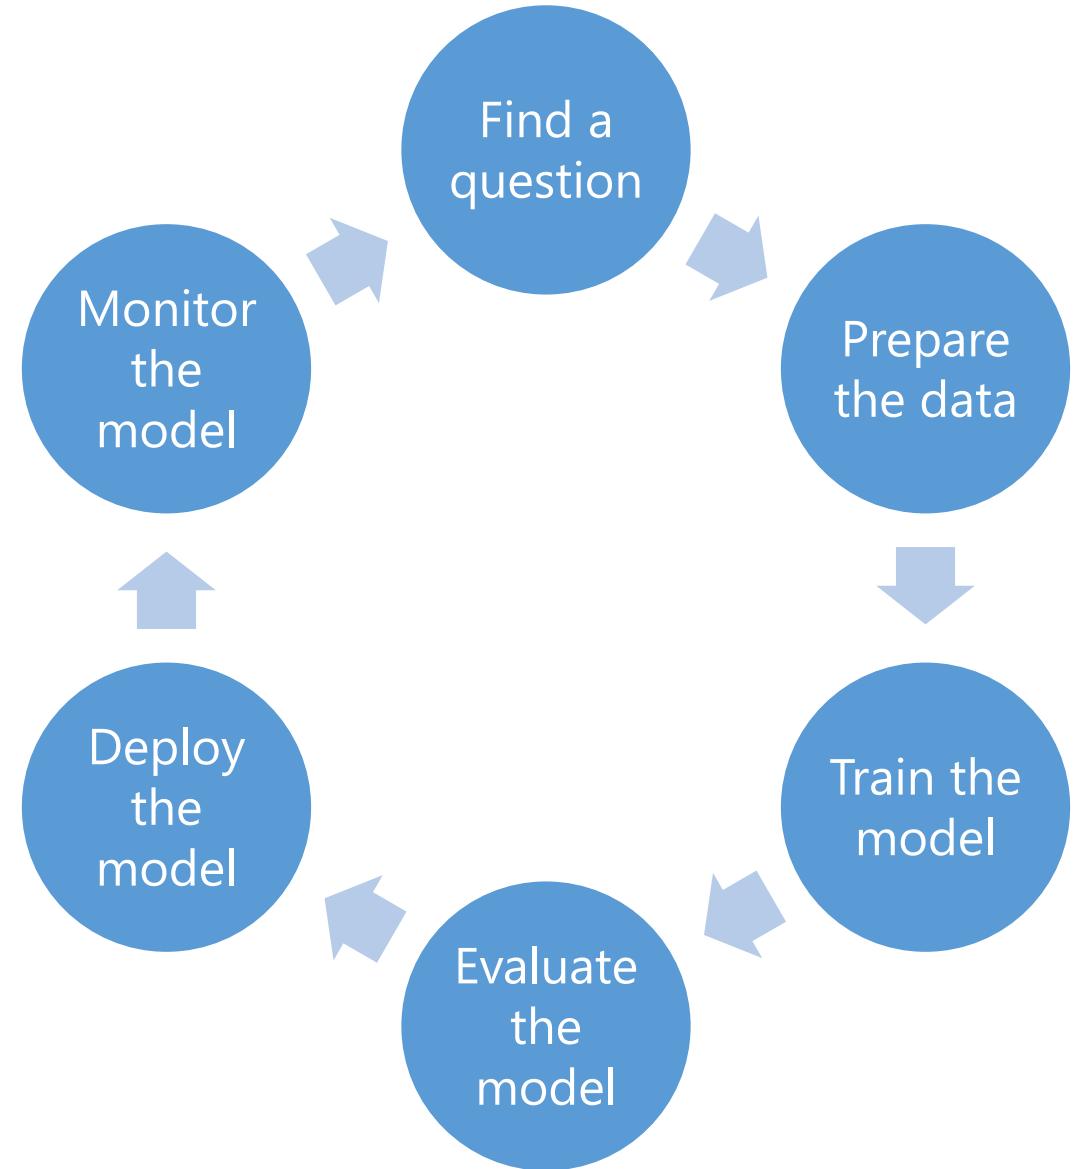
Prepare  
the data











Creating accurate and robust  
models is not easy

Find a  
question

Monitor  
the  
model

Prepare  
the data

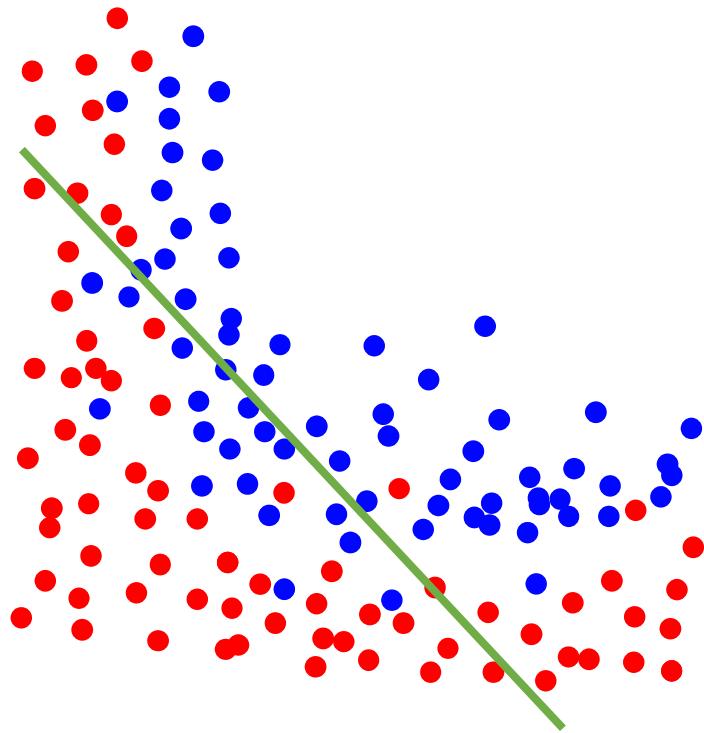
Deploy  
the  
model

Train  
the  
model

Evaluate  
the  
model

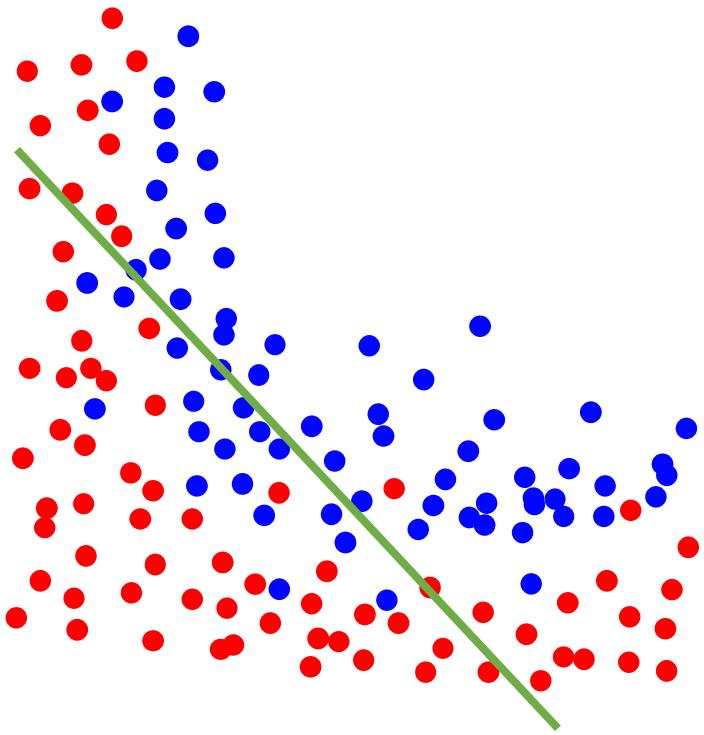
# Goodness of Fit

# Goodness of Fit

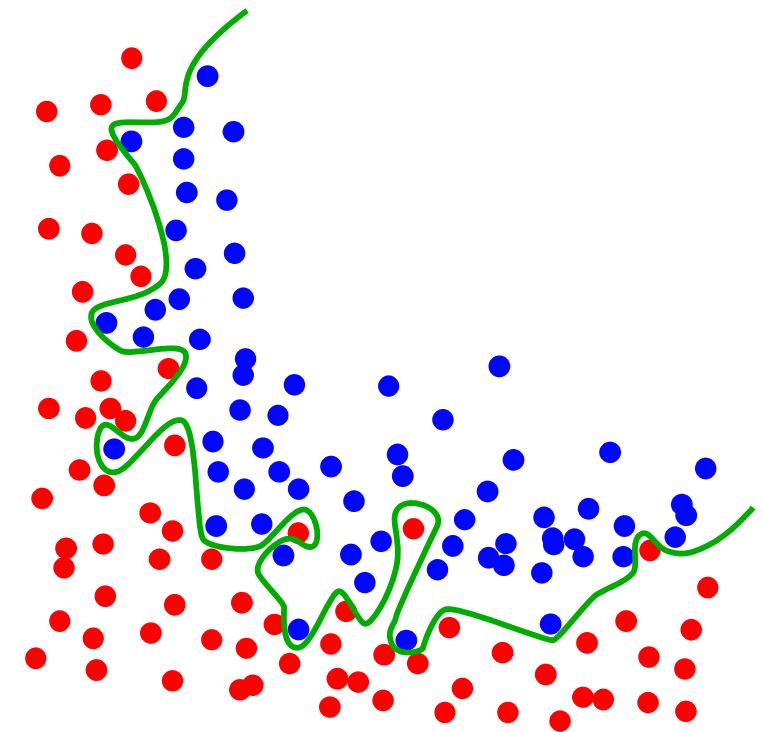


Underfit

# Goodness of Fit

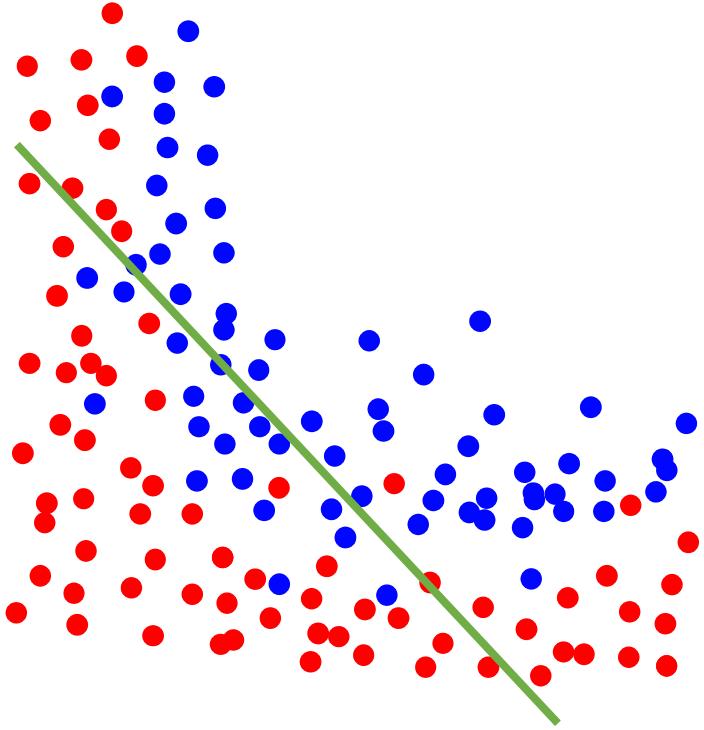


Underfit

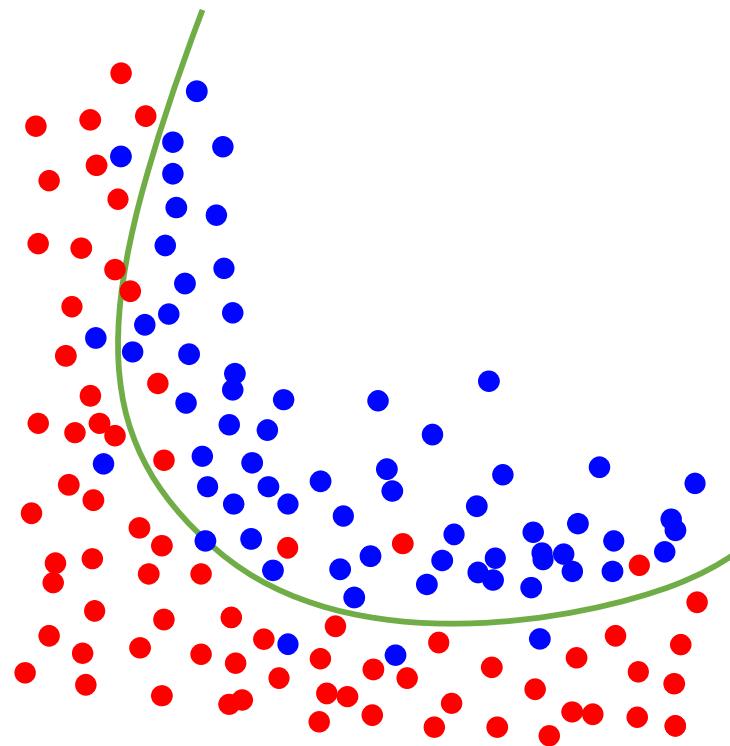


Overfit

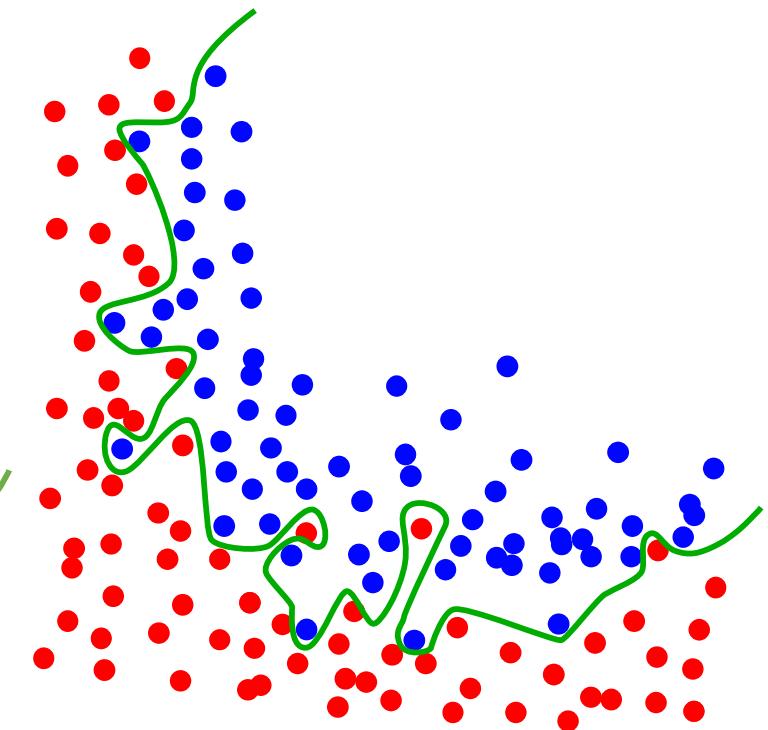
# Goodness of Fit



Underfit



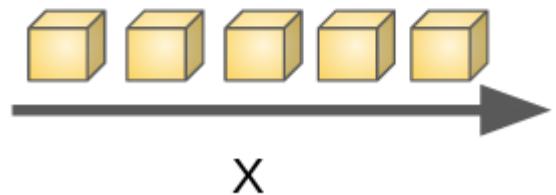
Good fit



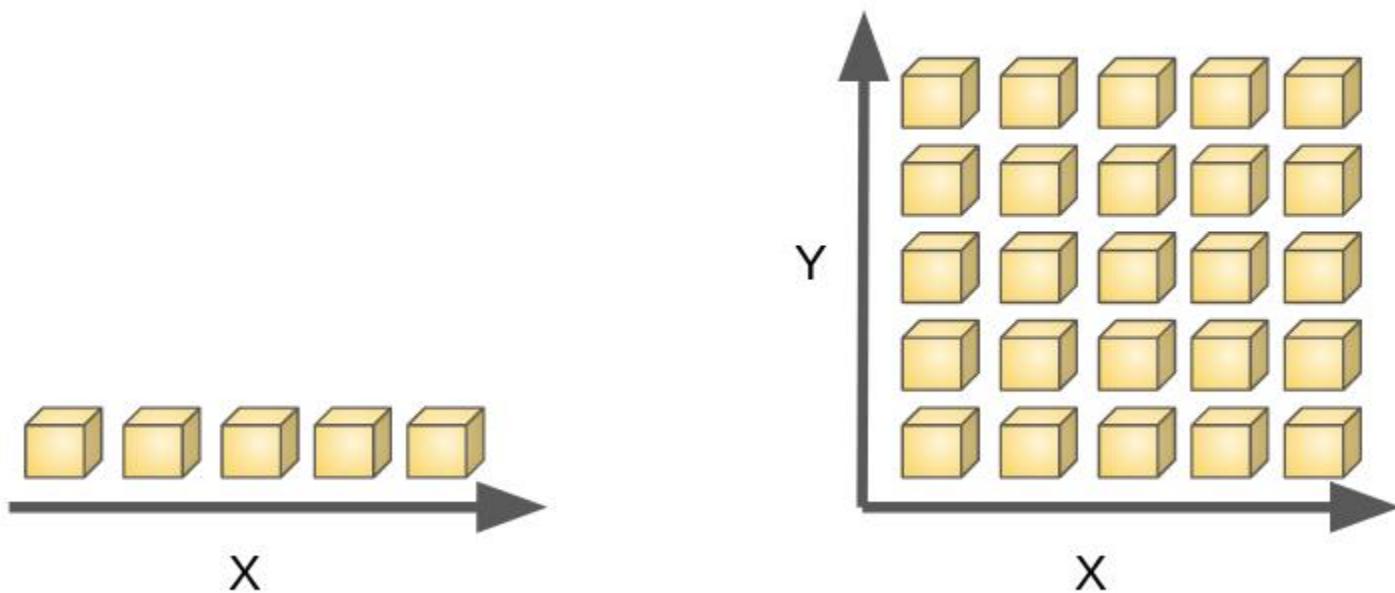
Overfit

# Curse of Dimensionality

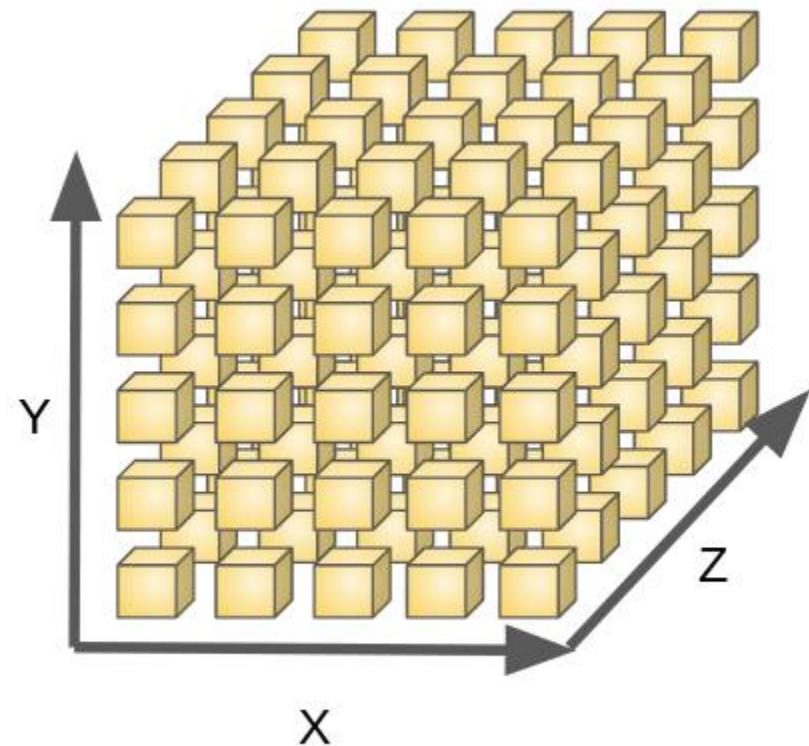
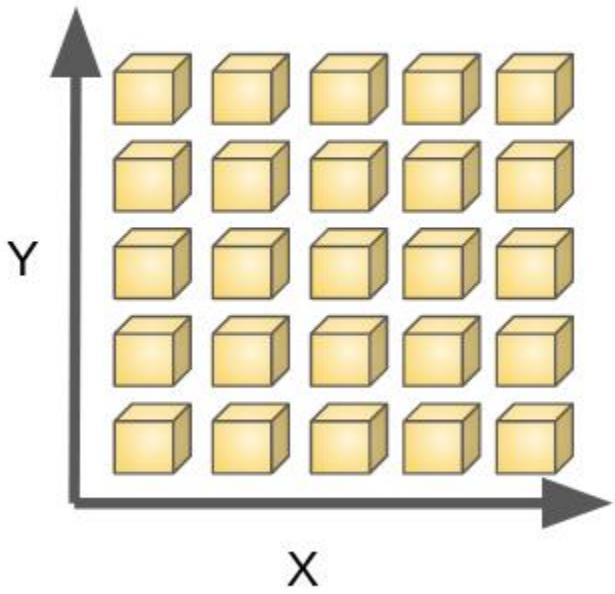
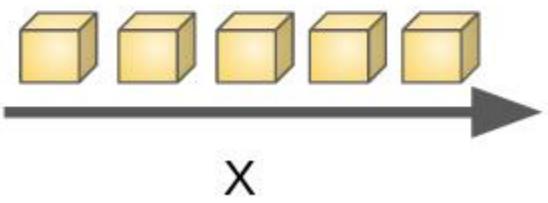
# Curse of Dimensionality



# Curse of Dimensionality



# Curse of Dimensionality



A police officer in uniform, wearing a cap with a badge, stands in a dark environment, holding a flashlight that illuminates the scene. Another officer is partially visible in the background. The scene is dimly lit, with the primary light source being the flashlight.

# Movie Break

# Demo 8 – ML in Practice

Goal: Predict survivors  
of the Titanic

# Lab 8A – ML in Practice (Easy)

Goal: Predict survivors  
of the Titanic

# Lab 8B – ML in Practice (Hard)

Goal: Predict risk in practice

# ML in Production

# How to Deploy to Production

- Deploy to web app (Shiny)
- Deploy to cloud (Azure ML)
- Deploy to server (ML Server)
- Deploy to any app (ONNX)



# Iris Species Predictor

Petal Length (cm)



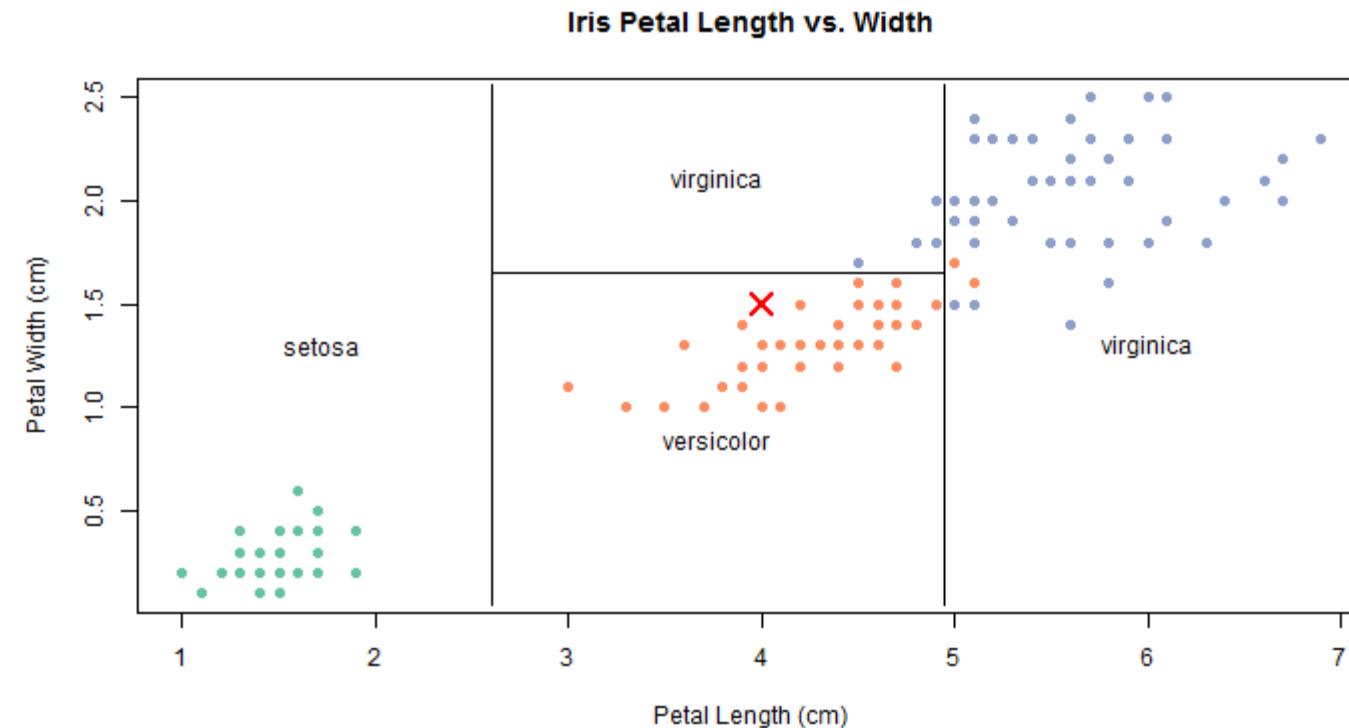
A horizontal slider for Petal Length (cm) ranging from 1 to 7. The value is set to 4.

Petal Width (cm)



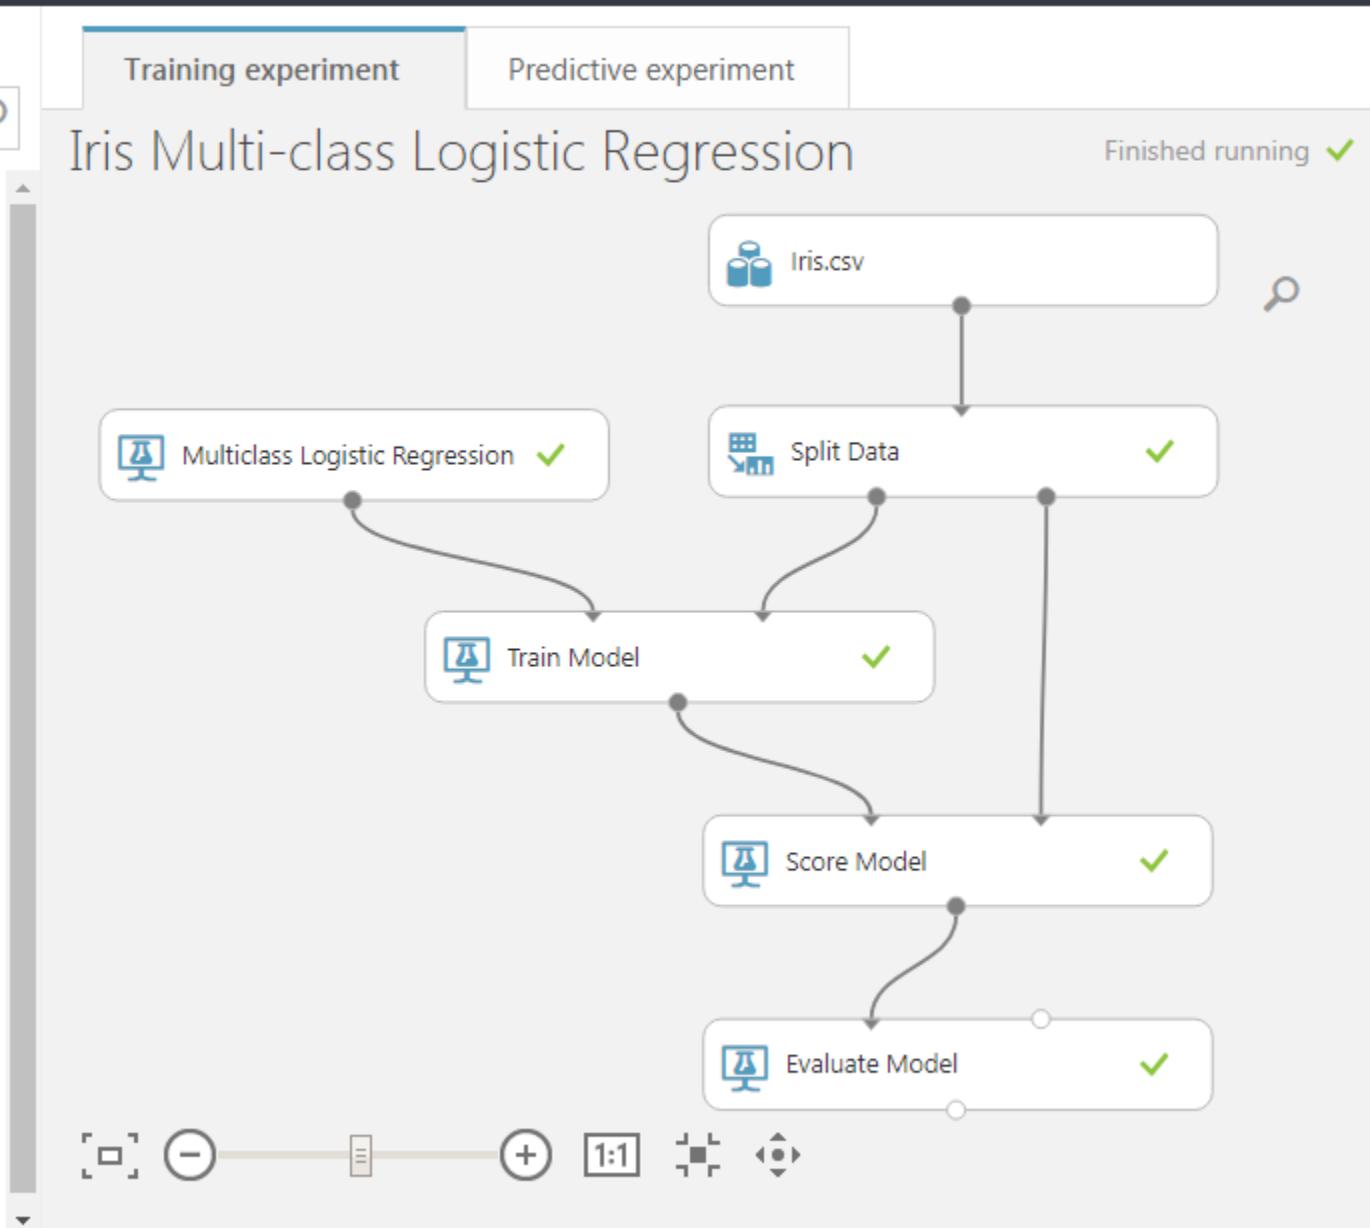
A horizontal slider for Petal Width (cm) ranging from 0 to 2.5. The value is set to 1.5.

The predicted species is versicolor



< Search experiment items 

- Saved Datasets
- Trained Models
- Data Format Conversions
- Data Input and Output
- Data Transformation
- Feature Selection
- Machine Learning
- OpenCV Library Modules
- Python Language Modules
- R Language Modules
- Statistical Functions
- Text Analytics
- Time Series
- Web Service
- Deprecated



Properties Project

Experiment Properties

START TIME	3/17/20...
END TIME	3/17/20...
STATUS CODE	Finished
STATUS DETAILS	None

Summary

Enter a few sentences describing your experiment (up to 140 characters).

Description

Enter the detailed description for your experiment.

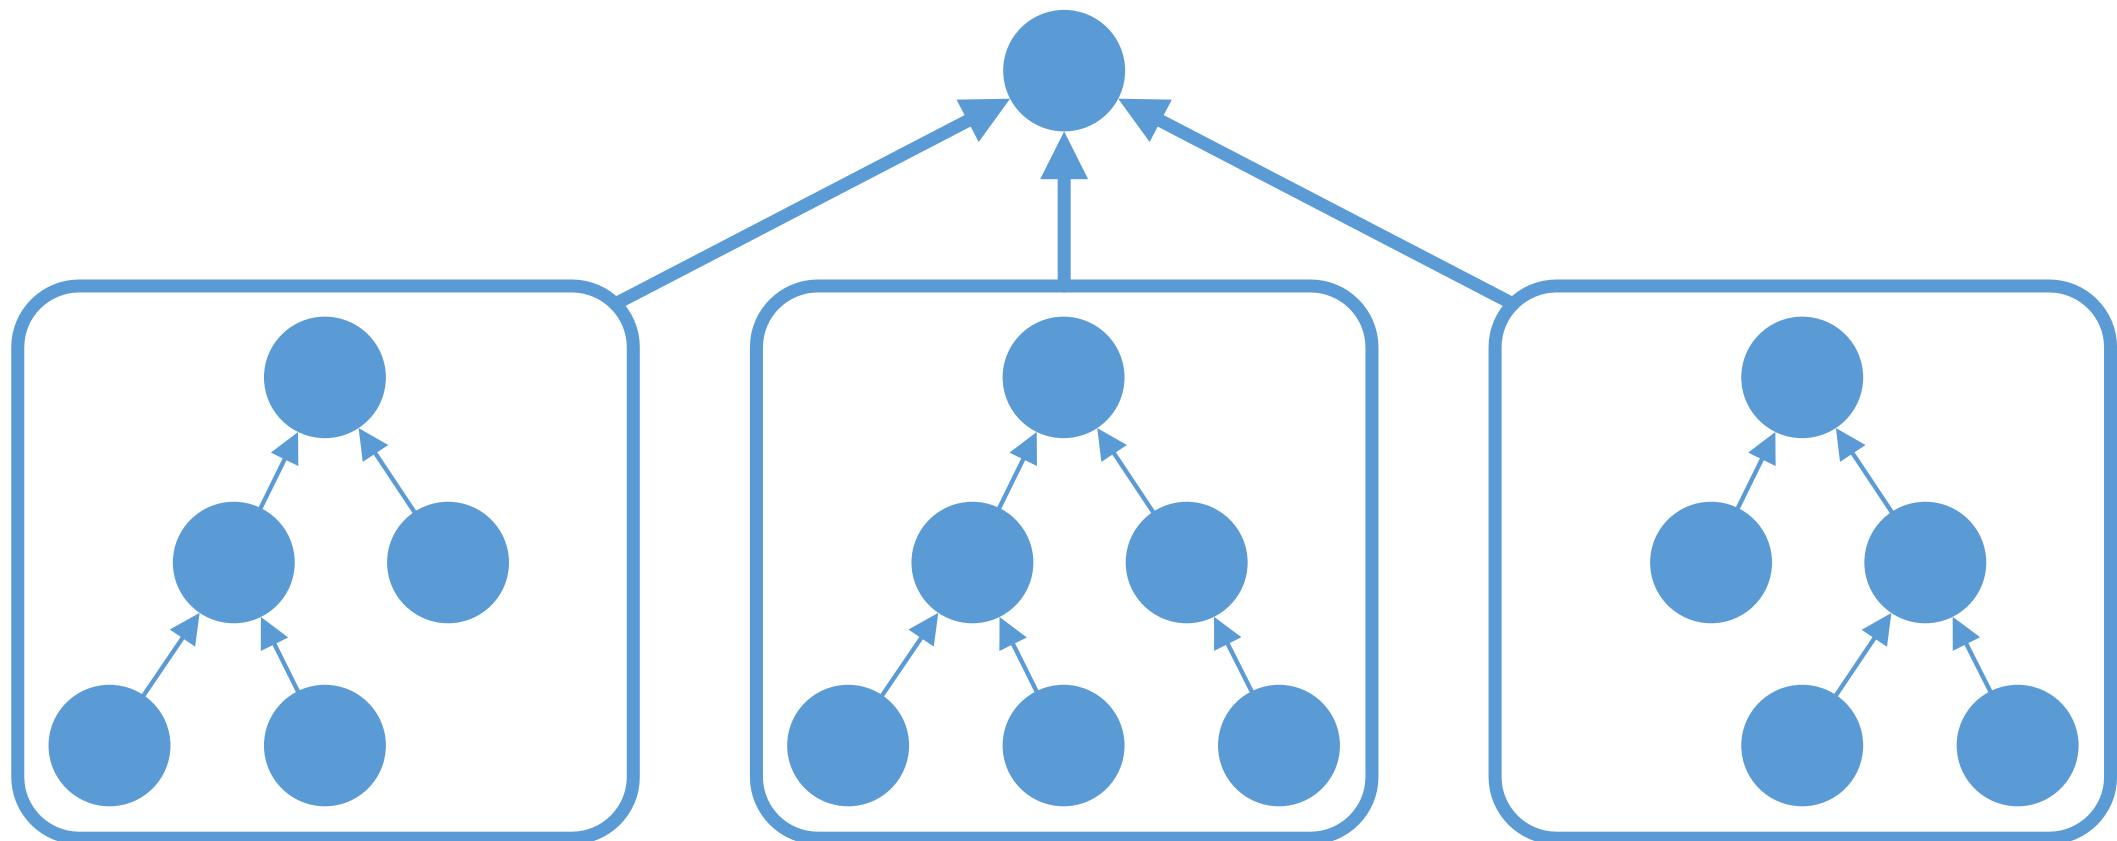
Quick Help

# Conclusion

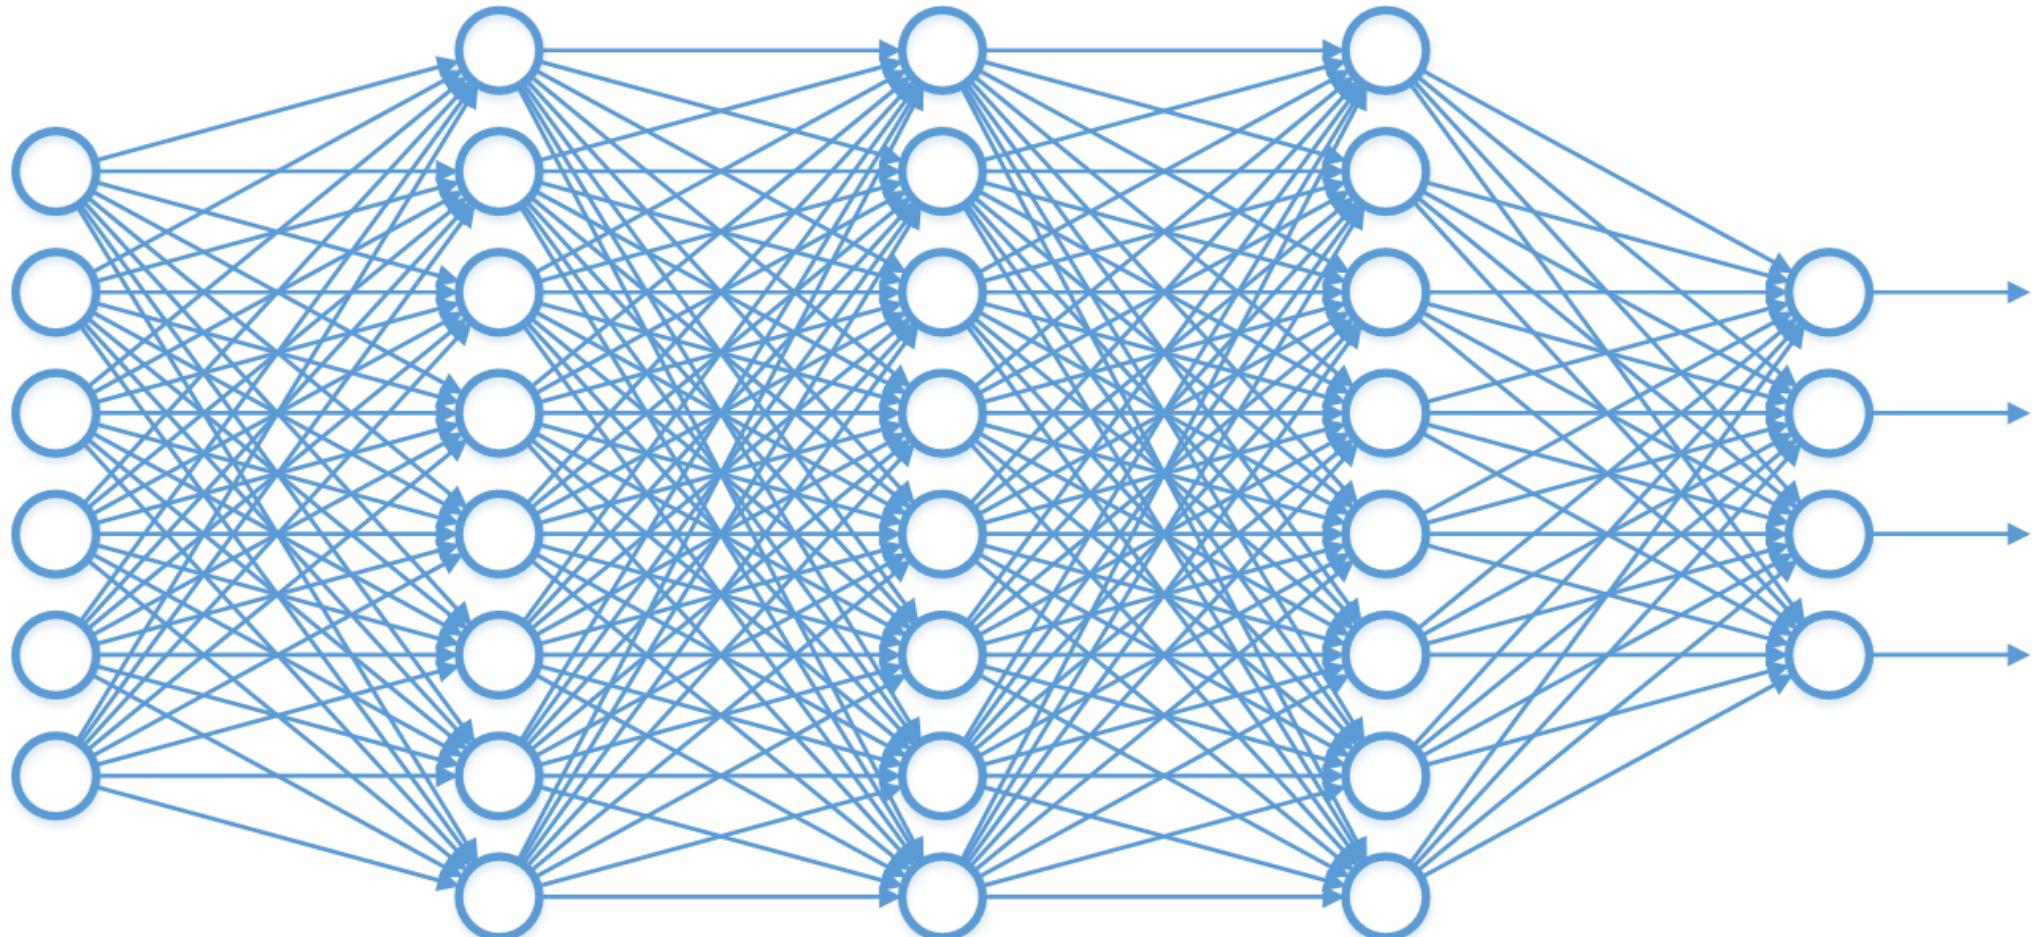


This is just the tip of the iceberg!

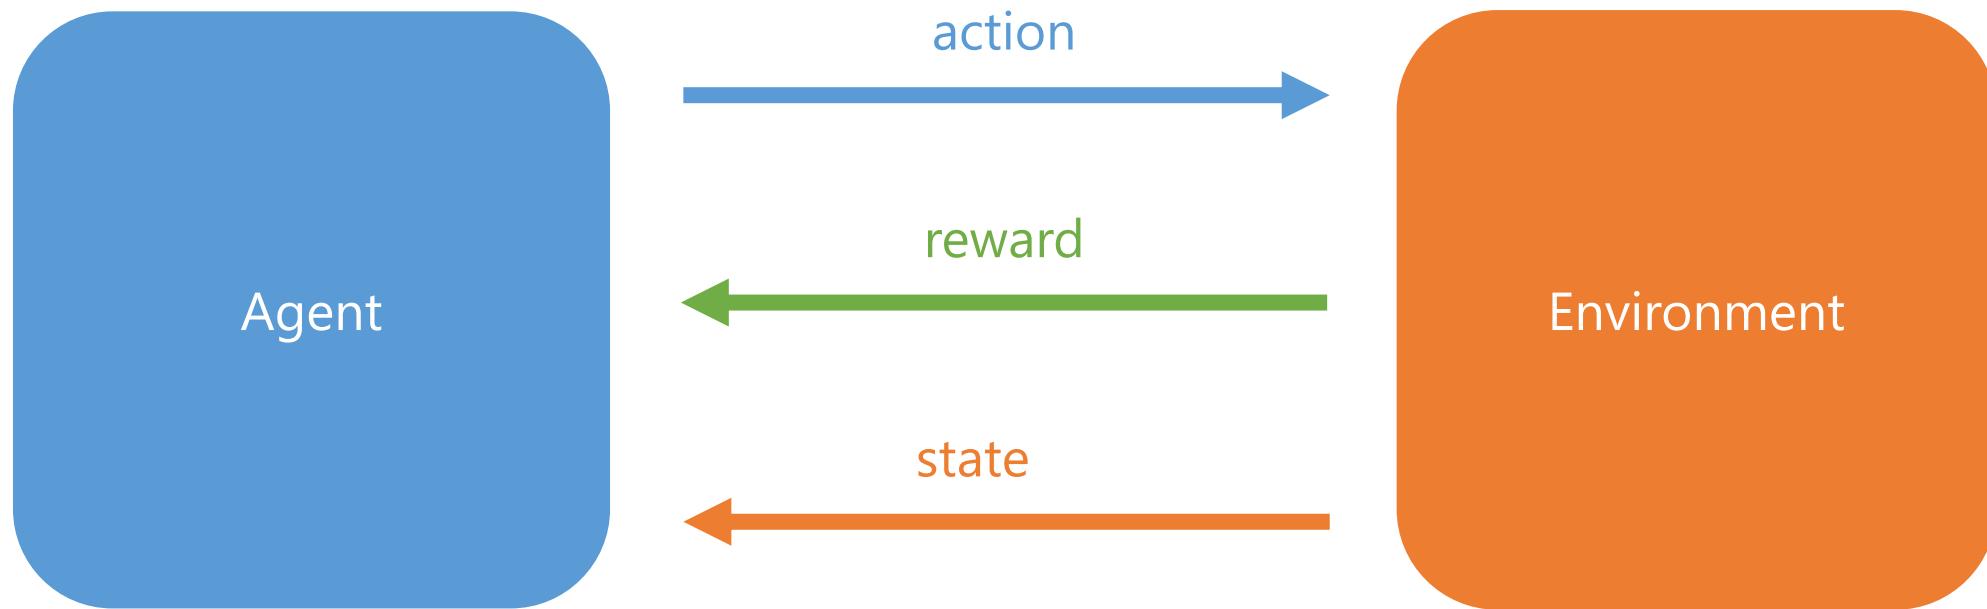
# Ensemble Learning



# Deep Learning



# Reinforcement Learning



Where do we go from here?

# Where to Go Next

Data Camp: <https://www.datacamp.com>

Pluralsight: <https://www.pluralsight.com>

Coursera: <https://www.coursera.org>

# Pluralsight Courses

Data Science with R

Data Science: The Big Picture

Deep Learning: The Big Picture

Exploratory Data Analysis with R

Data Visualization with R (3-part)



<https://www.pluralsight.com/authors/matthew-renze>

## News

### 2017-08-25 - Invitation to Speak at Devoxx Morocco

Very excited to announce that I've been invited to give a keynote in Casablanca at [Devoxx Morocco](#) in November. My keynote presentation will be on [Artificial Intelligence](#).



### 2017-08-16 - Invitation to Speak at Microsoft Ignite

I've been invited to speak at [Microsoft Ignite](#) in Orlando, Florida in September. This will be my first time speaking at Ignite. Talks will include both Data Science and Machine Learning with R.



Matthew is a data science consultant, author for [Pluralsight](#), international public speaker, a [Microsoft MVP](#), [ASPIInsider](#), and open-source software contributor.

### 2017-08-14 - Dev on Fire Interview

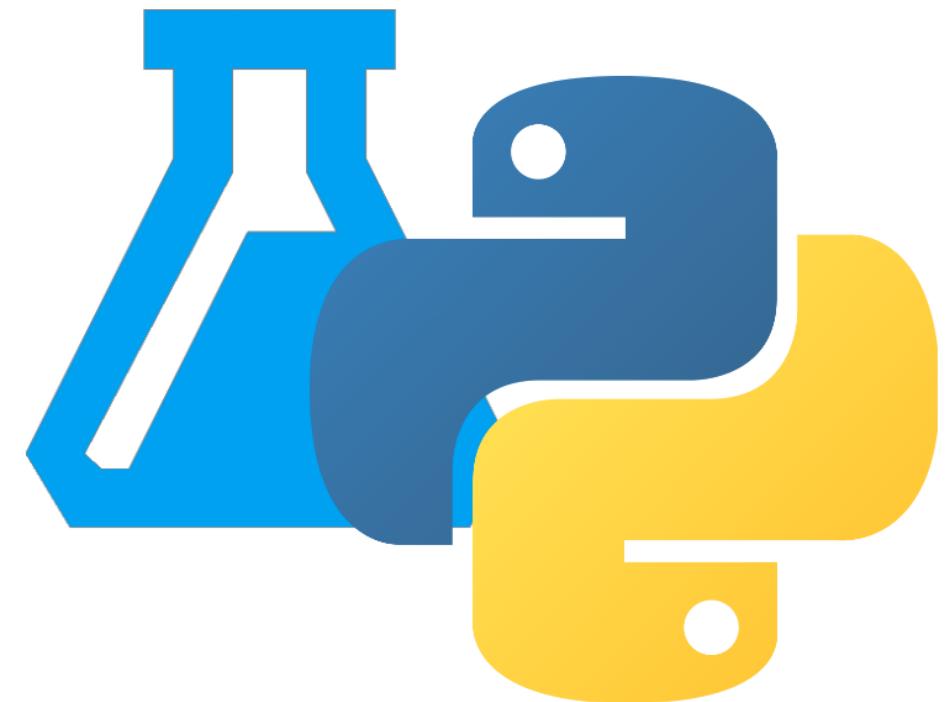
# Feedback

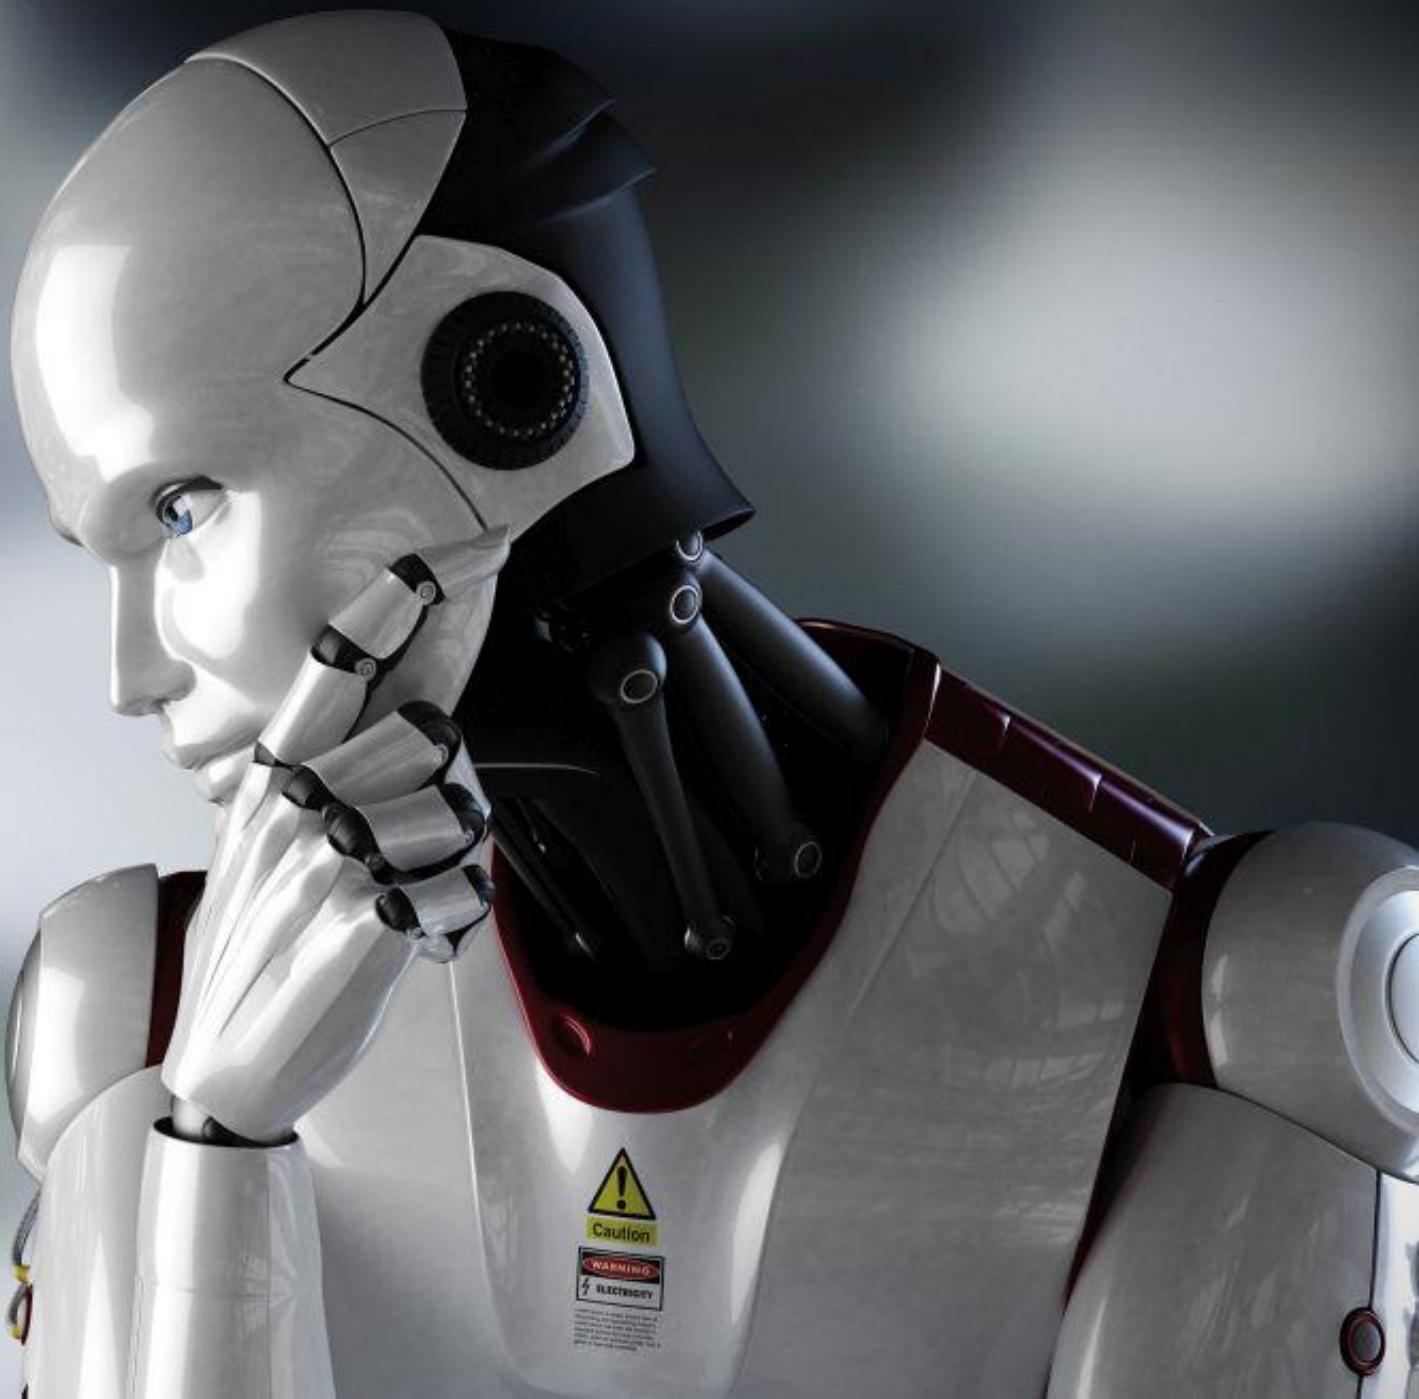
Very important to me!  
What did you like?  
What could I improve?

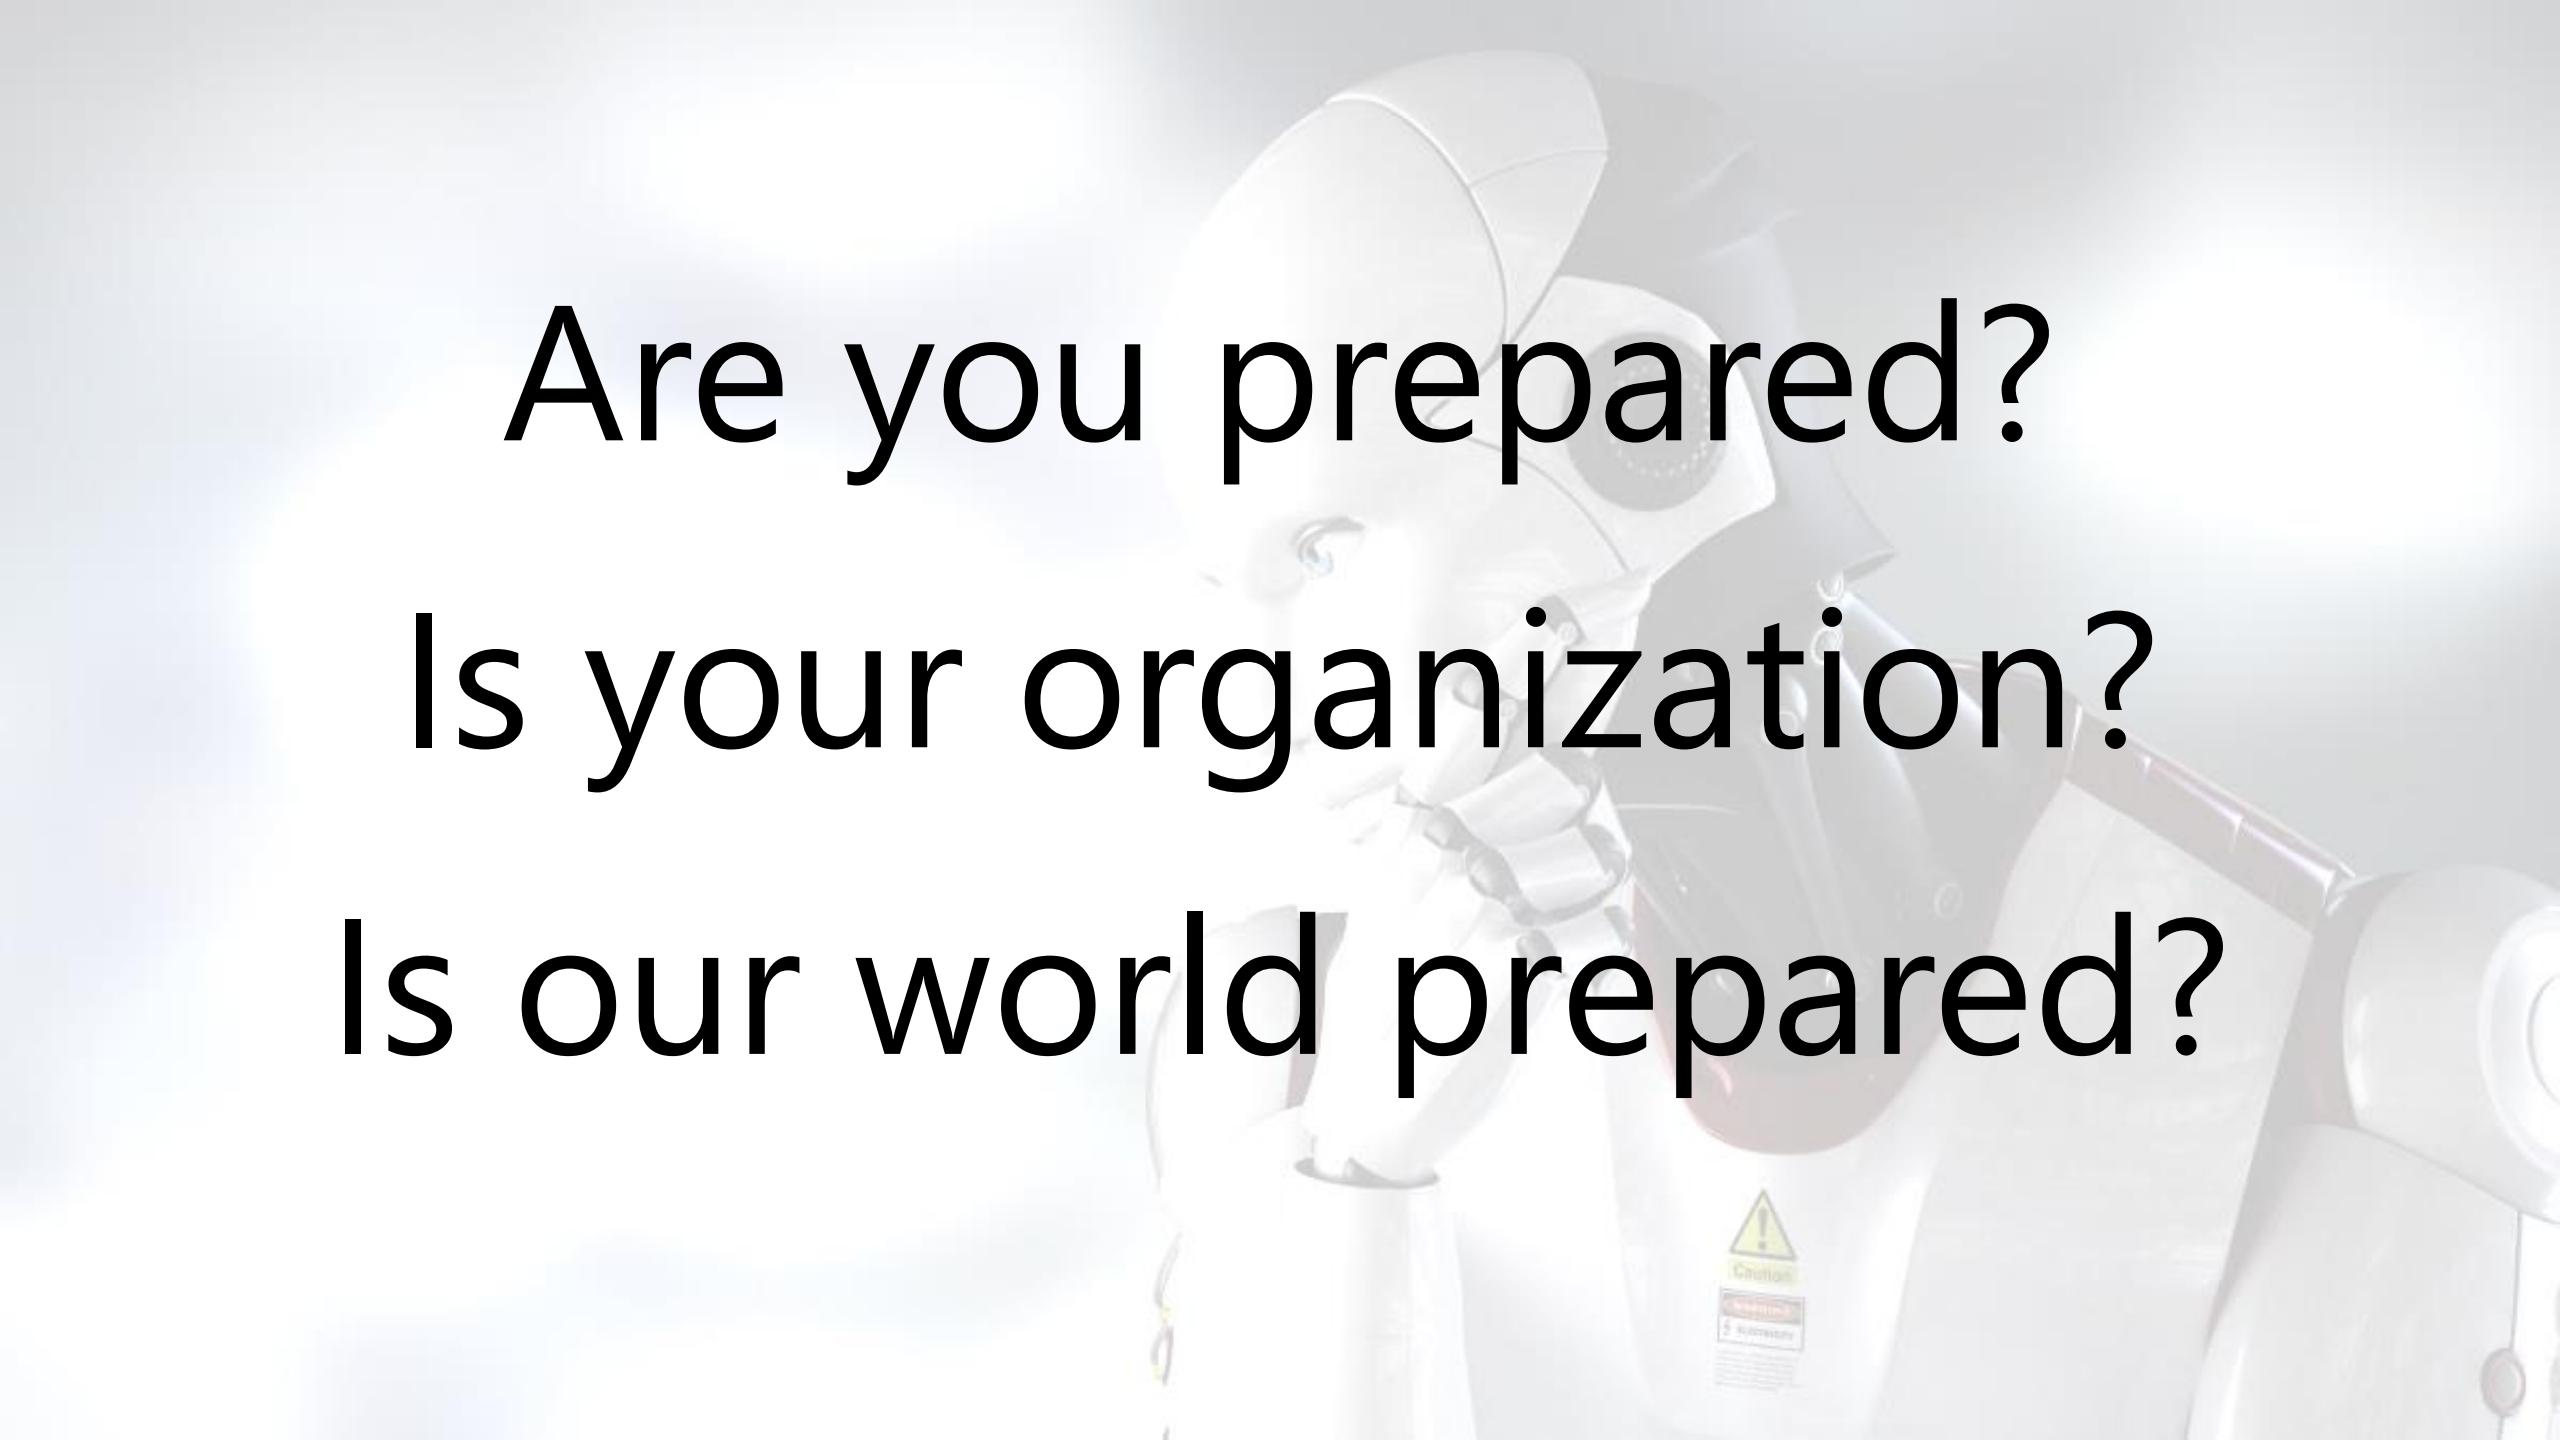


# Conclusion

1. Intro to ML and Python
2. Classification
3. Regression
4. Clustering
5. ML in Practice



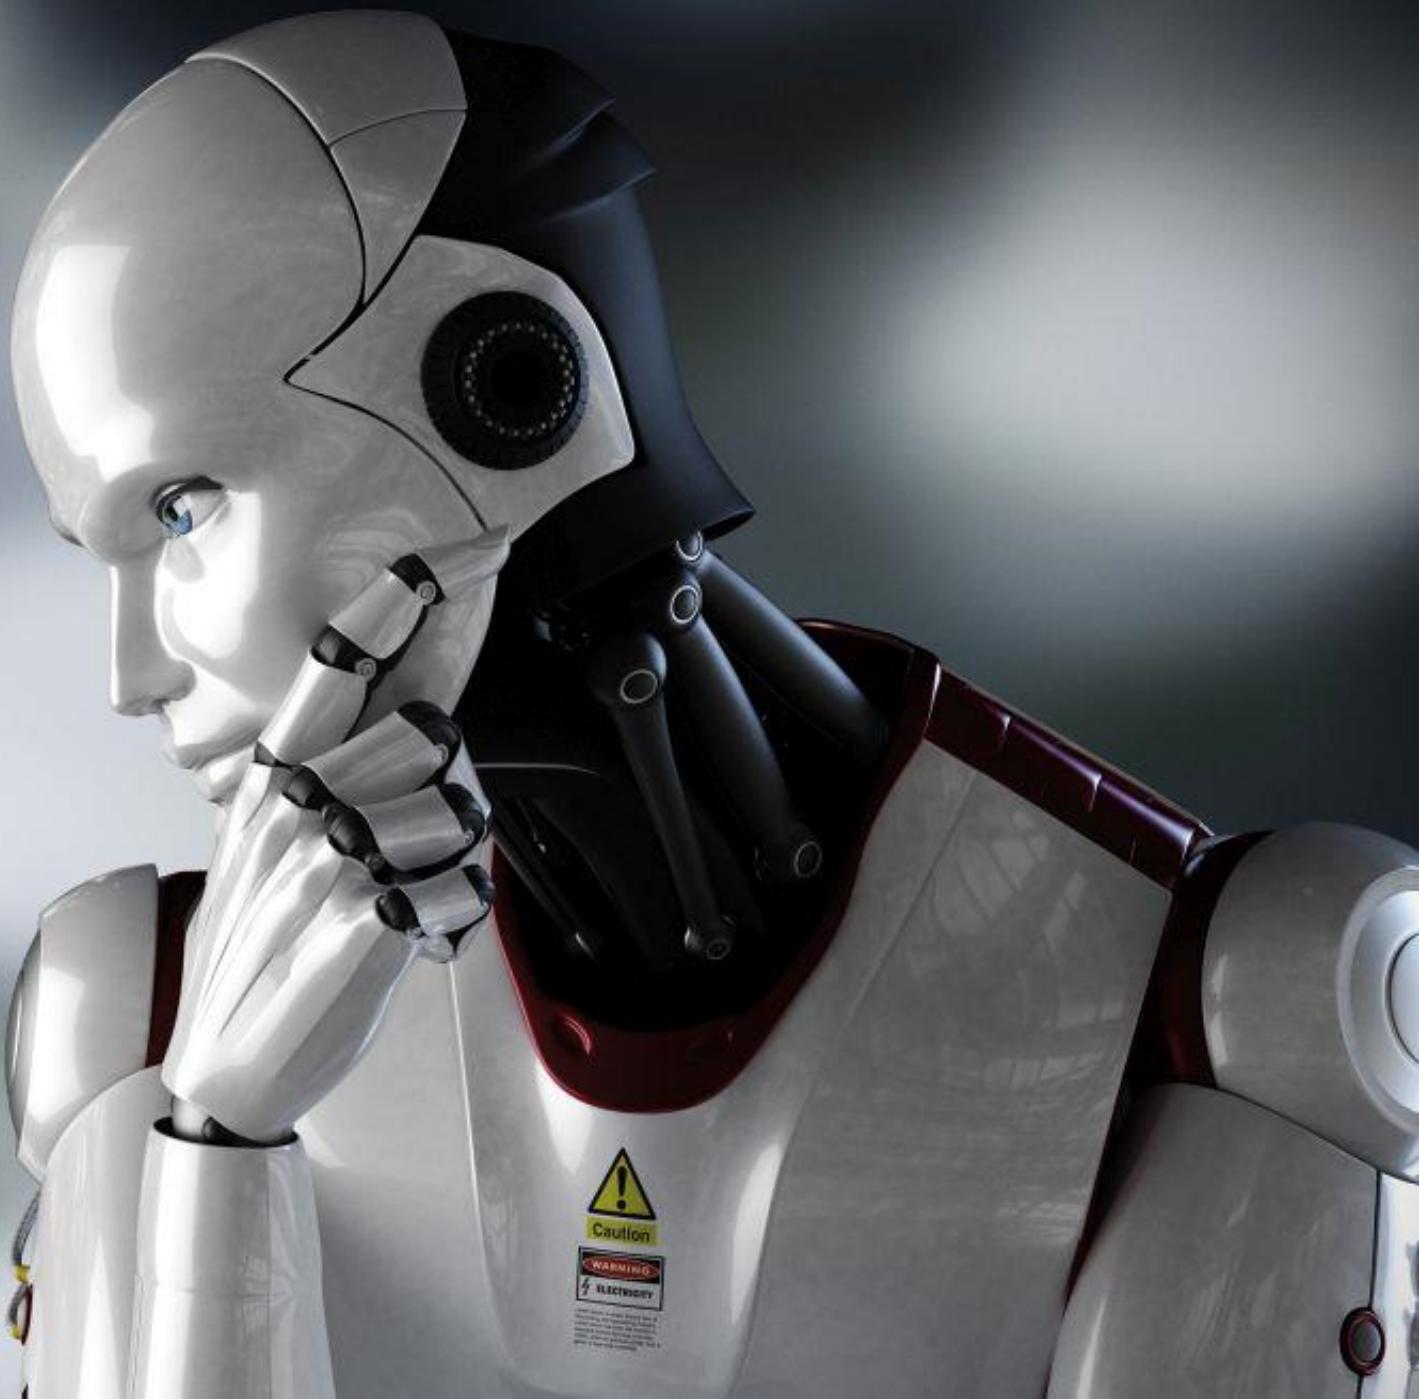




Are you prepared?

Is your organization?

Is our world prepared?



# Contact Info

Matthew Renze

Data Science Consultant  
Renze Consulting

Twitter: [@matthewrenze](https://twitter.com/matthewrenze)

Email: [info@matthewrenze.com](mailto:info@matthewrenze.com)

Website: [www.matthewrenze.com](http://www.matthewrenze.com)



Thank You! : )